

# Rensselaer 2020 REU Notebook

Alex Striff

May to July 2020

## Contents

<b>1</b>	<b>Project description (May 27, 2020)</b>	<b>3</b>
<b>2</b>	<b>Getting started (May 27, 2020)</b>	<b>3</b>
<b>3</b>	<b>Intensity-level entropy</b>	<b>3</b>
<b>4</b>	<b>Effect of smoothing on intensity-level entropy</b>	<b>4</b>
4.1	Natural image . . . . .	4
4.2	Random pixel values . . . . .	7
4.2.1	Beware: GIGO . . . . .	8
4.3	Comparing different levels of smoothing . . . . .	10
<b>5</b>	<b>Local metrics (May 28, 2020)</b>	<b>14</b>
5.1	Induced metrics . . . . .	14
<b>6</b>	<b>Kernels</b>	<b>15</b>
<b>7</b>	<b>Boxcar intensity-level entropy</b>	<b>16</b>
7.1	Standard deviation . . . . .	17
7.2	Intensity entropy . . . . .	19
7.3	Replace surprisal with other functions . . . . .	20
7.4	Intensity entropy on disjoint blocks . . . . .	22
<b>8</b>	<b>Fractal dimensions (May 29, 2020)</b>	<b>23</b>

<b>9</b>	<b>Fractal dimension regression</b>	<b>24</b>
9.1	Box-counting dimension . . . . .	25
9.2	Information dimension . . . . .	31
<b>10</b>	<b>Probability and inference (May 30, 2020)</b>	<b>36</b>
<b>11</b>	<b>Ising images (June 1, 2020)</b>	<b>37</b>
<b>12</b>	<b>Ising images</b>	<b>37</b>
12.1	Standard Ising (on a torus) . . . . .	38
12.2	Image-edge Ising . . . . .	39
12.3	Image-metric Ising . . . . .	43
12.3.1	Unrestricted swapping motion . . . . .	43
12.3.2	Nearest-neighbor swapping motion . . . . .	46
<b>13</b>	<b>Statistical Mechanics of Images (June 3, 2020)</b>	<b>48</b>
<b>14</b>	<b>Thermodynamic quantities for images from a microscopic model (June 4, 2020)</b>	<b>50</b>
14.1	Quantum filled-site model (FSM) . . . . .	50
14.2	Observables and thermodynamic state variables . . . . .	50
<b>15</b>	<b>Progress summary (from beginning) (June 6, 2020)</b>	<b>51</b>
<b>16</b>	<b>Simulations for canonical ensemble averages (June 10, 2020)</b>	<b>52</b>
<b>17</b>	<b>The Wang-Landau algorithm (density of states)</b>	<b>52</b>
17.1	The 2D Ising model . . . . .	54
17.1.1	Calculating canonical ensemble averages . . . . .	57
17.2	Thermal calculations on images . . . . .	60
17.2.1	Calculating canonical ensemble averages . . . . .	66
<b>18</b>	<b>Progress summary (June 12, 2020)</b>	<b>69</b>
<b>19</b>	<b>Description of MAXENT (June 13, 2020)</b>	<b>70</b>
<b>20</b>	<b>Maximum-entropy reconstruction</b>	<b>70</b>
20.1	Example: 1D Point from Gaussian . . . . .	71
20.2	Example: Image from PSF convolution (measurement) . . . . .	72

21	“Greedy” painting-like pictures (June 14, 2020)	76
22	Greedy Cubism	76

## 1 Project description

May 27, 2020 The aim of this REU project is to quantify the information present in images by the principled application of methods from statistical physics. The approach is to find a suitable notion of entropy which captures the salient features of particular kinds of images. We will consider a variety of features motivated by intuition or domain knowledge, and then move to machine learning as a tool for discovering other features.

## 2 Getting started

May 27, 2020 The initial goal is to characterize the most naïve calculation, which I’ll call the *intensity entropy*. This does *not* take into account the spatial correlation of pixels in an image.

## 3 Intensity-level entropy

Given a discrete random variable  $X$  with support  $\mathcal{X}$ , the *Shannon entropy* is

$$H = \sum_{x \in \mathcal{X}} -P(x) \ln P(x).$$

The *intensity-level entropy* is the Shannon entropy of the empirical distribution of intensity values.

```

1 import numpy as np

1 def shannon_entropy(h):
2     """The Shannon entropy in bits"""
3     return -sum(p*np.log2(p) if p > 0 else 0 for p in h)
4
5 def intensity_distribution(data):
6     """The intensity distribution of 8-bit `data`."""
7     hist, _ = np.histogram(data, bins=range(256+1), density=True)
8     return hist

```

```

9
10 def intensity_entropy(data):
11     """The intensity-level entropy of 8-bit image data"""
12     return shannon_entropy(intensity_distribution(data))
13
14 def intensity_expected(f, data):
15     """The intensity-distribution expected value of `f`."""
16     return sum(p*f(p) for p in intensity_distribution(data))

```

## 4 Effect of smoothing on intensity-level entropy

```

1 import numpy as np
2 import numpy.linalg as linalg
3 import matplotlib.pyplot as plt
4 from PIL import Image, ImageFilter, ImageOps
5 from src.utilities import *
6 from src.intensity_entropy import *

```

### 4.1 Natural image

```

1 img = ImageOps.grayscale(Image.open('test.jpg'))
2 scale = max(np.shape(img))
3 data = np.array(img)
4 img

```



```
1 intensity_entropy(img)
```

7.51132356216608

The problem with the intensity entropy is that it is usually near maximum (8 bits for these grayscale images).

```
1 def intensity_blur(img, scales, display=True):
2     scale = max(np.shape(img))
3
4     results = []
5     for k in scales:
6         simg = img.filter(ImageFilter.GaussianBlur(k * scale))
7         data = np.array(simg)
8         ihist, ibins = np.histogram(data, bins=range(256+1), density=True)
9         S = shannon_entropy(ihist)
10        if display:
11            hist = plt.hist(ibins[:-1], ibins, weights=ihist, alpha=0.5)
12            results.append((k, simg, hist, S))
13        else:
14            results.append((k, S))
15
```

```


16     if display:
17         plt.axvline(x=np.mean(np.array(img)))
18
19     return results

```

```

1 results = intensity_blur(img, np.linspace(0, 1.5, num=50), False)
2
3 plt.plot(*np.transpose(results), 'o-')
4 plt.ylim((0, 8))
5 plt.xlabel = "Smoothing"
6 plt.ylabel = "Intensity Entropy (bits)"


```



```

1 rings = [img for _, img, _, _ in intensity_blur(img, [0, 0.01, 0.05, 0.125, 0.25, 0.5])]
2 plt.show()


```



```

1 _, axarr = plt.subplots(1, len(rings))
2 for i, subimg in enumerate(rings):
3     axarr[i].imshow(subimg, cmap='gray')
4 plt.show()

```




## 4.2 Random pixel values

```

1 rsize = 250
2 randimg = Image.fromarray((256*np.random.rand(*2*[rsize])).astype('uint8'))
3 randimg

```




#### 4.2.1 Beware: GIGO


The boundary effects and discrete kernel of `ImageFilter.GaussianBlur` renders the data unreliable after the “minimum” of the intensity entropy with smoothing. This is immediately clear after even small smoothing for random pixel values, since there are no spatial correlations.

```
1 results = intensity_blur(randiimg, np.linspace(0, 0.3, num=75), False)
2
3 plt.plot(*np.transpose(results), 'o-')
4 plt.ylim((0, 8))
5 plt.xlabel = "Smoothing"
6 plt.ylabel = "Intensity Entropy (bits)"
```






```
1 rings = [img for _, img, _ in intensity_blur(randing, [0.01, 0.05, 0.25])]
1 plt.show()
```



```

1 _, axarr = plt.subplots(1, len(rings))
2 for i, subimg in enumerate(rings):
3     axarr[i].imshow(subimg, cmap='gray')
4 plt.show()

```



The rightmost image should be uniform: the renormalization emphasizes incorrect deviations. These are what keep the intensity entropy from vanishing.

### 4.3 Comparing different levels of smoothing

Is composing  $n$  Gaussian blurs with variance  $\sigma^2$  the same as doing one with variance  $n\sigma^2$  (considering the boundary effects and discrete kernel)?

```

1 nsmooths = 10
2 cimg = img
3 oneimg = cimg.filter(ImageFilter.GaussianBlur(np.sqrt(nsmooths)*2))
4 oneimg

```



```
1 nimg = cimg
2 for _ in range(nsmooths):
3     nimg = nimg.filter(ImageFilter.GaussianBlur(2))
4 nimg
```



Answer: **No**

The differences between results at different scales can be pretty wack.

```
1 Image.fromarray((255*rescale(np.array(nimg) - np.array(oneimg))).astype('uint8'))
```



```
1 smimg = img
2 smdiff = np.array(smimg.filter(ImageFilter.GaussianBlur(2))) -
  ↳ np.array(smimg.filter(ImageFilter.GaussianBlur(100)))
3 diffimg = Image.fromarray((255 * rescale(smdiff)).astype('uint8'))
4 diffimg
```



## 5 Local metrics

May 28, 2020 Given an image  $I : X \times Y \rightarrow \mathbb{Z}_n$ , we will now consider *local metrics* for the information it contains.

I want to be careful in understanding the statistical assumptions I am making, so I'll try to be explicit about distinguishing true distributions from empirical distributions, and how the assumptions behind postulating the existence of empirical distributions relate to the actual calculation being done. This should also aid in learning more solid probability theory.

### 5.1 Induced metrics

**Definition 1** (Lists). Given a set  $S$ , the collection of lists of elements from  $S$  is

$$\text{List}(S) = \bigcup_{n \in \mathbb{Z}_{\geq 0}} S^n,$$

where a list (tuple)  $s \in S^n$  is a map  $s : \mathbb{Z}_n \rightarrow S$  and  $|s| = n$ .

**Definition 2** (Image distributions). An *image distribution* is a map  $D$  that takes an image  $I$  and produces a random variable  $D(I) : \Omega \rightarrow E$ .

We are constructing empirical distributions from image data according to some map  $M : \text{Img} \rightarrow \text{List}(\Omega)$ , which produces the list of values  $V = M(I)$ . Then the probability of  $D(I)$  taking a value in a subset  $S \subseteq E$  is

$$P(X \in S) = \frac{1}{|V|} \sum_{s \in S} |V^{-1}(\{s\})|.$$

**Example 1.** The intensity-level entropy is a function of the *nonnegative* random variable from the image distribution of intensity values. That is, the map  $M$  takes an image and returns the list of its intensity values.

**Definition 3** (Induced image distributions). Given an image distribution  $D$ , and a subset  $S \subseteq \text{dom } I$ , we construct the *induced image distribution*  $D|_S$  by

$$D|_S(I) = D(I|_S).$$

**Definition 4** (Induced random variable). Given an image  $I$ , an image distribution  $D$  and collection of subsets  $\{S_i\}$  of  $\text{dom } I$ , a function  $H$  admits the random variables

$$H_i = (H \circ D|_{S_i})(I)$$

**Definition 5.** The *r-box* at  $(x, y)$  is  $B_r(x, y) = [x - r, x + r] \times [y - r, y + r]$ .

Given two real random variables  $A$  and  $B$  with joint PDF  $f_{A,B}(a, b)$ , the PDF of their sum is

$$f_{A+B}(c) = \int_{-\infty}^{\infty} da f_{A,B}(a, a - c) = \int_{-\infty}^{\infty} db f_{A,B}(b - c, b). \quad (1)$$

For independent  $A$  and  $B$ , EQ. 1 reduces to  $f_{A+B} = f_A * f_B$  over the marginals.

## 6 Kernels

Generalized to arbitrary functions on subregions of images.

```
1 import numpy as np
```

```

1 def box(x, y, r):
2     return np.s_[max(0, x-r) : x+r+1, max(0, y-r) : y+r+1]
3 def mapbox(r, f, a):
4     return np.reshape([f(a[box(*i, r)]) for i in np.ndindex(np.shape(a))], np.shape(a))
5 def mapboxes(rs, f, a):
6     return (mapbox(r, f, a) for r in rs)
7 def mapallboxes(f, a):
8     return mapboxes(range(max(np.shape(a))), f, a)
9
10 def mapblocks(h, w, f, a):
11     return np.array([[f(y) for y in np.array_split(x, w, axis=1)]
12                     for x in np.array_split(a, h)])

```

## 7 Boxcar intensity-level entropy

```

1 import numpy as np
2 import numpy.linalg as linalg
3 import matplotlib.pyplot as plt
4 from PIL import Image, ImageFilter, ImageOps
5 from src.utilities import *
6 from src.intensity_entropy import *
7 from src.kernels import *
8 plt.rcParams['image.cmap'] = 'inferno'

```


Let's compare the boxcar images for intensity entropy to those for a positive function on an image (the standard deviation) and for different functions of the induced intensity distribution.

```

1 img = ImageOps.grayscale(Image.open('test.jpg'))
2 scale = max(np.shape(img))
3 data = np.array(img)
4 plt.imshow(img);


```





## 7.1 Standard deviation

```
1 plt.imshow(mapbox(2, np.std, np.array(img)));
```

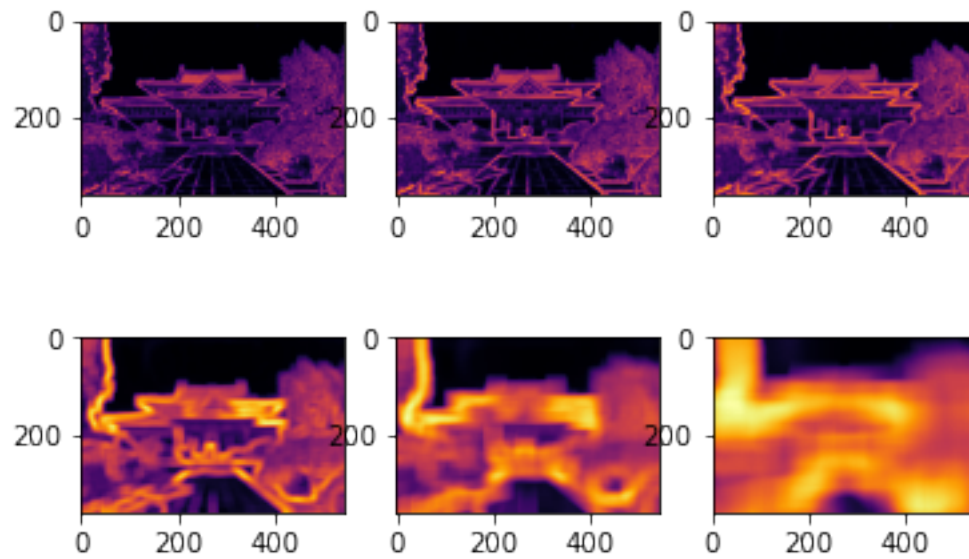


```

1 boxos = list(mapboxes([1,2,3,10,20,50], np.std, np.array(img)))


1 _, axarr = plt.subplots(2, np.ceil(len(boxos)/2).astype('int'))
2 for i, subimg in enumerate(boxos[:3]):
3     axarr[0,i].imshow(subimg)
4 for i, subimg in enumerate(boxos[3:]):
5     axarr[1,i].imshow(subimg)
6 plt.show()

```



## 7.2 Intensity entropy

```
plt.imshow(mapbox(2, intensity_entropy, np.array(img)));
```




```

1 boxSes = list(mapboxes([1,2,3,10,20,50], intensity_entropy, np.array(img)))

1 _, axarr = plt.subplots(2, np.ceil(len(boxSes)/2).astype('int'))
2 for i, subimg in enumerate(boxSes[:3]):
3     axarr[0,i].imshow(subimg)
4 for i, subimg in enumerate(boxSes[3:]):
5     axarr[1,i].imshow(subimg)
6 plt.show()

```



### 7.3 Replace surprisal with other functions


To what extent do the surprisal-related results depend upon the specific form of the *surprisal*  $x \mapsto -\log(x)$  in the expected value of the intensity distribution? We will replace the expected surprisal with the expected  $f$ , for different functions  $f$  on the empirical probabilities of a pixel taking some intensity.

Laurent:  $p \mapsto -1 + 1/p$ .

```


1 plt.imshow(mapbox(2, lambda I: intensity_expected(lambda p: -1 + 1/p if p > 0 else 0, I),
  ↪ np.array(img)));

```



Taylor:  $p \mapsto -(1 + p)$ .

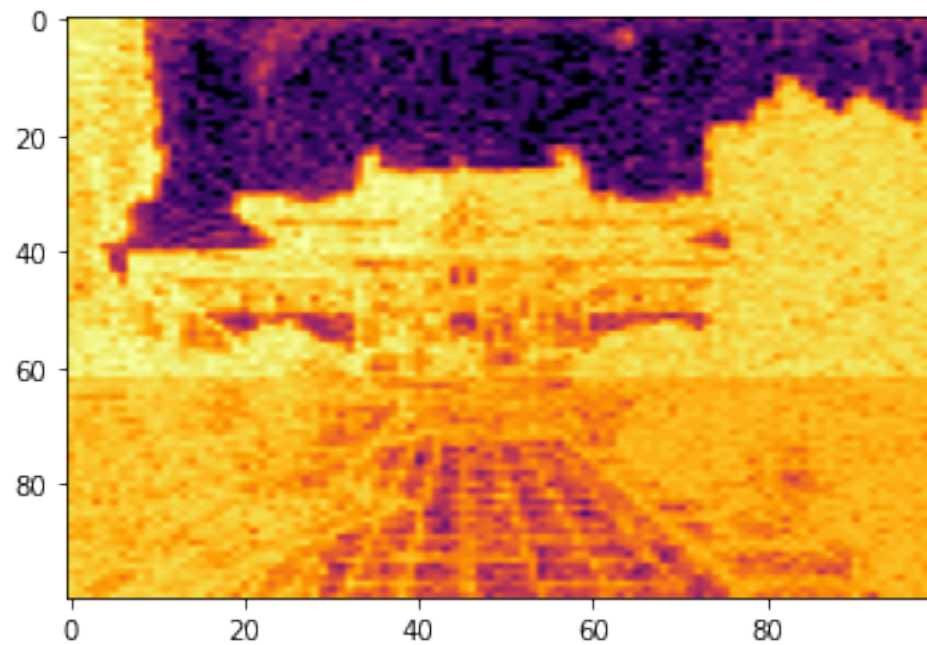
```
plt.imshow(mapbox(2, lambda I: intensity_expected(lambda p: -(1+p), I), np.array(img)));
```






## 7.4 Intensity entropy on disjoint blocks

```
1 plt.imshow(mapblocks(100, 100, intensity_entropy, np.array(img)),  
2             aspect=np.divide(*np.shape(img)));
```



```
1 plt.imshow(mapblocks(25, 25, intensity_entropy, np.array(img)),  
2             aspect=np.divide(*np.shape(img)));
```



## 8 Fractal dimensions

May 29, 2020 The previous results hint at characterizing the growth of the intensity entropy with different discretizations.

**Definition 6.** The *Rényi entropy of order  $\alpha \geq 0$*  of a discrete random variable  $X$  with support  $\mathcal{X}$  is

$$H_\alpha(X) = \frac{1}{1-\alpha} \log \sum_{x \in \mathcal{X}} P(x)^\alpha = \frac{\alpha}{1-\alpha} \log \|P\|_\alpha,$$

where  $\|P\|_\alpha$  denotes the  $\alpha$ -norm of the vector of probability values. The limit  $\alpha \rightarrow 1$  reproduces the Shannon entropy.

**Definition 7.** Given a real random variable  $X$ , define a discretized random variable

$$\langle X \rangle_\varepsilon = \frac{\lfloor \varepsilon X \rfloor}{\varepsilon}.$$

Then the *generalized dimension* of  $X$  is

$$d_\alpha(X) = \lim_{\varepsilon \rightarrow 0} \frac{H_\alpha(\langle X \rangle_\varepsilon)}{\log \varepsilon} = \lim_{\varepsilon \rightarrow 0} \frac{\alpha}{1-\alpha} \log (\|\langle X \rangle_\varepsilon\|_\alpha - \varepsilon).$$

The case  $\alpha \rightarrow 1$  is the *information dimension* of  $X$ . The generalized dimension may be estimated from linear regression of  $H_\alpha(\langle X \rangle_\varepsilon)$  with  $\log \varepsilon$  as the independent variable.

## 9 Fractal dimension regression

```
1 import numpy as np
2 import numpy.linalg as linalg
3 import matplotlib.pyplot as plt
4 from PIL import Image, ImageFilter, ImageOps
5 from scipy import interpolate
6 from scipy import integrate
7 from src.intensity_entropy import *
8 from src.kernels import *
9 plt.rcParams['image.cmap'] = 'inferno'

1 img = ImageOps.grayscale(Image.open('test.jpg'))
2 scale = max(np.shape(img))
3 data = np.array(img)
4 img
```






## 9.1 Box-counting dimension

```
1 def boxdim(data):
2     es = np.linspace(2, min(np.shape(data)))
3     boxes = [np.log(np.sum(mapblocks(
4         ε, ε, lambda x: 1 if np.any(x) else 0, data))) for ε in es]
5     loges = np.log(es)
6     endes = loges[[0, -1]]
7     dimfit = np.polyfit(np.log(es), boxes, 1) # [slope, intercept]
8     plt.plot(endes, dimfit[0]*endes + dimfit[1])
9     plt.plot(loges, boxes, '+')
10    plt.xlabel('Scale (ln ε)')
11    plt.ylabel('Box count (ln N)')
12    return dimfit[0]
```

```
1 boxdim(data)
```

2.0087040269581435




```
1 sky = data.copy()
2 sky[sky < 128+32] = 0
3 Image.fromarray(sky)
```




```
1 boxdim(sky)
```

```
1.7877778191348215
```




```
1 nosky = data.copy()
2 nosky[nosky < 128+64] = 0
3 Image.fromarray(nosky)
```



```
1 boxdim(nosky)
```

```
1.6214794967487127
```




```
1 dots = data.copy()
2 dots[nosky < 128+64+16] = 0
3 Image.fromarray(dots)
```



```
1 boxdim(dots)
```

```
1.2821025677557252
```



## 9.2 Information dimension

```

1 def discretize(f, a, b, ε, N=20):
2     return [integrate.simps(f(np.linspace(c - ε/2, c + ε/2, N)), dx=ε / (N - 1))
3             for c in np.arange(a + ε/2, b, ε)]

1 def infodim(dist, s=1e-5):
2     l = len(dist)
3     spl = interpolate.splrep(range(l), dist, s=s)
4     f = lambda x: interpolate.splev(x, spl)
5
6     εs = 1 / np.linspace(10, 1)
7     loges = -np.log2(εs)
8     endes = loges[[0, -1]]
9     entropies = [shannon_entropy(discretize(f, 0, 1, ε)) for ε in εs]
10    dimfit, cov = np.polyfit(loges, entropies, 1, cov='unscaled')
11
12    plt.plot(endes, dimfit[0]*endes + dimfit[1])
13    plt.plot(loges, entropies, '+')
14    plt.xlabel('Scale (lg ε)')
15    plt.ylabel('Shannon entropy (bits)')

```

```

16
17     return dimfit[0], cov[0,0]

```

The Gaussian distribution

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

is continuous, so its information dimension is 1.


```

1 def gaussian(μ, σ, x):
2     return np.exp(-(x - μ)**2 / (2*σ**2)) / (σ*np.sqrt(2*np.pi))

1 infodim((10/256) * gaussian(0, 1, np.linspace(-5, 5, 256)))

(1.0014184290221988, 0.015707497682893923)

```




The rectified Gaussian distribution  $g(x) = \Theta(x)f(x) + \delta(x)/2$  is half-continuous, so its information dimension is  $1/2$ .

```

1 dist = np.concatenate([[0]*256, (5/256)*gaussian(0, 1, np.linspace(0, 5, 256))])
2 plt.plot(dist);


```






```
1 infodim(dist)
```

```
(0.4979088715795226, 0.012313889825394398)
```




Now that we've validated infodim, what does it say about the intensity distribution of an image?

```
1 dist = intensity_distribution(img)
2 plt.plot(dist);
```



```
1 infodim(dist)
```

```
(0.98265843047326, 0.015707497682893923)
```



## 10 Probability and inference

May 30, 2020 Let's look at a simple inference problem before considering images. This example illustrates how the approach founded on probability theory differs from the naïve statistical approach usually taken by physicists.

**Example 2** (Biased coin tosses). Consider tossing a biased coin  $N$  times to obtain  $n$  heads. What is the probability  $p'$  that the next coin toss comes up heads?

The temptation is to claim  $n/N$  as the probability, but this is *incorrect* if we want to allow all consistent biases. The problem with this solution is that the most probable bias is assumed to be the true bias.

The probability of getting  $m$  heads if a single head has probability  $p$  is

$$P(m | p) = \binom{N}{m} p^m (1 - p)^{N-m}.$$

We have no other information, so we assume that all of the biases are equally likely. This means that  $P(p)$  is constant (the uniform prior). The distribution of

biases  $p$  given the observation of  $m$  heads is then

$$P(p | m) = \frac{P(m | p)P(p)}{P(m)} = \frac{P(m | p)P(p)}{\int_0^1 d\tilde{p} P(m | \tilde{p})P(\tilde{p})} = \frac{P(m | p)}{\int_0^1 d\tilde{p} P(m | \tilde{p})}.$$

We compute that

$$P(m) = \binom{N}{m} \int_0^1 dp p^m (1-p)^{N-m} = \binom{N}{m} \frac{m!(N-m)!}{(N+1)!} = \frac{1}{N+1},$$

so the next coin toss is heads with probability

$$\begin{aligned} p' &= \int_0^1 dp P(\text{head} | n, p) P(p | n) = \int_0^1 dp p P(p | n) \\ &= \int_0^1 dp p (N+1) \binom{N}{n} p^n (1-p)^{N-n} = \frac{n+1}{N+2}. \end{aligned}$$

For  $n = 3$  and  $N = 10$ ,  $p' = 0.33$ . This is a more conservative estimate than  $p' = 0.30$  from the most probable bias.

## 11 Ising images

*June 1, 2020* What happens if we apply a model from statistical physics to an image?

## 12 Ising images

```


1 import numpy as np
2 import numpy.linalg as linalg
3 import matplotlib.pyplot as plt
4 from PIL import Image, ImageFilter, ImageOps
5 import imageio
6 plt.rcParams['image.cmap'] = 'gray'

1 from ipywidgets import IntProgress
2 from IPython.display import display
3 import time
```

## 12.1 Standard Ising (on a torus)

In grayscale for fun.


```
1 def neighbors(a, i, j):
2     return np.hstack([a[:,j].take([i-1,i+1], mode='wrap'),
3                       a[i,:].take([j-1,j+1], mode='wrap')])
4
5 def energy(img, i, j):
6     return -1 + np.sum(np.abs(img[i, j] - neighbors(img, i, j)))
7
8 def isingstep( $\beta$ , img):
9     w, h = np.shape(img)
10    i = np.random.randint(w)
11    j = np.random.randint(h)
12    E0 = energy(img, i, j)
13    img[i, j] *= -1
14    E1 = energy(img, i, j)
15    P = np.exp(- $\beta$ *(E1 - E0)) if E1 > E0 else 1
16    if np.random.rand() > P: # Restore old
17        img[i, j] *= -1
18    return img
19
20 img = 2*np.random.rand(50, 50) - 1
21 plt.imshow(img);
```



```

1  n = 100000
2  for i in range(n):
3      isingstep(3 * (np.pi / 2) / np.arctan(n - i), img)
4  plt.imshow(img);

```



## 12.2 Image-edge Ising

```

1  edges = Image.open("ising-edges.png")
2  edata = np.array(edges) > 128
3  edges

```



```

1 def eenergy(img, edges, i, j):
2     """Edge-modified Ising energy:  $\theta$  on edge."""
3     if edges[i, j]:
4         return  $\theta$ 
5     w, h = np.shape(img)
6     c = img[i, j]
7     l = img[i-1, j] if i > 0 else img[w-1, j]
8     r = img[i+1, j] if i < w-1 else img[0, j]
9     t = img[i, j-1] if j > 0 else img[i, h-1]
10    b = img[i, j+1] if j < h-1 else img[i, 0]
11    return -img[i, j] * (1 + r + t + b)
12
13 def nenergy(img, edges, i, j):
14     """Neighbor-modified Ising energy:  $\theta$  interactions with edges."""
15     if edges[i, j]:
16         return  $\theta$ 
17
18     w, h = np.shape(img)
19     c = img[i, j]
20     l = r = t = b =  $\theta$ 
21     if i > 0:
22         l = img[i-1, j] if not edges[i-1, j] else  $\theta$ 
23     else:
24         l = img[w-1, j] if not edges[w-1, j] else  $\theta$ 
25
26     if i < w - 1:
27         r = img[i+1, j] if not edges[i+1, j] else  $\theta$ 
28     else:
29         r = img[0, j] if not edges[0, j] else  $\theta$ 
30
31     if j > 0:
32         t = img[i, j-1] if not edges[i, j-1] else  $\theta$ 
33     else:
34         t = img[i, h-1] if not edges[i, h-1] else  $\theta$ 
35
36     if j < h - 1:
37         b = img[i, j+1] if not edges[i, j+1] else  $\theta$ 
38     else:
39         b = img[i, 0] if not edges[i, 0] else  $\theta$ 
40
41    return -img[i, j] * (1 + r + t + b)
42
43 def eisingstep( $\beta$ , img, edges):
44     w, h = np.shape(img)
45     i = np.random.randint(w)

```




```

46     j = np.random.randint(h)
47     E0 = nenergy(img, edges, i, j)
48     img[i, j] *= -1
49     E1 = nenergy(img, edges, i, j)
50     P = np.exp(-β*(E1 - E0)) if E1 > E0 else 1
51     if np.random.rand() > P: # Restore old
52         img[i, j] *= -1
53     return img
54
55 def frame(writer, data):
56     writer.append_data((255 * ((eimg + 1) / 2)).astype('uint8'))

1  img = Image.open("ising-letters.png")
2  eimg = -1 + 2 * (np.array(img) / 255)
3  plt.imshow(eimg);

```



movie.gif: Full neighbor Ising.

```

1  n = 1000000
2  f = IntProgress(min=0, max=1 + (n-1) // 1000) # instantiate the bar
3  display(f)
4  with imageio.get_writer('movie.gif', mode='I') as writer:
5      frame(writer, eimg)
6      for i in range(n):


```

```

7         eisingstep(0.5 * (np.pi / 2) / np.arctan(n - i), eimg, edata)
8         if i % 1000 == 0:
9             f.value += 1
10            frame(writer, eimg)
11 plt.imshow(eimg);

```

IntProgress(value=0, max=1000)




imovie.gif: Normal Ising.

```

1  n = 1000000
2  img = eimg
3  with imageio.get_writer('imovie.gif', mode='I') as writer:
4      frame(writer, img)
5      for i in range(n):
6          isingstep(0.5 * (np.pi / 2) / np.arctan(n - i), img)
7          if i % 1000 == 0:
8              frame(writer, img)
9  plt.imshow(img);

```



## 12.3 Image-metric Ising

```

1 # def takewrap(a, i, j, xs=np.arange(-1, 2), ys=np.arange(-1, 2)):
2 def takewrap(a, i, j, xs=np.arange(0, 1), ys=np.arange(0, 1)):
3     return np.array([x for v in a.take(xs+i, axis=0, mode='wrap')
4                       for x in v.take(ys+j, mode='wrap')])

```

### 12.3.1 Unrestricted swapping motion

Swapping preserves the intensity distribution.

```

1 def sienergy(img, init, i, j):
2     """Inversion-symmetric image energy"""
3     eq = takewrap(img, i, j) == takewrap(init, i, j)
4     return -np.abs(np.sum(2*eq - 1))
5
6 def ienergy(img, init, i, j):
7     """Image energy based on 3x3 block deviation"""
8     return np.abs(init[i, j] - img[i, j])
9
10 def swisingstep(beta, img, edges):
11     w, h = np.shape(img)


```

```

12     i0 = np.random.randint(w)
13     i1 = np.random.randint(w)
14     j0 = np.random.randint(h)
15     j1 = np.random.randint(h)
16     E0 = ienergy(img, edges, i0, j0) + ienergy(img, edges, i1, j1)
17     img[i0, j0], img[i1, j1] = img[i1, j1], img[i0, j0]
18     E1 = ienergy(img, edges, i0, j0) + ienergy(img, edges, i1, j1)
19     P = np.exp(-β*(E1 - E0)) if E1 > E0 else 1
20     if np.random.rand() > P: # Restore old
21         img[i0, j0], img[i1, j1] = img[i1, j1], img[i0, j0]
22     return img
23
24 def nnisingstep(β, img, edges):
25     w, h = np.shape(img)
26     i0 = np.random.randint(w)
27     i1 = int((i0 + np.sign(np.random.rand() - 1/2)) % w)
28     j0 = np.random.randint(h)
29     j1 = int((j0 + np.sign(np.random.rand() - 1/2)) % h)
30     E0 = ienergy(img, edges, i0, j0) + ienergy(img, edges, i1, j1)
31     img[i0, j0], img[i1, j1] = img[i1, j1], img[i0, j0]
32     E1 = ienergy(img, edges, i0, j0) + ienergy(img, edges, i1, j1)
33     P = np.exp(-β*(E1 - E0)) if E1 > E0 else 1
34     if np.random.rand() > P: # Restore old
35         img[i0, j0], img[i1, j1] = img[i1, j1], img[i0, j0]
36     return img

1  img = Image.open("barbara.png")
2  eimg = -1 + 2 * (np.array(img) / 255)
3  initimg = eimg.copy()
4  plt.imshow(initimg);

```



swmmovie.gif: Image metric Ising (arbitrary swaps with ienergy).

```

1  n = 2000000
2  f = IntProgress(min=0, max=(1 + (n-1) // 1000)) # instantiate the bar
3  display(f)
4  with imageio.get_writer('swmmovie.gif', mode='I') as writer:
5      frame(writer, eimg)
6      for i in range(n):
7          k = i/n
8          swisingstep(3, eimg, initimg)
9          if i % 1000 == 0:
10             f.value += 1
11             frame(writer, eimg)
12     # for i in range(n):
13     #     k = i/n
14     #     swisingstep(4*(1 - k) + 1e-3*k, eimg, initimg)
15     #     if i % 1000 == 0:
16     #         f.value += 1
17     #         frame(writer, eimg)
18     # for i in range(n):
19     #     k = i/n
20     #     swisingstep(1e-3*(1 - k) + 4*k, eimg, initimg)
21     #     if i % 1000 == 0:
22     #         f.value += 1


```

```

23 #         frame(writer, eimg)
24
25 plt.imshow(eimg);

IntProgress(value=0, max=2000)

```




### 12.3.2 Nearest-neighbor swapping motion

```

1 img = Image.open("barbara.png")
2 eimg = -1 + 2 * (np.array(img) / 255)
3 initimg = eimg.copy()
4 plt.imshow(initimg);

```



nnmovie.gif: Image metric Ising (neighborly swaps with ienergy).

```

1  n = 2000000
2  f = IntProgress(min=0, max=3*(1 + (n-1) // 1000)) # instantiate the bar
3  display(f)
4  with imageio.get_writer('nnmovie.gif', mode='I') as writer:
5      frame(writer, eimg)
6      for i in range(n):
7          k = i/n
8          nnisingstep(5*(1 - k) + 1e-4*k, eimg, initimg)
9          if i % 1000 == 0:
10             f.value += 1
11             frame(writer, eimg)
12     for i in range(n):
13         nnisingstep(1e-4, eimg, initimg)
14         if i % 1000 == 0:
15             f.value += 1
16             frame(writer, eimg)
17     for i in range(n):
18         k = i/n
19         nnisingstep(1e-4*(1 - k) + 5*k, eimg, initimg)
20         if i % 1000 == 0:
21             f.value += 1
22             frame(writer, eimg)


```

```

23
24 plt.imshow(eimg);

```

```
IntProgress(value=0, max=6000)
```



## 13 Statistical Mechanics of Images

*June 3,* Given the qualitative success of the image-metric based Ising images, we consider  
*2020* generalizations.

**Definition 8.** An  $N$ -element image space over metric spaces  $(K, d)$  and  $(P, a)$  is the space  $\text{Img} = (P \times K)^N$ . A corresponding *image* is an element of  $\text{Img}$ .

The space  $P$  determines the spatial arrangement of the image, and is usually two-dimensional Euclidean space. We usually consider the subset of an image space where the  $P$ -coordinates are fixed, in a grid layout. The space  $K$  determines the qualities of an image at a point in  $P$ . This is usually a color or intensity space, and in practical applications is a machine integer like  $128 \in \mathbb{Z}_{256}$ .

**Definition 9.** An *image system* on an image space  $\text{Img}$  consists of a *ground image*  $I_0 \in \text{Img}$  and a *dispersion relation*  $E : \mathbb{R} \rightarrow \mathbb{R}$ . This defines the *energy* of an image



$I$  as

$$E(I) = \sum_{(p_0, k_0) \in I_0} \sum_{k \in \{k: (p_0, k) \in I\}} E(d(k_0, k)).$$

**Example 3.** For usual images in  $((\mathbb{Z}_n \times \mathbb{Z}_m) \times \mathbb{Z}_{256})^N$ , where  $N = nm$  and the positions of  $I_0$  and  $I$  coincide (indexed by  $i$  and  $j$ ), we have

$$E(I) = \sum_{i=1}^n \sum_{j=1}^m E(d(k_0^{ij}, k^{ij})) = \sum_{i=1}^n \sum_{j=1}^m \varepsilon \left| k_0^{ij} - k^{ij} \right|^1,$$

with typical choices of  $E$  and  $d$ .

In the binary case ( $K = \mathbb{Z}_2$ ), we have  $N$  independent two-level systems.

**Example 4** (Grayscale images). Consider a pixel of a ground grayscale image, with integer value  $k_0 \in 0, \dots, K-1$  for even  $K$ . There are then

$$2g = 2 \begin{cases} k_0, & k_0 < K/2 \\ K - k_0 - 1, & \text{else} \end{cases}$$

energy values that occur twice, and  $K - 2g$  energy values that occur once (like  $|x|$  on an interval like  $[-3, 8]$ ). Thus the partition function for this single pixel is

$$\begin{aligned} Z_g &= \sum_{k=-g}^{K-g-1} e^{-\beta \varepsilon |k|} = 1 + \sum_{k=1}^g e^{-\beta \varepsilon k} + \sum_{k=1}^{K-g-1} e^{-\beta \varepsilon k} \\ &= 1 + \frac{e^{-\beta g \varepsilon} (e^{\beta g \varepsilon} - 1)}{e^{\beta \varepsilon} - 1} + \frac{e^{-\beta (K-g-1) \varepsilon} (e^{\beta (K-g-1) \varepsilon} - 1)}{e^{\beta \varepsilon} - 1} \end{aligned}$$

and the partition function for the whole image is

$$Z = \prod_{g=0}^{-1+K/2} Z_g^{NP(g)},$$

where  $NP(g)$  is the number of pixels in the ground image with the given  $g$ -value. We then see that

$$\ln Z = \sum_{g=0}^{-1+K/2} NP(g) \ln Z_g = N \langle \ln Z_g \rangle_G,$$

where  $G$  is the random variable that takes the value  $g$  with probability  $P(g)$ . It then follows that  $\langle E/N \rangle = \langle E_g \rangle_G$  and  $S/N = \langle S_g \rangle_G$  as usual for extensive variables.

## 14 Thermodynamic quantities for images from a microscopic model

*June 4, 2020* Since we are thinking of images as statistical entities, what is the corresponding microscopic model? Given such a model, what quantities do we consider in thermal equilibrium, and how can we understand different ensembles?

### 14.1 Quantum filled-site model (FSM)

**Definition 10** (FSM). We define a lattice model corresponding to a *ground image*  $I_0$  as follows. Each pixel with value  $k_0 \in K \subseteq \mathbb{Z}$  in the image corresponds to a *site*, which is a discrete system with  $K$  levels. The energy of level  $k$  is  $E_{k_0}(k) = \varepsilon|k - k_0|^r$ .<sup>1</sup> We usually have  $r = 1$  or  $2$ . We suppose that the levels are filled by fermions that interact according to the Hamiltonian

$$H = \sum_{k \in K} \sum_i V c_{ik}^\dagger c_{ik} - \sum_{\ell \in \mathcal{N}_k} \sum_{j \in \mathcal{N}_i} t_{\ell} c_{ik}^\dagger c_{j\ell}.$$

For  $K \subseteq \mathbb{Z}$ ,  $\mathcal{N}_k = k + \{-1, 0, 1\}$ .

### 14.2 Observables and thermodynamic state variables

Several observables of the FSM are of interest:

- **Pixel colors.** The occupations  $n_k$  of different levels at a *single* site induce a distribution on  $K$ . In equilibrium, the mean level

$$\langle k \rangle \equiv \frac{\sum_{k \in K} k n_k}{\sum_{k \in K} n_k}$$

should be near  $k_0$ , since the energy of a level is symmetric about  $k_0$ . As the temperature increases, so will the variance of the mean level. On the flip side, does *varying*  $k_0$  for many pixels quasistatically (changing the ground image) do work on the system? Yes, but is this consistent with what we expect?

- **Color distribution.** The net occupations  $m_k$  of different levels across *all* sites induce a distribution on  $K$ . In the special case of gray images (so levels are intensity), the entropy of the induced random variable is intensity

---

<sup>1</sup>If we want to consider colors, then  $K$  is a metric space and we replace  $k - k_0$  with the metric.

entropy that we have studied previously. This distribution is stationary when different levels cannot interact, but is it so at finite temperature?

- **Opacity.** The net occupancy of a *site* could be connected to its opacity. In equilibrium, this should be similar across all sites. Then regions with *no* particles during nonequilibrium processes make sense. The picture of a gas with fluctuating density that emits light comes to mind. When at maximum opacity, the gas in that region cannot be compressed further, and cannot accept more particles. The canonical density properties of a photon gas (like energy density) might be a good reason to choose the particles to be bosons.
- **Number of pixels.** We could vary the total number of pixels different ways. One way is to have a continuous ground image, and choose different grid discretizations. Another way is to have a large ground image grid and vary the zoom level. It would be sensible to combine these sorts of transformations with pixel color transformations, since they include translation and rotation as special cases. This seems most most similar to varying the volume of a gas. Including opacity makes fast adiabatic piston motion volume changes like  $V \mapsto 2V$  make sense.
- **Number of particles.** Depending on if we allow interactions between levels, it may be appropriate to consider chemical potentials. Either for all particles, or for each color. Could this be conjugate to opacity?
- **The usual.** Given that quasistatic transformation of the other quantities does work the way we expect, we can consider the usual response variables like heat capacities and compressibilities. There is also the thermodynamic entropy.

## 15 Progress summary (from beginning)

*June 6,* Over the last two weeks, I calculated some metrics on images and did some basic  
*2020* simulations. The most important metric was the intensity entropy, which is used as the “entropy” of an image in the maximum entropy method (MEM) of image reconstruction used by astronomers. This was calculated for a whole image and locally in different regions of an image. On the topic of scaling, fractal dimensions were explored. The box-counting dimension was computed for different

images, and the information dimension of the intensity distribution was considered. I also did readings on probability theory and machine learning, since the usual frequentist approach that experimental physicists take is not applicable. Variants of Ising models were simulated for images, which led to the postulation of a microscopic model for varying images (the FSM), which is similar to a Hubbard model. The implications of this approach remain to be explored.

## 16 Simulations for canonical ensemble averages

June 10,  
2020

- Wang-Landau
- Parallel tempering (replica exchange MCMC)

## 17 The Wang-Landau algorithm (density of states)

We determine thermodynamic quantities from the partition function by obtaining the density of states from a simulation.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy import interpolate
```

A Wang-Landau algorithm, with quantities as logarithms and with monte-carlo steps proportional to  $f^{-1/2}$  (a “Zhou-Bhat schedule”).

```
1 def flat(H, tol = 0.2):
2     """Determines if an evenly-spaced histogram is approximately flat."""
3     Hμ = np.mean(H)
4     Hf = np.max(H)
5     H0 = np.min(H)
6     return Hf / (1 + tol) < Hμ < H0 / (1 - tol)
7 # def flat(H, tol = 0.2):
8 #     """Determines if an evenly-spaced histogram is approximately flat."""
9 #     Hμ = np.mean(H)
10 #     return not np.any(H < (1 - tol) * Hμ) and np.all(H ≠ 0)

1 def wanglandau(system,
2                 M = 1_00_000, # Monte carlo step scale
3                 ε = 1e-8,     # f tolerance
```

```

4         logf0 = 1,      # Initial log f
5         N = 8**2 + 1,   # Number of energy bins
6         E0 = -2 * 8**2, # Minimum energy
7         Ef = 2 * 8**2   # Maximum energy
8     ):
9         # Initial values
10        logf = logf0
11        logftol = np.log(1 + ε)
12        Es = np.linspace(E0, Ef, N)
13        S = np.zeros(N) # Set all initial g's to 1
14        H = np.zeros(N, dtype=int)
15        # Linearly bin the energy
16        i = max(0, min(N - 1, int(round((N - 1) * (system.E - E0) / (Ef - E0)))))
17
18        # Logging
19        mciters = 0
20        fiter = 0
21        ΔE = (Ef - E0) / (N - 1)
22        fiters = int(np.ceil(np.log2(logf0) - np.log2(logftol)))
23        print("ΔE = {}".format(ΔE))
24
25        while logftol < logf:
26            H[:] = 0
27            logf /= 2
28            iters = 0
29            niters = int((M + 1) * np.exp(-logf / 2))
30            fiter += 1
31            while not flat(H[2:-2]) and iters < niters: # Ising-specific histogram
32                # while not flat(H) and iters < niters:
33                system.propose()
34                Ev = system.Ev
35                j = max(0, min(N - 1, int(round((N - 1) * (Ev - E0) / (Ef - E0)))))
36                if E0 - ΔE/2 ≤ Ev ≤ Ef + ΔE/2 and (S[j] < S[i] or np.random.rand() < np.exp(S[i] -
37                    ↪ S[j])):
38                    system.accept()
39                    i = j
40                    H[i] += 1
41                    S[i] += logf
42                    iters += 1
43                mciters += iters
44                print("f: {} / {} \t {} / {}".format(fiter, fiters, iters, niters))
45
46        print("Done: {} total MC iterations.".format(mciters))
47        return Es, S, H

```

## 17.1 The 2D Ising model

```
1 class Ising:
2     def __init__(self, n):
3         self.n = n
4         self.spins = np.sign(np.random.rand(n, n) - 0.5)
5         self.E = self.energy()
6         self.Ev = self.E
7     def neighbors(self, i, j):
8         return np.hstack([self.spins[:,j].take([i-1,i+1], mode='wrap'),
9                             self.spins[i,:].take([j-1,j+1], mode='wrap')])
10    def energy(self):
11        return -0.5 * sum(np.sum(s * self.neighbors(i, j))
12                            for (i, j), s in np.ndenumerate(self.spins))
13    def propose(self):
14        i, j = np.random.randint(self.n), np.random.randint(self.n)
15        self.i, self.j = i, j
16        dE = 2 * np.sum(self.spins[i, j] * self.neighbors(i, j))
17        self.dE = dE
18        self.Ev = self.E + dE
19    def accept(self):
20        self.spins[self.i, self.j] *= -1
21        self.E = self.Ev
```

Note that this class-based approach adds some overhead. For speed, instances of Ising should be inlined into the simulation method.

```
1 isingn = 8
2 sys = Ising(isingn)
3 Es, S, H = wanglandau(sys);
```

```
 $\Delta E = 4.0$ 
f: 1 / 27 (55874 / 77880)
f: 2 / 27 (32200 / 88250)
f: 3 / 27 (72964 / 93942)
f: 4 / 27 (70127 / 96924)
f: 5 / 27 (78111 / 98450)
f: 6 / 27 (43058 / 99222)
f: 7 / 27 (99611 / 99611)
f: 8 / 27 (99805 / 99805)
f: 9 / 27 (99903 / 99903)
f: 10 / 27 (99952 / 99952)
f: 11 / 27 (99976 / 99976)
```

```

f: 12 / 27 (99988 / 99988)
f: 13 / 27 (99994 / 99994)
f: 14 / 27 (99997 / 99997)
f: 15 / 27 (99999 / 99999)
f: 16 / 27 (100000 / 100000)
f: 17 / 27 (100000 / 100000)
f: 18 / 27 (100000 / 100000)
f: 19 / 27 (100000 / 100000)
f: 20 / 27 (100000 / 100000)
f: 21 / 27 (100000 / 100000)
f: 22 / 27 (100000 / 100000)
f: 23 / 27 (100000 / 100000)
f: 24 / 27 (100000 / 100000)
f: 25 / 27 (100000 / 100000)
f: 26 / 27 (100000 / 100000)
f: 27 / 27 (100000 / 100000)
Done: 2451559 total MC iterations.

```


The energies at indices 1 and -1 are not occupied in the Ising model.

```

1 Es, S, H = Es[2:-2], S[2:-2], H[2:-2]


1 plt.plot(Es / isingn**2, S)
2 plt.xlabel("E / N")
3 plt.ylabel("log g(E) + C");

```



```
1 plt.plot(Es / isingn**2, H)
2 plt.xlabel("E / N")
3 plt.ylabel("Visits");
```






### 17.1.1 Calculating canonical ensemble averages

```

1  gspl = interpolate.splrep(Es, S, s=2*np.sqrt(2))
2  gs = np.exp(interpolate.splev(Es, gspl) - min(S))

1  plt.plot(Es / isingn**2, S)
2  plt.plot(Es / isingn**2, interpolate.splev(Es, gspl))
3  plt.xlabel("E / N")
4  plt.ylabel("log g(E) + C");

```



Translate energies to have minimum zero so that  $Z$  is representable.

```

1 nEs = Es - min(Es)

1 Z = lambda beta: np.sum(gEs * np.exp(-beta * nEs))

```

Ensemble averages

```

1 beta_s = [np.exp(k) for k in np.linspace(-3, 1, 200)]
2 E_mu = lambda beta: np.sum(nEs * gEs * np.exp(-beta * nEs)) / Z(beta)
3 E2 = lambda beta: np.sum(nEs**2 * gEs * np.exp(-beta * nEs)) / Z(beta)
4 CV = lambda beta: (E2(beta) - E_mu(beta)**2) * beta**2
5 F = lambda beta: -np.log(Z(beta)) / beta
6 Sc = lambda beta: beta * E_mu(beta) + np.log(Z(beta))


```

Heat capacity

```


1 plt.plot(np.log(beta_s), [CV(beta) for beta in beta_s])
2 plt.xlabel("ln beta")
3 plt.ylabel("Heat capacity")
4 plt.show()

```



## Entropy

```
1 plt.plot(np.log(βs), [Sc(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("S(β) + C")
4 plt.show()
```



## 17.2 Thermal calculations on images

```

1 class StatisticalImage:
2     def __init__(self, I0):
3         self.I0 = I0
4         self.I = I0.copy()
5         self.w, self.h = np.shape(I0)
6         self.E = self.energy()
7         self.Ev = self.E
8     def energy(self):
9         return sum(x0 - x if x < x0 else x - x0
10                for x, x0 in zip(self.I.flat, self.I0.flat))
11    def propose(self):
12        i, j = np.random.randint(self.w), np.random.randint(self.h)
13        self.i, self.j = i, j
14        x0 = self.I0[i, j]
15        x = self.I[i, j]
16        r = 16
17        dx = np.random.randint(-min(r, x), min(r, 255 - x) + 1)
18        x1 = x + dx
19        dE = (x0 - x1 if x1 < x0 else x1 - x0) - (x0 - x if x < x0 else x - x0)
20        self.dx = dx

```

```

21         self.dE = dE
22         self.Ev = self.E + dE
23     def accept(self):
24         self.I[self.i, self.j] += self.dx
25         self.E = self.Ev

1  Ls = range(1, 11, 2)
2  wlresults = [wanglandau(StatisticalImage(128 * np.ones((L, L), dtype=int)),
3                                     M=1_000_000, N=127*L*L + 1, E0=0, Ef=127*L*L)
4               for L in Ls]

```

$\Delta E = 1.0$

```

f: 1 / 27  (3903 / 778801)
f: 2 / 27  (9528 / 882497)
f: 3 / 27  (9083 / 939414)
f: 4 / 27  (10139 / 969234)
f: 5 / 27  (26254 / 984497)
f: 6 / 27  (27029 / 992218)
f: 7 / 27  (34230 / 996102)
f: 8 / 27  (32353 / 998049)
f: 9 / 27  (38764 / 999024)
f: 10 / 27 (16429 / 999512)
f: 11 / 27 (38150 / 999756)
f: 12 / 27 (64131 / 999878)
f: 13 / 27 (129645 / 999939)
f: 14 / 27 (25586 / 999970)
f: 15 / 27 (53326 / 999985)
f: 16 / 27 (31720 / 999993)
f: 17 / 27 (33121 / 999997)
f: 18 / 27 (24031 / 999999)
f: 19 / 27 (49997 / 1000000)
f: 20 / 27 (49427 / 1000000)
f: 21 / 27 (42771 / 1000000)
f: 22 / 27 (34318 / 1000000)
f: 23 / 27 (39775 / 1000000)
f: 24 / 27 (26611 / 1000000)
f: 25 / 27 (52471 / 1000000)
f: 26 / 27 (26318 / 1000000)
f: 27 / 27 (21238 / 1000000)

```

Done: 950348 total MC iterations.

$\Delta E = 1.0$   
f: 1 / 27 (778801 / 778801)  
f: 2 / 27 (882497 / 882497)  
f: 3 / 27 (693591 / 939414)  
f: 4 / 27 (969234 / 969234)  
f: 5 / 27 (984497 / 984497)  
f: 6 / 27 (992218 / 992218)  
f: 7 / 27 (885823 / 996102)  
f: 8 / 27 (998049 / 998049)  
f: 9 / 27 (999024 / 999024)  
f: 10 / 27 (999512 / 999512)  
f: 11 / 27 (999756 / 999756)  
f: 12 / 27 (999878 / 999878)  
f: 13 / 27 (999939 / 999939)  
f: 14 / 27 (999970 / 999970)  
f: 15 / 27 (999985 / 999985)  
f: 16 / 27 (999993 / 999993)  
f: 17 / 27 (999997 / 999997)  
f: 18 / 27 (999999 / 999999)  
f: 19 / 27 (1000000 / 1000000)  
f: 20 / 27 (1000000 / 1000000)  
f: 21 / 27 (1000000 / 1000000)  
f: 22 / 27 (1000000 / 1000000)  
f: 23 / 27 (1000000 / 1000000)  
f: 24 / 27 (1000000 / 1000000)  
f: 25 / 27 (1000000 / 1000000)  
f: 26 / 27 (1000000 / 1000000)  
f: 27 / 27 (1000000 / 1000000)  
Done: 26182763 total MC iterations.  
 $\Delta E = 1.0$   
f: 1 / 27 (778801 / 778801)  
f: 2 / 27 (882497 / 882497)  
f: 3 / 27 (939414 / 939414)  
f: 4 / 27 (969234 / 969234)  
f: 5 / 27 (984497 / 984497)  
f: 6 / 27 (992218 / 992218)  
f: 7 / 27 (996102 / 996102)  
f: 8 / 27 (998049 / 998049)

```

f: 9 / 27 (999024 / 999024)
f: 10 / 27 (999512 / 999512)
f: 11 / 27 (999756 / 999756)
f: 12 / 27 (999878 / 999878)
f: 13 / 27 (999939 / 999939)
f: 14 / 27 (999970 / 999970)
f: 15 / 27 (999985 / 999985)
f: 16 / 27 (999993 / 999993)
f: 17 / 27 (999997 / 999997)
f: 18 / 27 (999999 / 999999)
f: 19 / 27 (1000000 / 1000000)
f: 20 / 27 (1000000 / 1000000)
f: 21 / 27 (1000000 / 1000000)
f: 22 / 27 (1000000 / 1000000)
f: 23 / 27 (1000000 / 1000000)
f: 24 / 27 (1000000 / 1000000)
f: 25 / 27 (1000000 / 1000000)
f: 26 / 27 (1000000 / 1000000)
f: 27 / 27 (1000000 / 1000000)
Done: 26538865 total MC iterations.
ΔE = 1.0
f: 1 / 27 (778801 / 778801)
f: 2 / 27 (882497 / 882497)
f: 3 / 27 (939414 / 939414)
f: 4 / 27 (969234 / 969234)
f: 5 / 27 (984497 / 984497)
f: 6 / 27 (992218 / 992218)
f: 7 / 27 (996102 / 996102)
f: 8 / 27 (998049 / 998049)
f: 9 / 27 (999024 / 999024)
f: 10 / 27 (999512 / 999512)
f: 11 / 27 (999756 / 999756)
f: 12 / 27 (999878 / 999878)
f: 13 / 27 (999939 / 999939)
f: 14 / 27 (999970 / 999970)
f: 15 / 27 (999985 / 999985)
f: 16 / 27 (999993 / 999993)
f: 17 / 27 (999997 / 999997)

```

```

f: 18 / 27 (999999 / 999999)
f: 19 / 27 (1000000 / 1000000)
f: 20 / 27 (1000000 / 1000000)
f: 21 / 27 (1000000 / 1000000)
f: 22 / 27 (1000000 / 1000000)
f: 23 / 27 (1000000 / 1000000)
f: 24 / 27 (1000000 / 1000000)
f: 25 / 27 (1000000 / 1000000)
f: 26 / 27 (1000000 / 1000000)
f: 27 / 27 (1000000 / 1000000)
Done: 26538865 total MC iterations.
ΔE = 1.0
f: 1 / 27 (778801 / 778801)
f: 2 / 27 (882497 / 882497)
f: 3 / 27 (939414 / 939414)
f: 4 / 27 (969234 / 969234)
f: 5 / 27 (984497 / 984497)
f: 6 / 27 (992218 / 992218)
f: 7 / 27 (996102 / 996102)
f: 8 / 27 (998049 / 998049)
f: 9 / 27 (999024 / 999024)
f: 10 / 27 (999512 / 999512)
f: 11 / 27 (999756 / 999756)
f: 12 / 27 (999878 / 999878)
f: 13 / 27 (999939 / 999939)
f: 14 / 27 (999970 / 999970)
f: 15 / 27 (999985 / 999985)
f: 16 / 27 (999993 / 999993)
f: 17 / 27 (999997 / 999997)
f: 18 / 27 (999999 / 999999)
f: 19 / 27 (1000000 / 1000000)
f: 20 / 27 (1000000 / 1000000)
f: 21 / 27 (1000000 / 1000000)
f: 22 / 27 (1000000 / 1000000)
f: 23 / 27 (1000000 / 1000000)
f: 24 / 27 (1000000 / 1000000)
f: 25 / 27 (1000000 / 1000000)
f: 26 / 27 (1000000 / 1000000)

```




```
f: 27 / 27 (1000000 / 1000000)
Done: 26538865 total MC iterations.
```

```
1 import pickle
2 with open('wlresults.pickle', 'wb') as f:
3     pickle.dump(list(Ls), f)
4     pickle.dump(wlresults, f)
```


We plot the results for  $L = 5$ , since it is in between the “slow” behavior of  $L = 3$  and the “fast” (phase transition?) behavior of  $L = 7$ .

```
1 L = Ls[2]
2 Es, S, H = wlresults[2]

1 plt.plot(Es / L**2, S)
2 plt.xlabel("E / N")
3 plt.ylabel("log g(E) + C");
```



```
1 plt.plot(Es / L**2, H)
2 plt.xlabel("E / N")
3 plt.ylabel("Visits");
```



### 17.2.1 Calculating canonical ensemble averages

```

1  gspl = interpolate.splrep(Es, S, s=0*np.sqrt(2))
2  gs = np.exp(interpolate.splev(Es, gspl) - min(S))

```

Translate energies to have minimum zero so that  $Z$  is representable.

```

1  nEs = Es - min(Es)

1  Z = lambda beta: np.sum(gs * np.exp(-beta * nEs))

```

Ensemble averages

```

1  beta_s = [np.exp(k) for k in np.linspace(-7, 3, 500)]
2  E_mu = lambda beta: np.sum(nEs * gs * np.exp(-beta * nEs)) / Z(beta)
3  E2 = lambda beta: np.sum(nEs**2 * gs * np.exp(-beta * nEs)) / Z(beta)
4  CV = lambda beta: (E2(beta) - E_mu(beta)**2) * beta**2
5  F = lambda beta: -np.log(Z(beta)) / beta
6  Sc = lambda beta: beta * E_mu(beta) + np.log(Z(beta))


```

Helmholtz free energy

```

1 plt.plot(np.log(βs), [F(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("Helmholtz free energy")
4 plt.show()

```




### Heat capacity

```

1 plt.plot(np.log(βs), [CV(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("Heat capacity")
4 plt.show()

```




## Entropy

```

1 plt.plot(np.log(βs), [Sc(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("S(β) + C")
4 plt.show()

```



## 18 Progress summary

*June 12, 2020* This week, I investigated different methods for obtaining thermodynamic quantities from simulations. The issue is that quantities like entropy and the Helmholtz free energy depend on global properties of the phase space (the probability or density of a microstate), and thus cannot be constructed as cumulants of microstates during a simulation. A related issue is the improbability of “tunneling” across energy barriers when taking the usual temperature-weighted steps.

These considerations motivate histogram-based methods, like the Wang-Landau algorithm that I implemented. Instead of operating in the canonical ensemble, we take a biased random walk on energies so that the result is a flat histogram of visited energies. The density of states from this process may then be used to compute a canonical partition function. Modifications where we keep a joint density of states with respect to another variable make other ensembles accessible.

Further progress with the Wang-Landau algorithm was slowed by a discrepancy in the steps needed between Wang and Landau’s results and mine for the 32 by 32 Ising ferromagnet. Their original paper claims a 0.035 % average error in

the density of states after only 700 000 total spin flips.<sup>2</sup> This is far fewer than the spin flips I needed, and another paper that looks at the scaling of the tunneling time (spin flips to go from ground to anti-ground state) might corroborate this.<sup>3</sup> Their tunneling times are all well above 1 720 000 (an eighth of their  $\tau_{\text{exact}}$ ) for the same simulation. Success in the Wang-Landau algorithm requires visiting all energies many times, so execution takes several tunneling times. We are still trying to resolve this discrepancy, as well as other vague details from the original paper, like the possible choice of energy bins for continuous systems and unspecified edge behavior for energy intervals.

I have looked at other histogram methods like WHAM, as well as parallel tempering (replica exchange MCMC). Both parallel tempering and the Wang-Landau algorithm are attractive because they are easily parallelizable. I have also come back to considering the MAXENT/MEM image reconstruction technique and the issues of entropy and feature representation again.

## 19 Description of MAXENT

*June 13, 2020* I describe the basic idea of the MAXENT reconstruction technique, and implement the most naïve approach. This makes it clear that a better optimization algorithm is needed. The idea here is less to make it work than to get an idea of how it works, and then use a better software package later on if we do end up using MAXENT.

## 20 Maximum-entropy reconstruction

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy import signal, misc
```

Given a measurement  $D$  of a system  $I_0$ , we wish to reconstruct  $I_0$  from  $D$ . We take the Bayesian approach and find the  $I$  that maximizes  $P(I|D) \propto P(D|I)P(I)$ , where the likelihood is related to the error in  $D$  due to  $I$  and the prior is  $P(I) = \exp(\lambda S(I))$ , where  $S$  is some notion of the entropy of an image. If  $D = T(I_0)$ , then  $P(D|I) = \exp(-E(T(I), D))$ , for some metric  $E$  on  $D$ 's. From the form of the objective  $\ln P(D|I) + \ln(I) = -E(T(I), D) + \lambda S(I)$ , we see that the entropy acts as a regularizer.

---

<sup>2</sup>10.1103/PhysRevLett.86.2050

<sup>3</sup>10.1103/PhysRevLett.92.097201

The astronomers have both  $I$  and  $D$  as intensity lists  $\mathbb{Z}_N \rightarrow \mathbb{R}_{\geq 0}$ , where

$$S(I) = \sum_i I_i \ln I_i$$

and

$$E(D', D) = \sum_i \exp \left( -\frac{(D'_i - D_i)^2}{2\sigma_i} \right)$$

(with empirical  $\sigma_i$ ). The transformation  $T$  is convolution with the point spread function of the telescope.

To perform the optimization, we need not only the functions  $S$  and  $E$ , but also a procedure to modify candidate images.

```

1 def maxent_objective(D, I, λ): # For minimization
2     return D.E(I.transform()) - λ*I.S

1 # Greedy for now, but can be something like simulated annealing
2 def maxent(D, I, λ = 1, N = 1_000_000, ε = 1e-8):
3     f0 = np.inf
4     f = maxent_objective(D, I, λ)
5     i = 0
6     while ε < f0 - f and i < N:
7         I.propose()
8         fv = maxent_objective(D, I, λ)
9         if fv < f: # Greedy
10             f0, f = f, fv
11             I.accept()
12             i += 1
13         if i % (N // 20) == 0:
14             print("Maxent: {} / {}".format(i, N))
15
16     print("Maxent: i: {} / {}, Δf: {}".format(i, N, f0 - f))
17     return I, i, f

```

## 20.1 Example: 1D Point from Gaussian

```

1 class Gaussian:
2     def __init__(self, μ=0, σ=1):
3         self.μ = μ
4         self.σ = σ
5     def E(self, Gv):
6         return (self.μ - Gv.μ)**2 + (self.σ - Gv.σ)**2
7

```

```

8  class Point:
9      def __init__(self, x=0):
10         self.x = x
11         self.dx = 0
12         self.S = self.entropy()
13     def propose(self):
14         self.dx = np.random.rand() - 0.5
15         self.S = self.entropy()
16     def accept(self):
17         self.x += self.dx
18         self.dx = 0 # Idempotence
19     def entropy(self):
20         return -(self.x - 10)**2 # Opposite from true max
21     def transform(self):
22         return Gaussian( $\mu$  = self.x + self.dx)

1  D = Gaussian(0, 1)
2  I = Point(9)
3  I0, _, _ = maxent(D, I,  $\lambda$  = 1, N = 1_000_000,  $\epsilon$  = 1e-4)
4  I0.x

```

Maxent: i: 1000000 / 1000000,  $\Delta f$ : 0.011299906795997572

4.268883578482481

Results are sort of near 5. The optimization is terrible, but you get the idea: the entropy shifts the best point away from zero.

## 20.2 Example: Image from PSF convolution (measurement)

```

1  class DImage:
2      def __init__(self, I):
3          self.I = I
4      def E(self, Dv):
5          return np.sum((self.I - Dv.I)**2)
6
7  class IImage:
8      def __init__(self, I):
9          self.I = I

```




```


10     self.w, self.h = np.shape(I)
11     self.i, self.j = 0, 0
12     self.I0 = 0
13     self.Iv = 0
14     n = int(np.sqrt(self.w * self.h))
15     self.G = signal.windows.gaussian(n // 10, n // 50)
16     self.S = self.entropy()
17     def propose(self):
18         self.i, self.j = np.random.randint(self.w), np.random.randint(self.h)
19         I0 = self.I[self.i, self.j]
20         self.I0 = I0
21         self.Iv = np.random.randint(256)
22         self.S = self.entropy()
23     def accept(self):
24         self.I[self.i, self.j] = self.Iv
25     def entropy(self):
26         return -np.sum(np.log(self.I + 1)) + self.I0*np.log(self.I0 + 1) -
           ↪ self.Iv*np.log(self.Iv + 1)
27     def transform(self):
28         Iv = self.I.copy()
29         Iv[self.i, self.j] = self.Iv
30         return DImage(signal.sepfir2d(Iv, self.G, self.G))

1  I = IImage(misc.ascent())[250:350,250:350])
2  plt.imshow(I.I, cmap='gray');

```



```
1 plt.imshow(I.transform().I, cmap='gray');
```




```

1 I = IImage(misc.ascent()[250:350,250:350])
2 Iguess = IImage(128 * np.ones((I.w, I.h), dtype=int))
3 I0, _, _ = maxent(I.transform(), Iguess,  $\lambda = 1e-9$ , N = 1_000_000,  $\epsilon = 1e-20$ ) # Just do the max
    ↪ iterations

Maxent: 50000 / 1000000
Maxent: 100000 / 1000000
Maxent: 150000 / 1000000
Maxent: 200000 / 1000000
Maxent: 250000 / 1000000
Maxent: 300000 / 1000000
Maxent: 350000 / 1000000
Maxent: 400000 / 1000000
Maxent: 450000 / 1000000
Maxent: 500000 / 1000000
Maxent: 550000 / 1000000
Maxent: 600000 / 1000000
Maxent: 650000 / 1000000
Maxent: 700000 / 1000000
Maxent: 750000 / 1000000
Maxent: 800000 / 1000000
Maxent: 850000 / 1000000
Maxent: 900000 / 1000000
Maxent: 950000 / 1000000
Maxent: 1000000 / 1000000
Maxent: i: 1000000 / 1000000,  $\Delta f$ : 4.847227362683043

1 plt.imshow(I0.I, cmap='gray');

```



## 21 “Greedy” painting-like pictures

*June 14, 2020* We noticed that swapping neighboring pixels in the previous simulations produced images that looked somewhat like paintings. There are other ways to make images look like paintings, and this is one of them. We randomly draw a shape of fixed size on the image. The color of the shape is the mean color of the image in the rectangle where the shape is drawn. This is done many times, and the size of the shapes gradually decreases. I call this “greedy” because choosing the mean color is the locally optimal color. This method has the advantage of being similar to the process of painting, which is why it is presented instead of the usual “oilify” filters in image processing software. An improvement would try to draw shapes that reflect the shapes of the subject matter, and perhaps add other texture. These kinds of considerations might require more global strategies.

## 22 Greedy Cubism

Draw an image by greedily drawing cubes.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from PIL import Image, ImageOps, ImageDraw
4 from scipy import signal, misc


1 img = Image.open("barbara.png")

1 def greedy_cubes(img, N):
2     xs, ys = img.size
3     art = Image.new(img.mode, (xs, ys))
4     draw = ImageDraw.Draw(art)
5     rmax = int(np.sqrt(xs*ys) / 10)
6     r = rmax
7     ε = 10
8     for i in range(N):
9         x = np.random.randint(xs)
10        y = np.random.randint(ys)
11        [np.mean(c) for c in cimg.split()]
12        r = int(rmax * (1 - (i/(N+1))**2)) + 1
13        box = [x - r, y - r, x + r, y + r]
14        color = tuple(int(np.round(np.mean(c))) for c in img.crop(box=box).split())
15        draw.ellipse(box, fill=color)
16    return art

1 art = greedy_cubes(img, 10000)

1 art

```



```
1 cimg = Image.open('test.jpg')  
2 cimg
```



```
1 art = greedy_cubes(cimg, 1000000)
2 art
```

