

# Understanding visual complexity with statistical physics

Alex Striff & Vincent Meunier

Turbulence

Noise

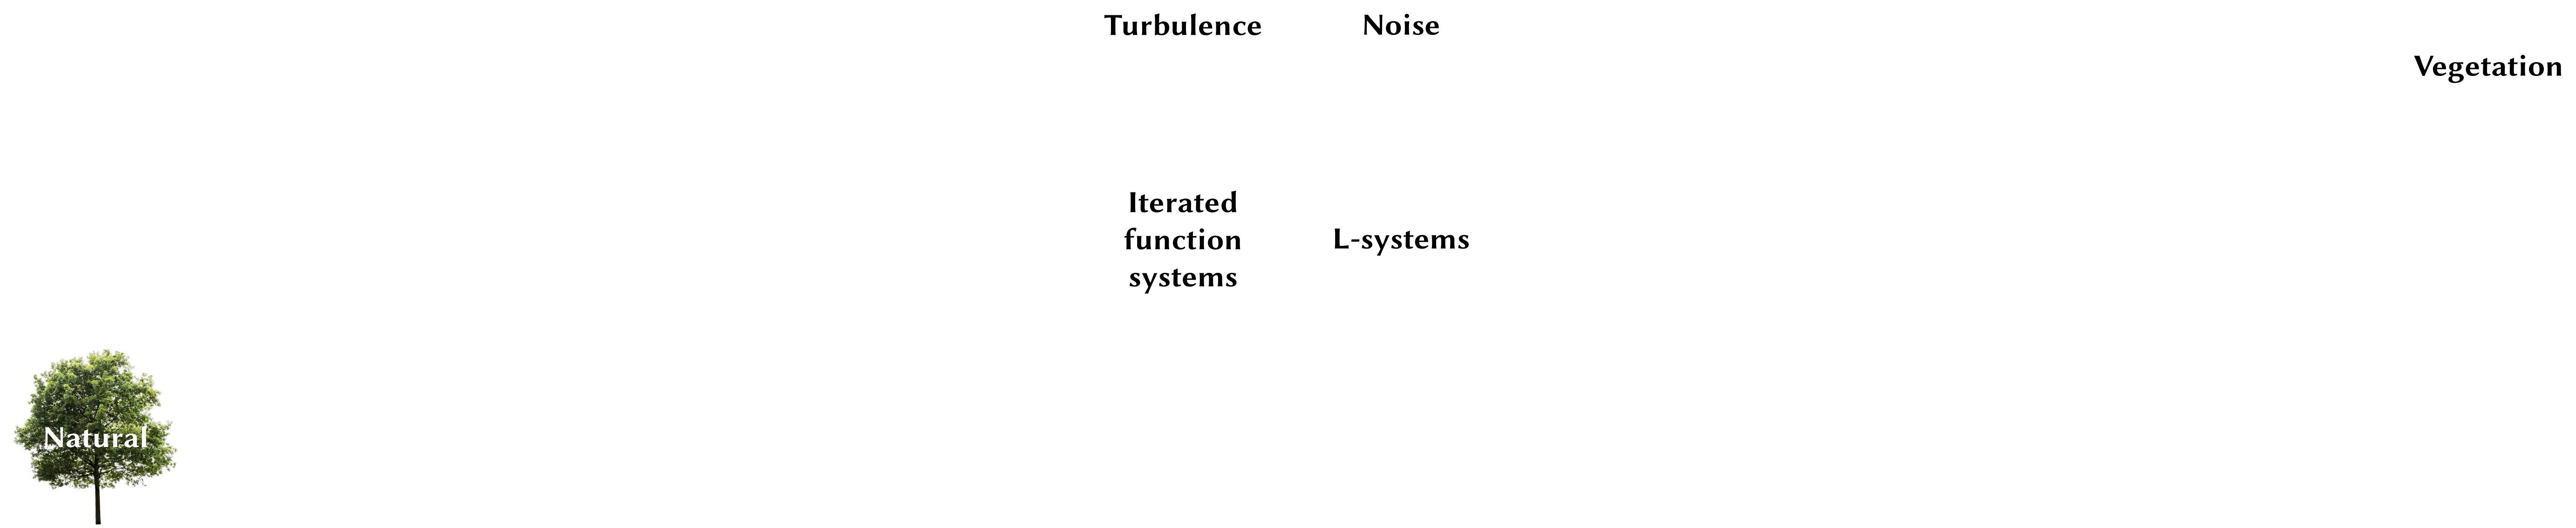
Vegetation

Iterated  
function  
systems

L-systems

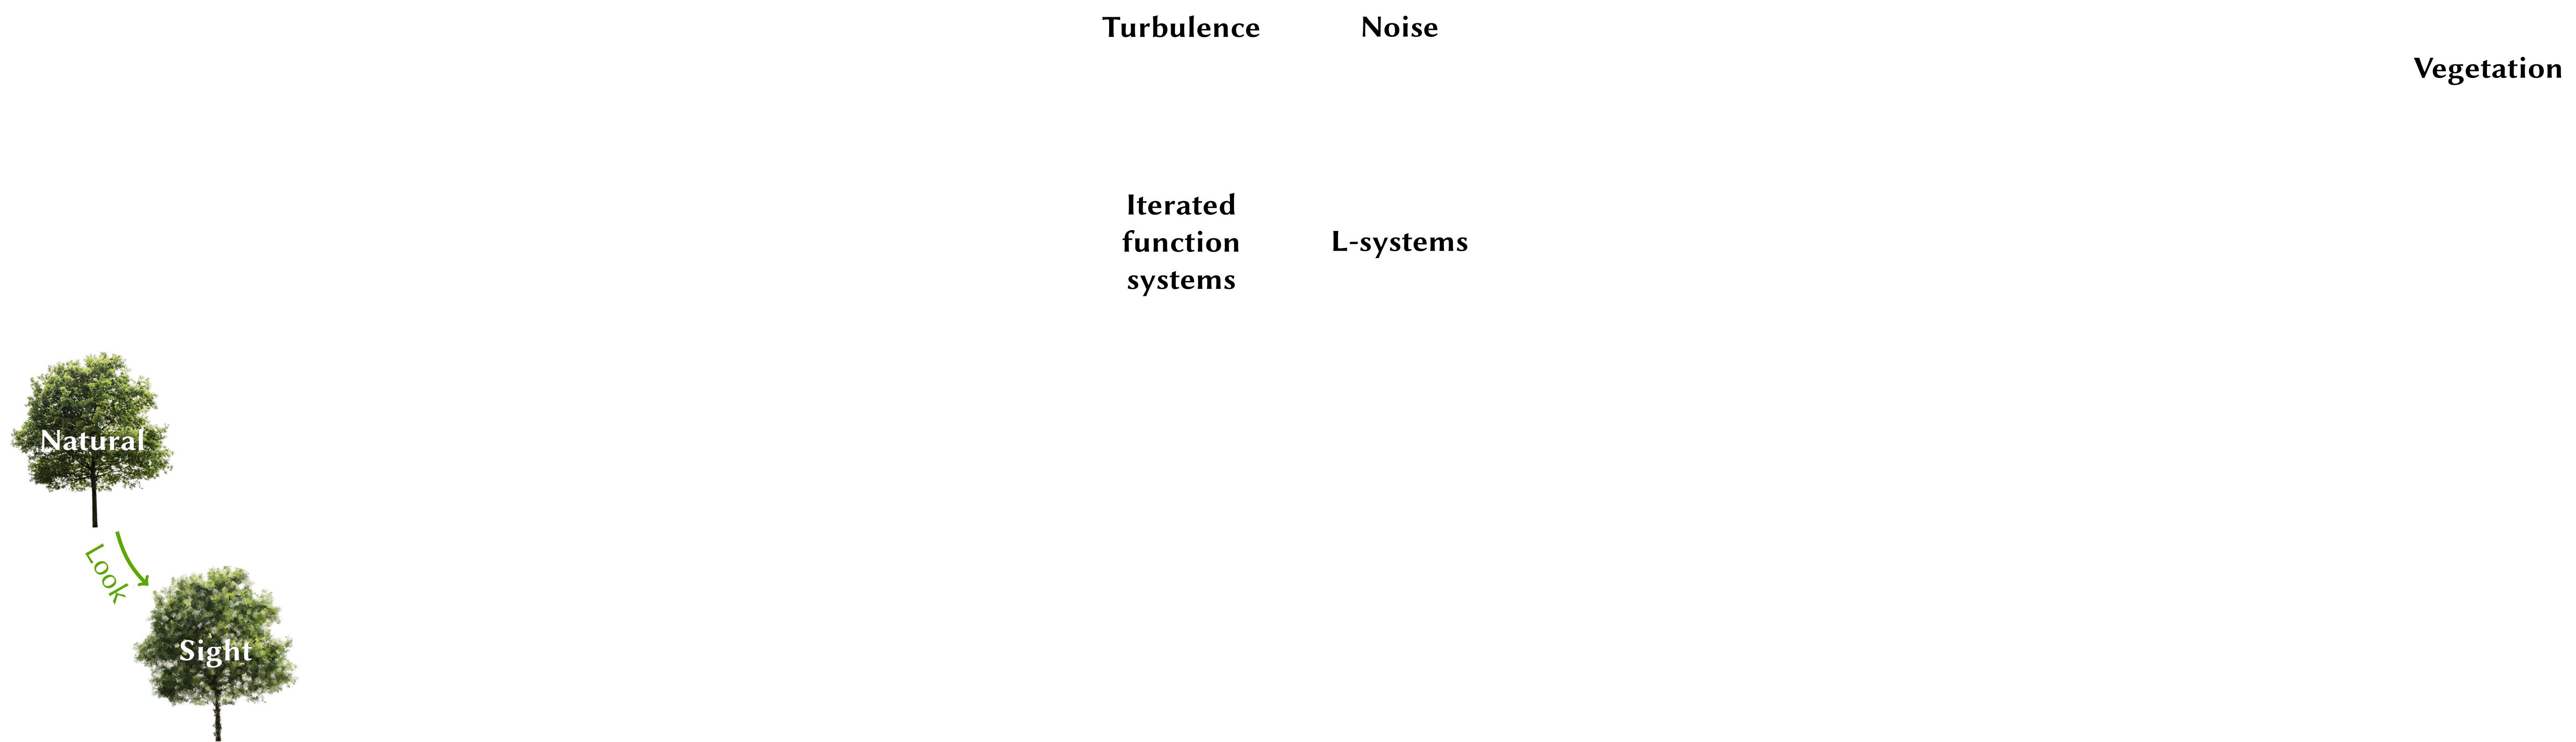
# Understanding visual complexity with statistical physics

Alex Striff & Vincent Meunier



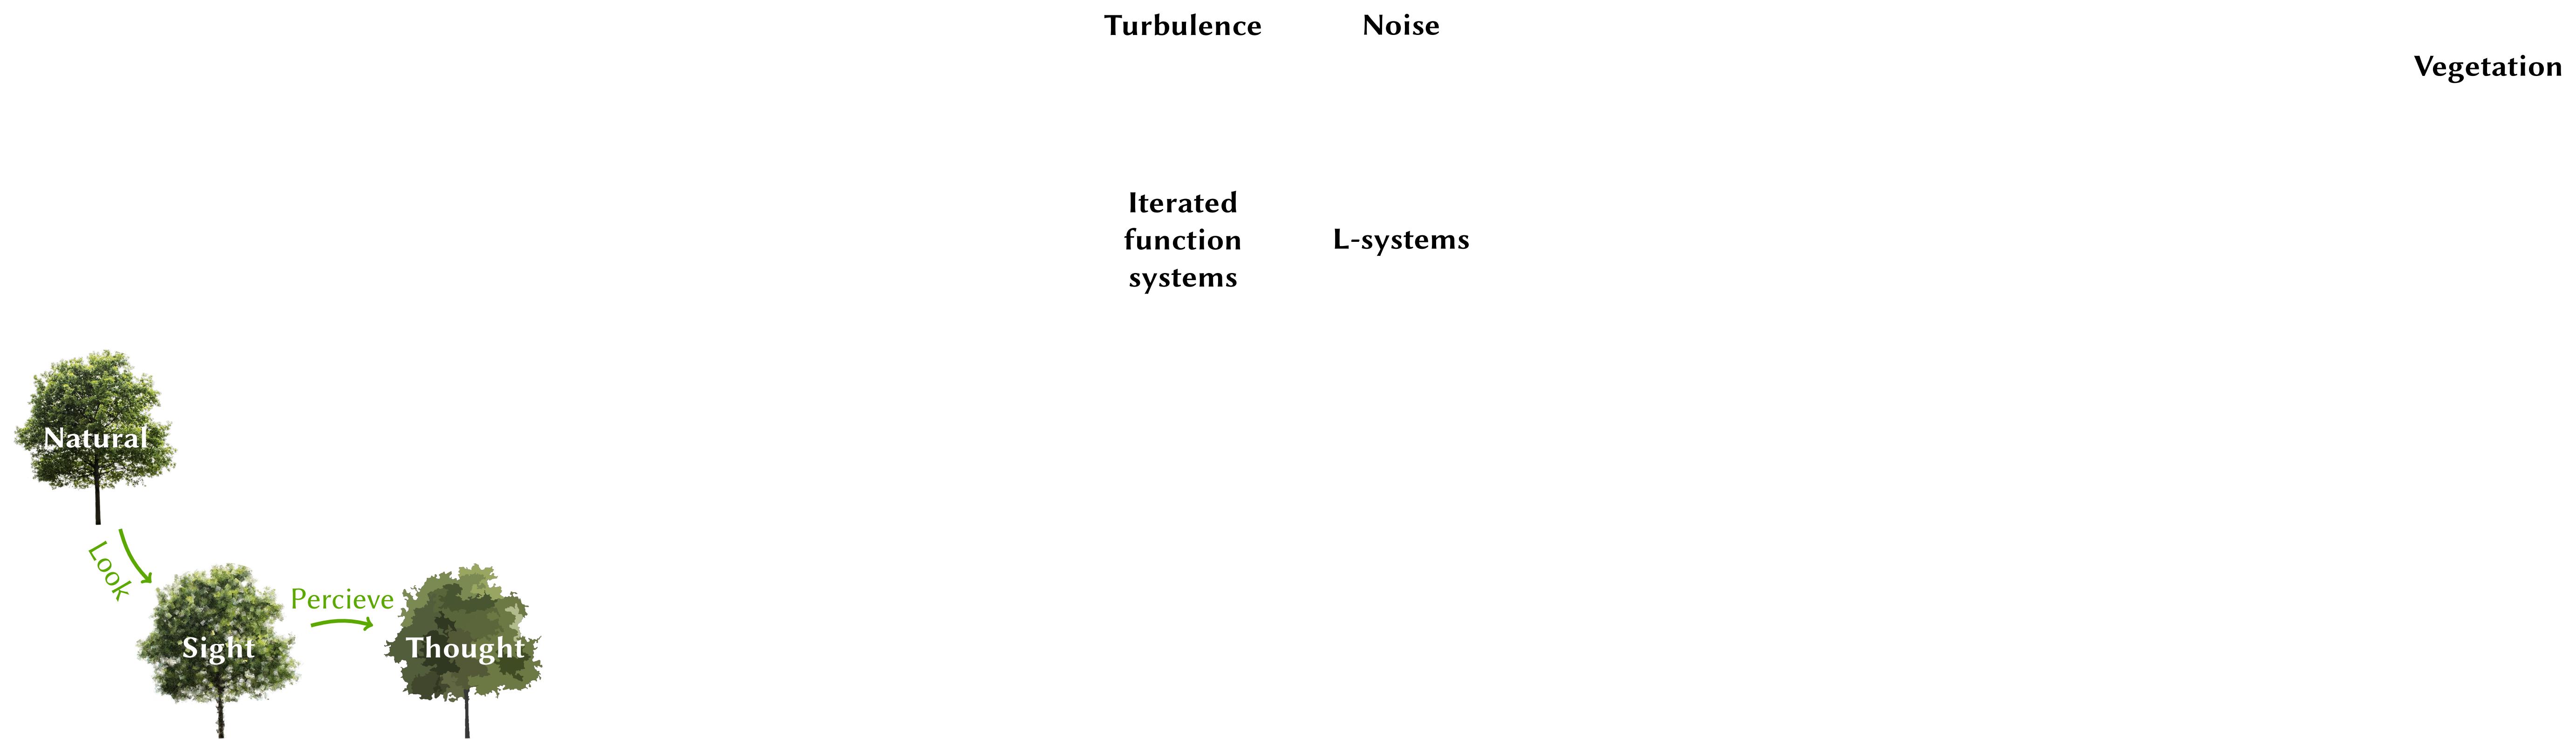
# Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier



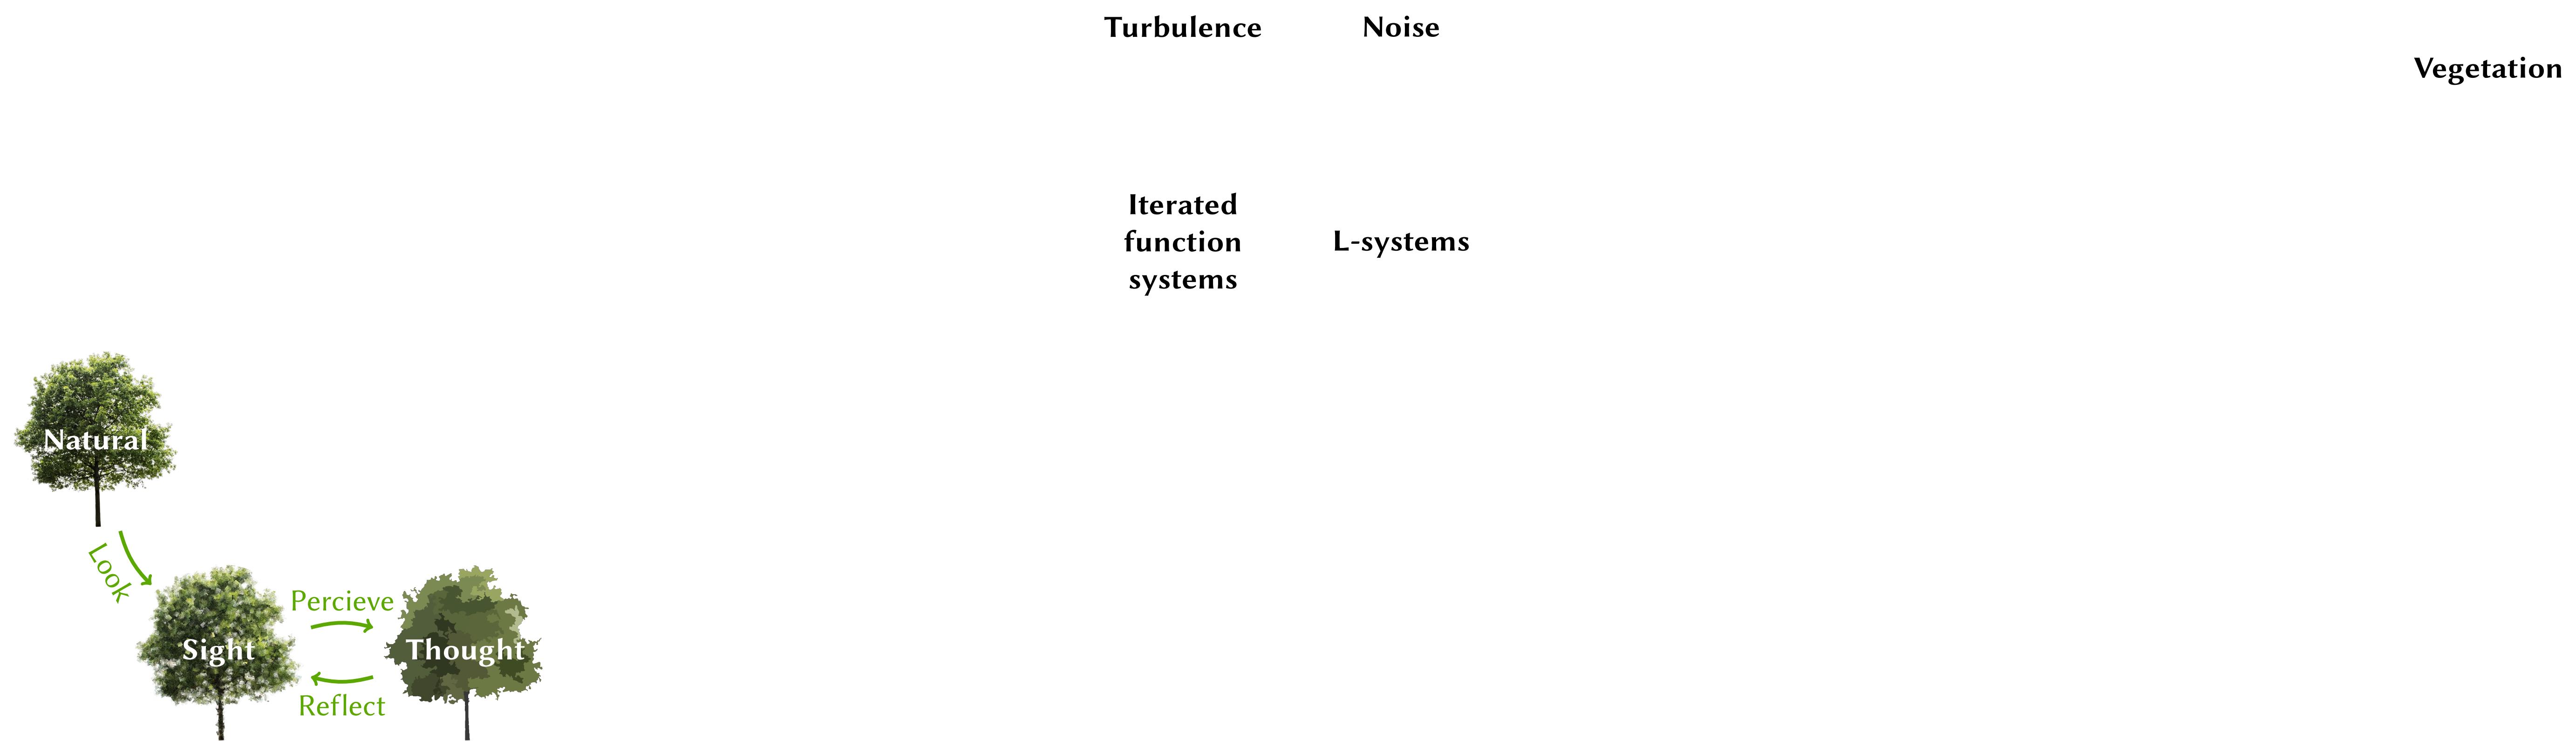
# Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier



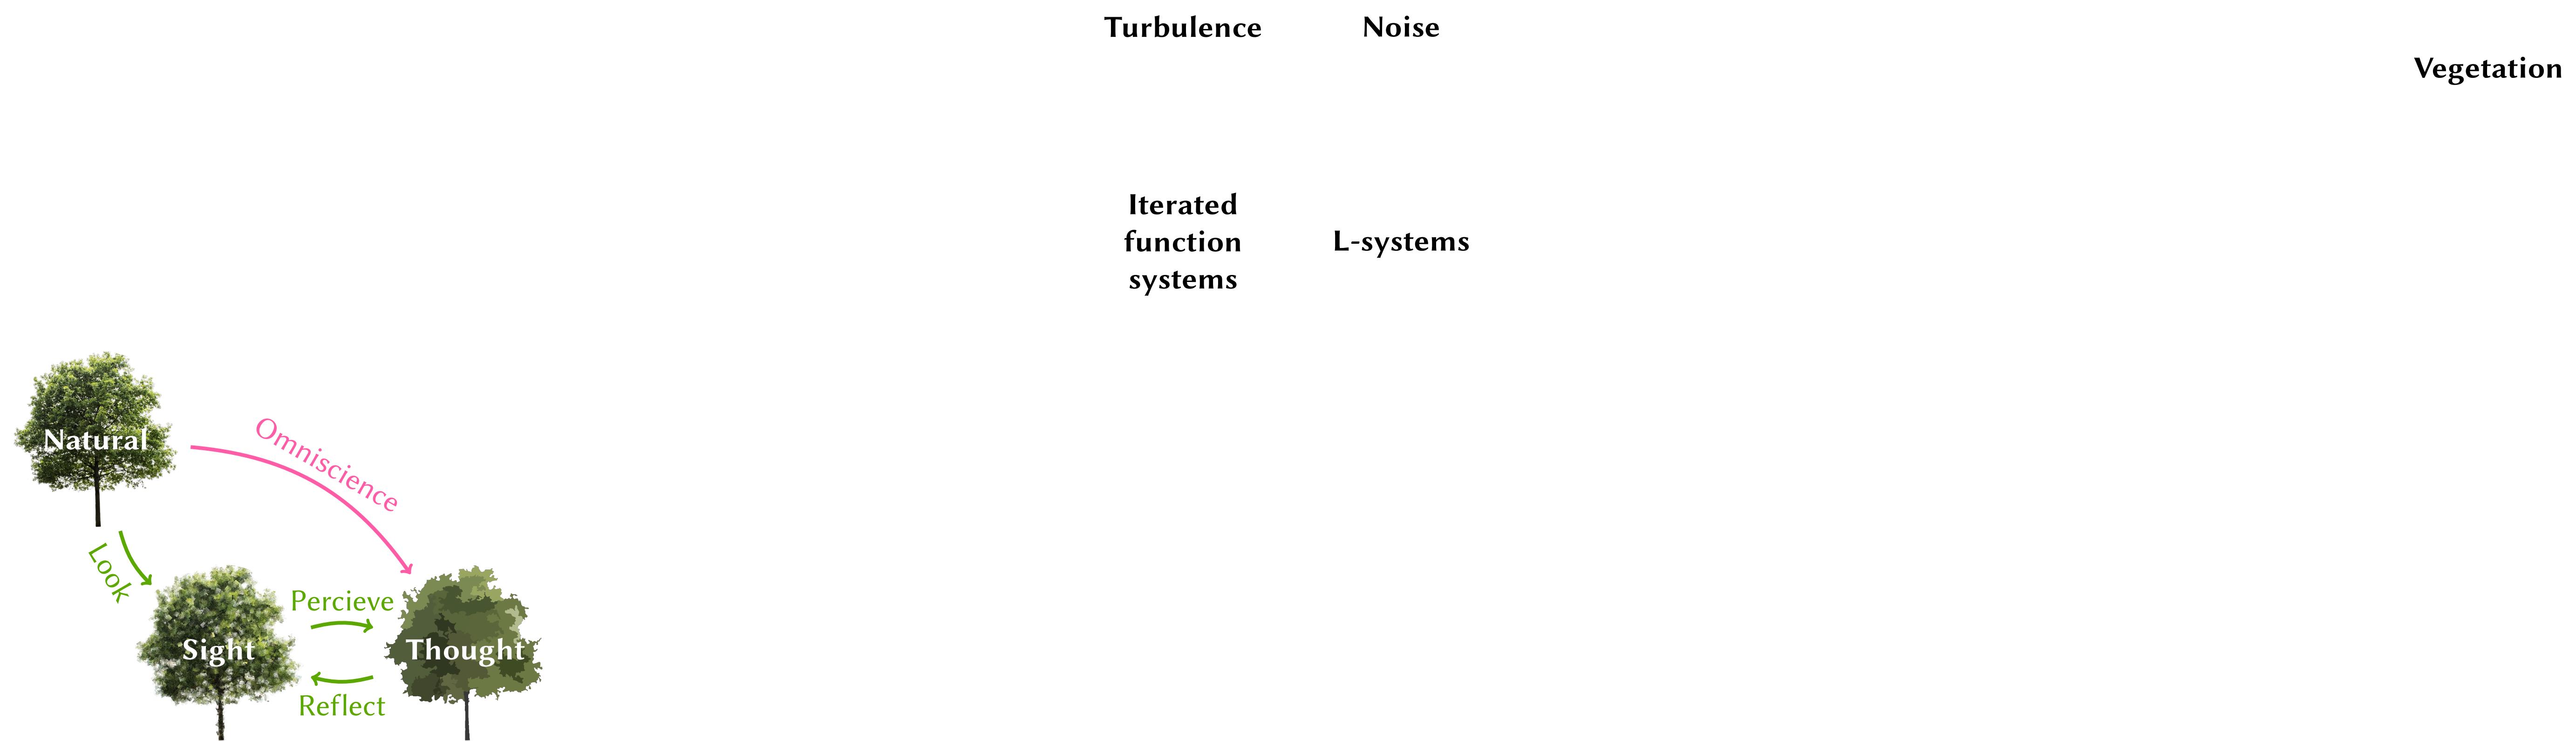
# Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier



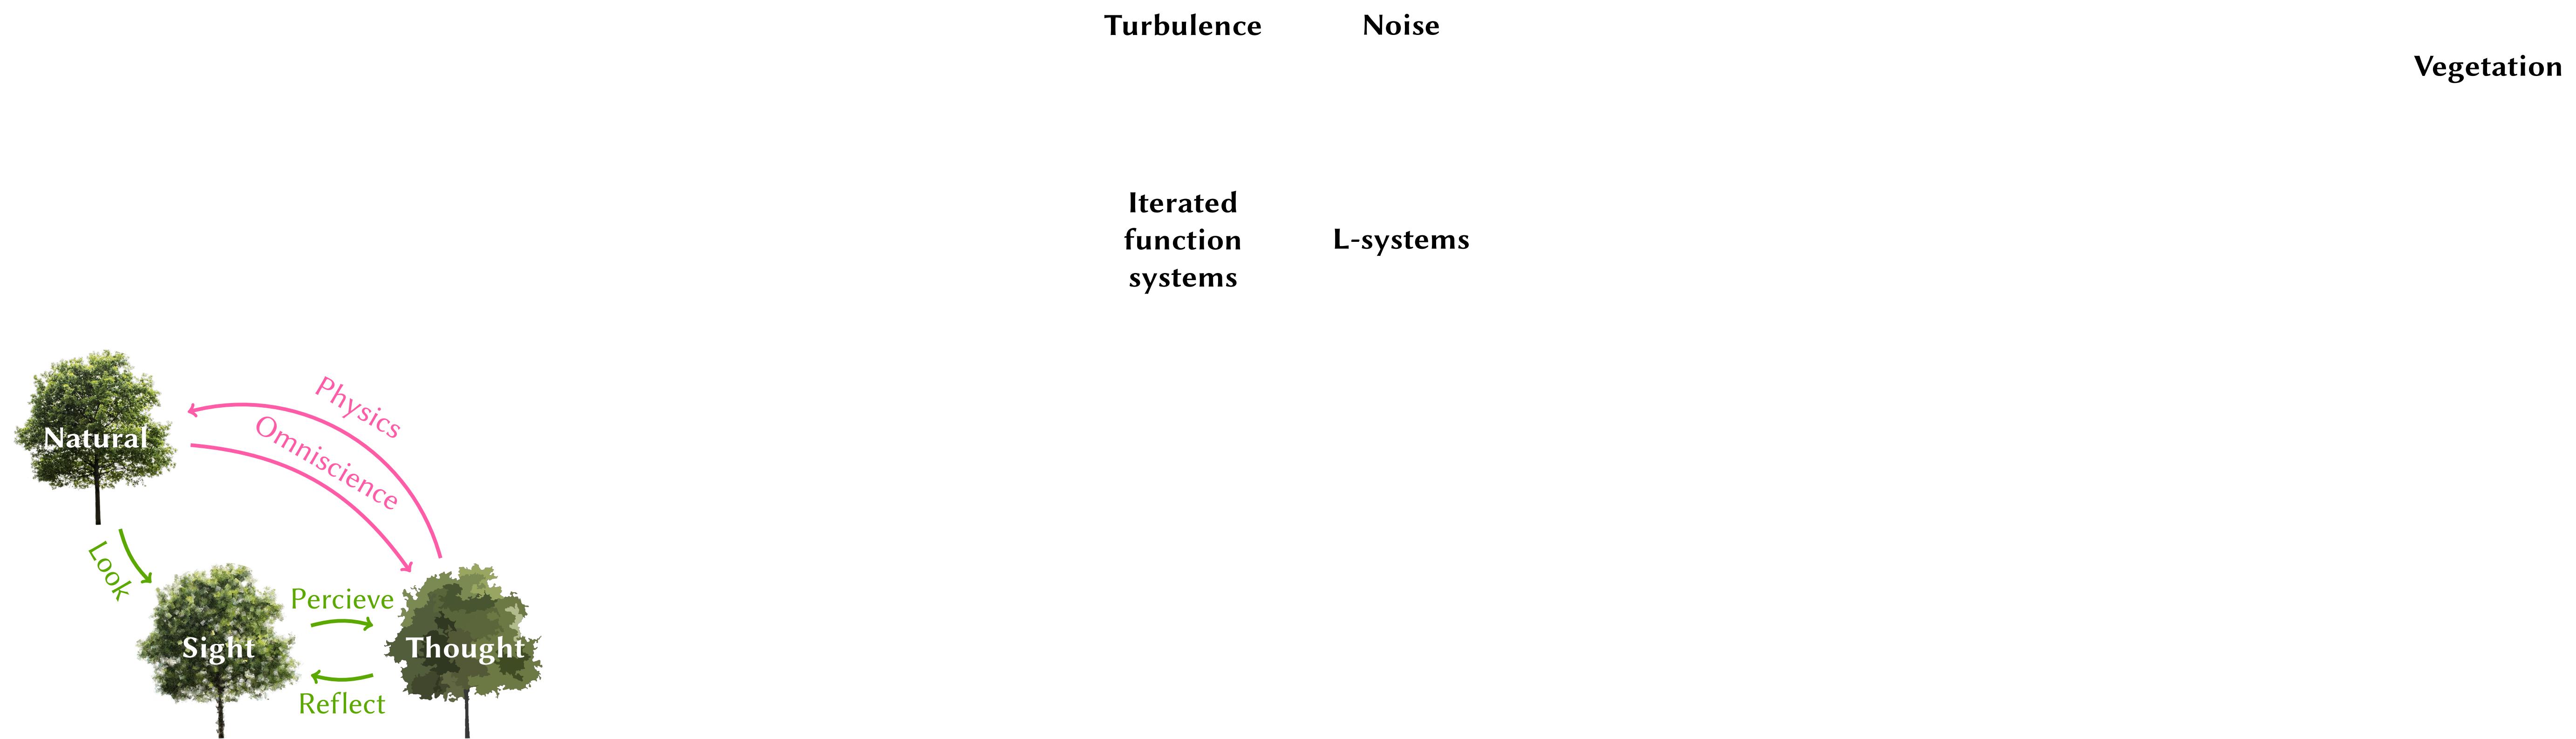
## Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier



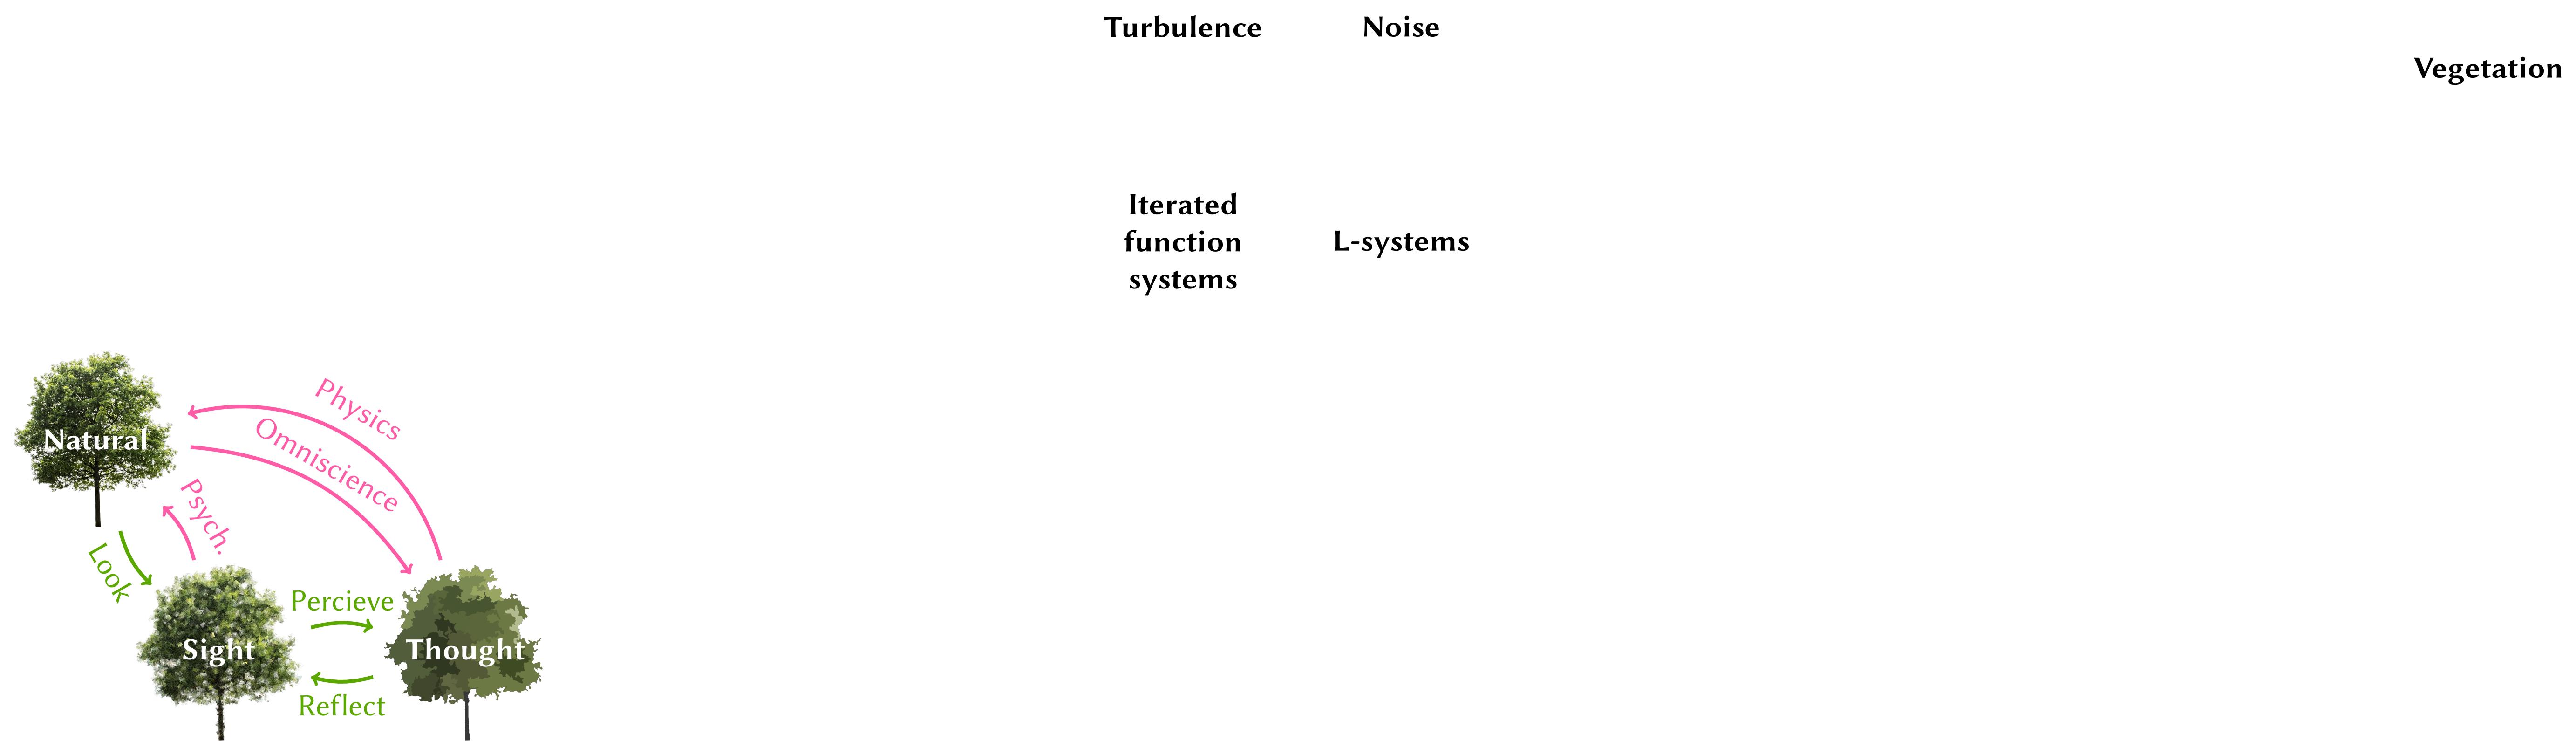
## Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier



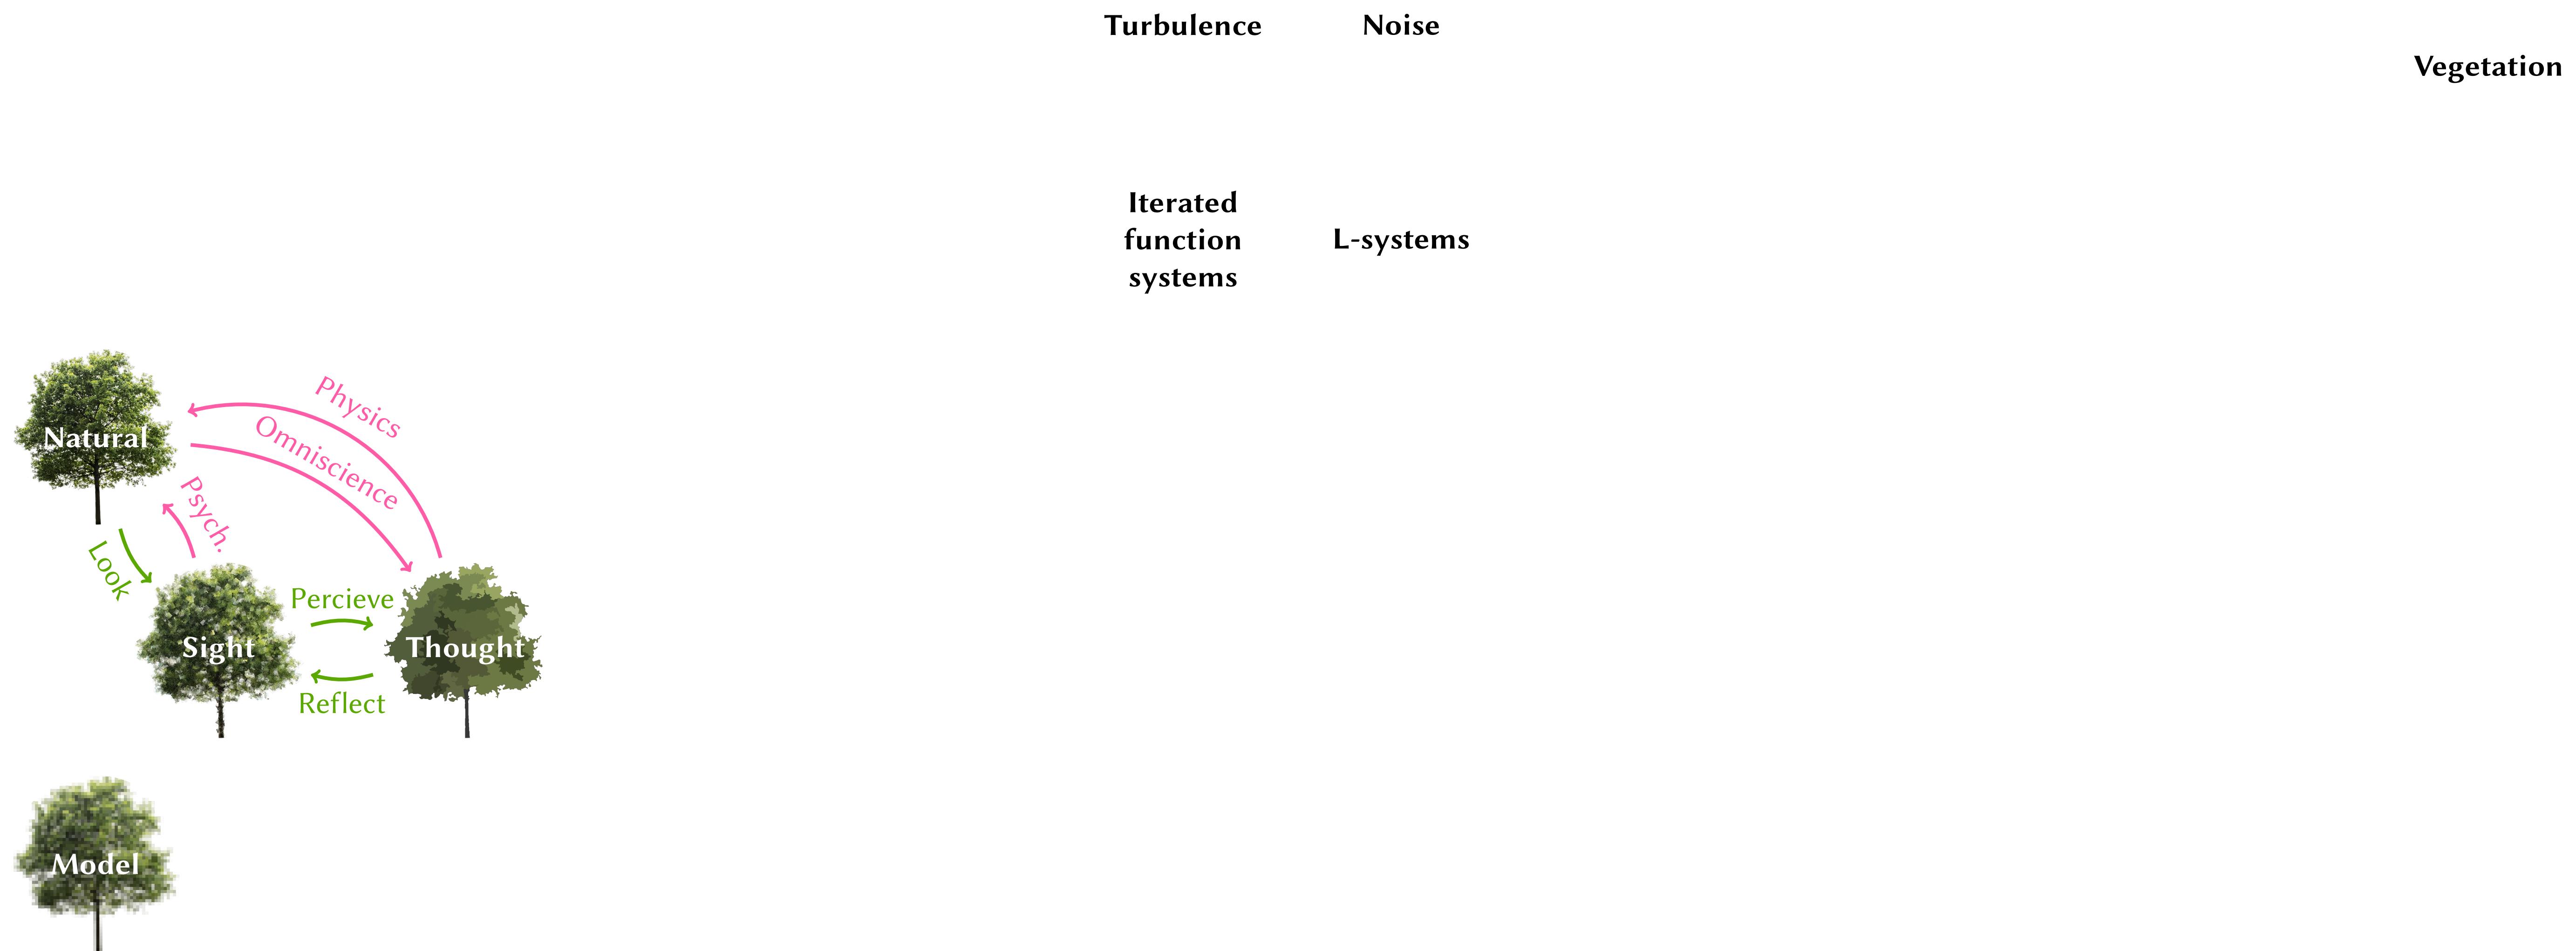
## Understanding visual complexity with statistical physics

Alex Striff &amp; Vincent Meunier

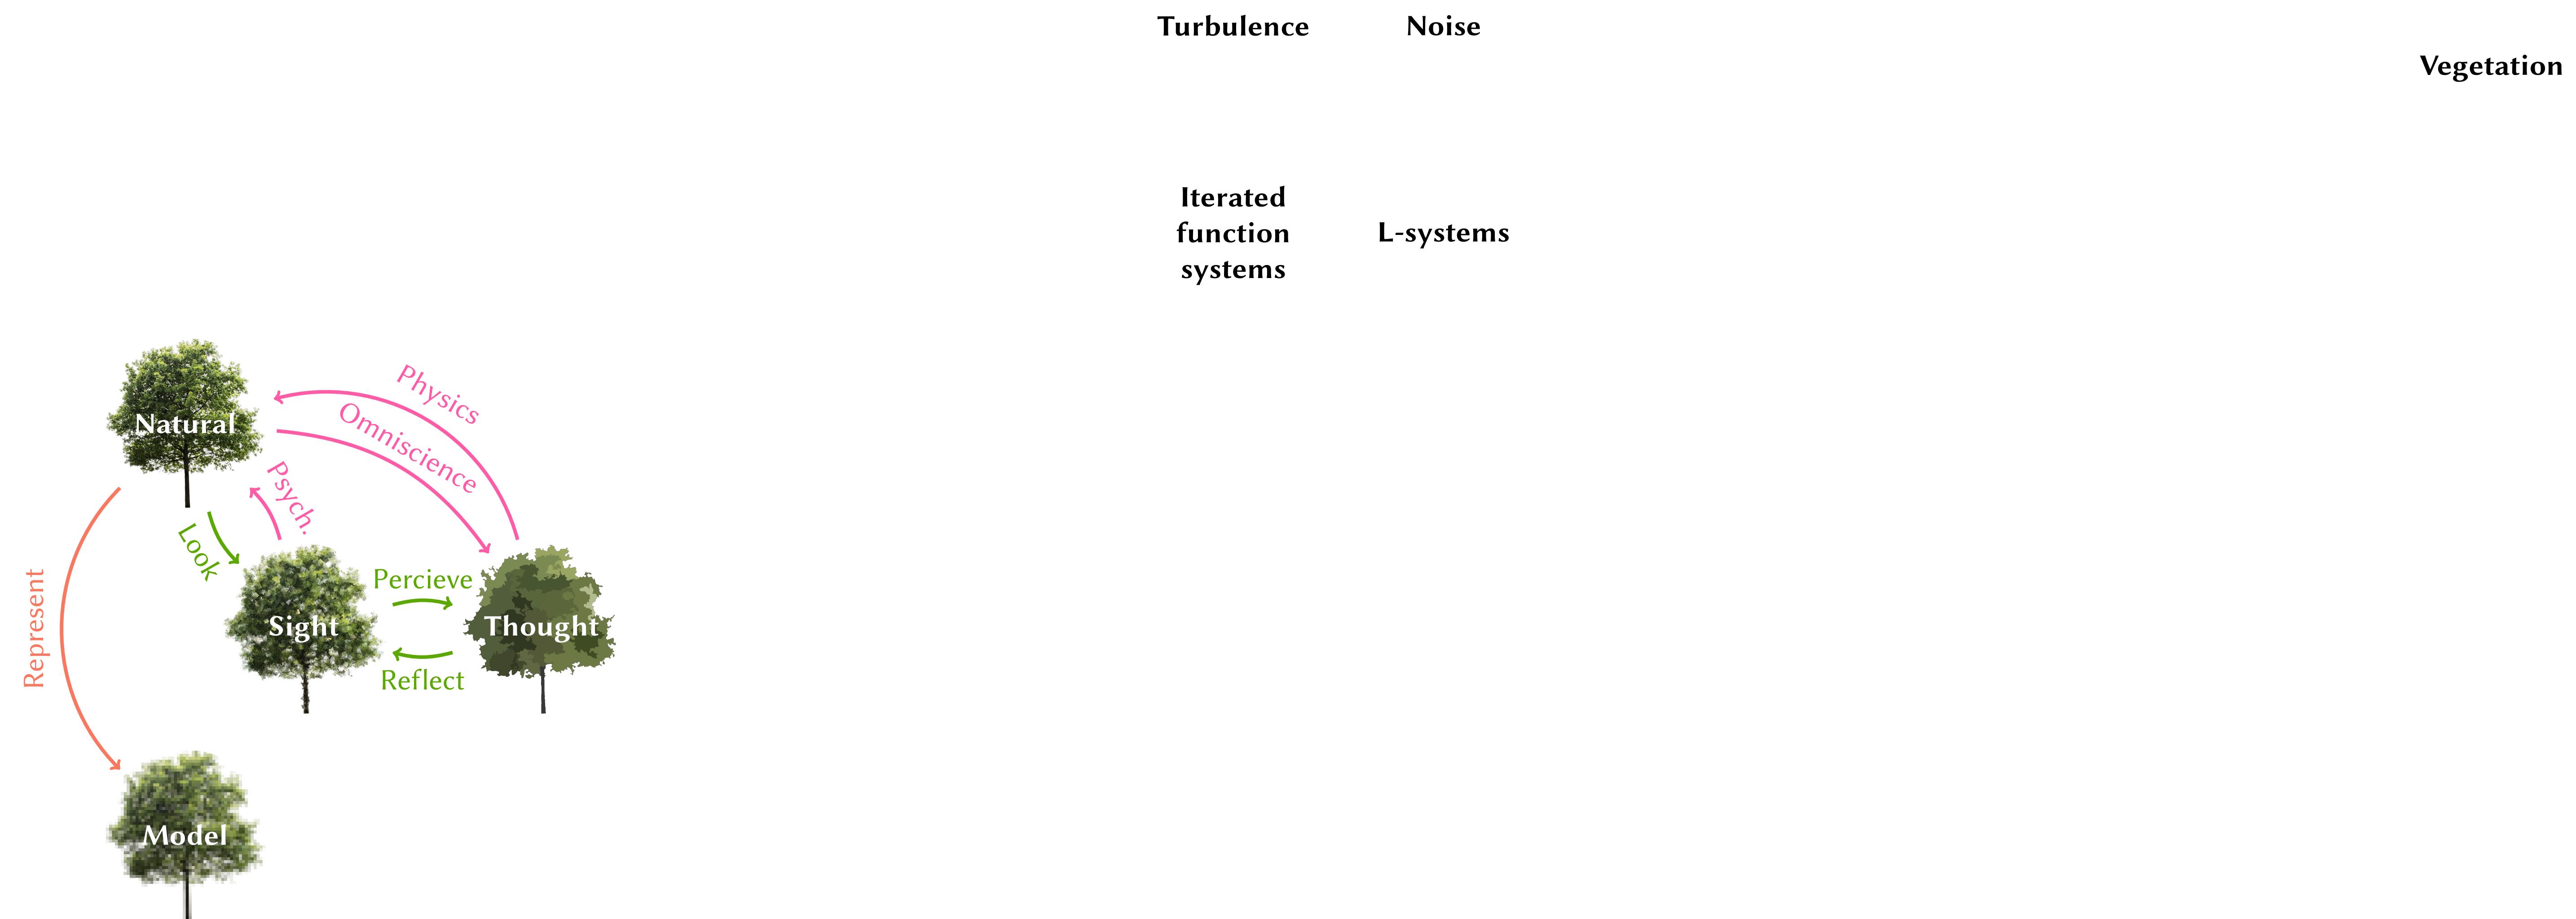


## Understanding visual complexity with statistical physics

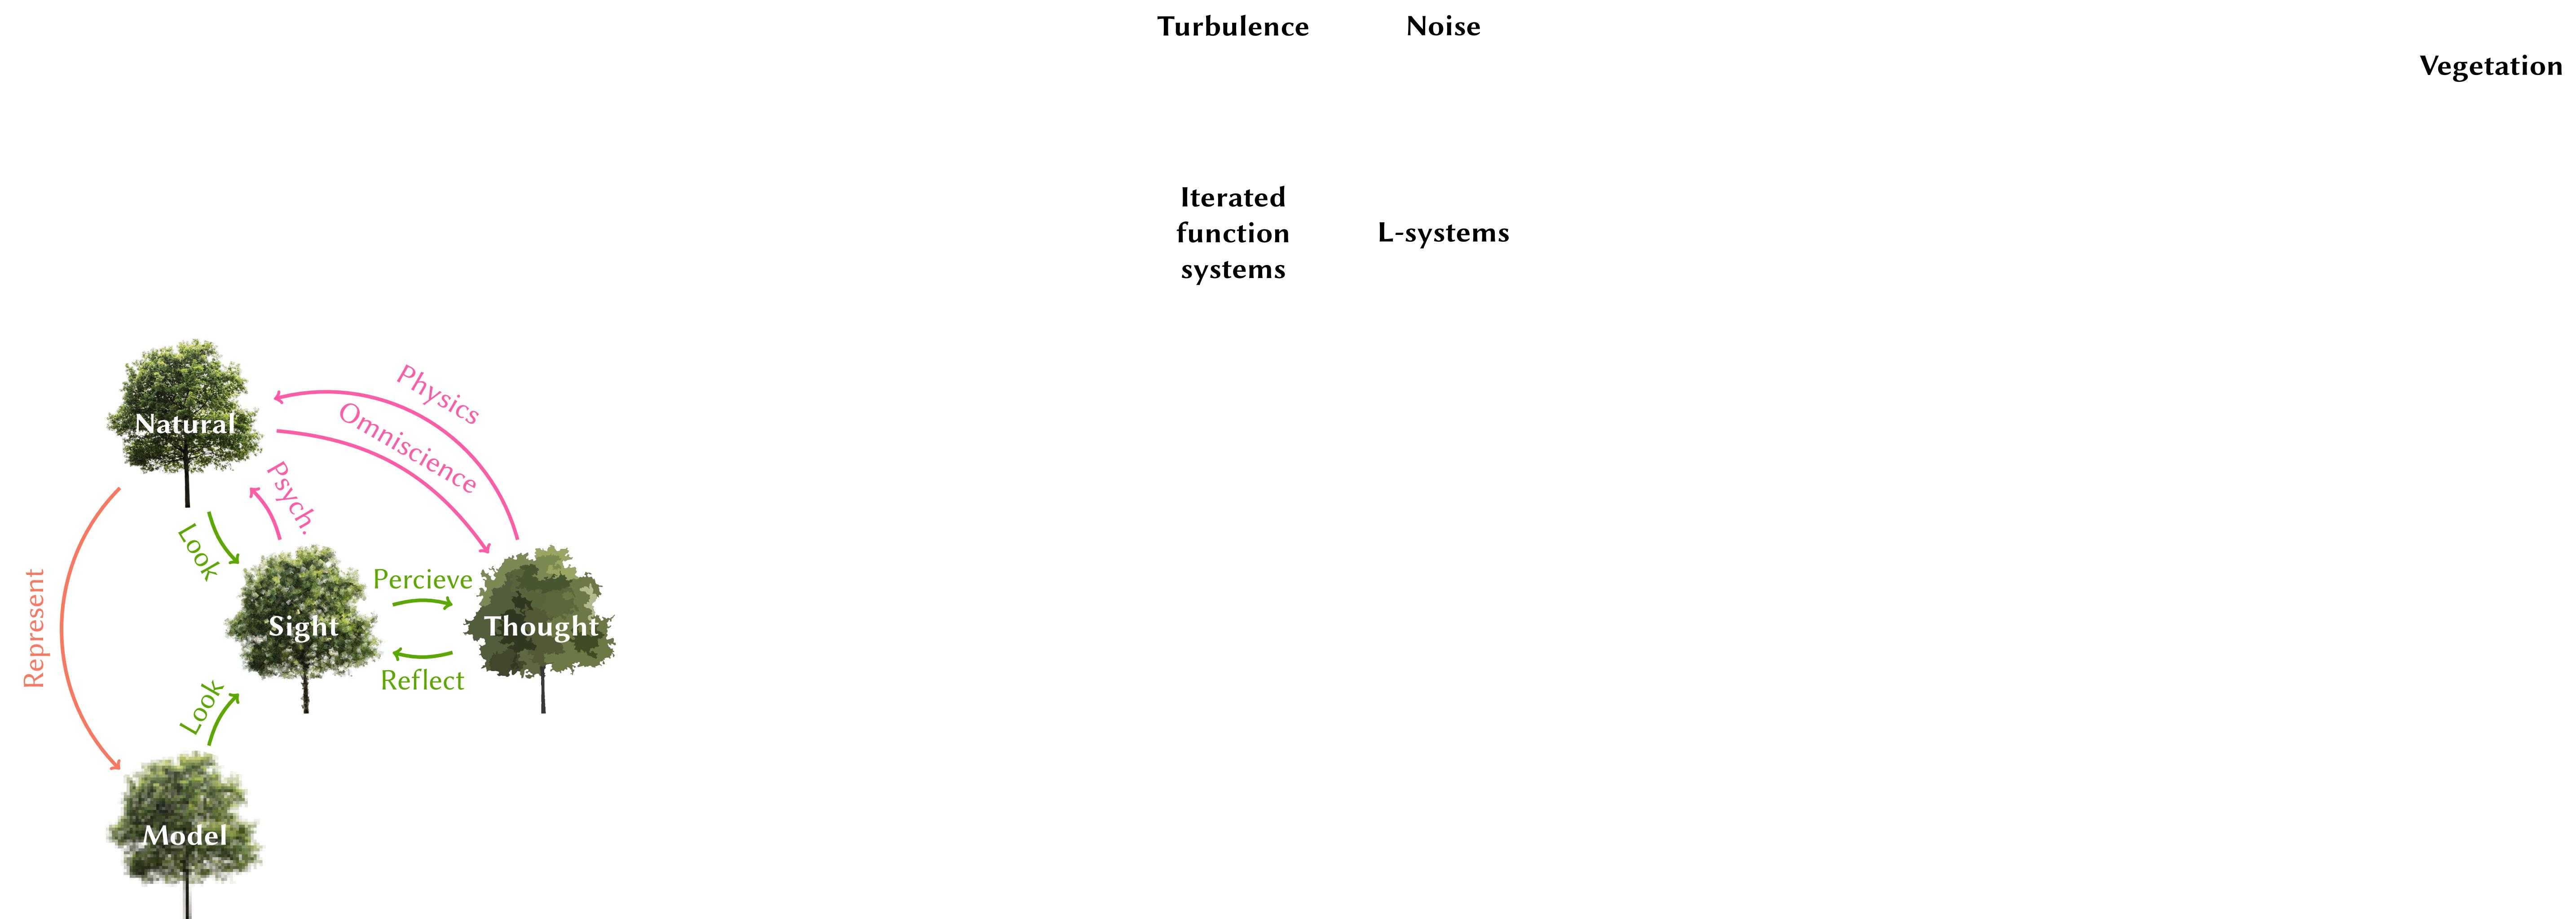
Alex Striff &amp; Vincent Meunier



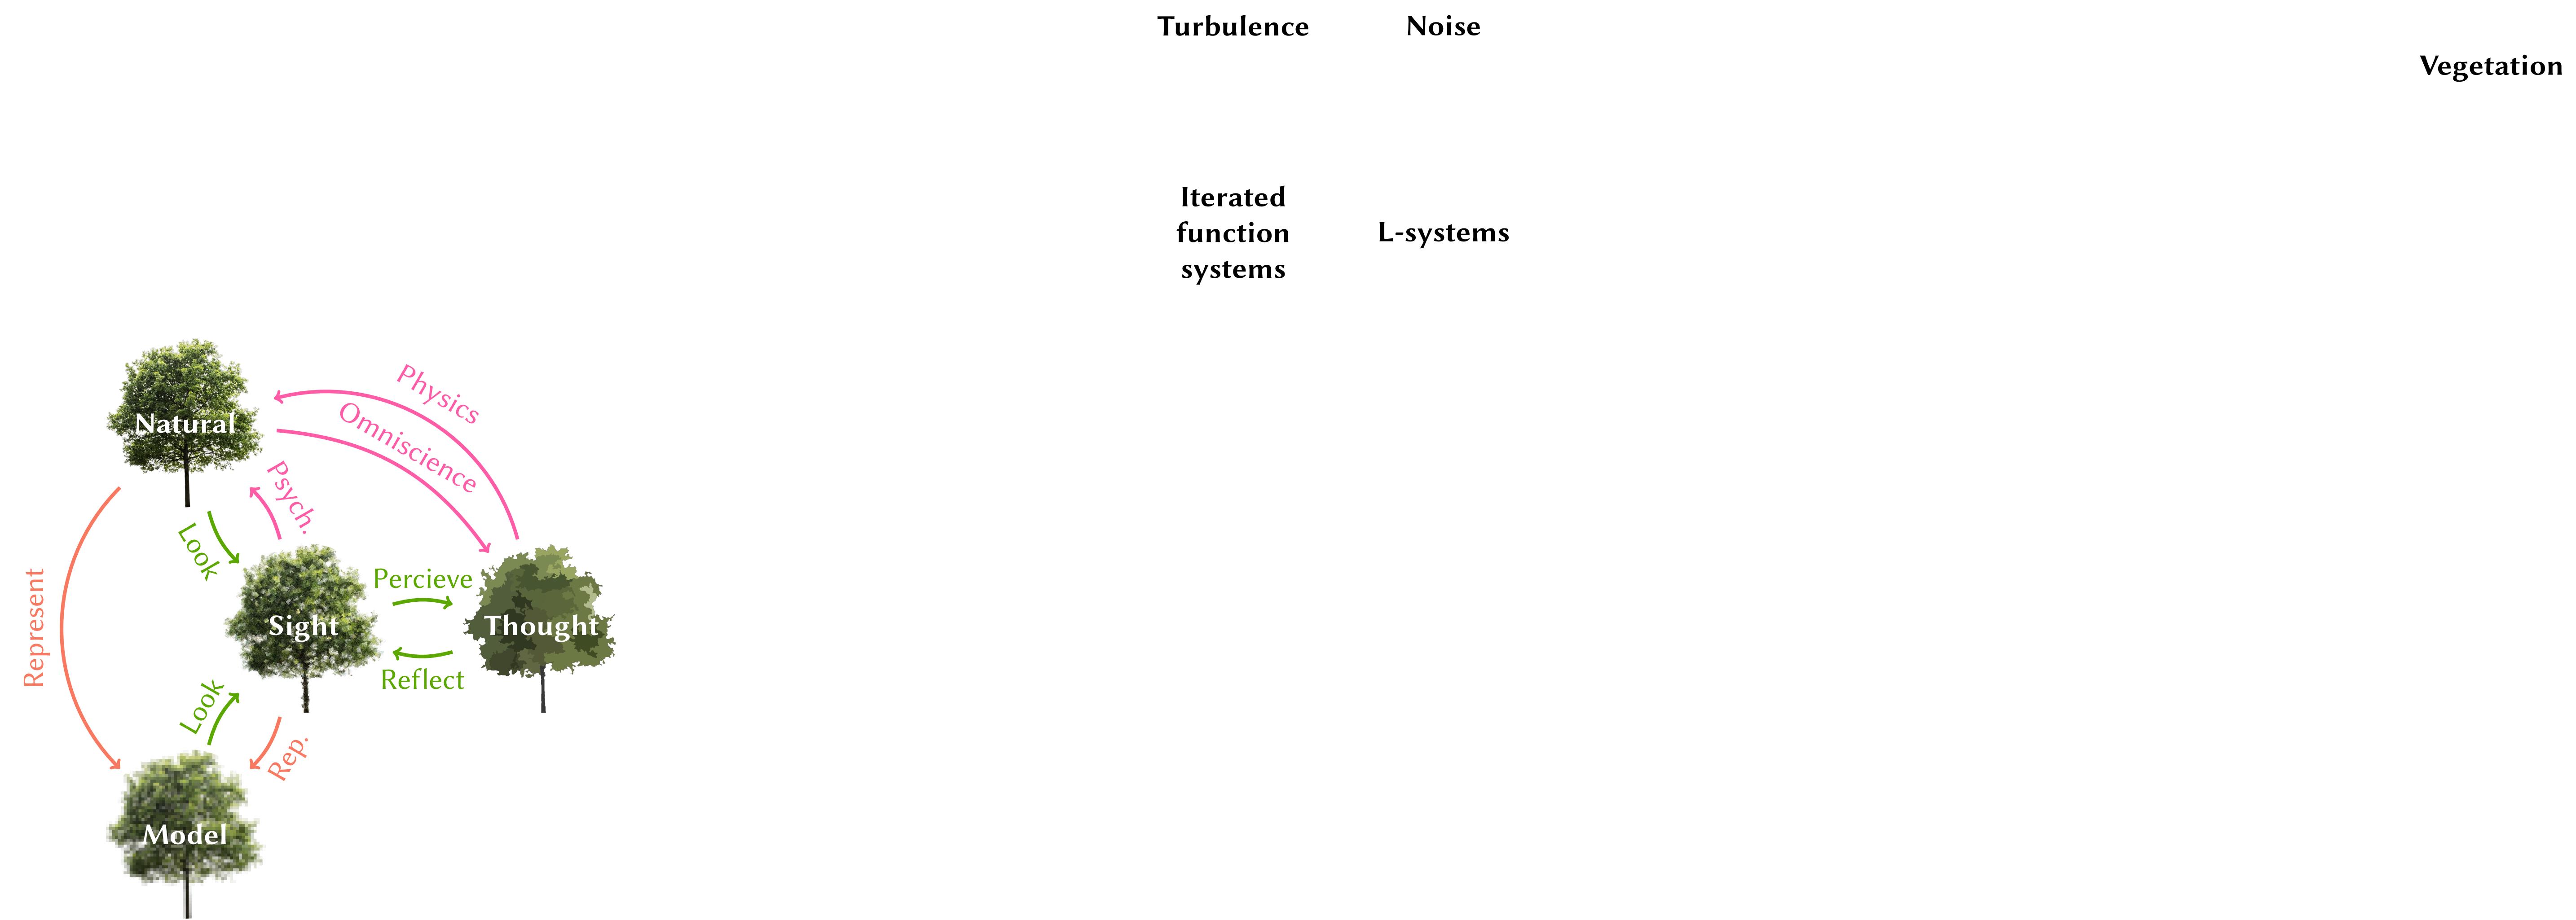
## Understanding visual complexity with statistical physics



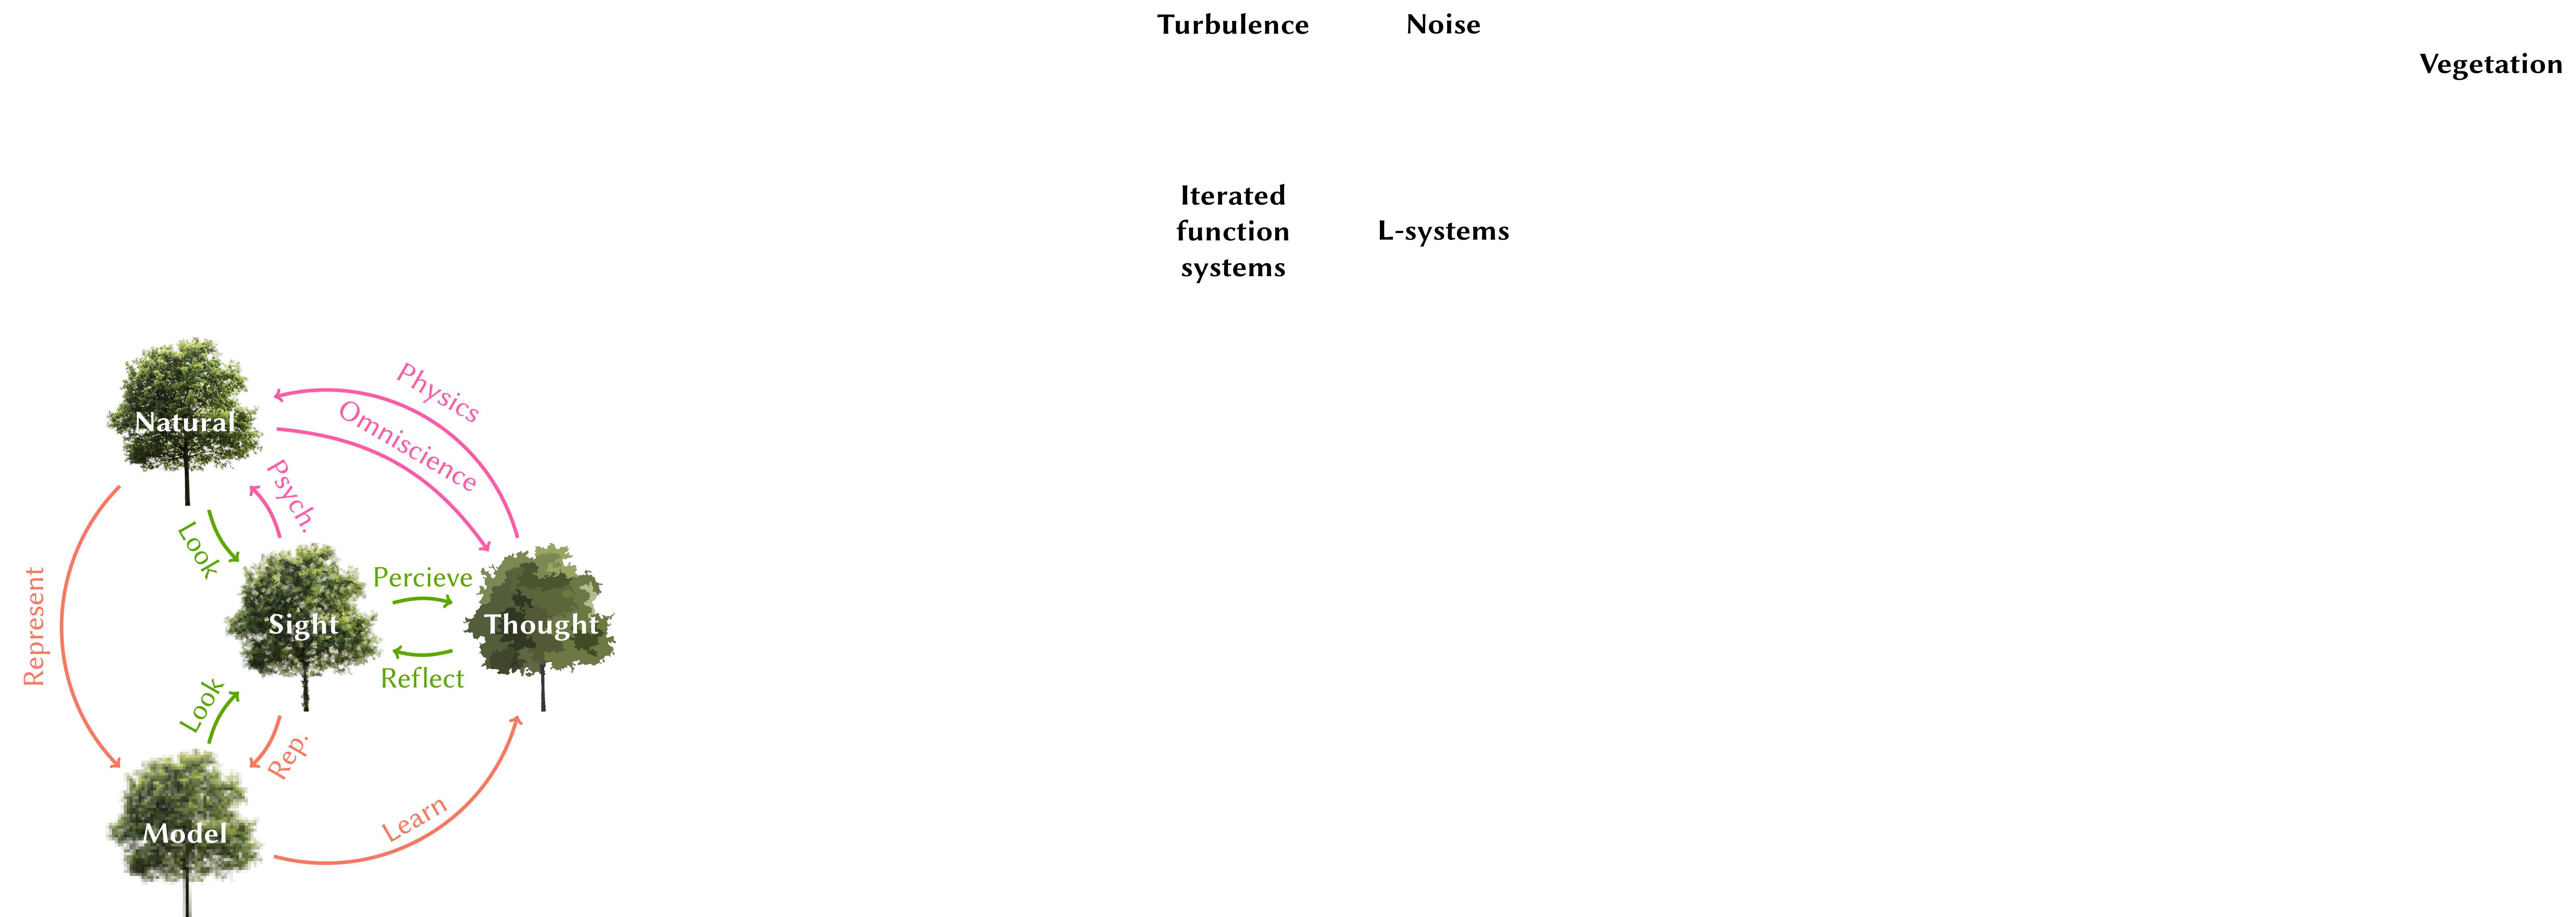
## Understanding visual complexity with statistical physics



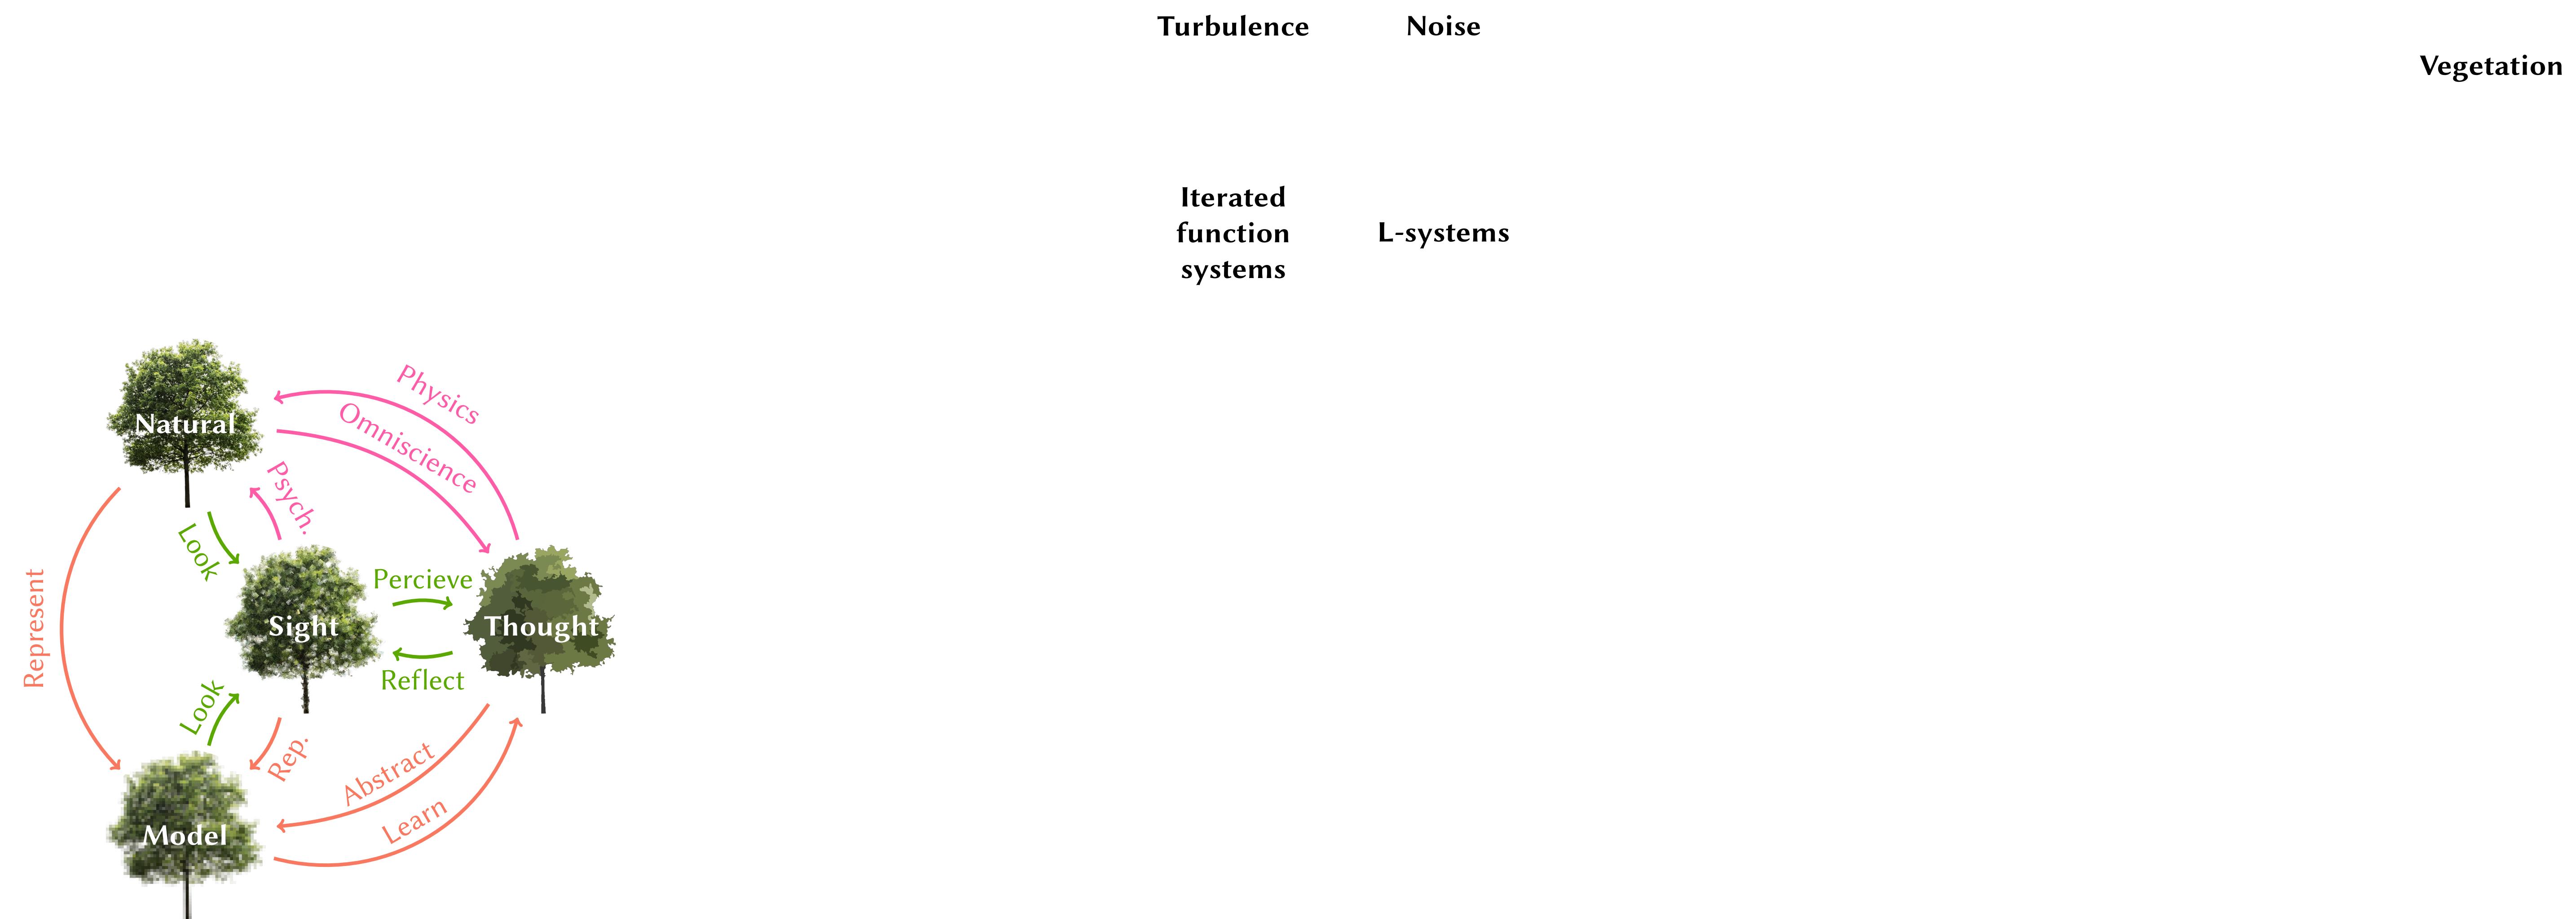
## Understanding visual complexity with statistical physics



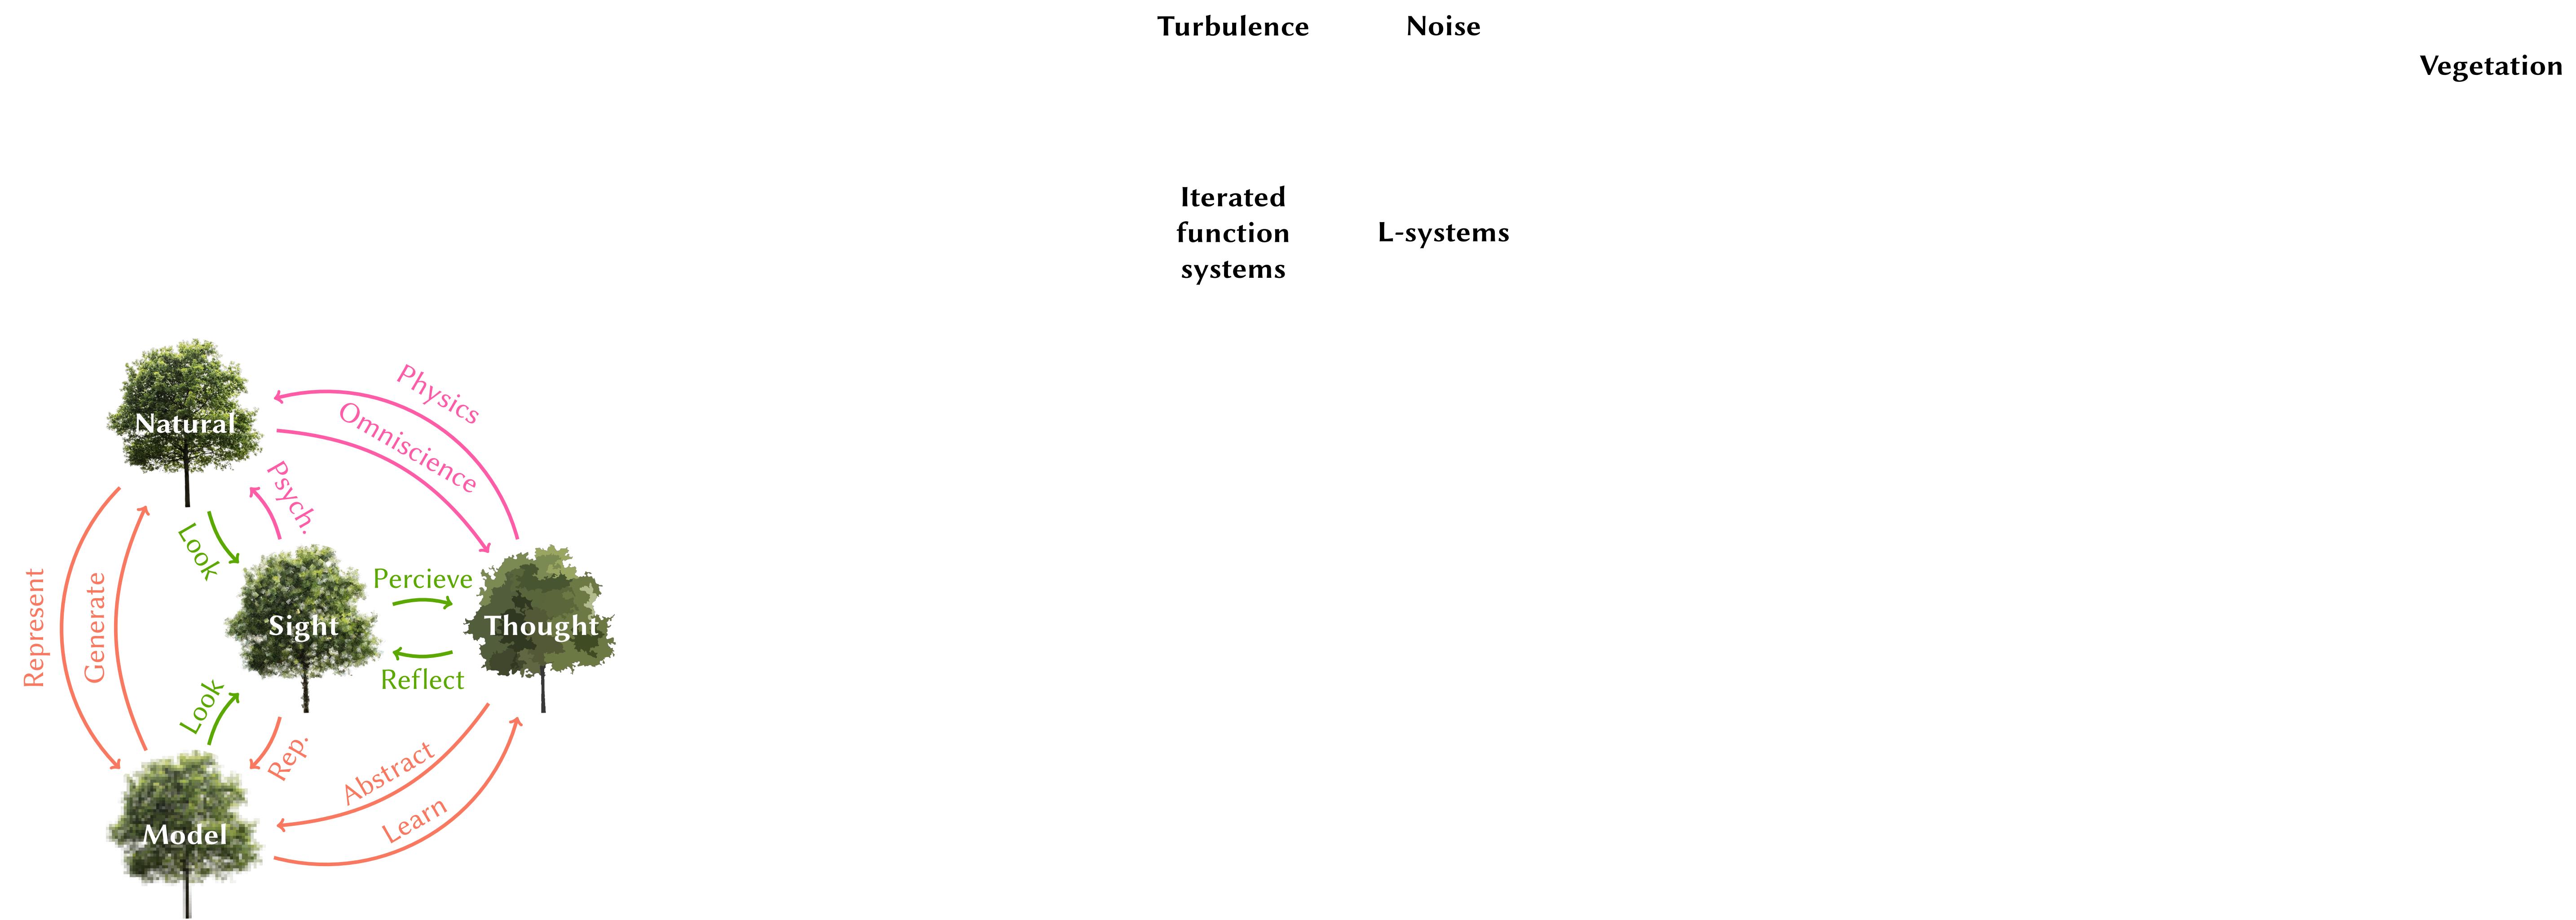
## Understanding visual complexity with statistical physics

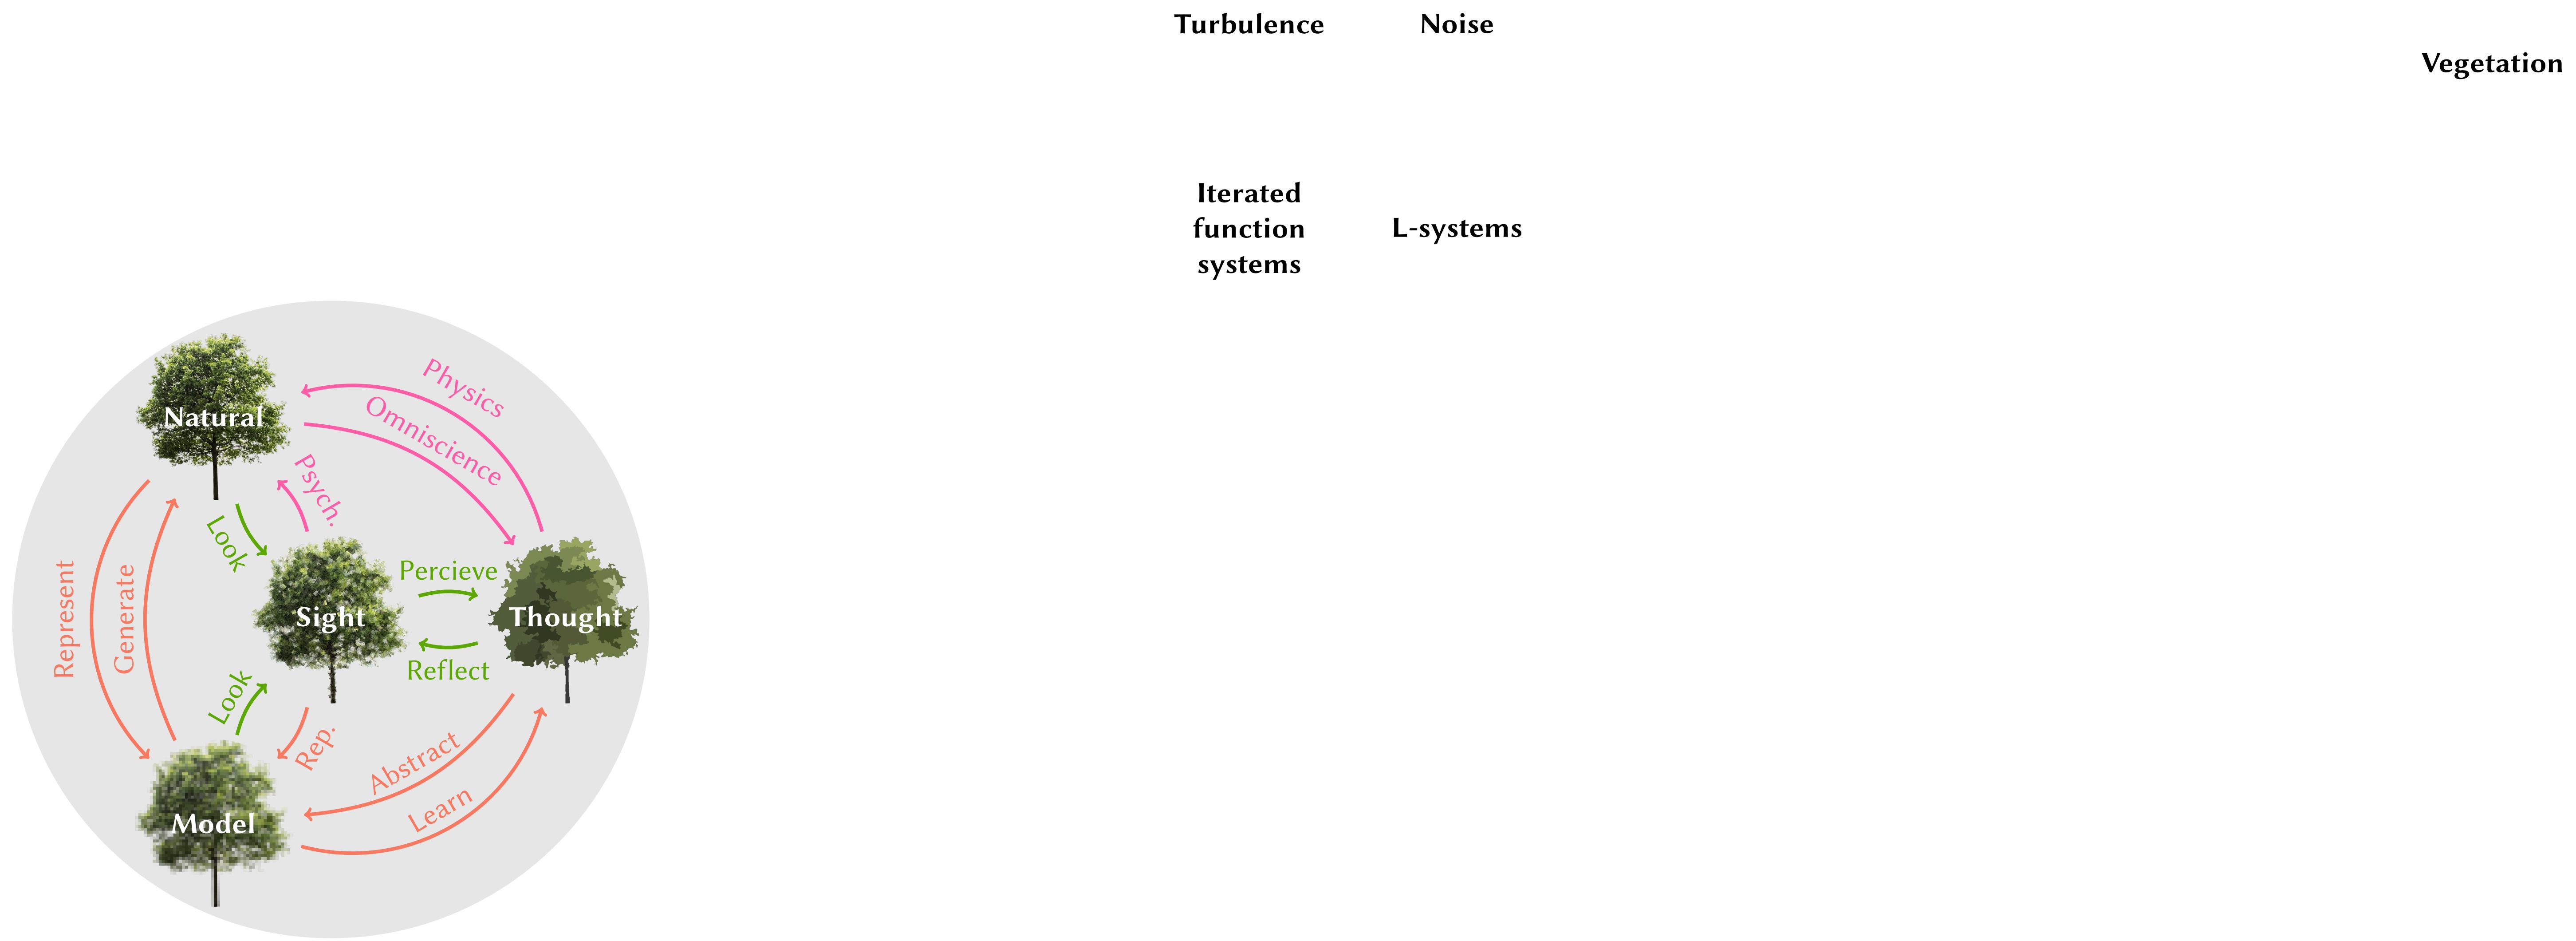


## Understanding visual complexity with statistical physics



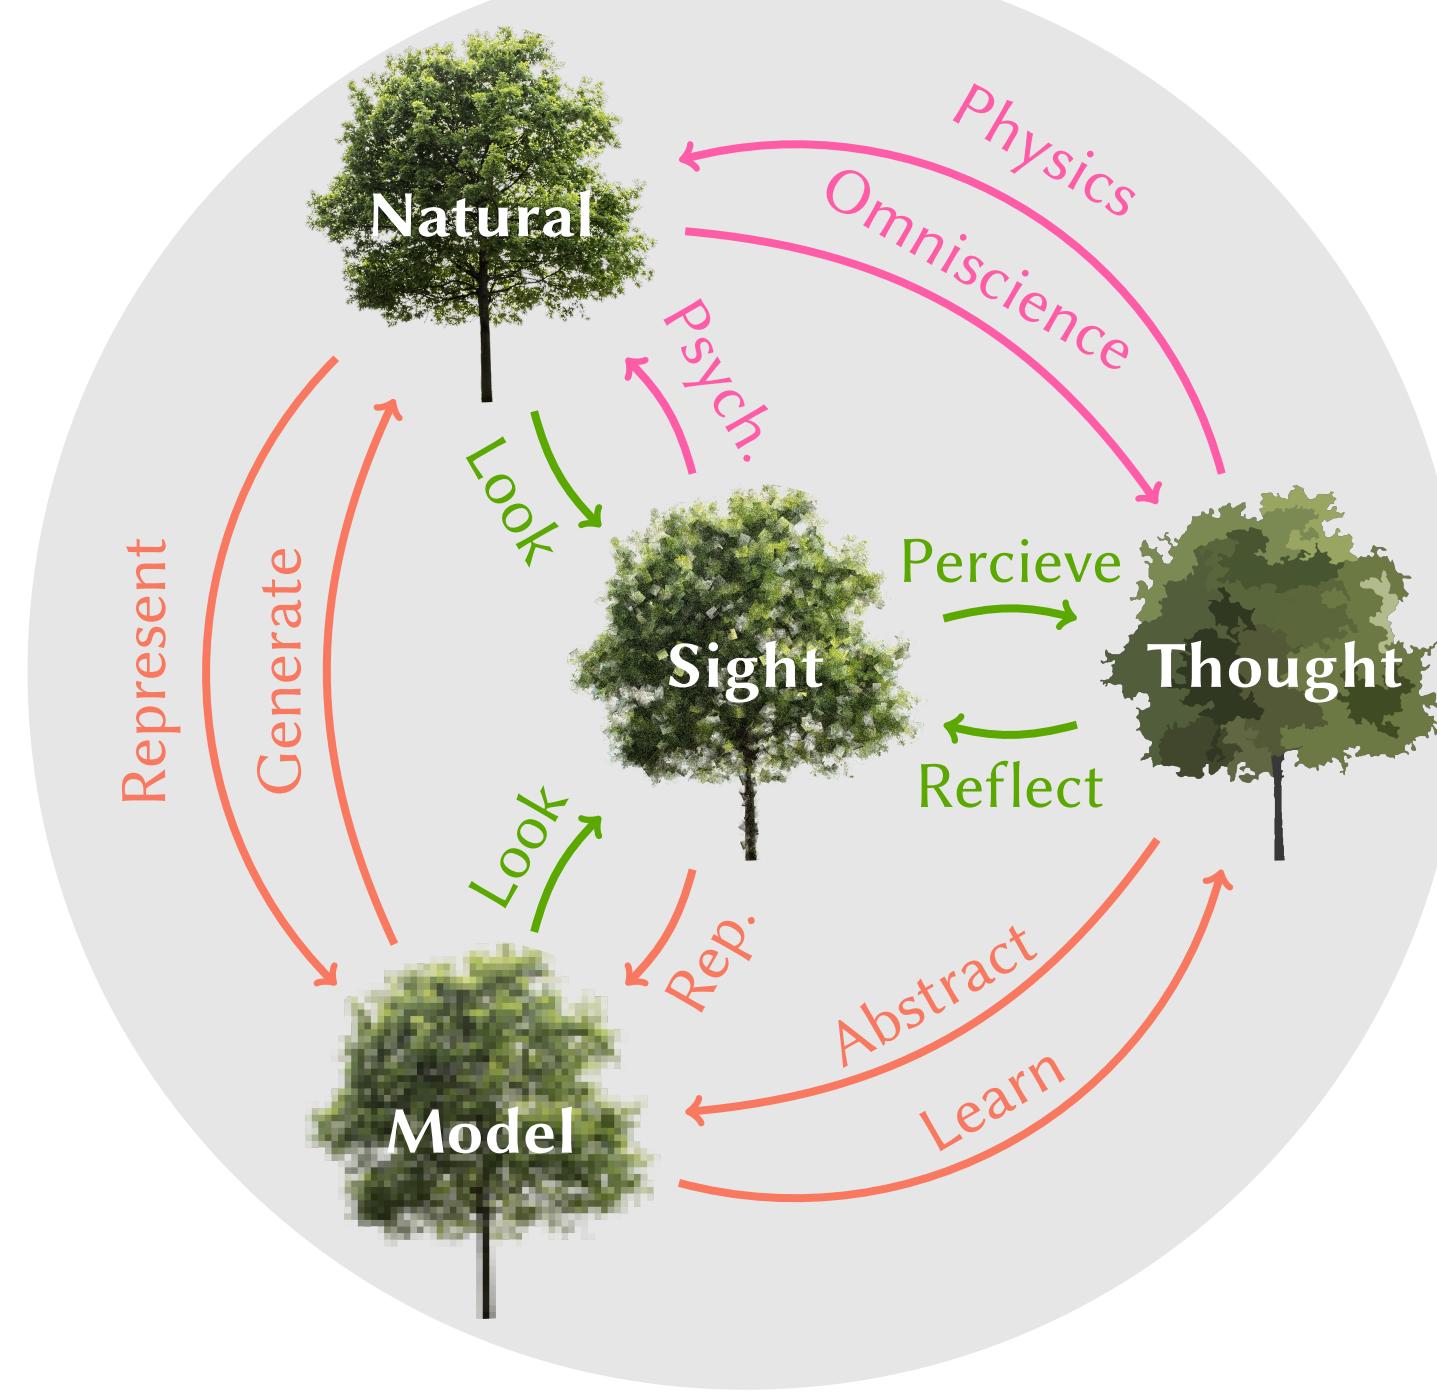
## Understanding visual complexity with statistical physics



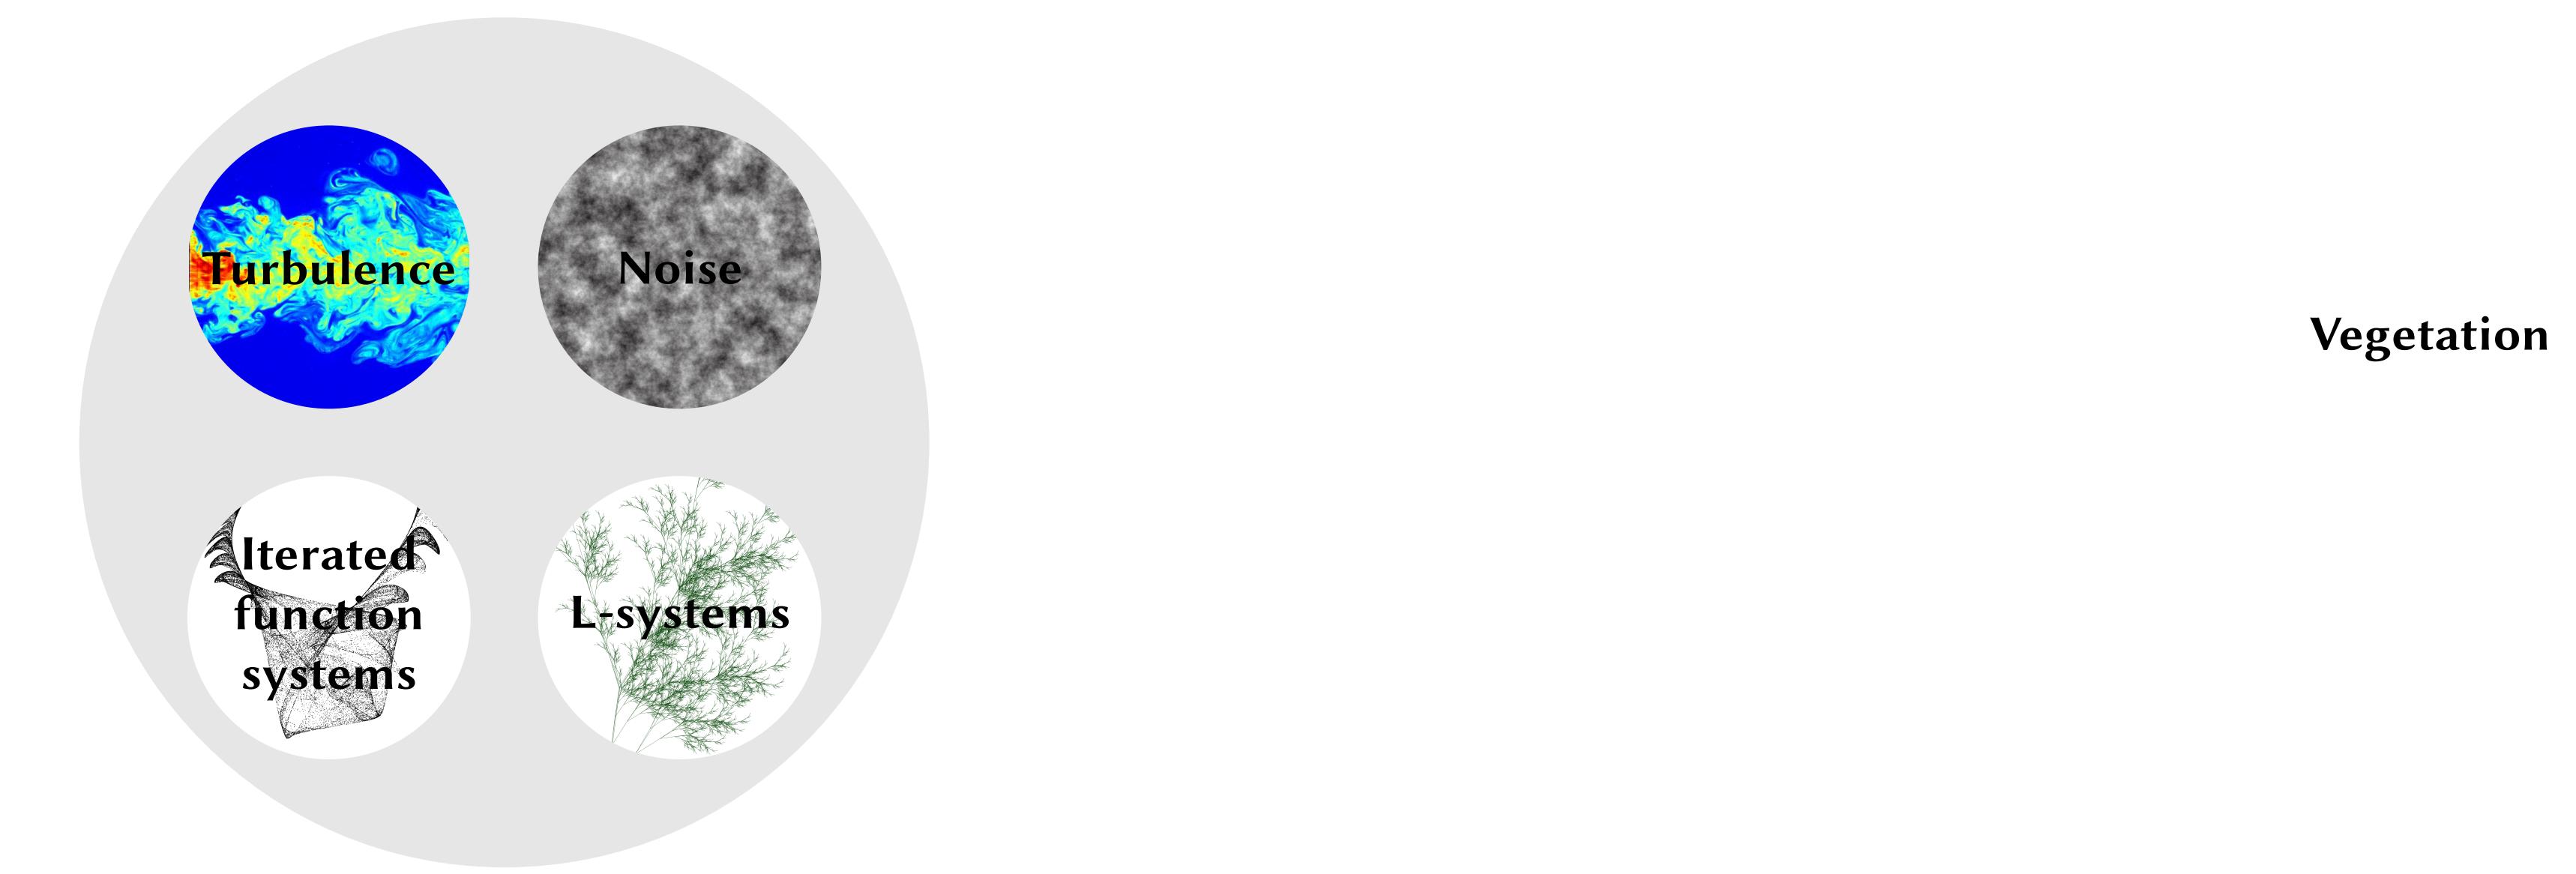


We cannot **directly percieve** the true nature of the world,  
but we can form models to **indirectly percieve** nature.  
A model must represent nature, while still being something we can readily look at and use to guide thought.  
**What are good models of how nature looks to us?**

# Understanding visual complexity with statistical physics

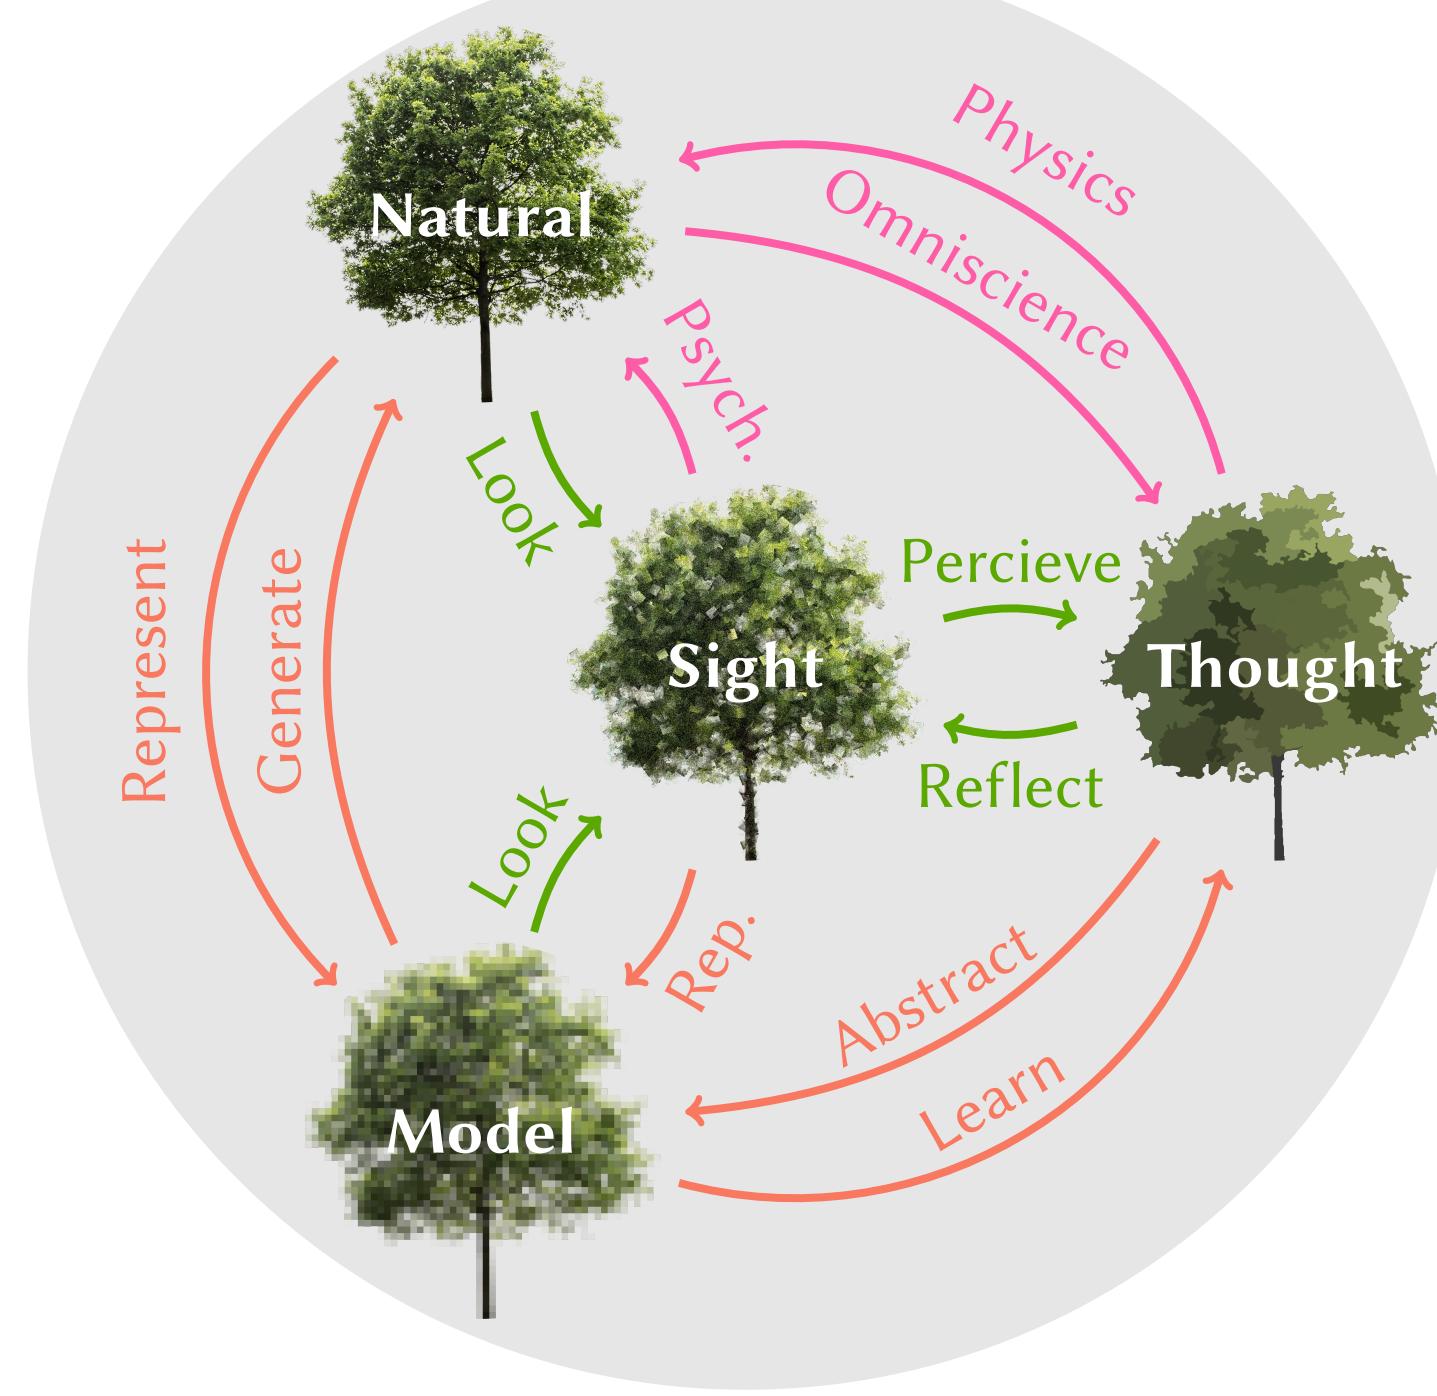


We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

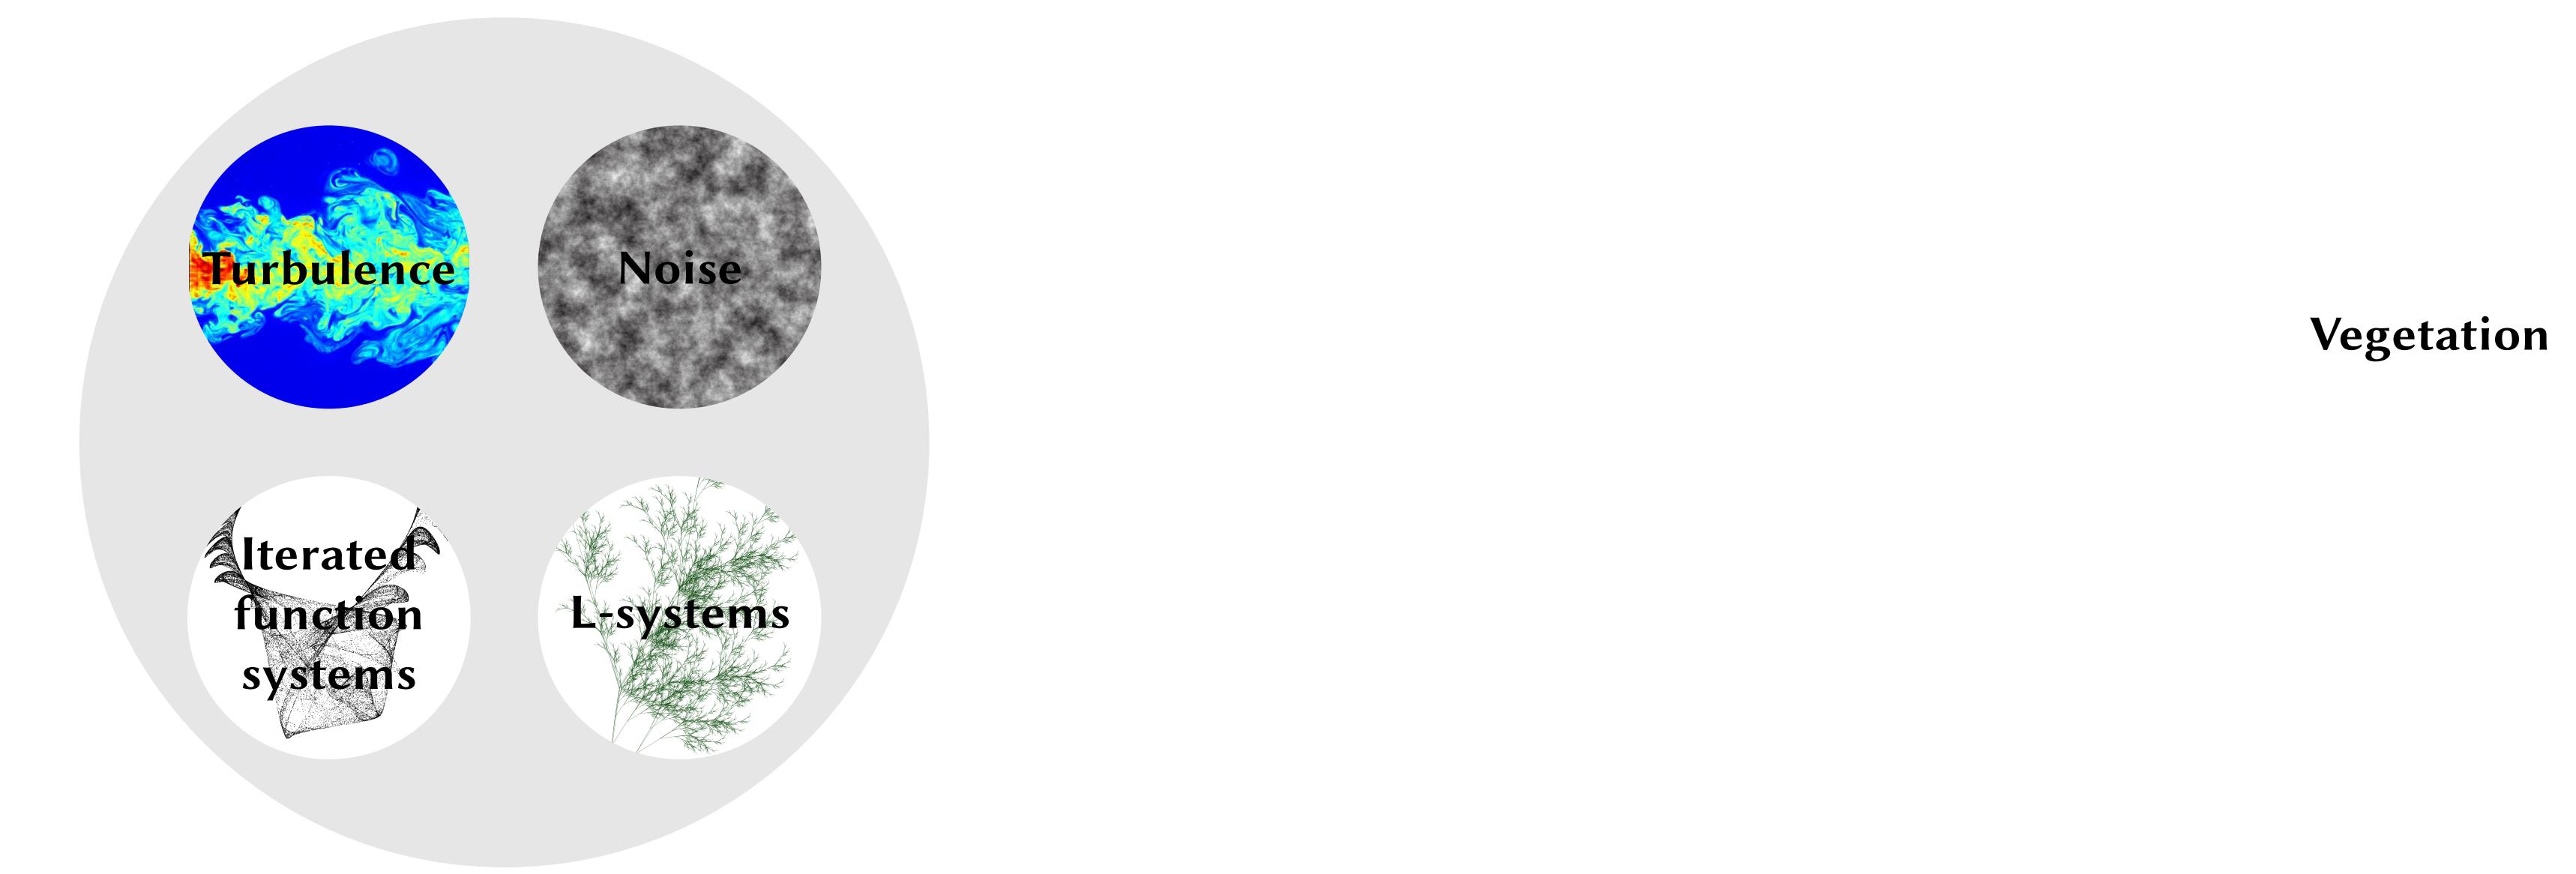


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

# Understanding visual complexity with statistical physics

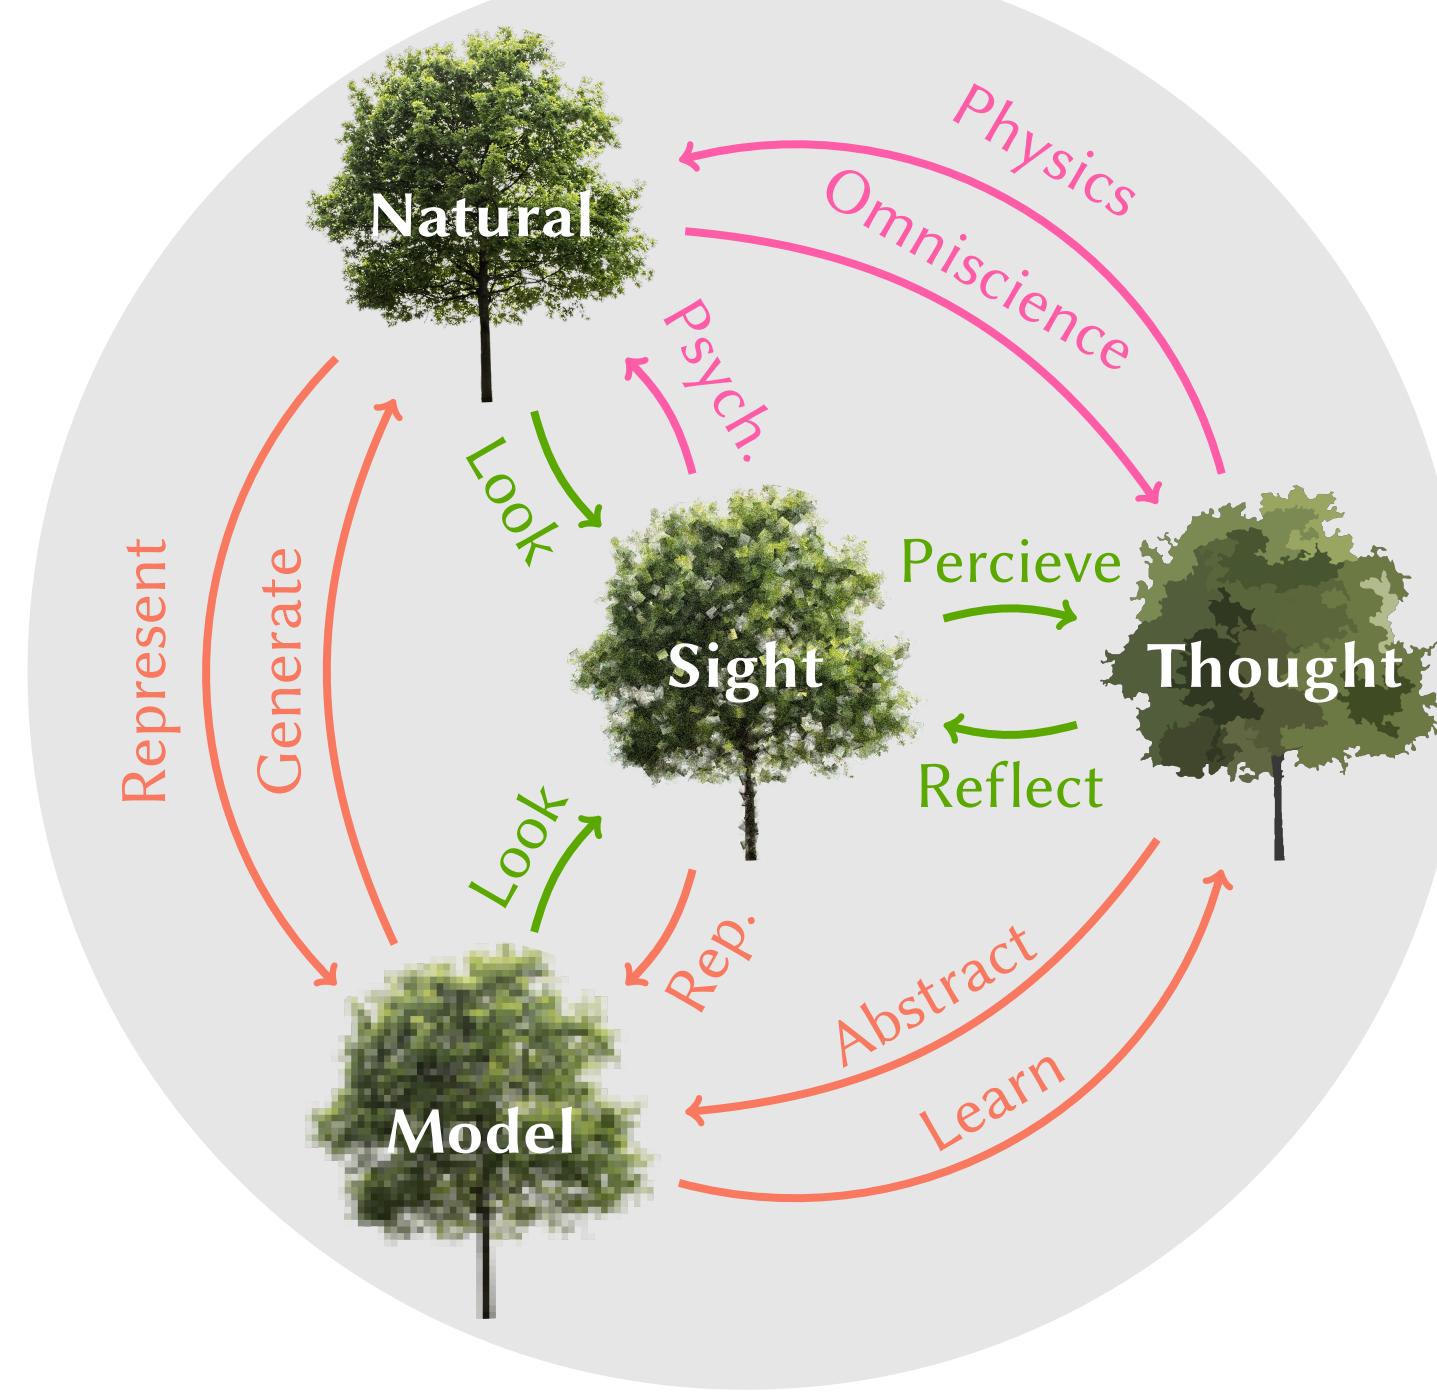


We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

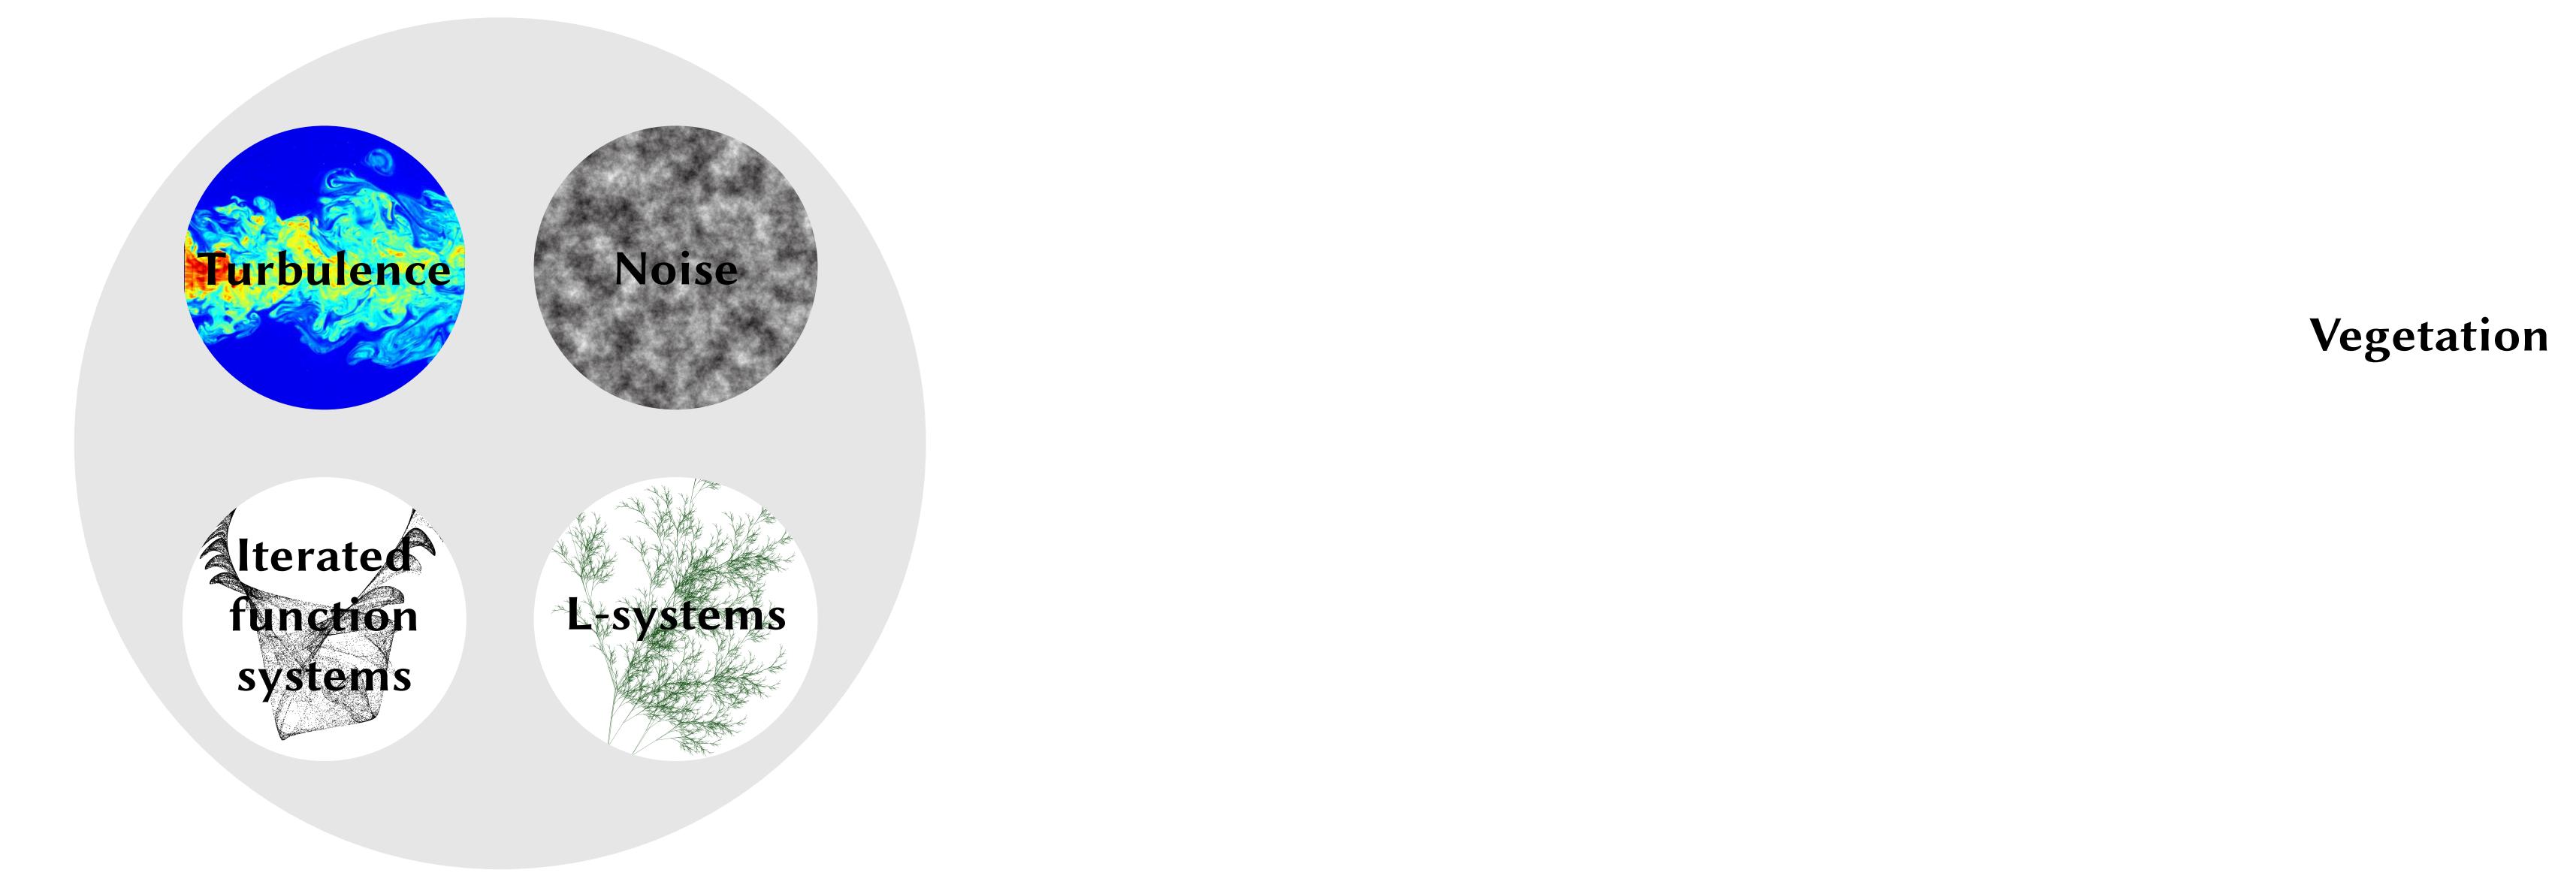


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

# Understanding visual complexity with statistical physics

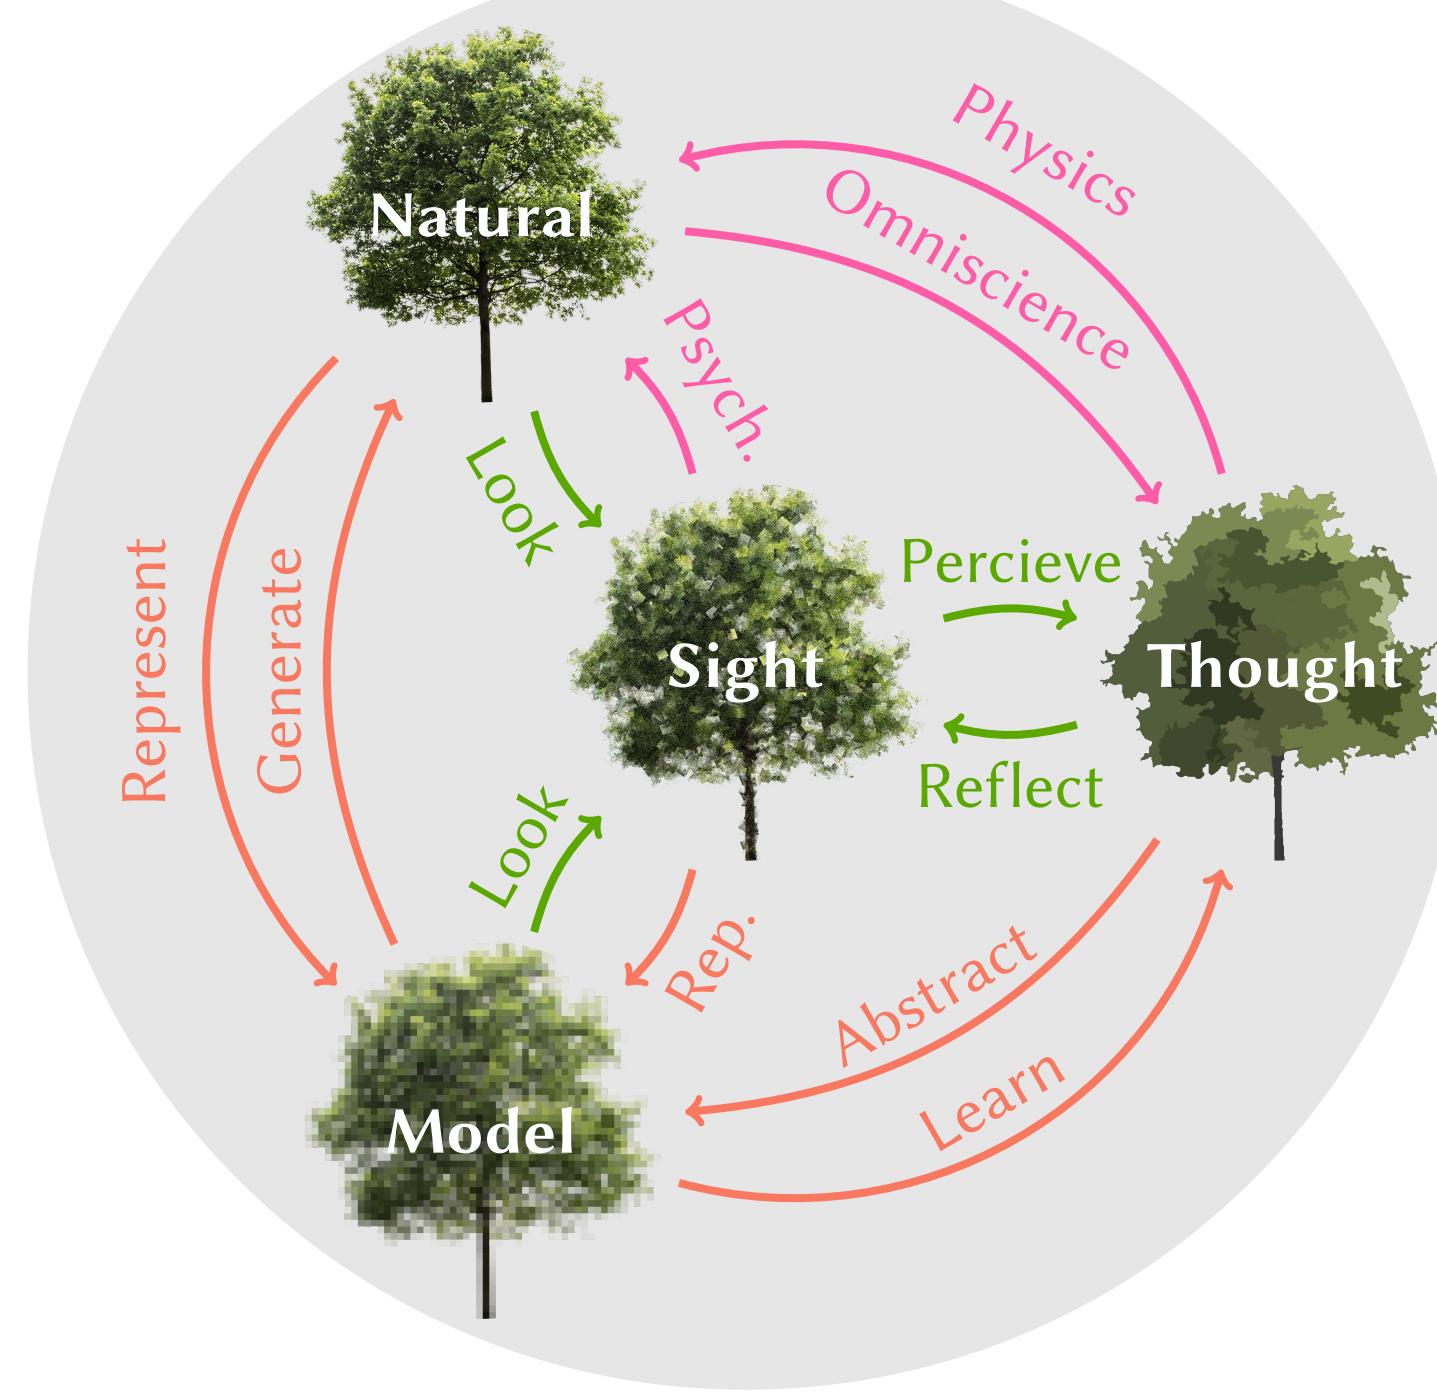


We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

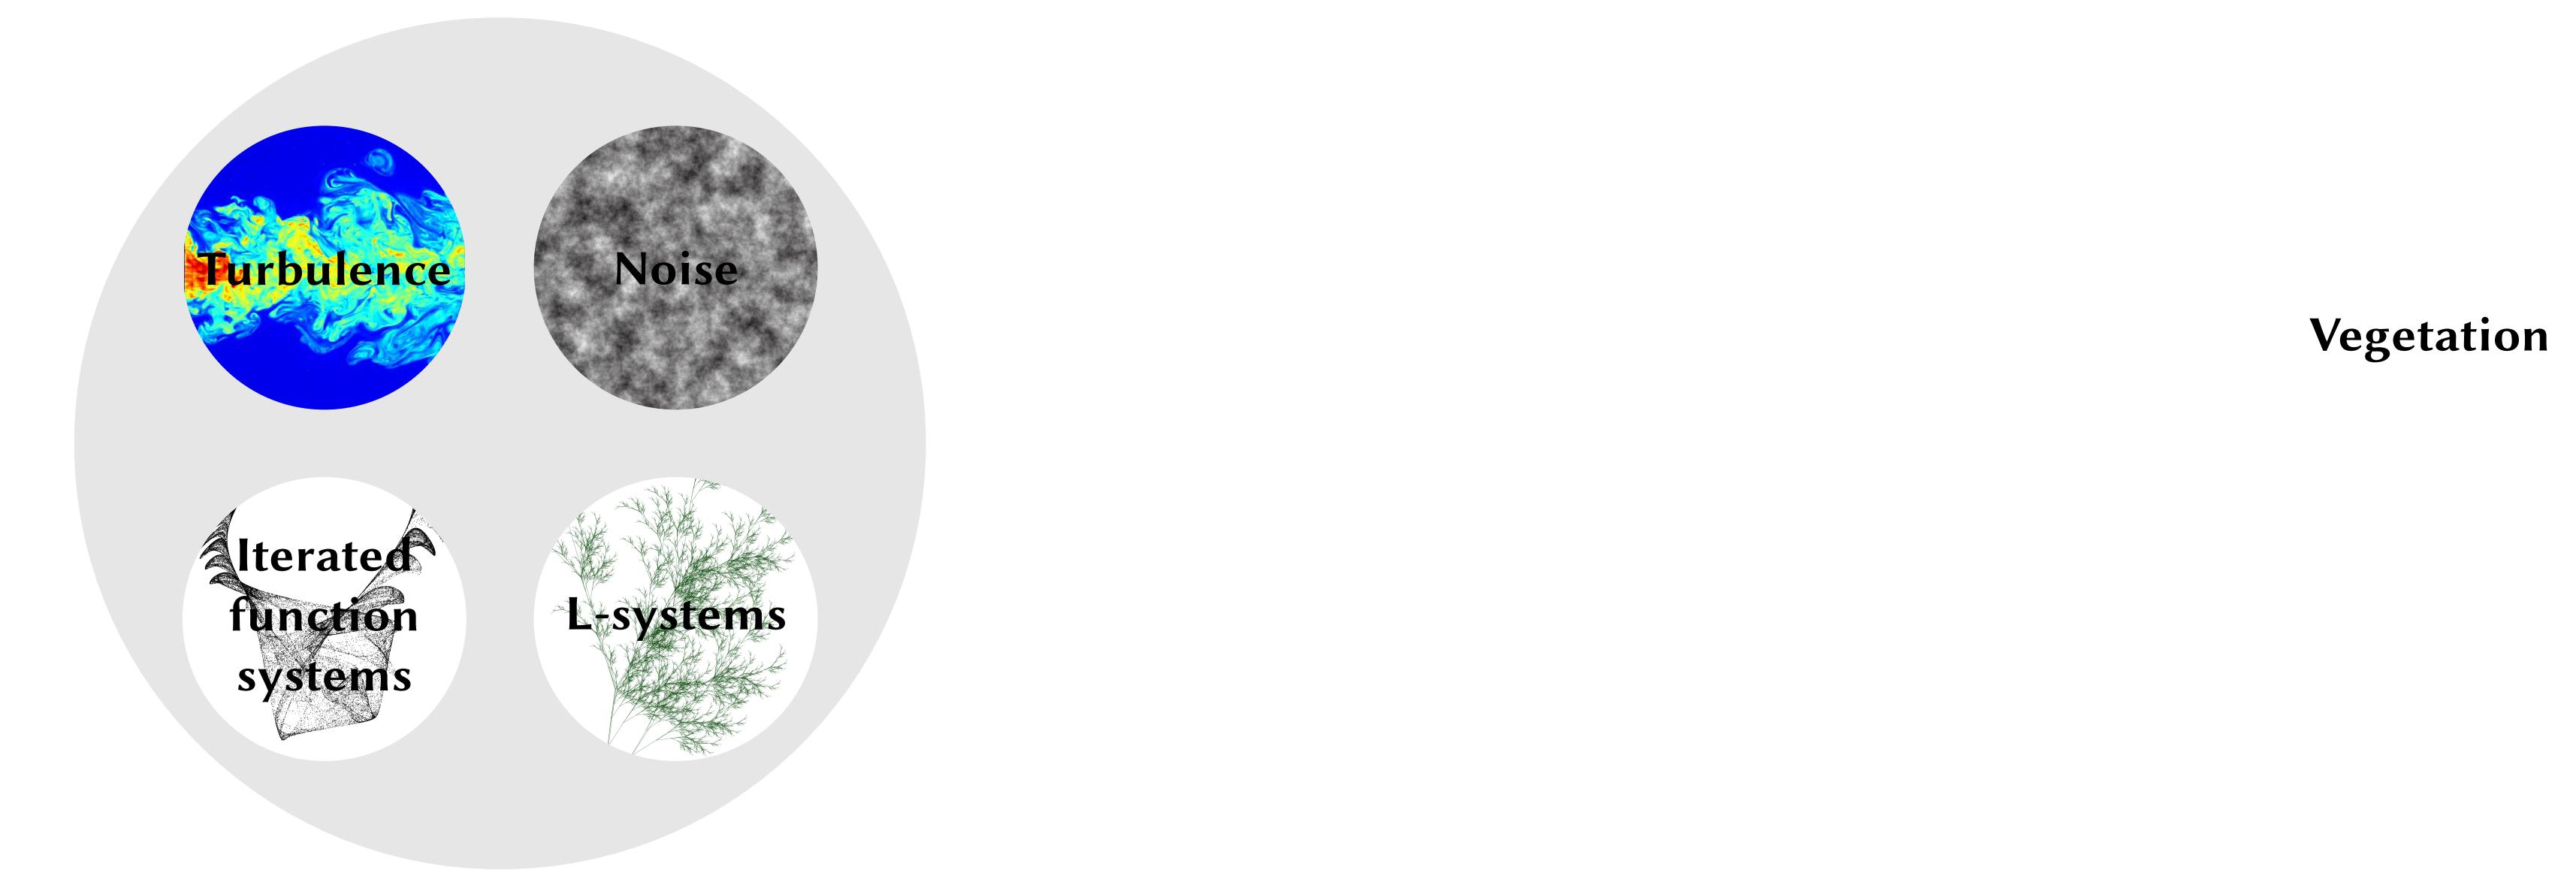


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

# Understanding visual complexity with statistical physics

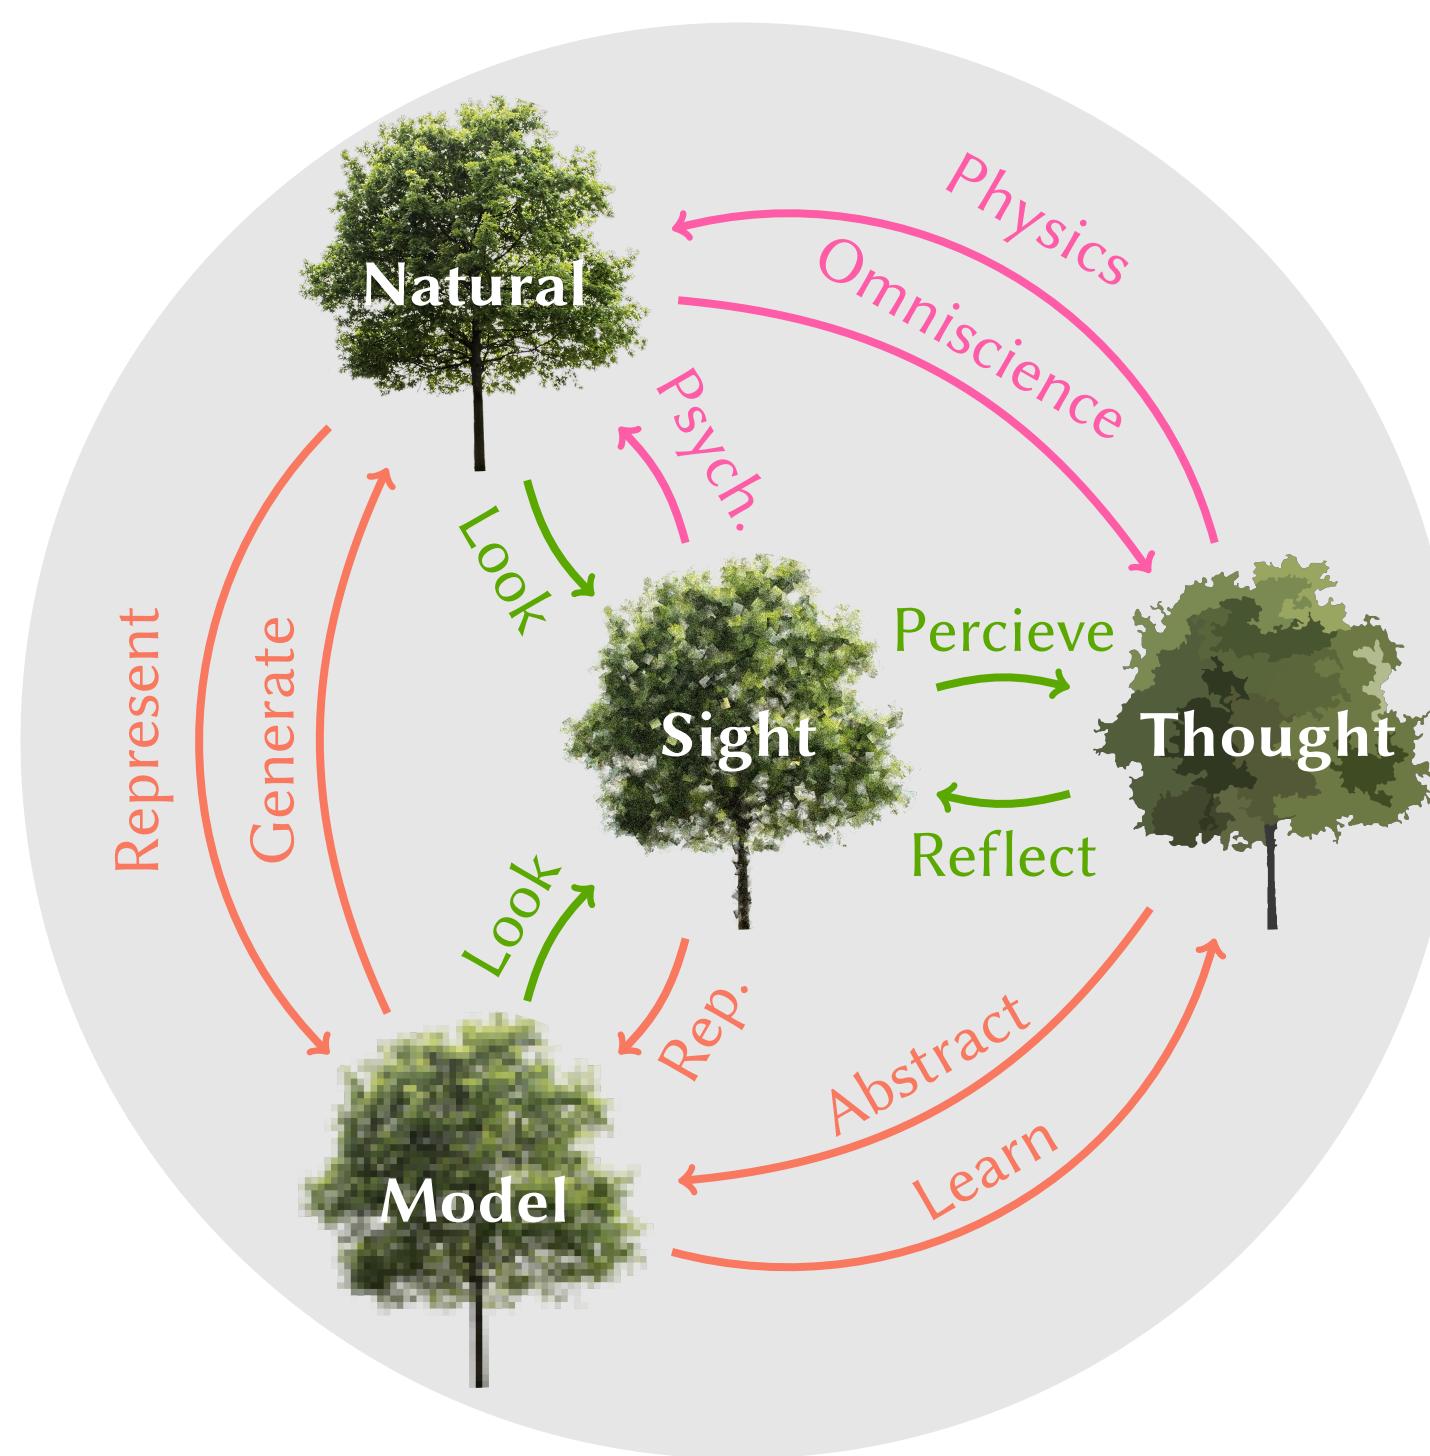


We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

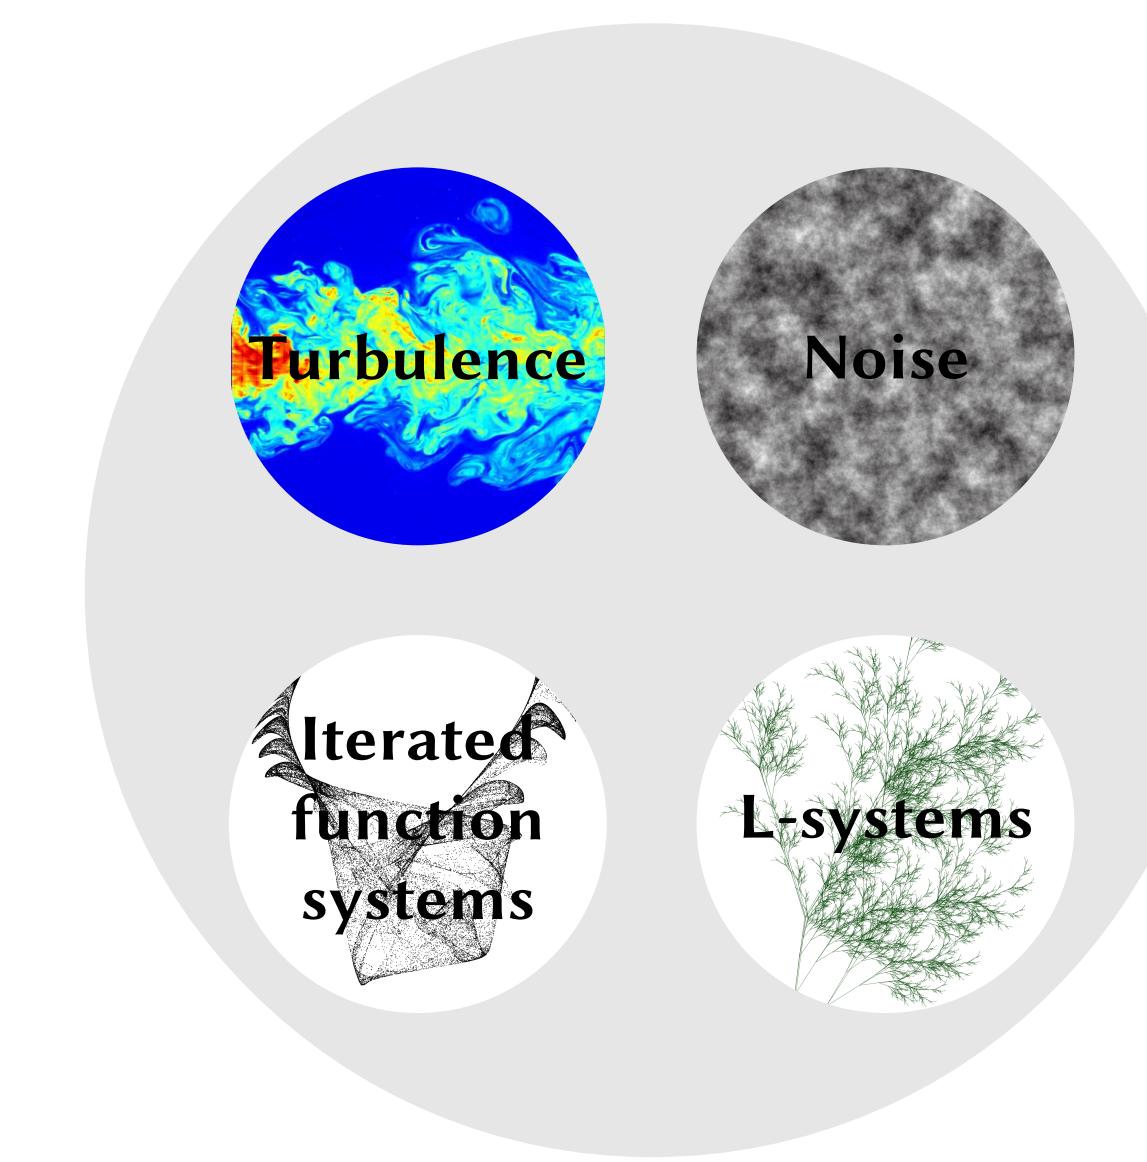


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

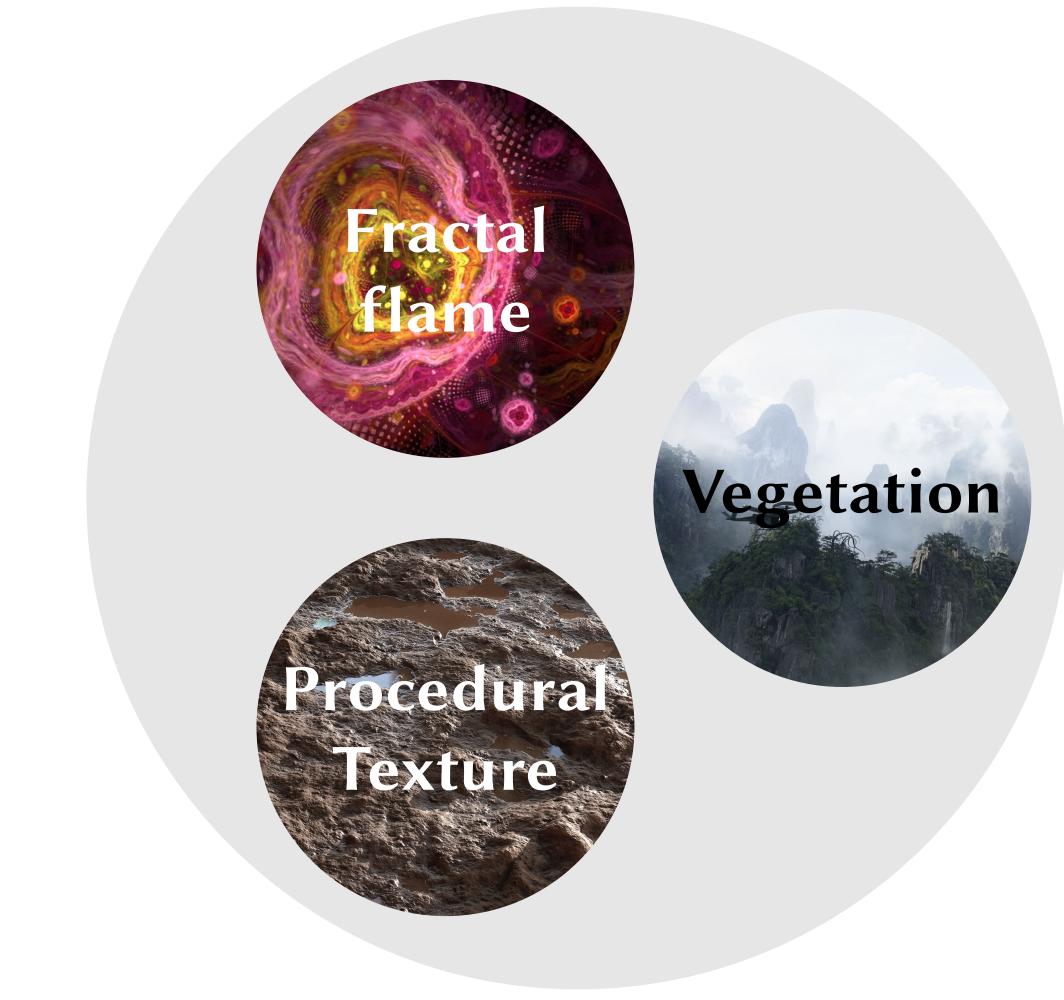
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**



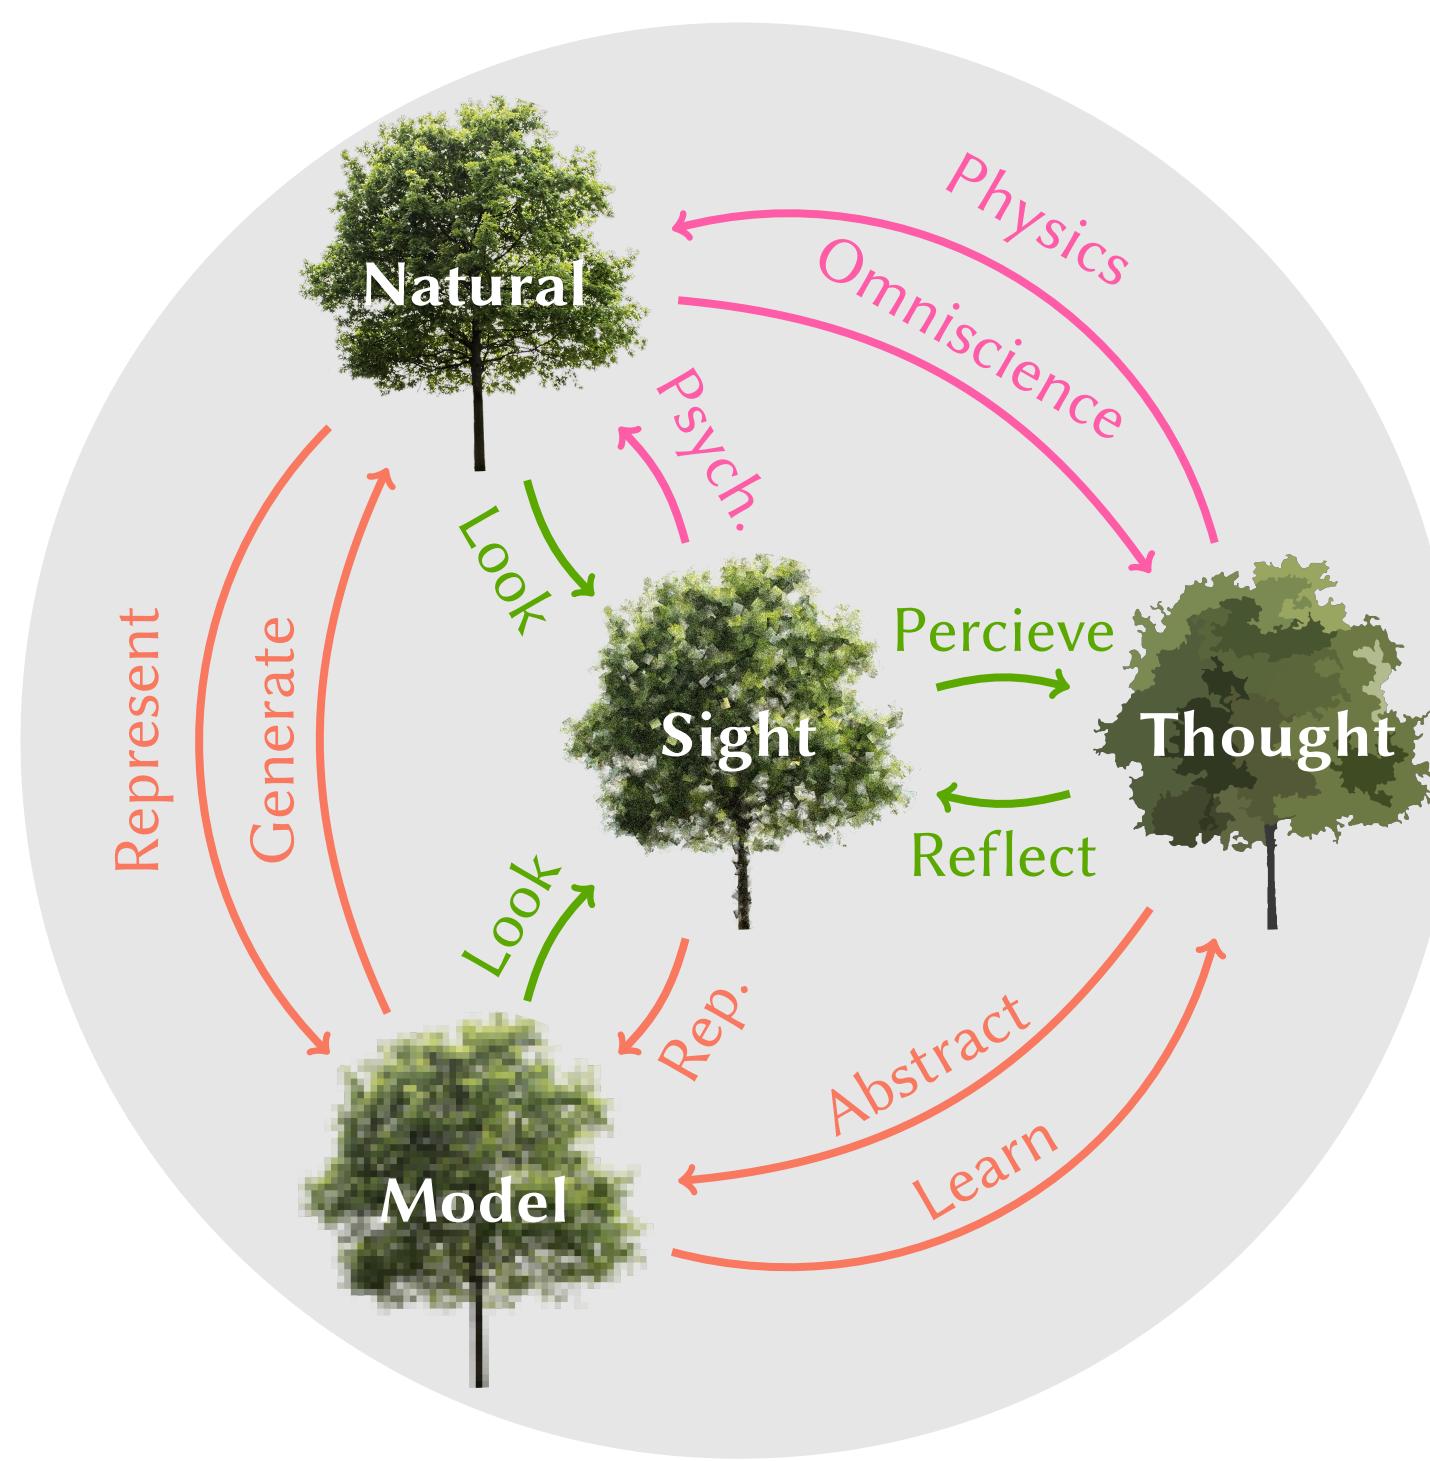
**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



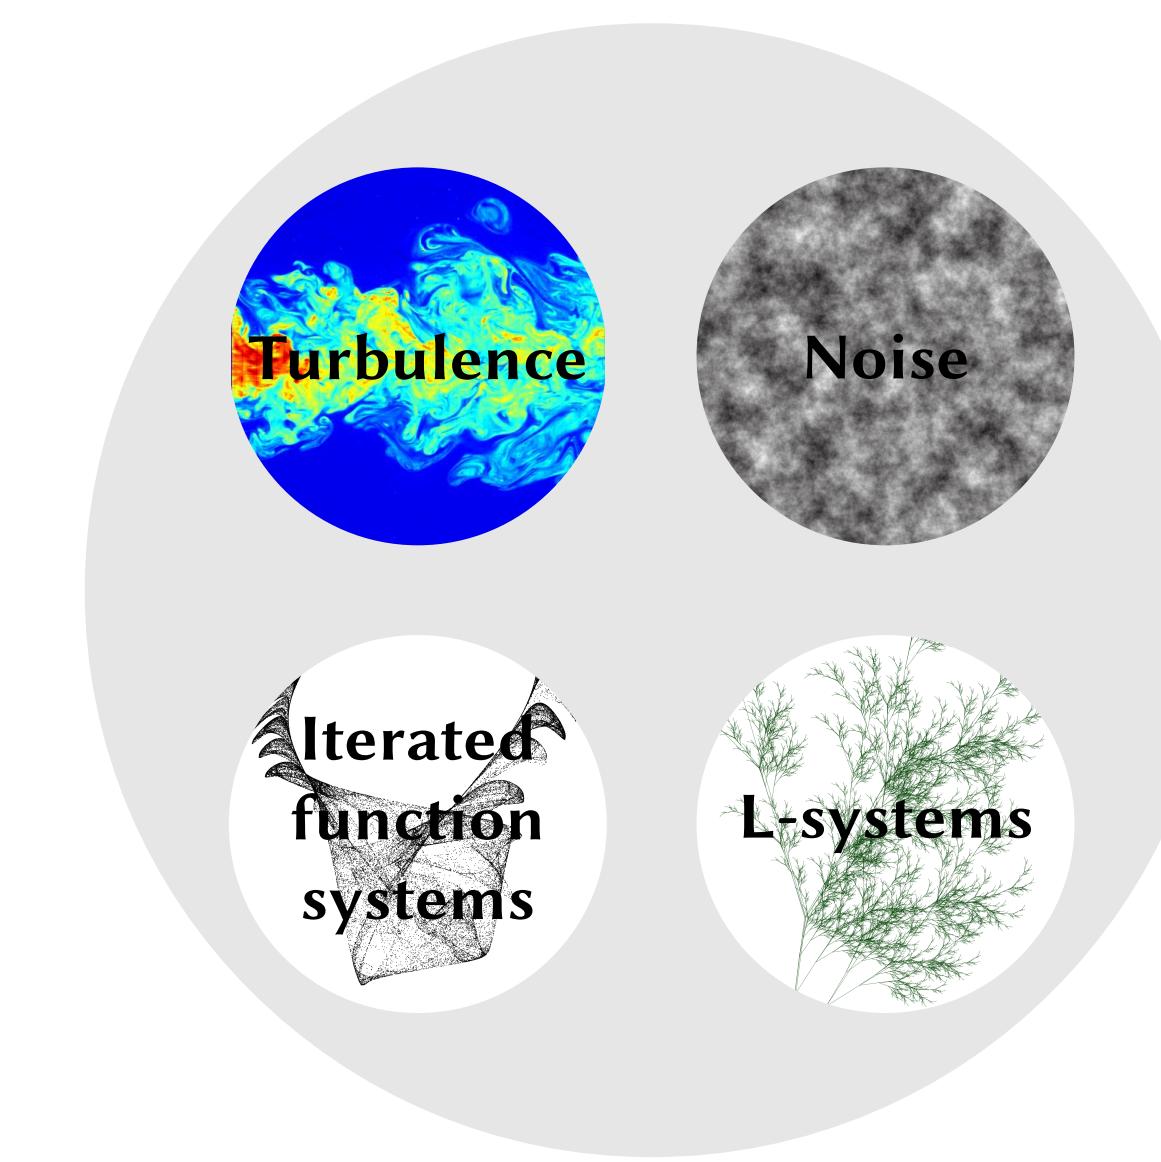
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

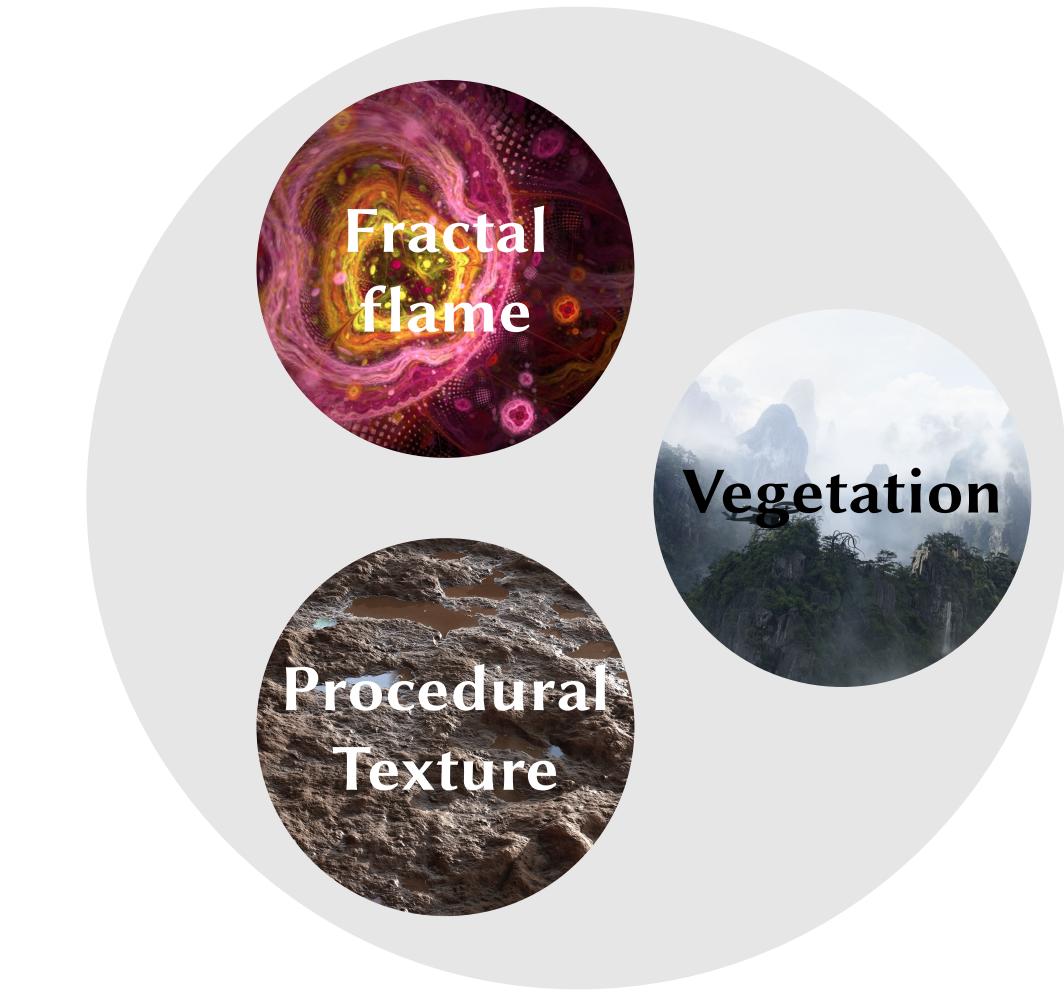
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**



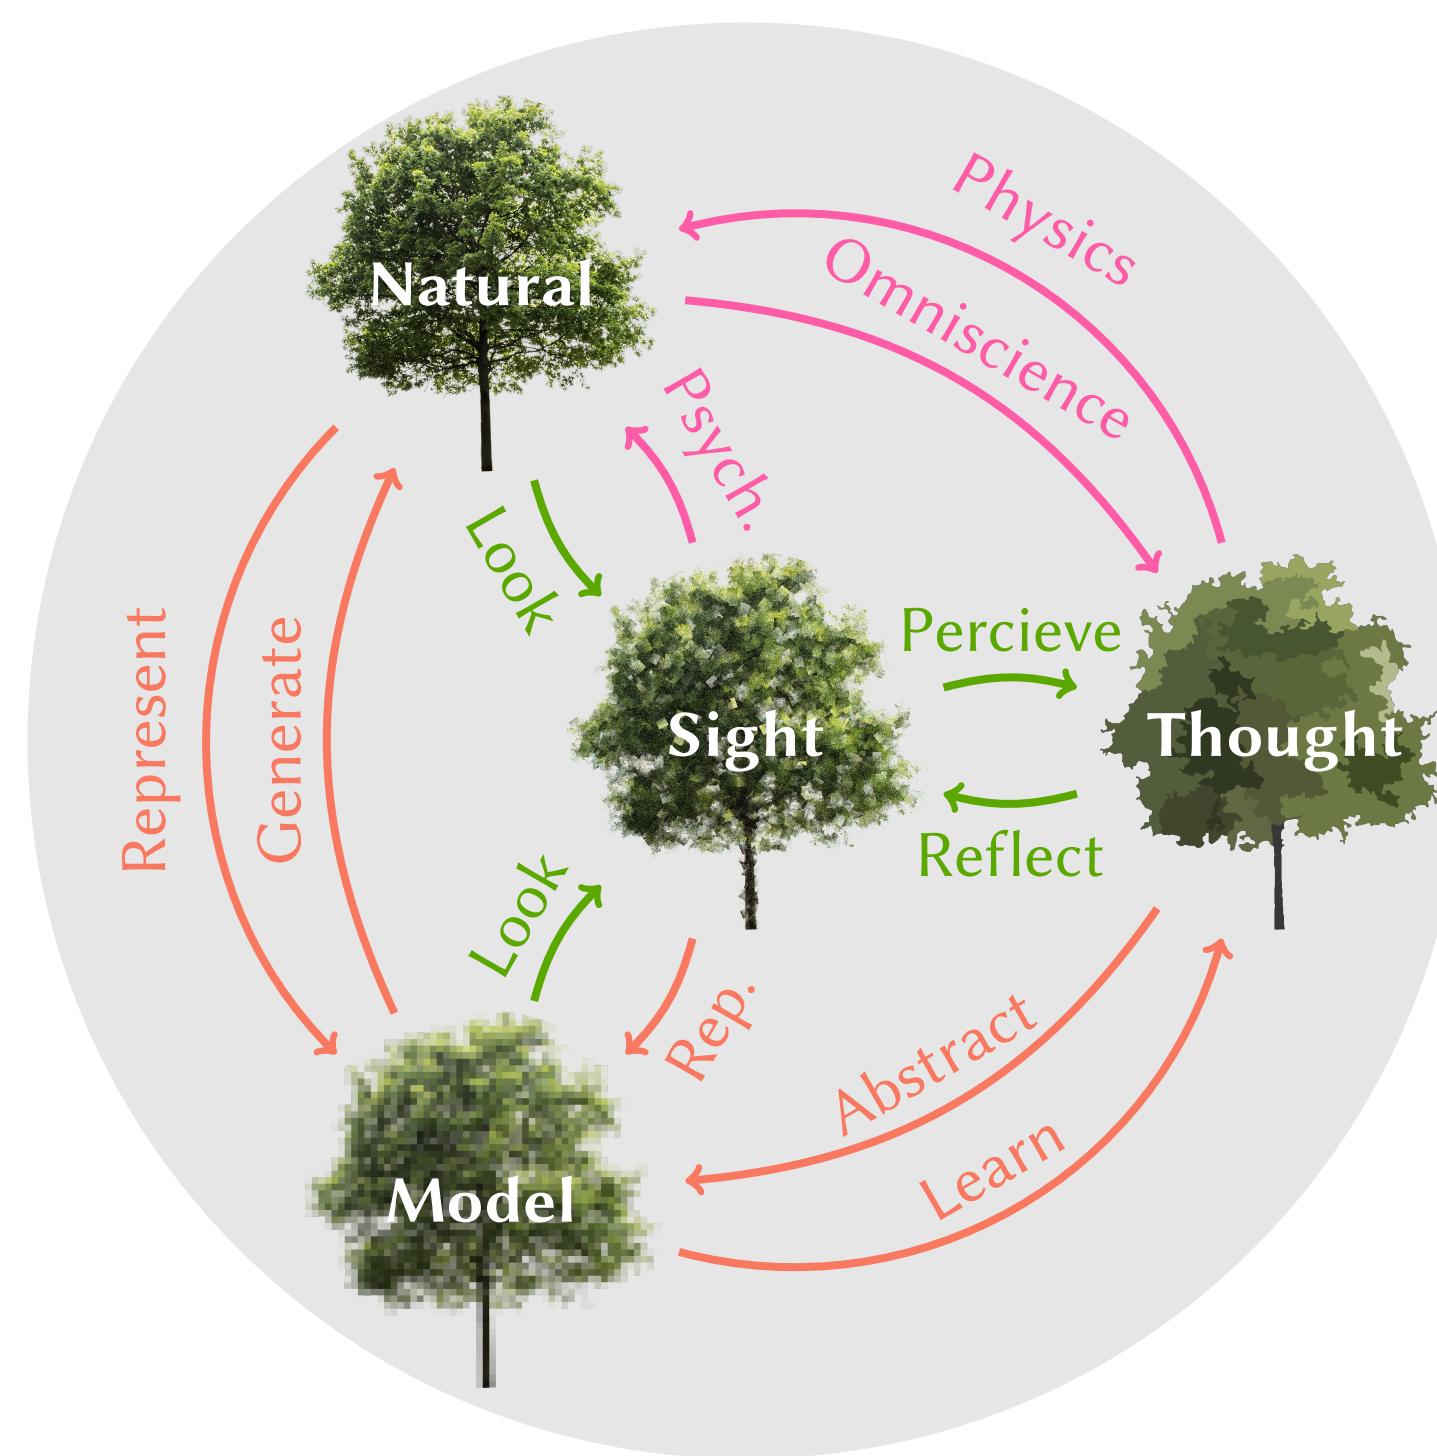
**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



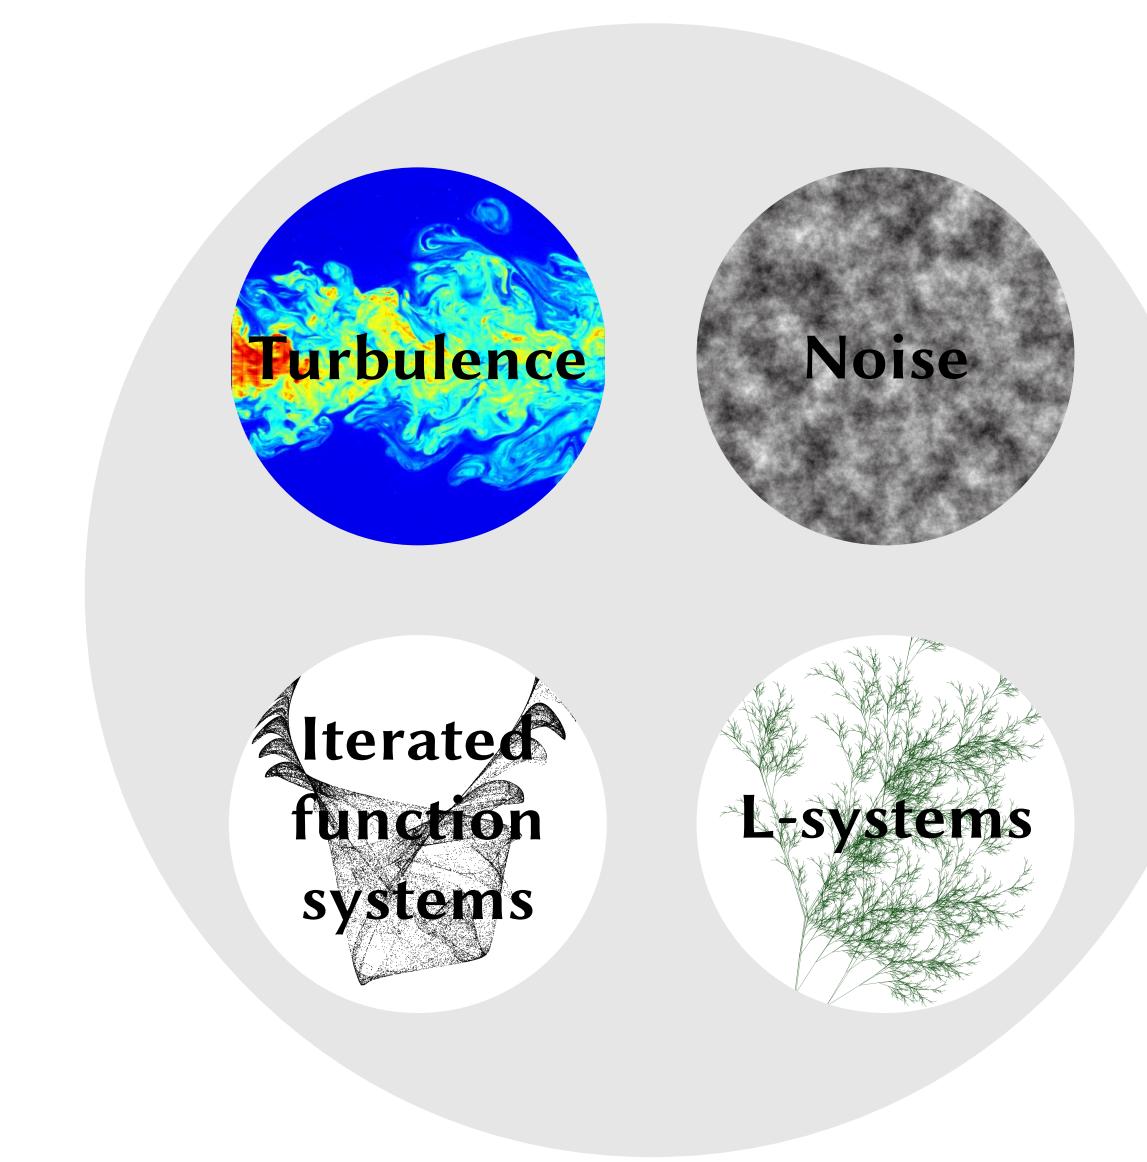
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

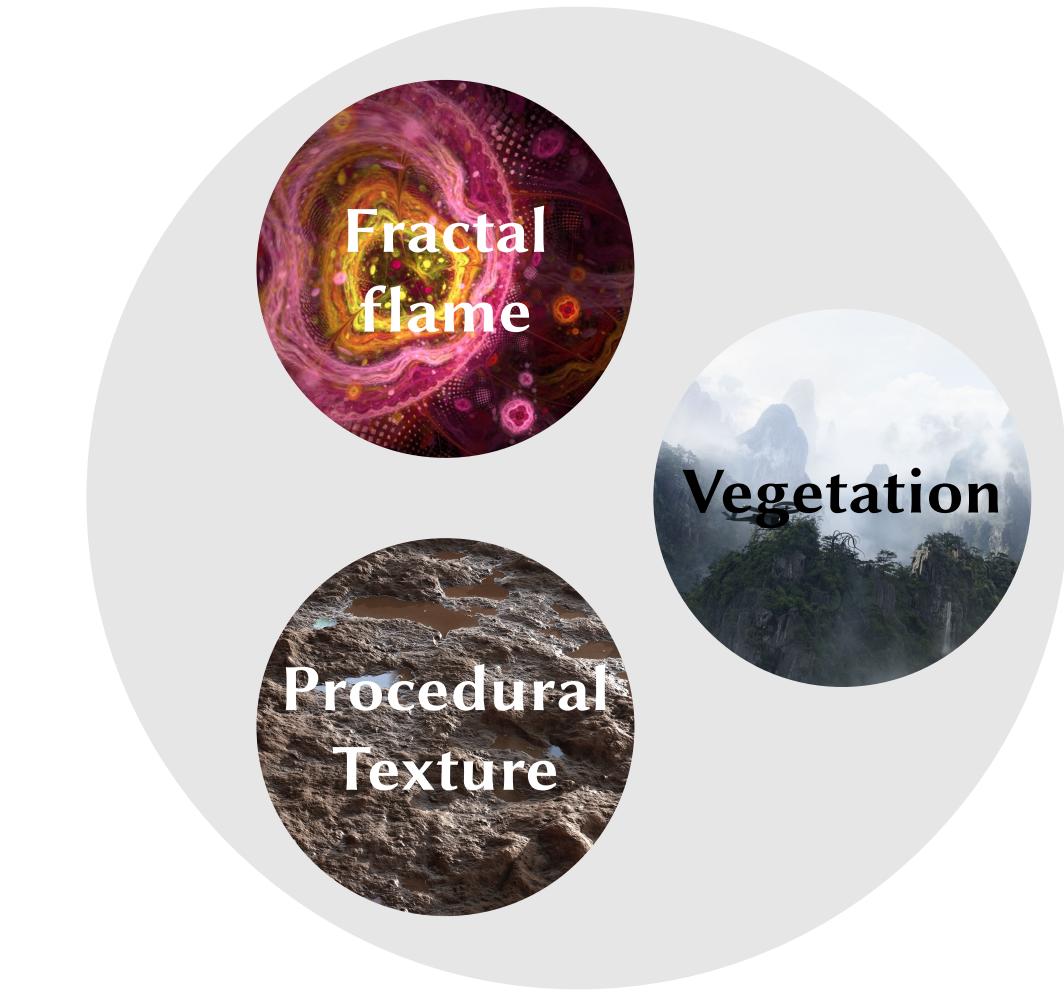
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**



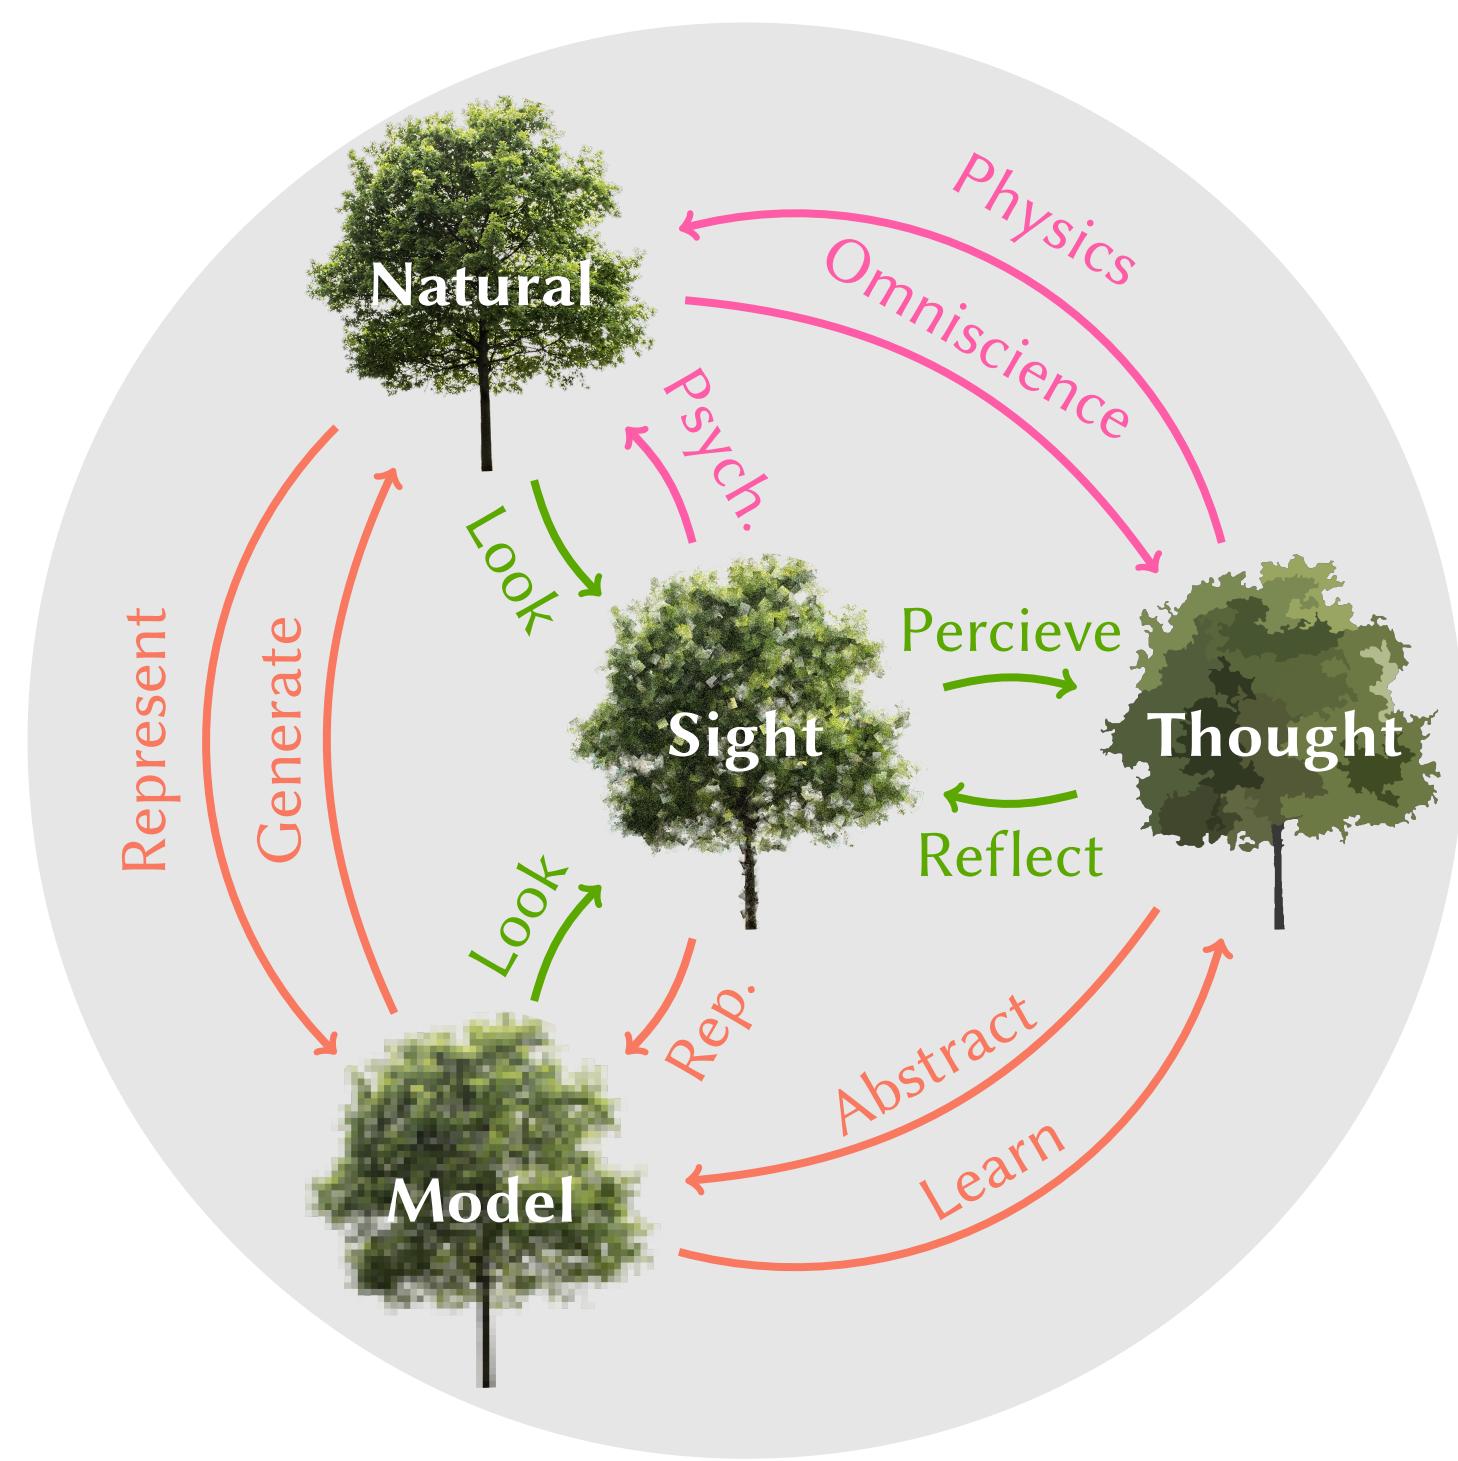
**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



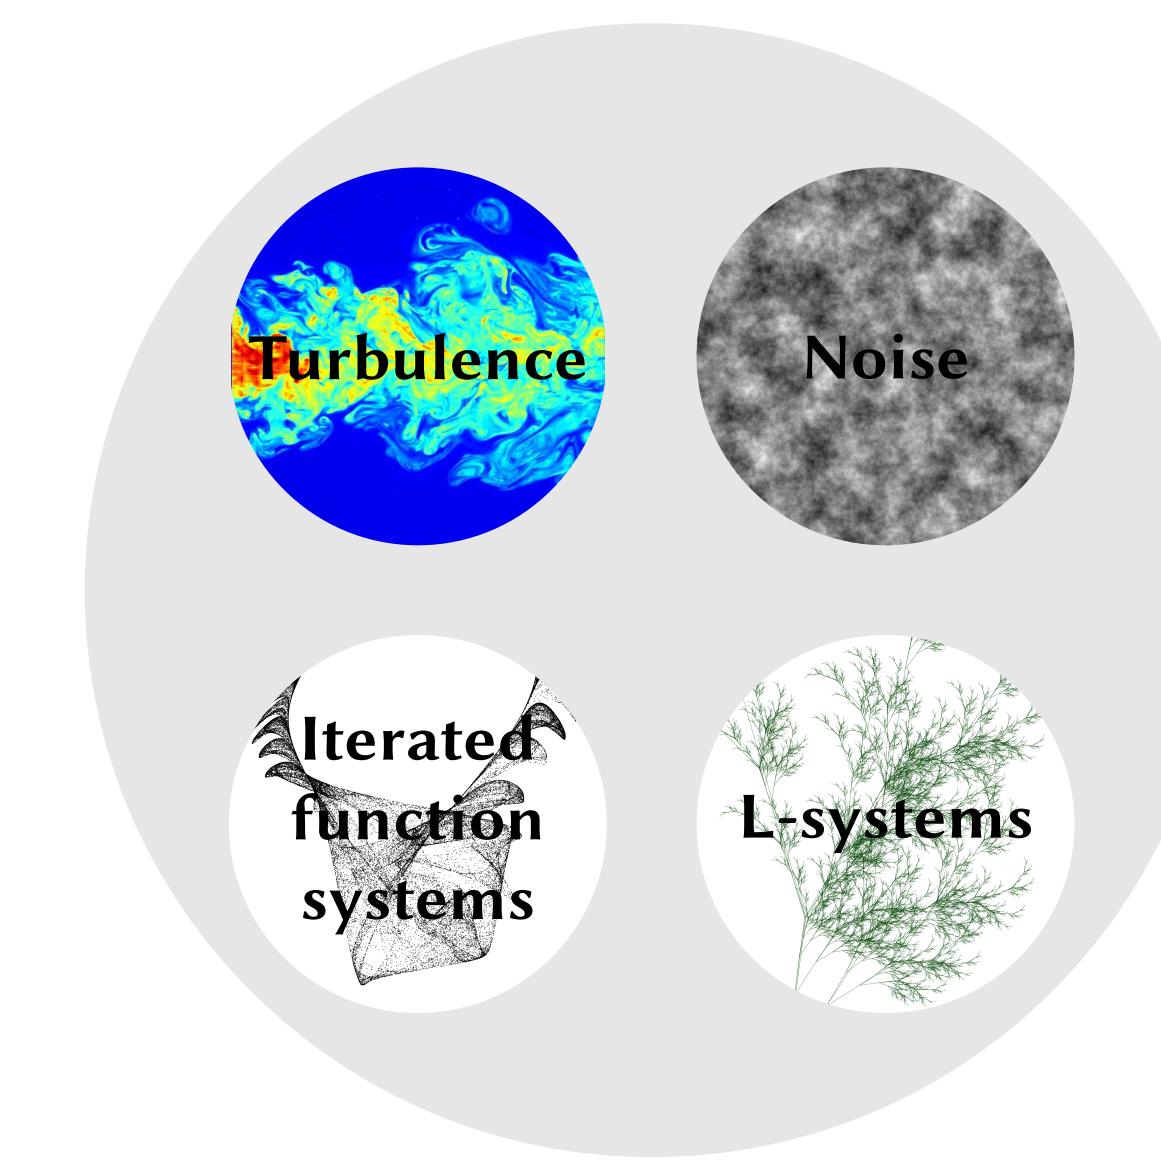
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

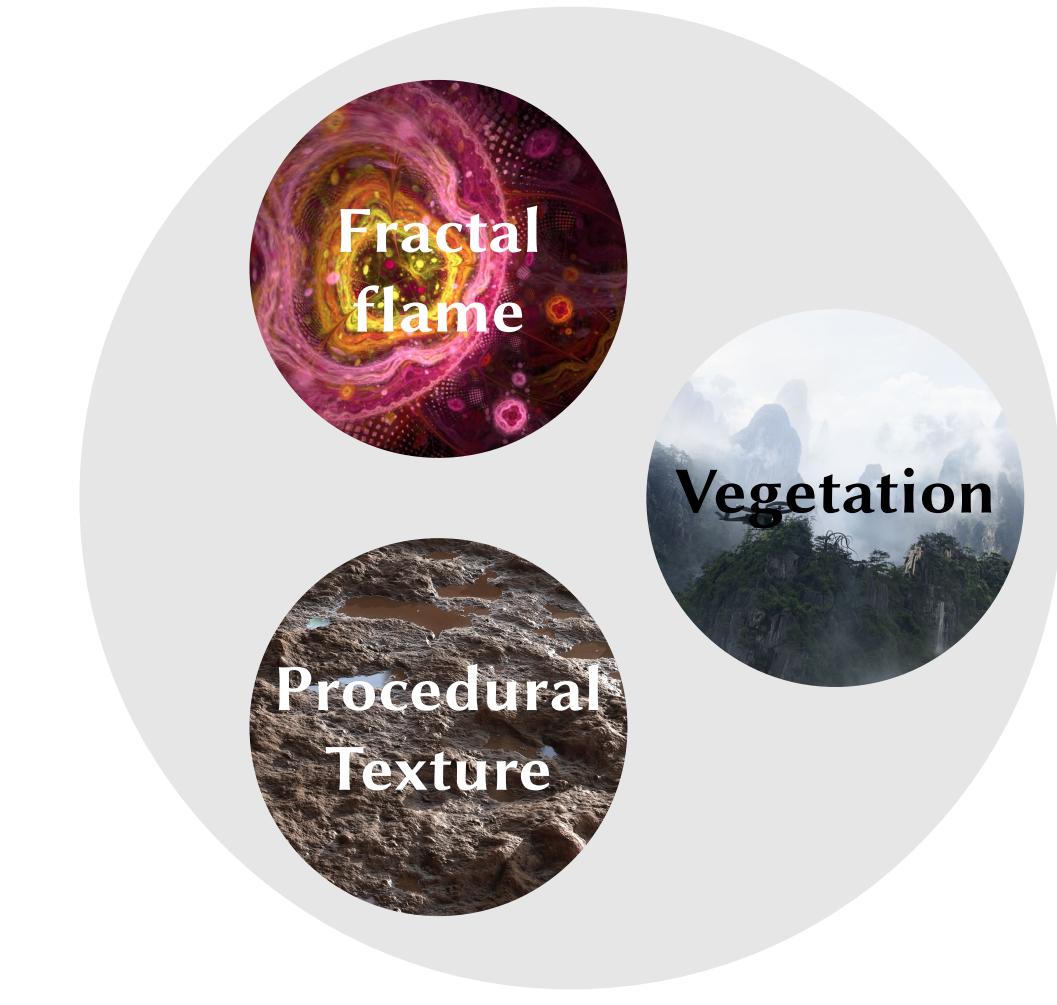
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

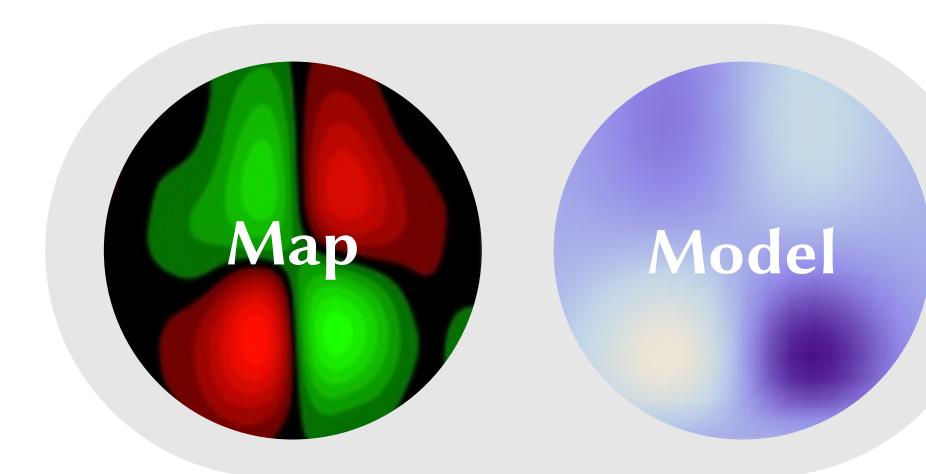


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



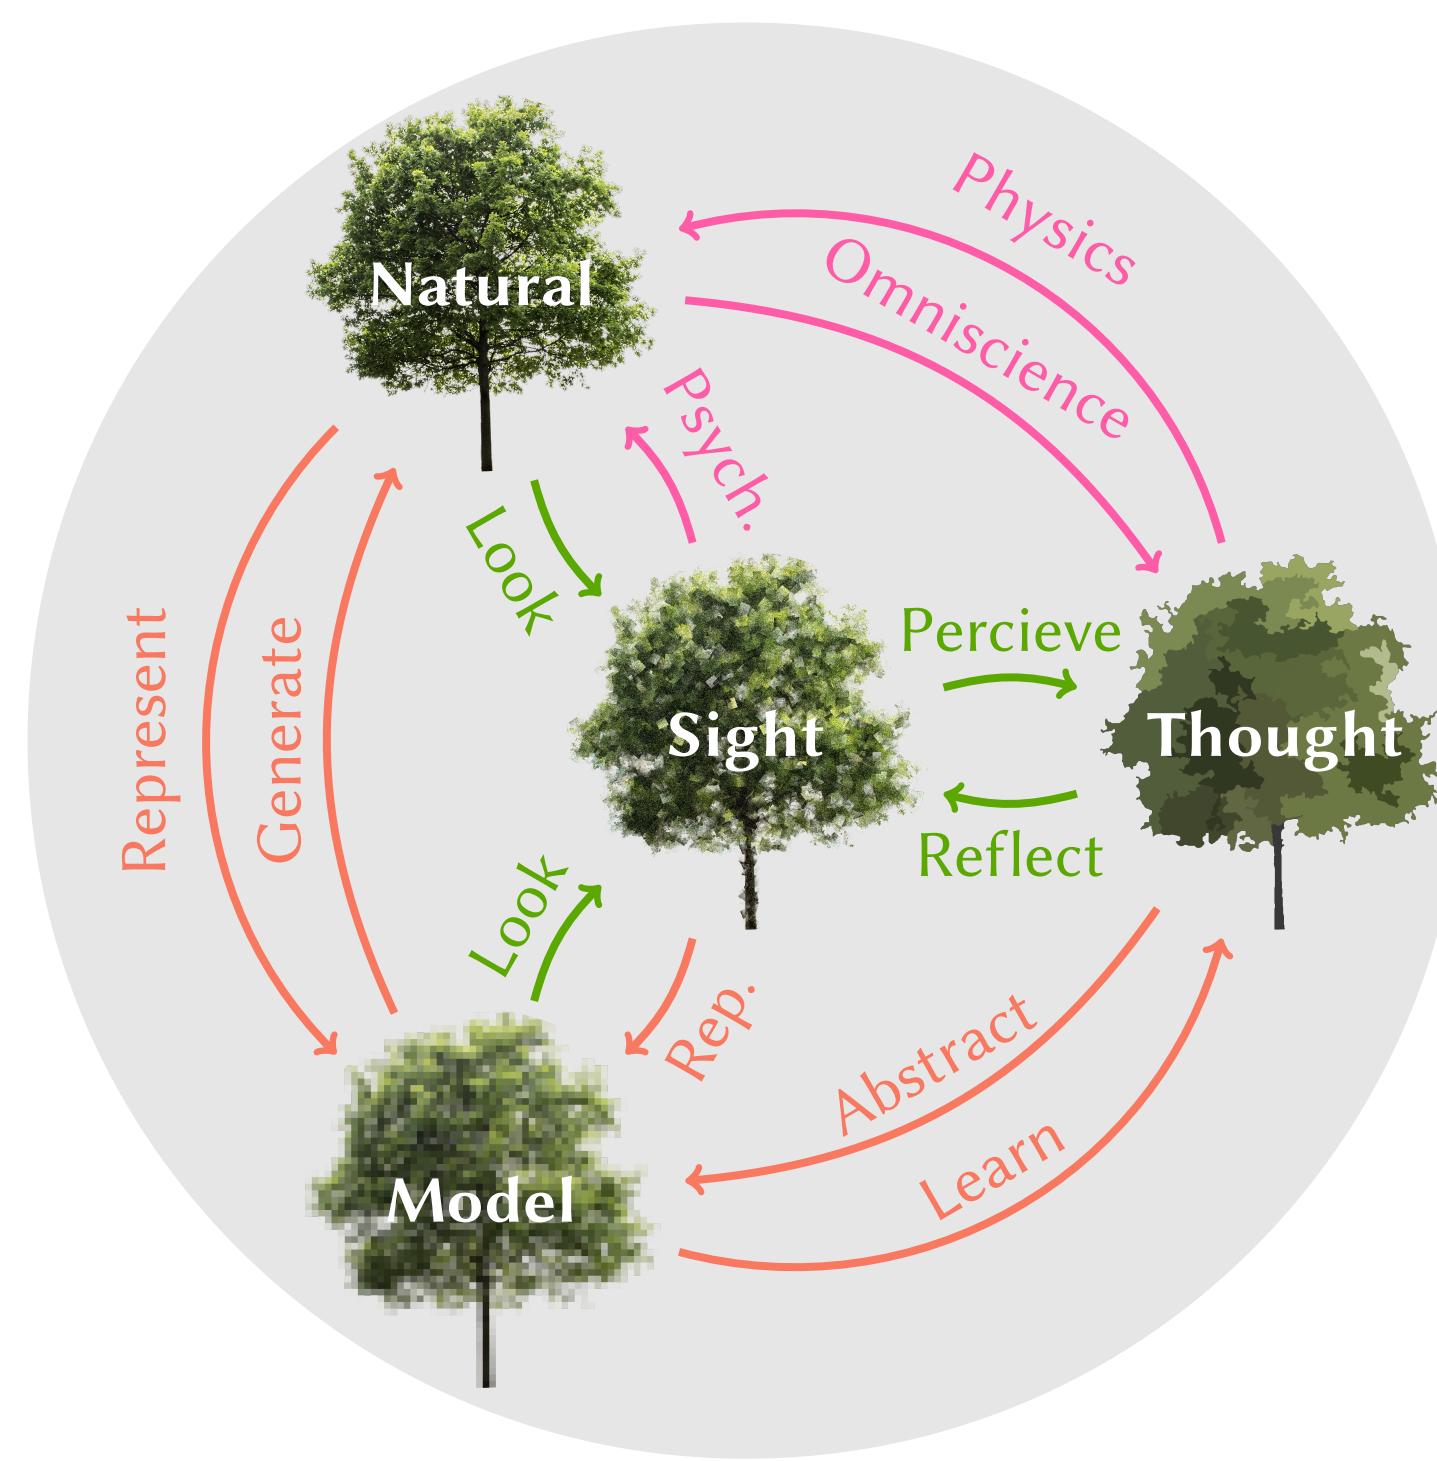
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

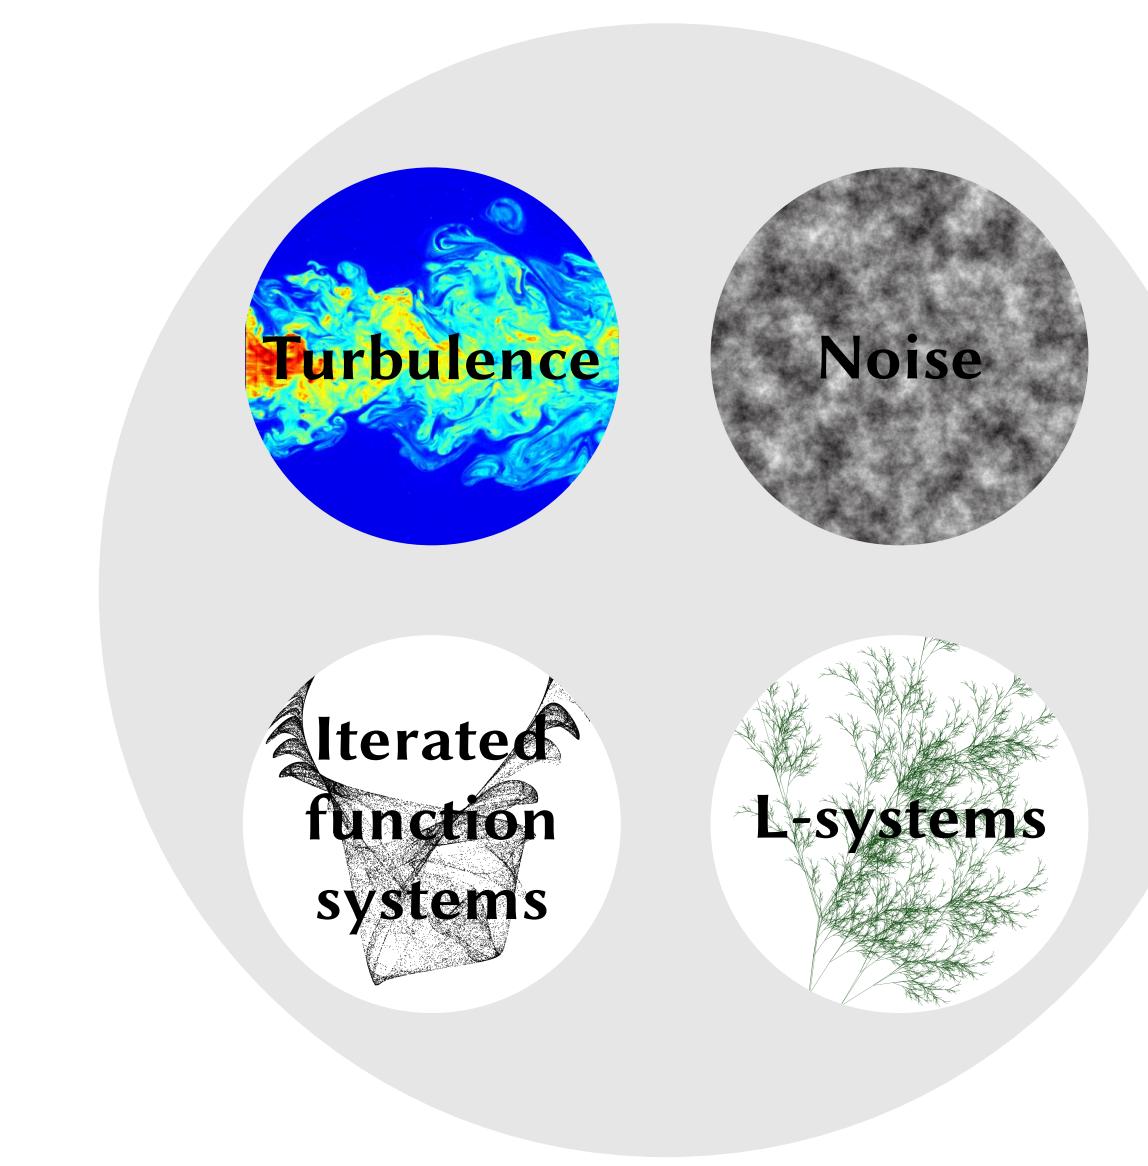


**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.

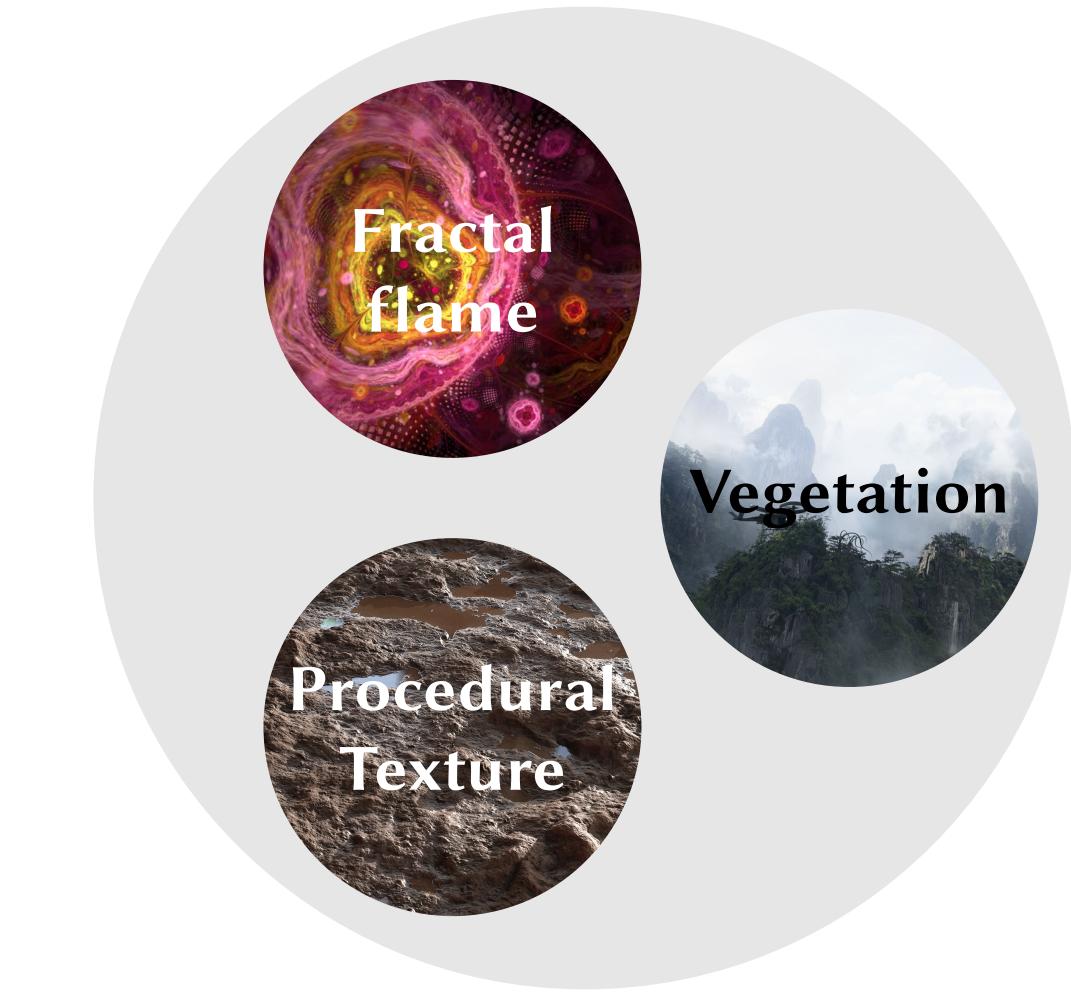
# Understanding visual complexity with statistical physics



We cannot **directly percieve** the true nature of the world, but we can form models to **indirectly percieve** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

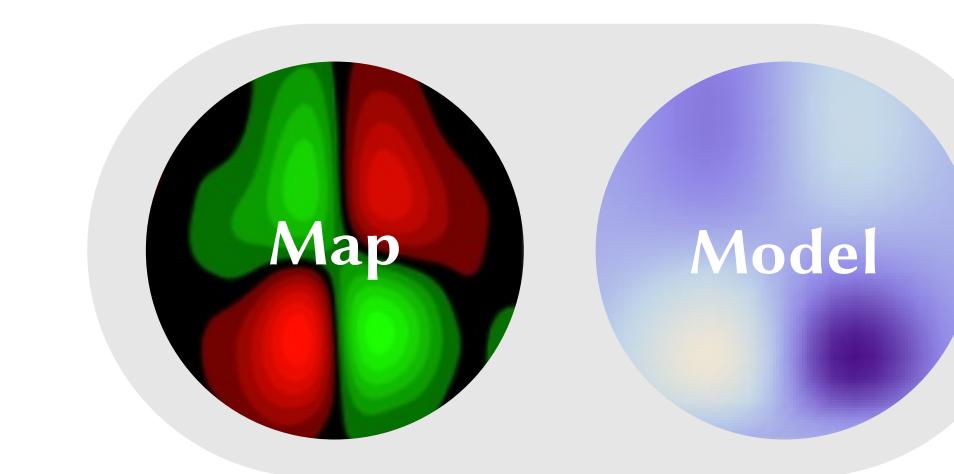


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



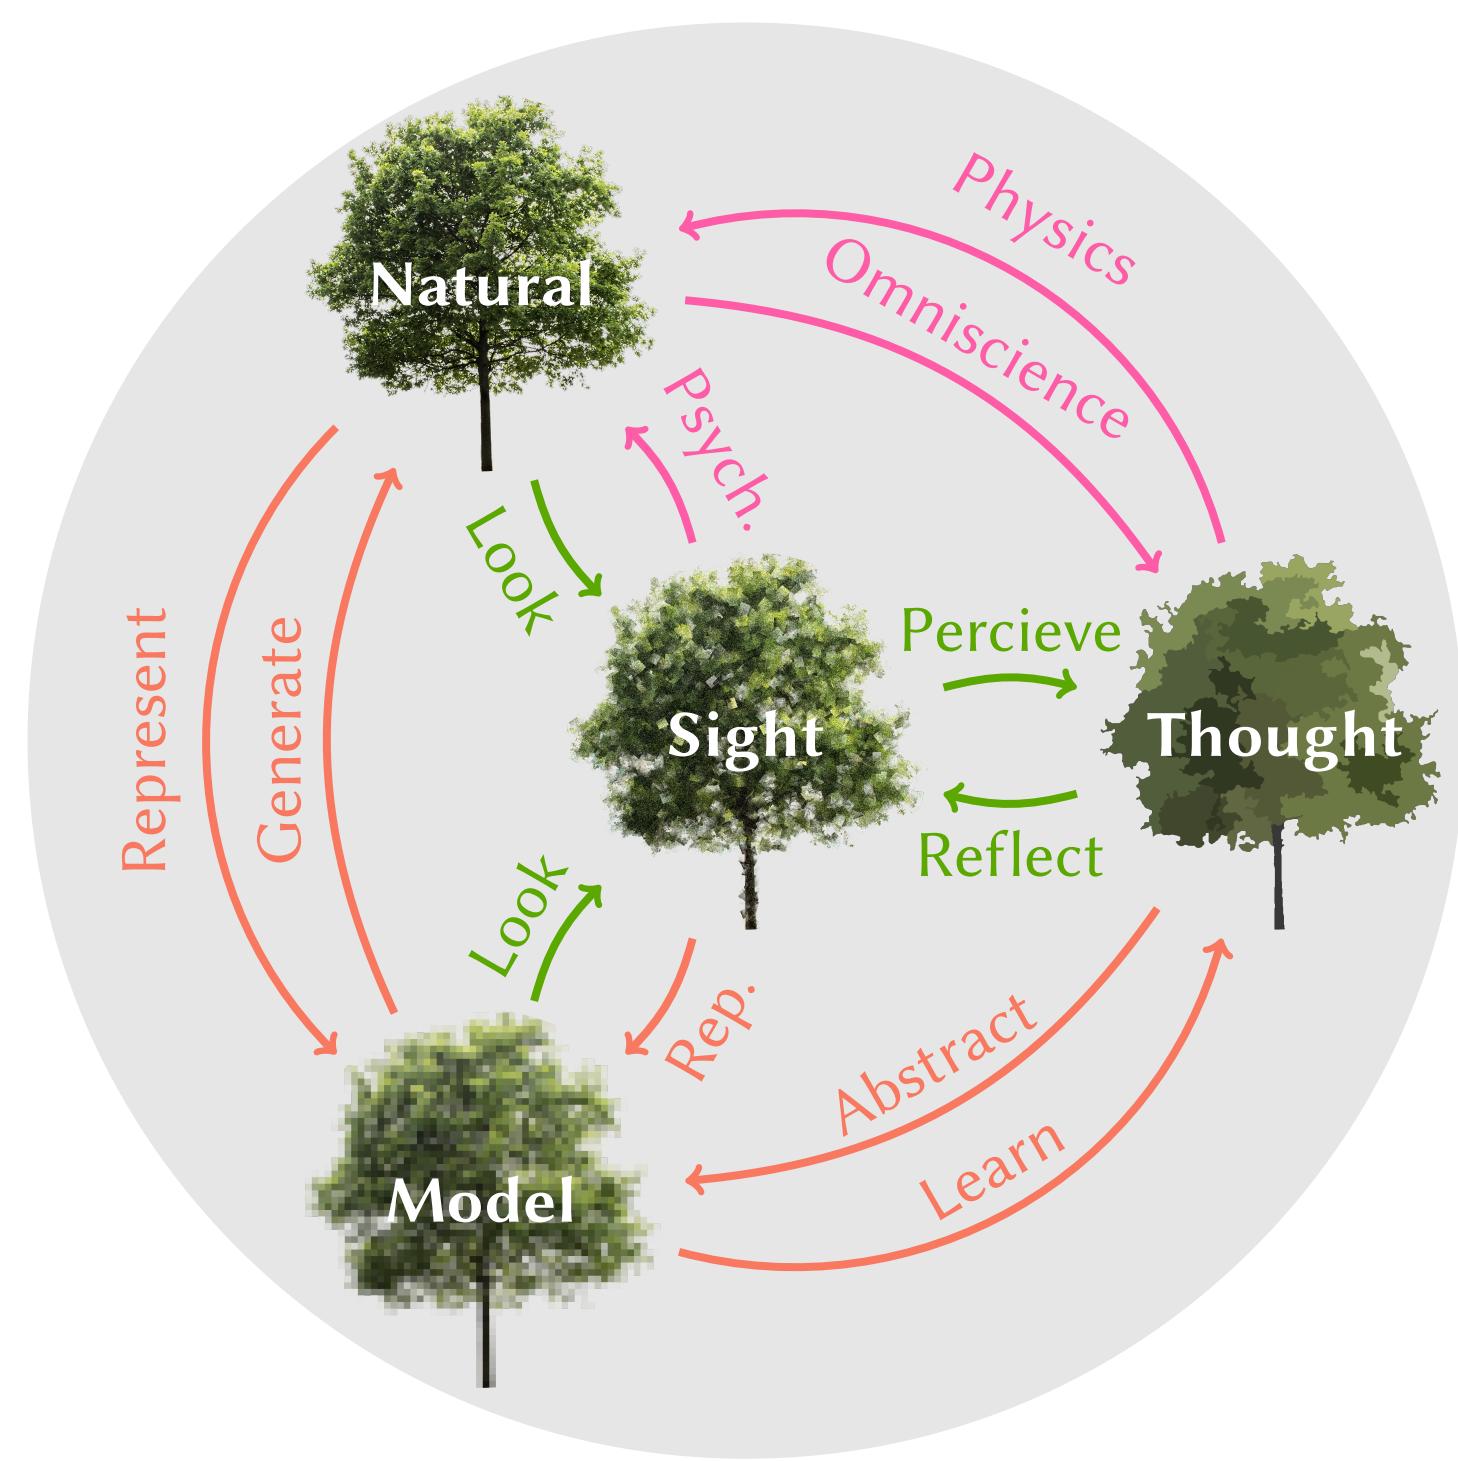
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

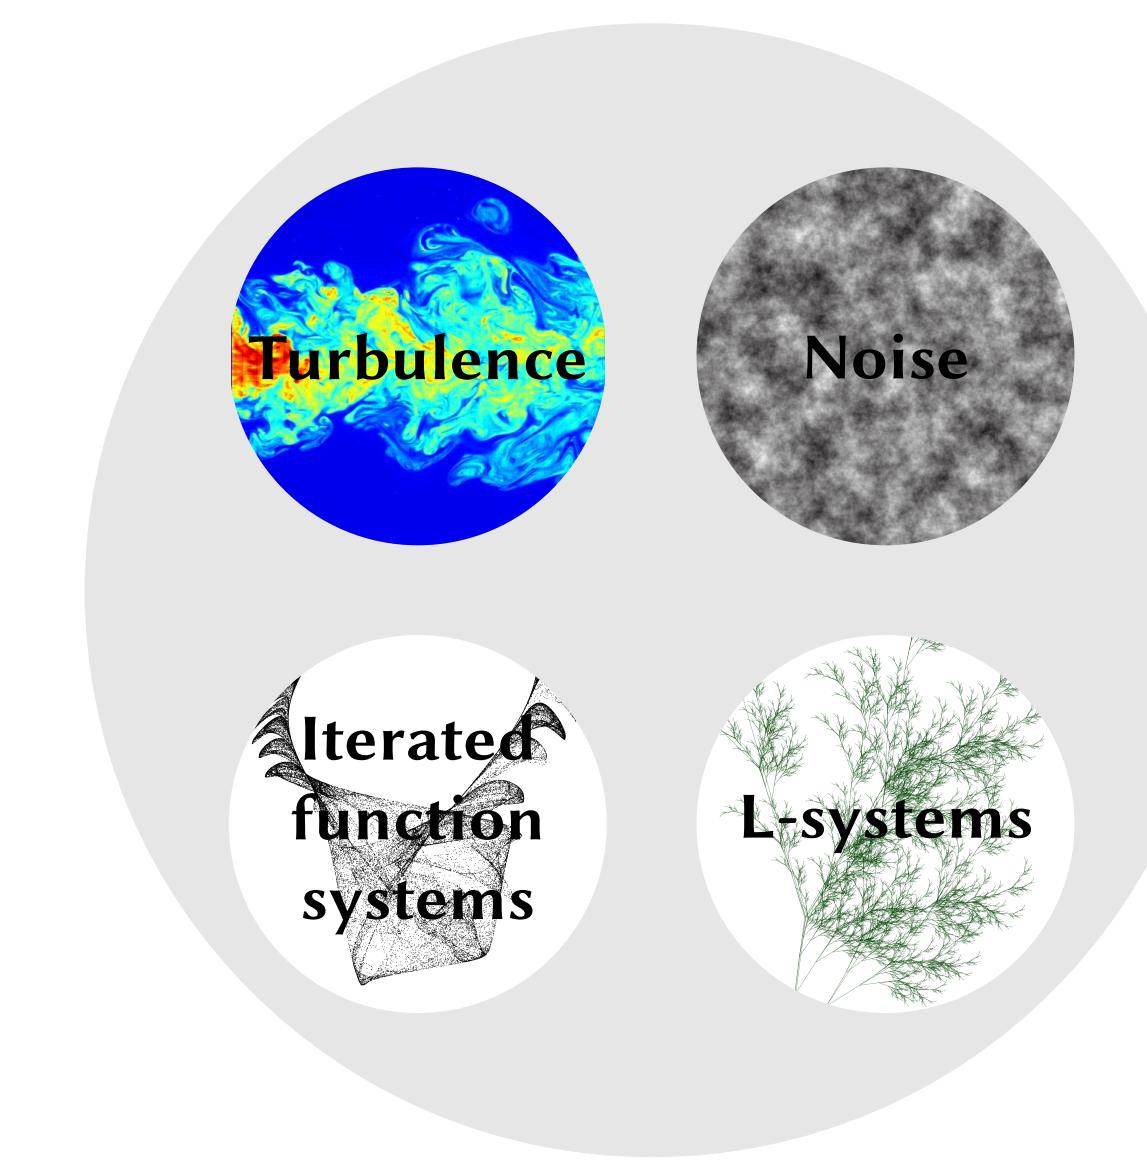


**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.

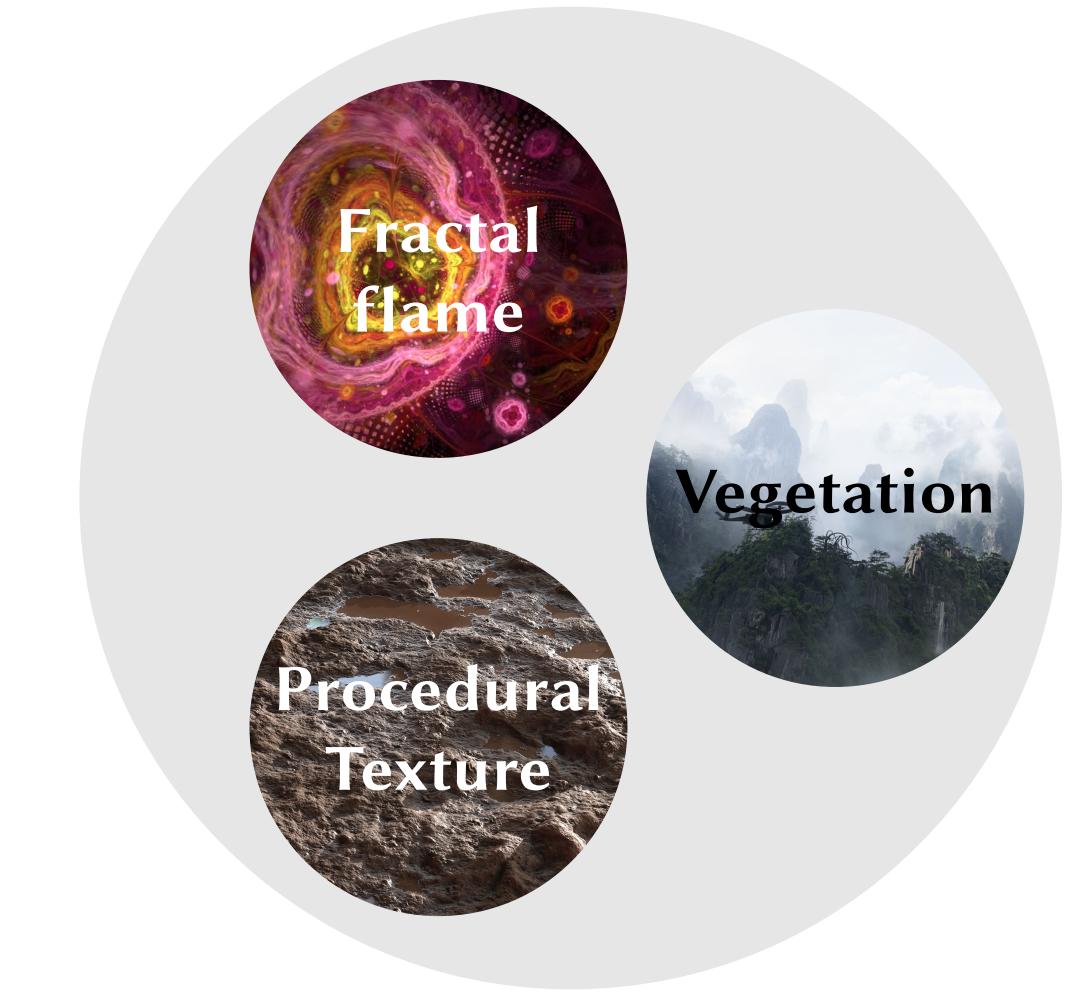
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

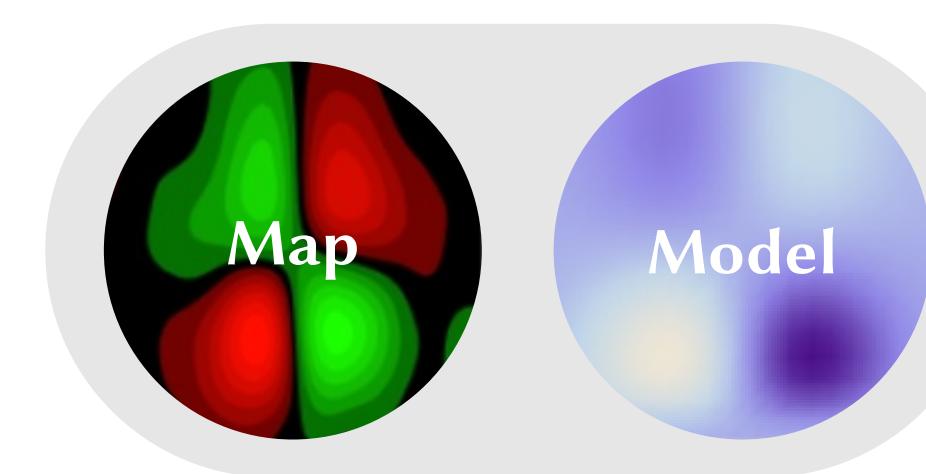


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



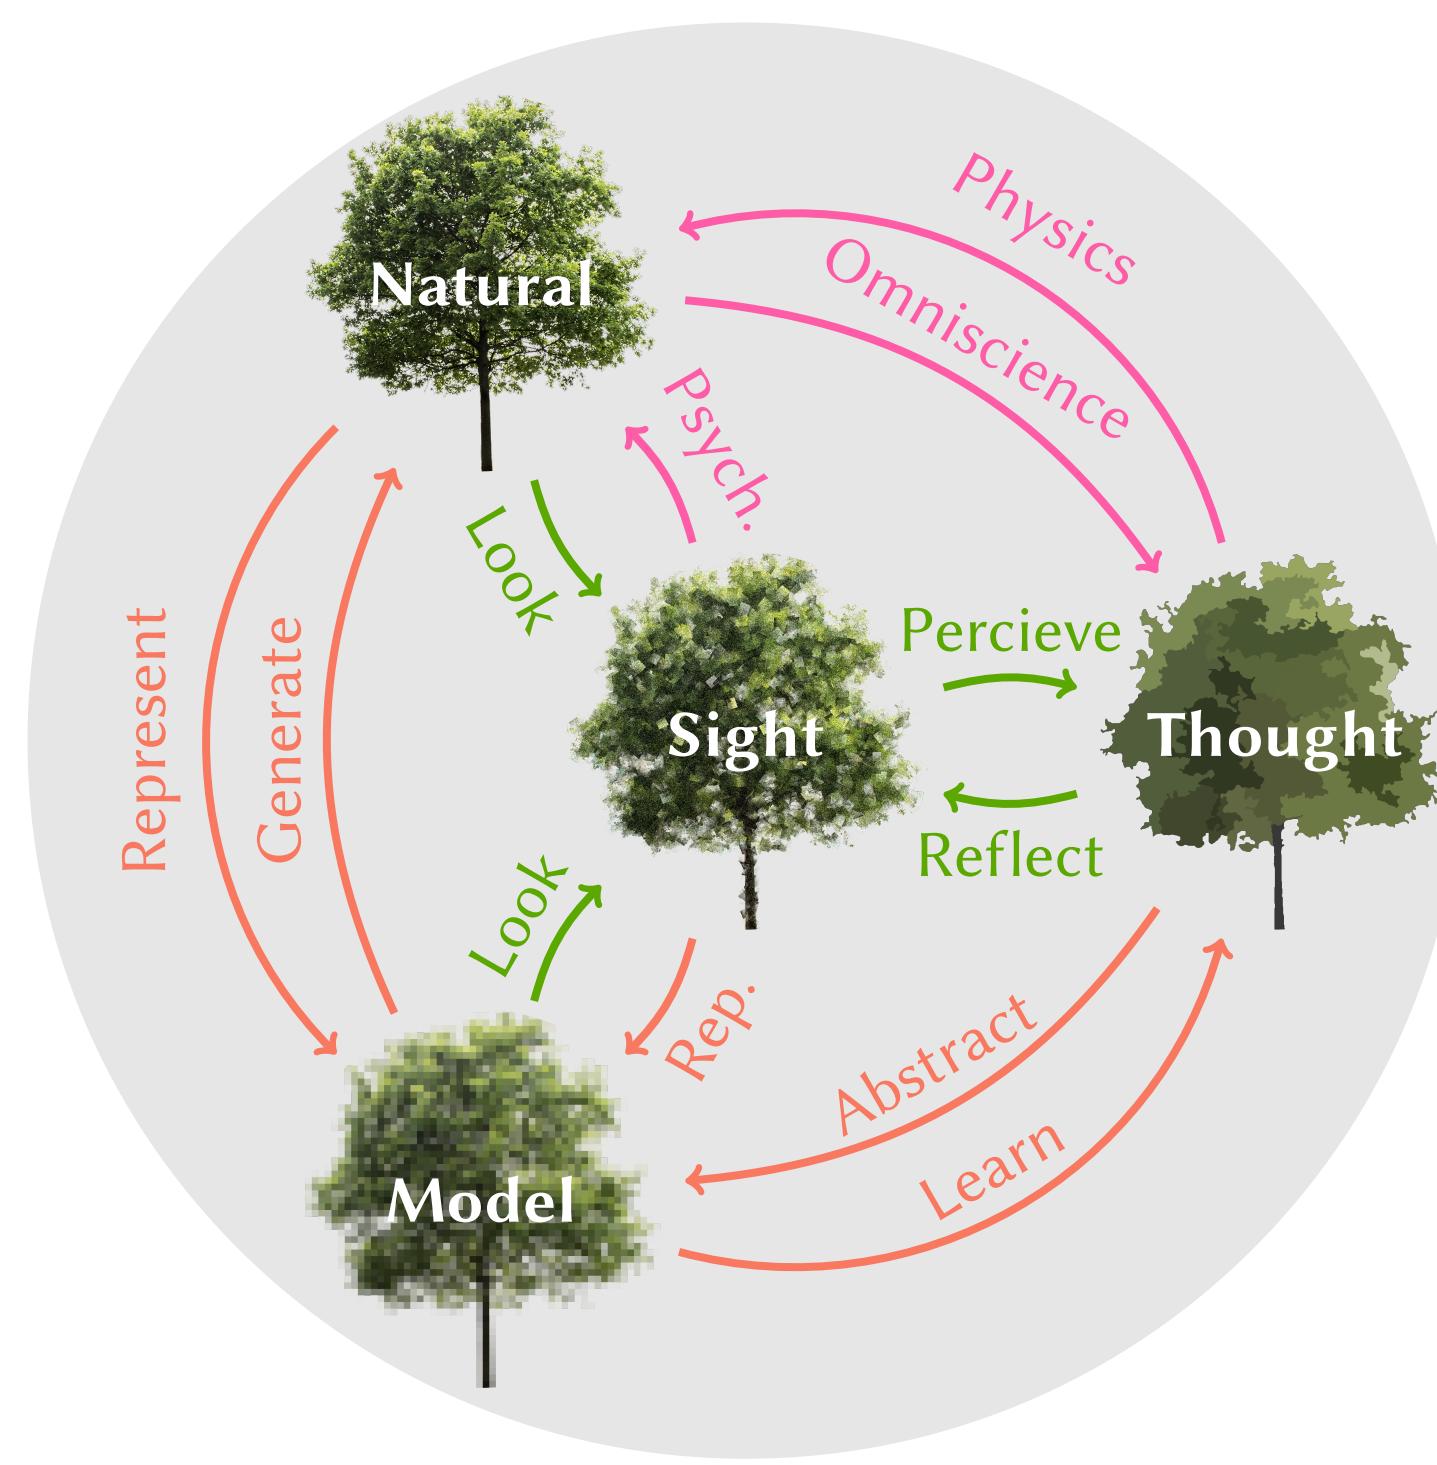
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

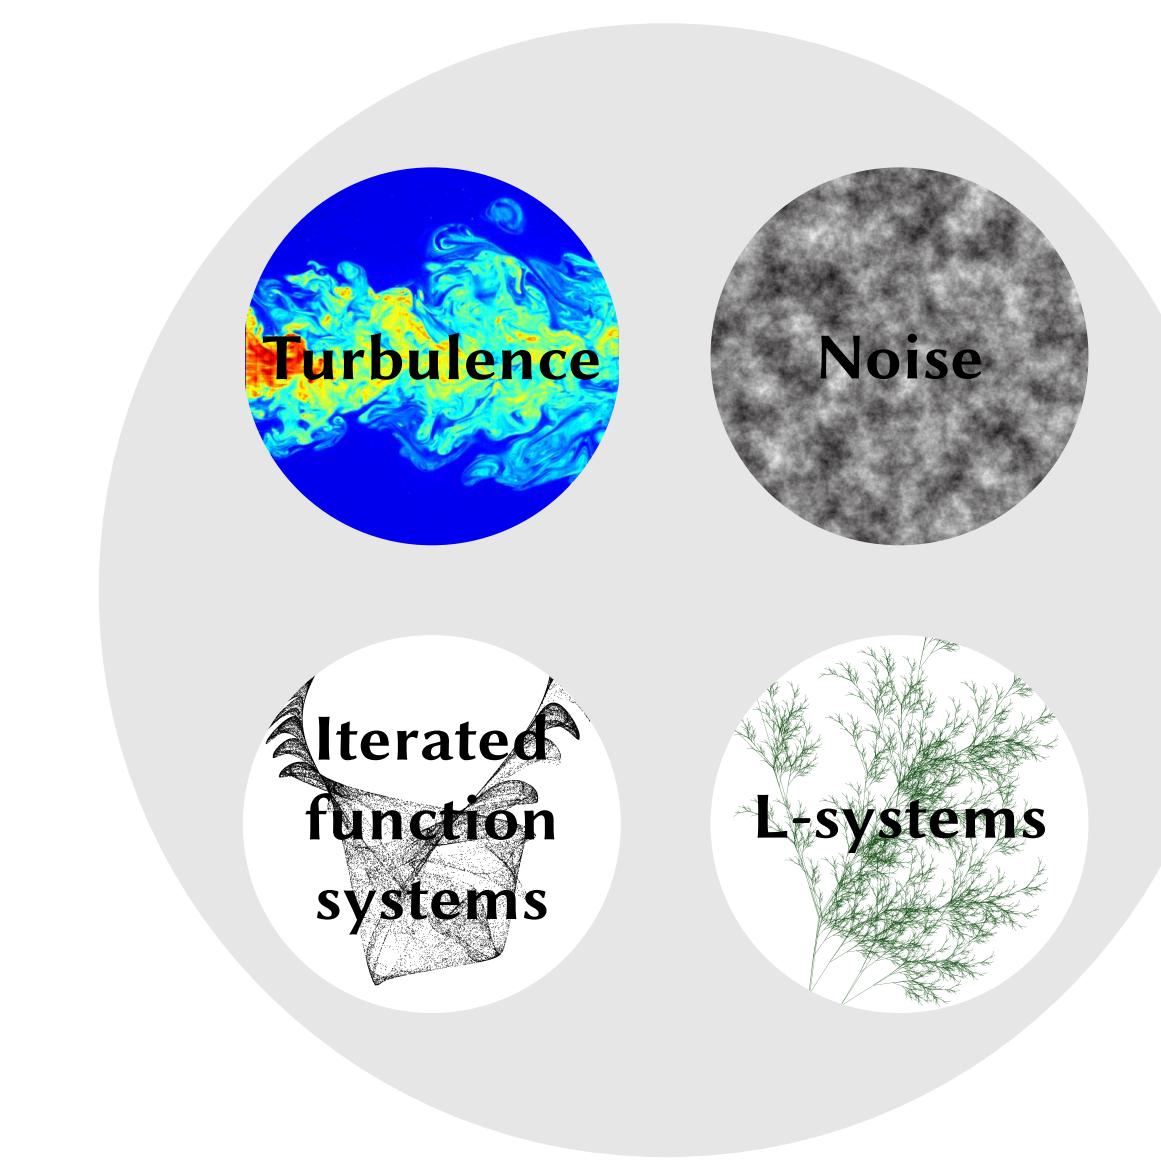


**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.

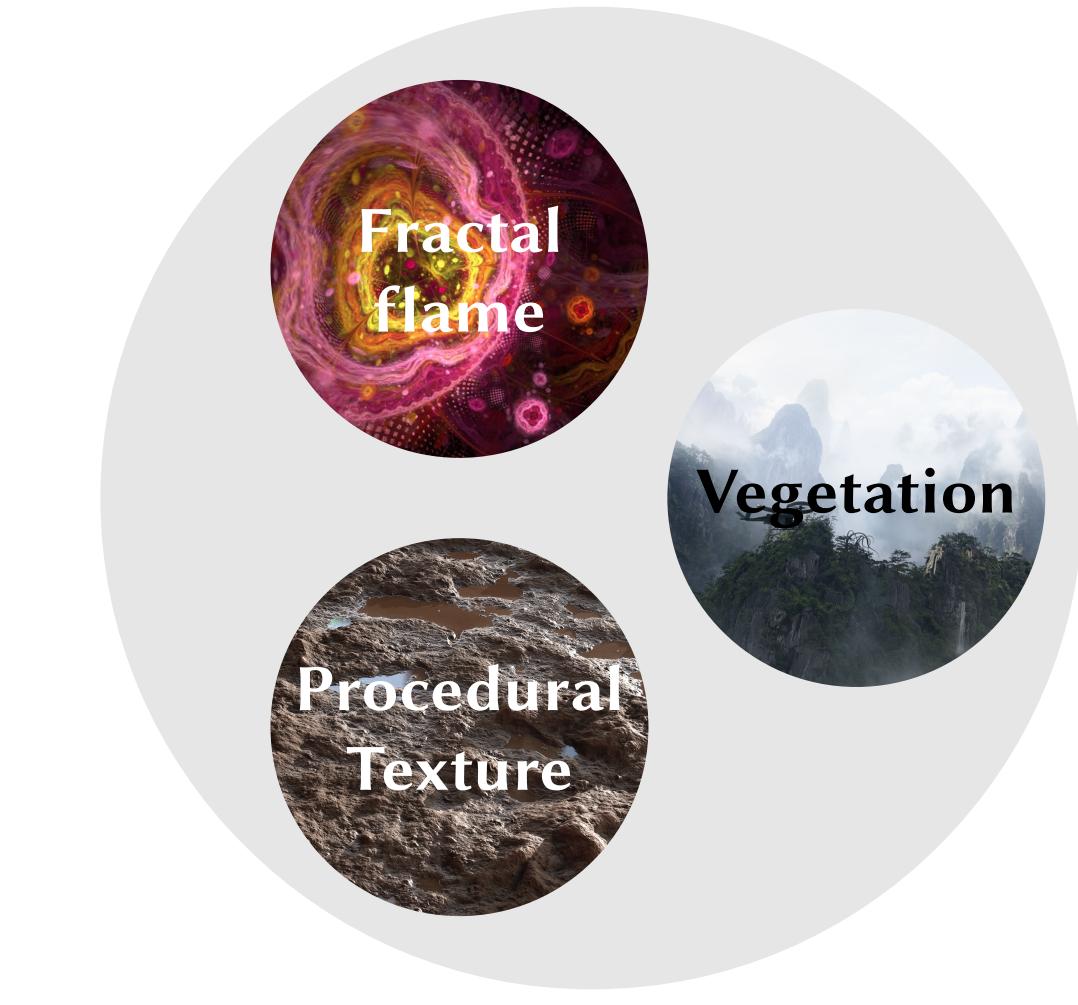
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

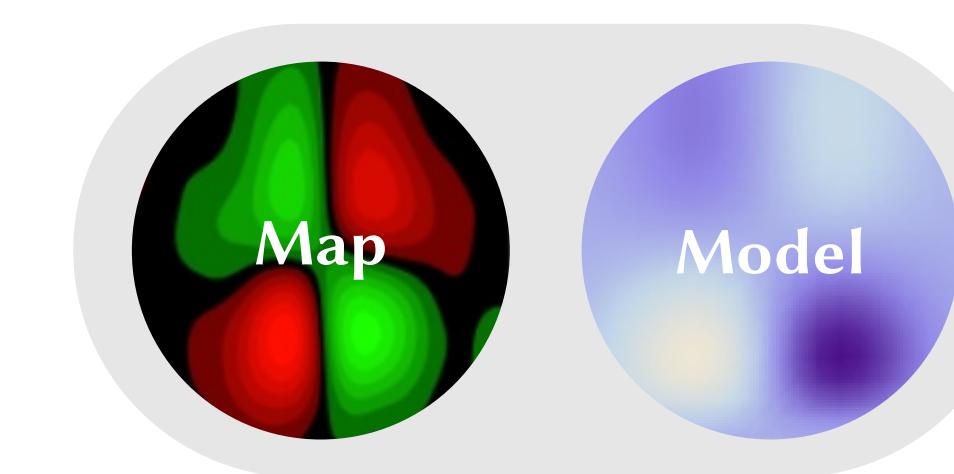


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.



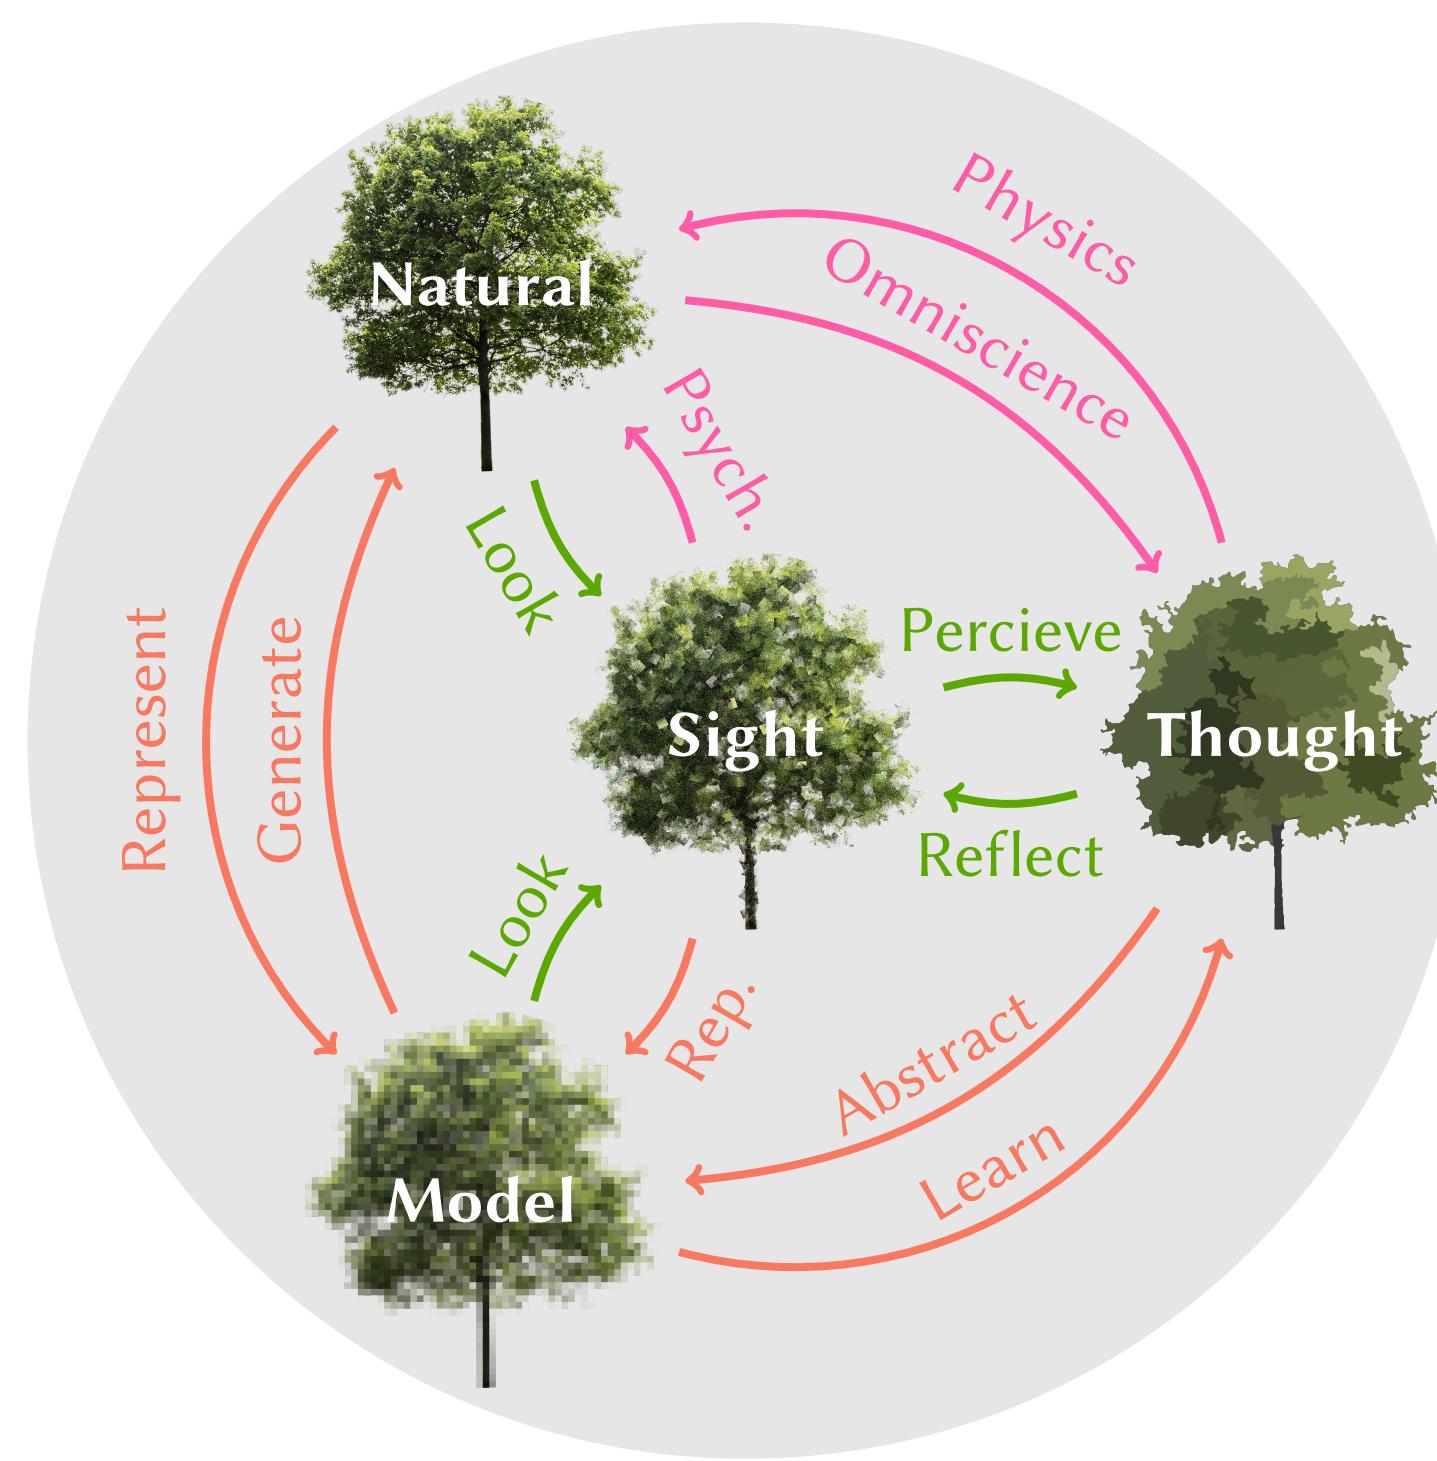
These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

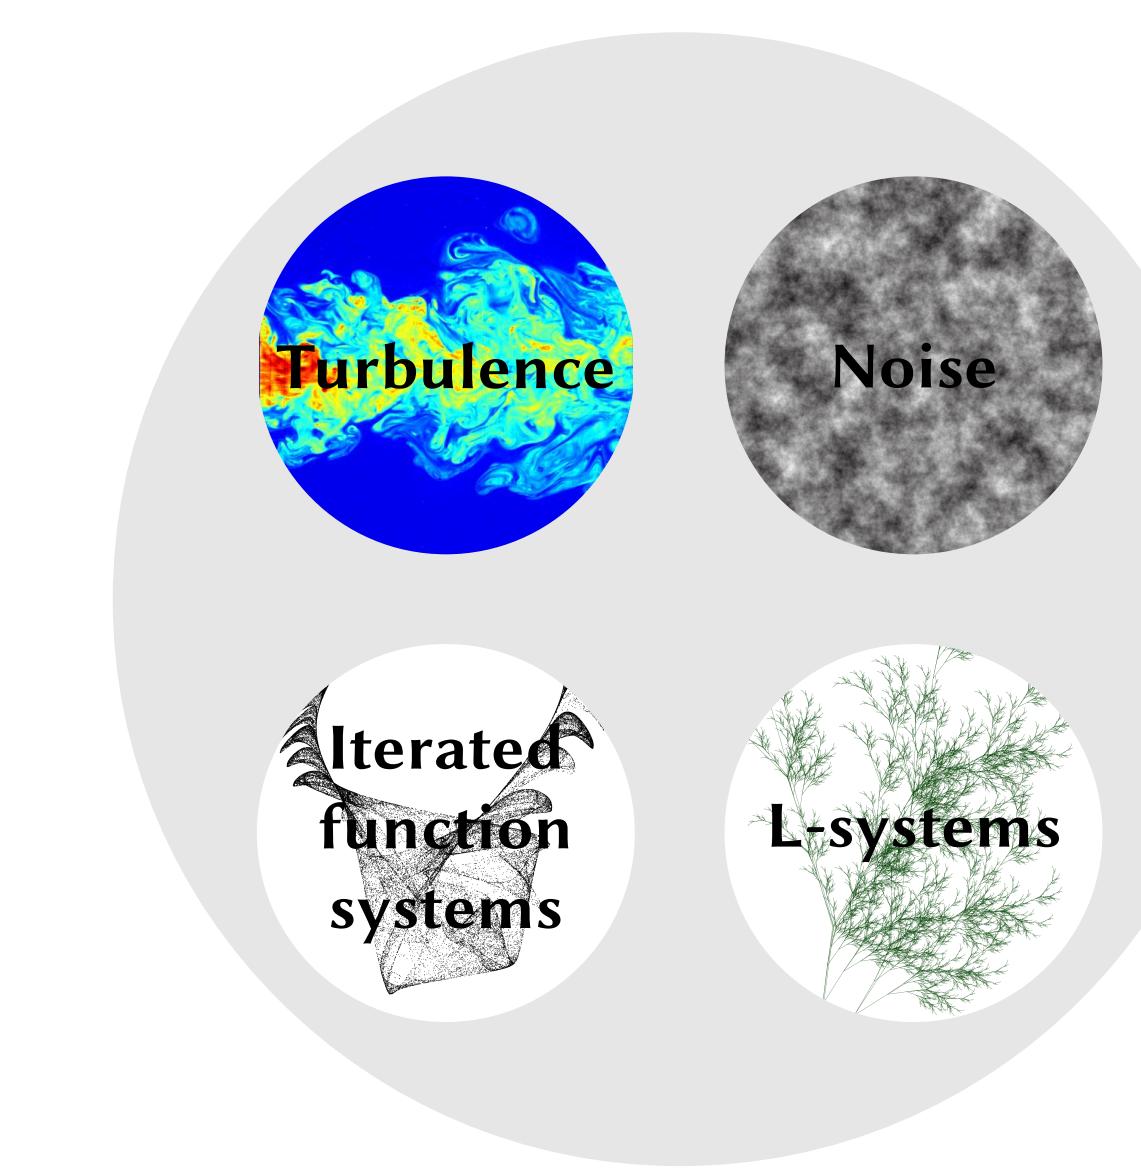


**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.

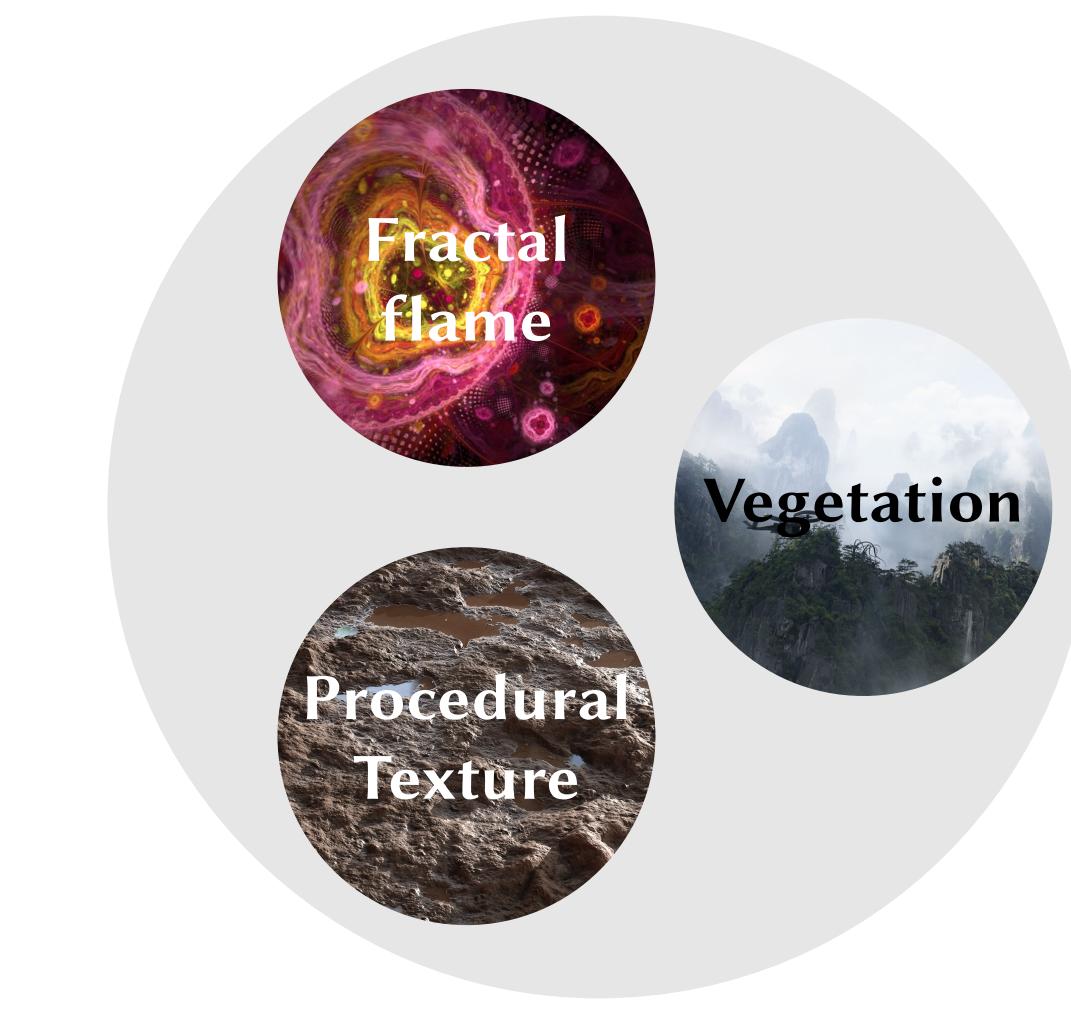
# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought. **What are good models of how nature looks to us?**

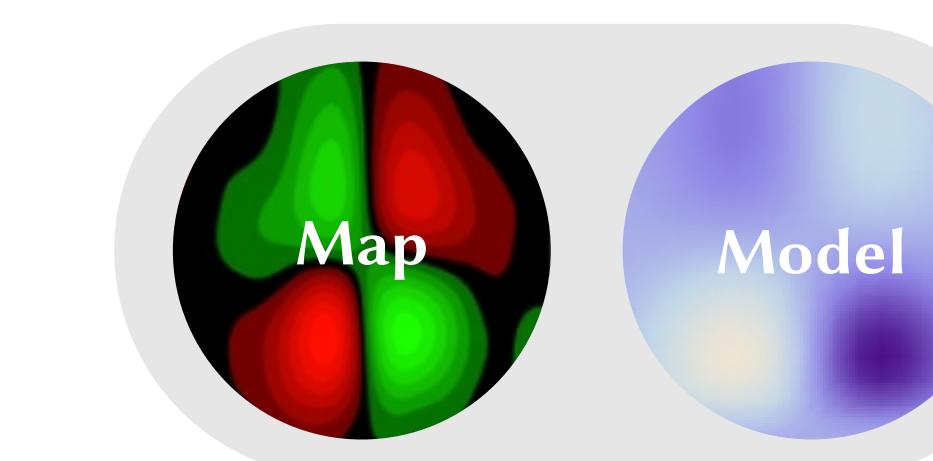


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

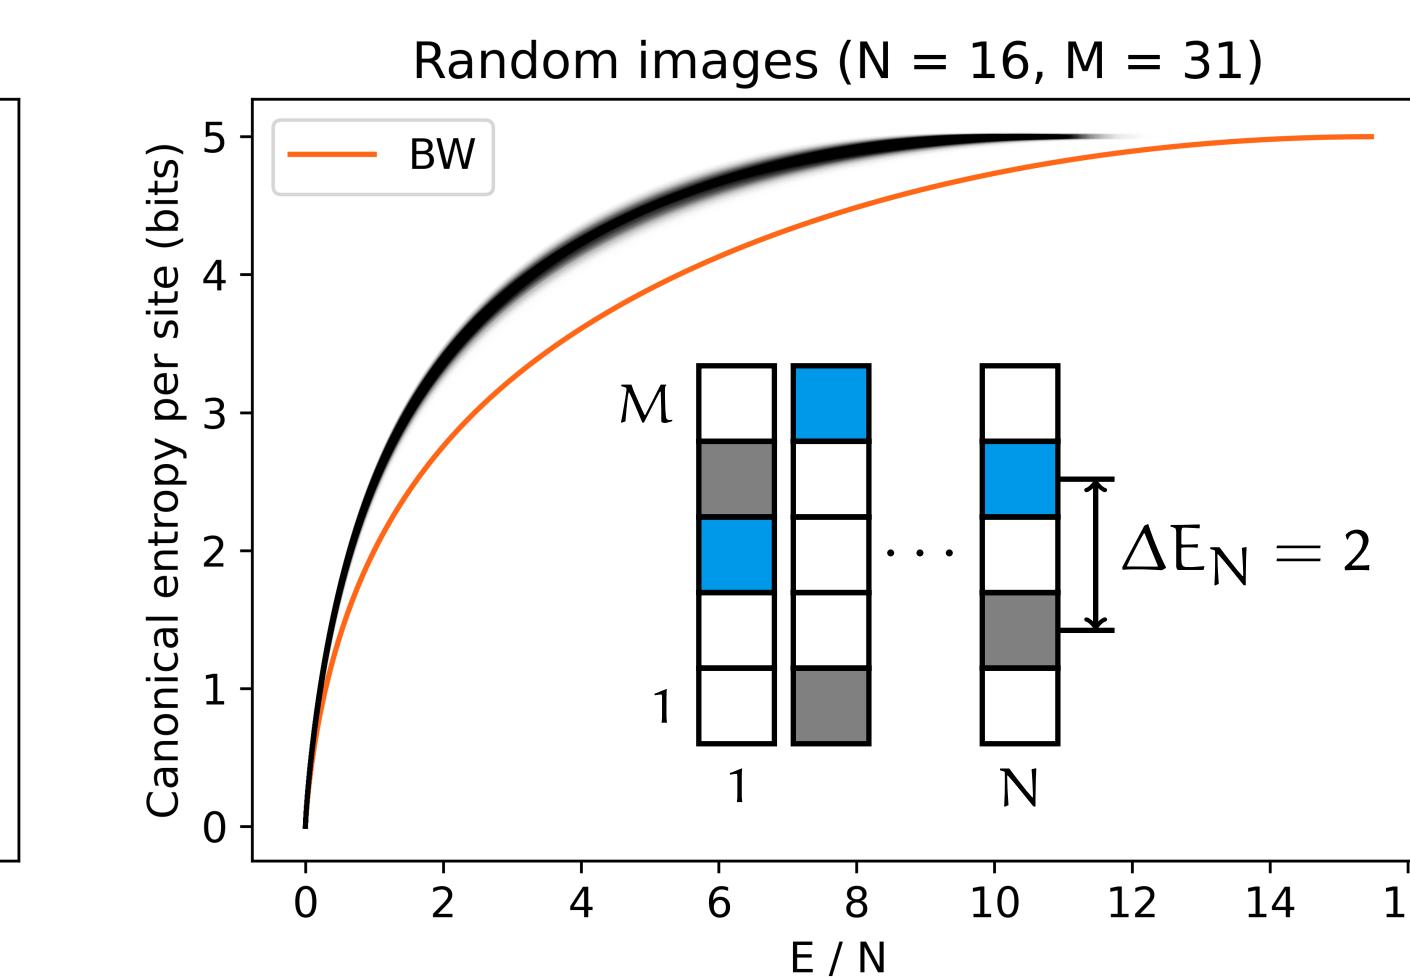
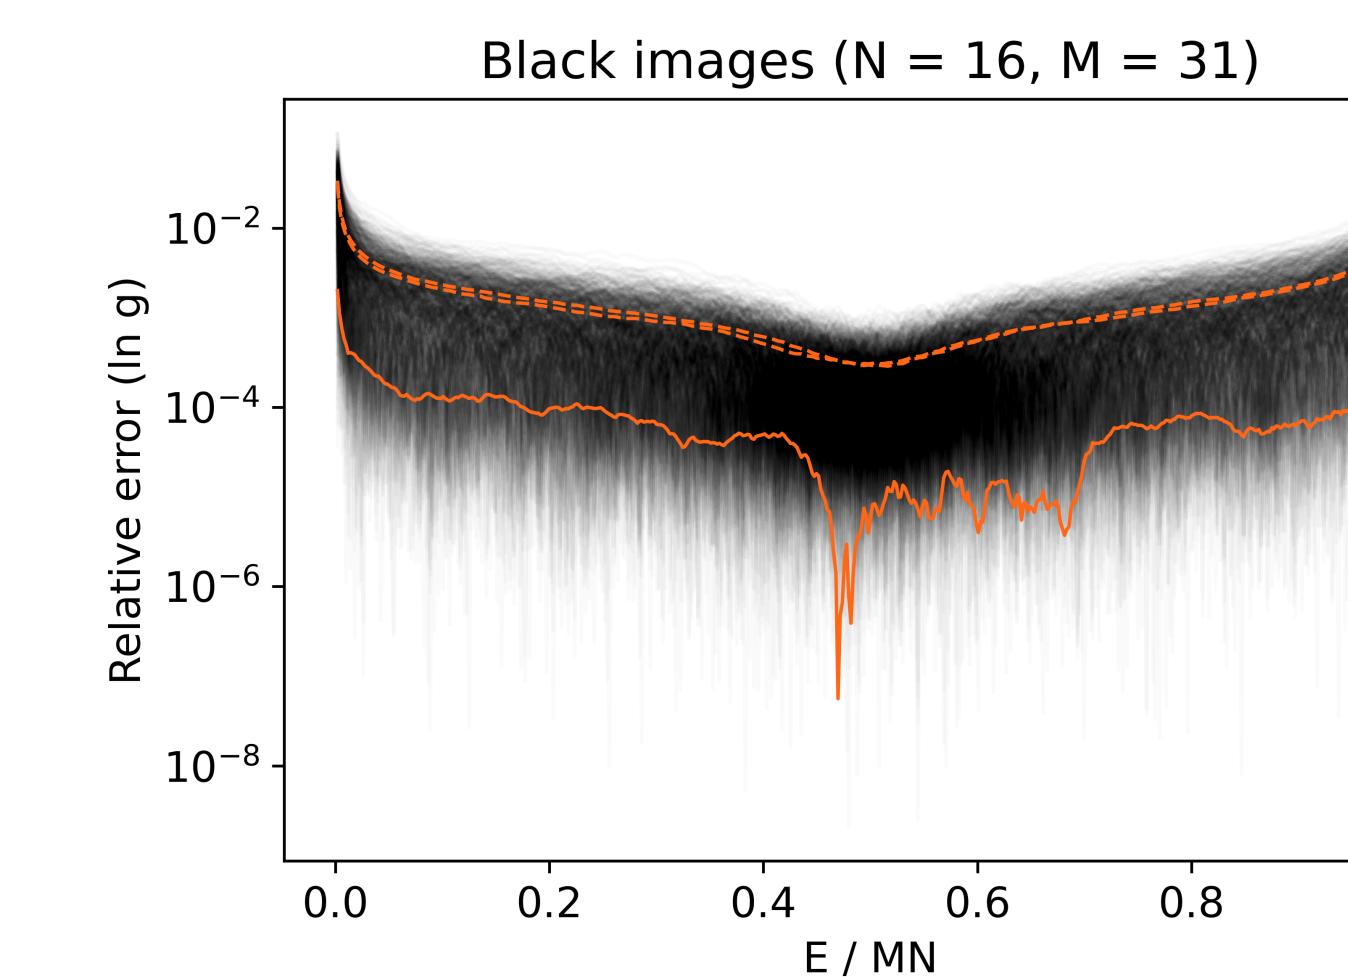


These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.

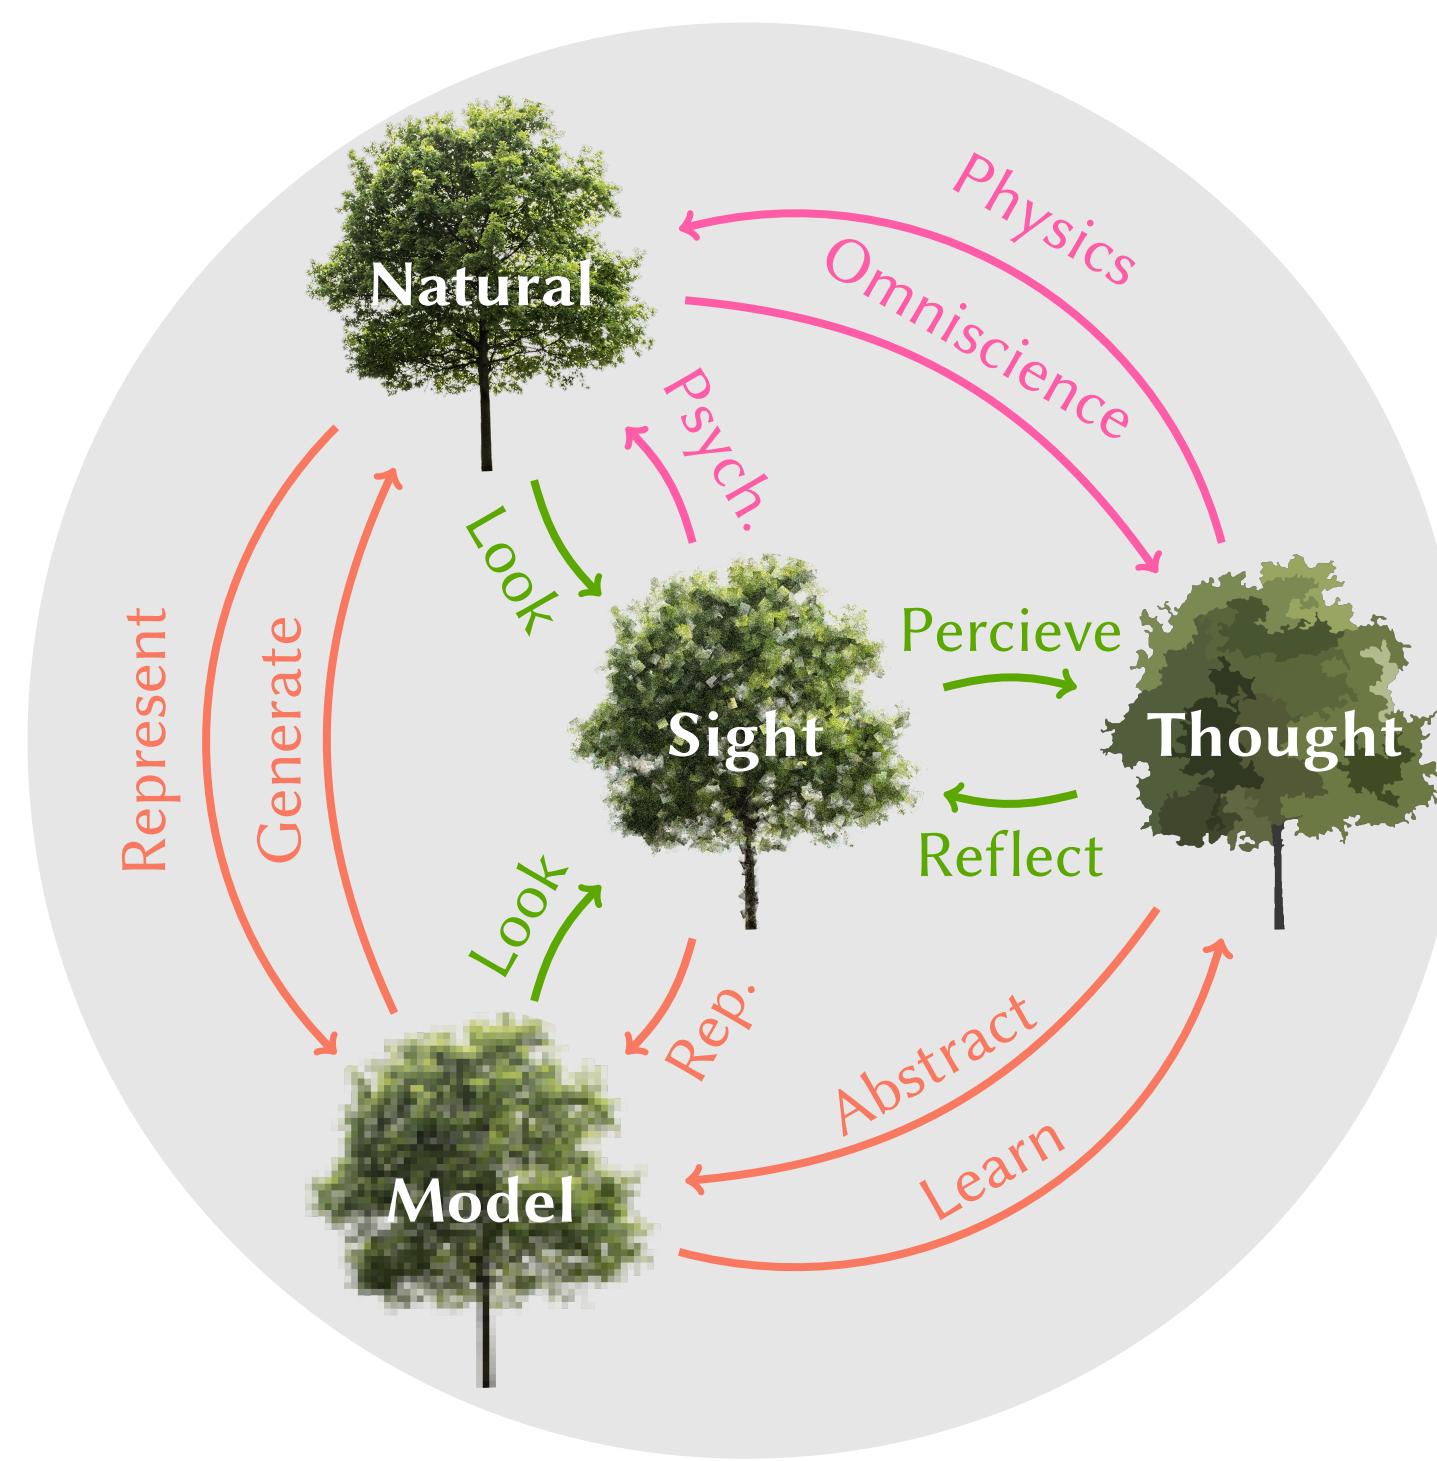


**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.



For a notion of lost information in an image, we may perform simulations of **images as discrete-level systems**, where the **energy** is the sum of the differences in pixel values. The *Wang-Landau* algorithm was implemented for images, and confirmed to have the expected convergence properties on **black or white (BW)** images (*left*). Simulations for 1024 random grayscale images (*right*) reveal how entropy scales with how different the pixel values are (average energy).

# Understanding visual complexity with statistical physics



We cannot **directly perceive** the true nature of the world, but we can form models to **indirectly perceive** nature. A model must represent nature, while still being something we can readily look at and use to guide thought.

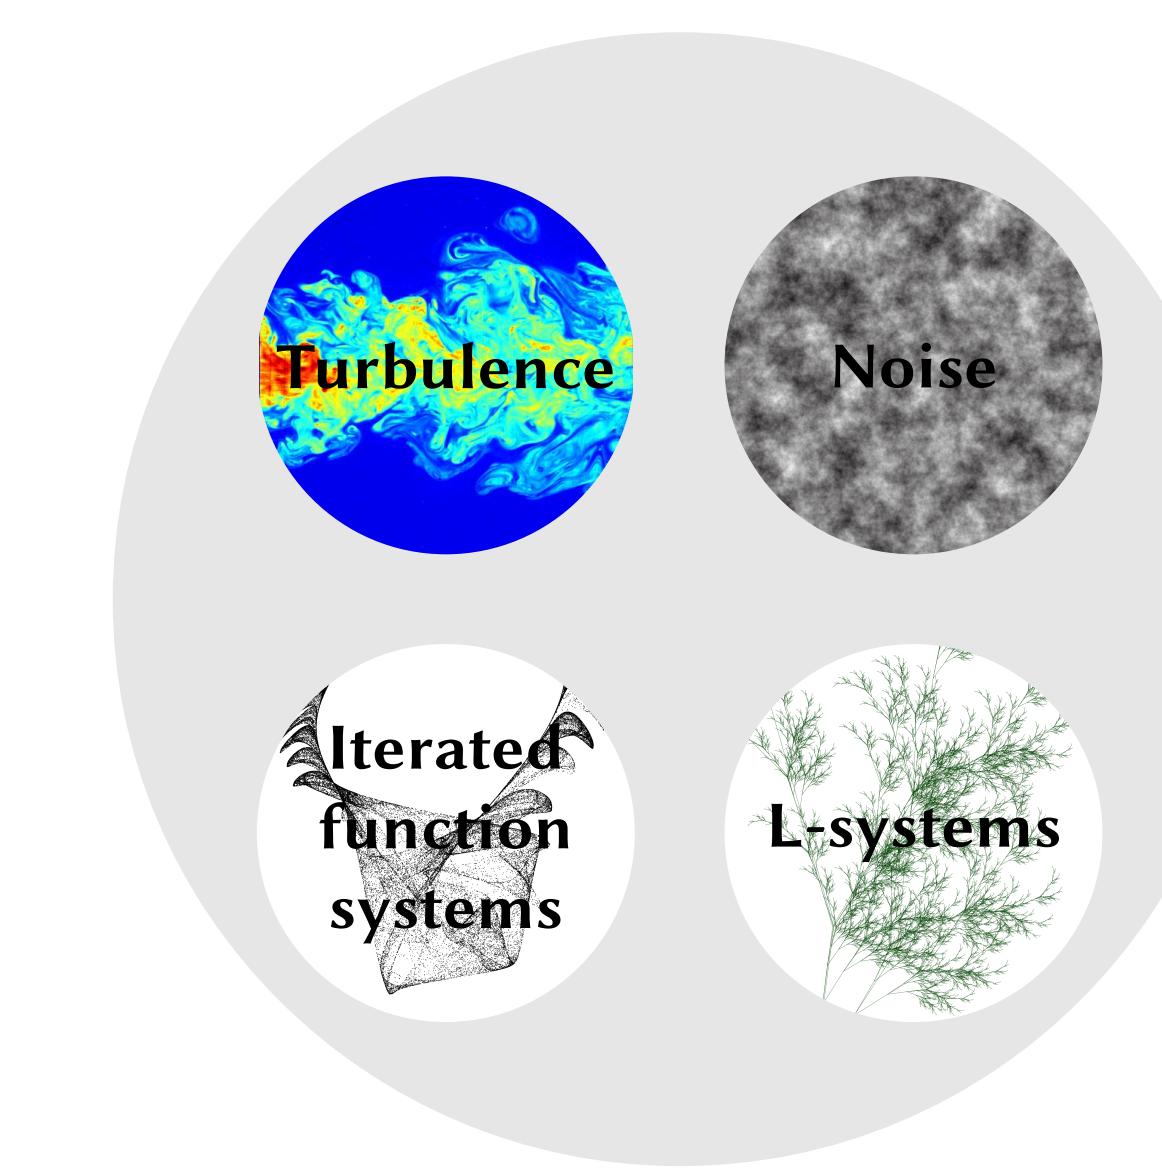
**What are good models of how nature looks to us?**

## Statistical Physics and Information Theory

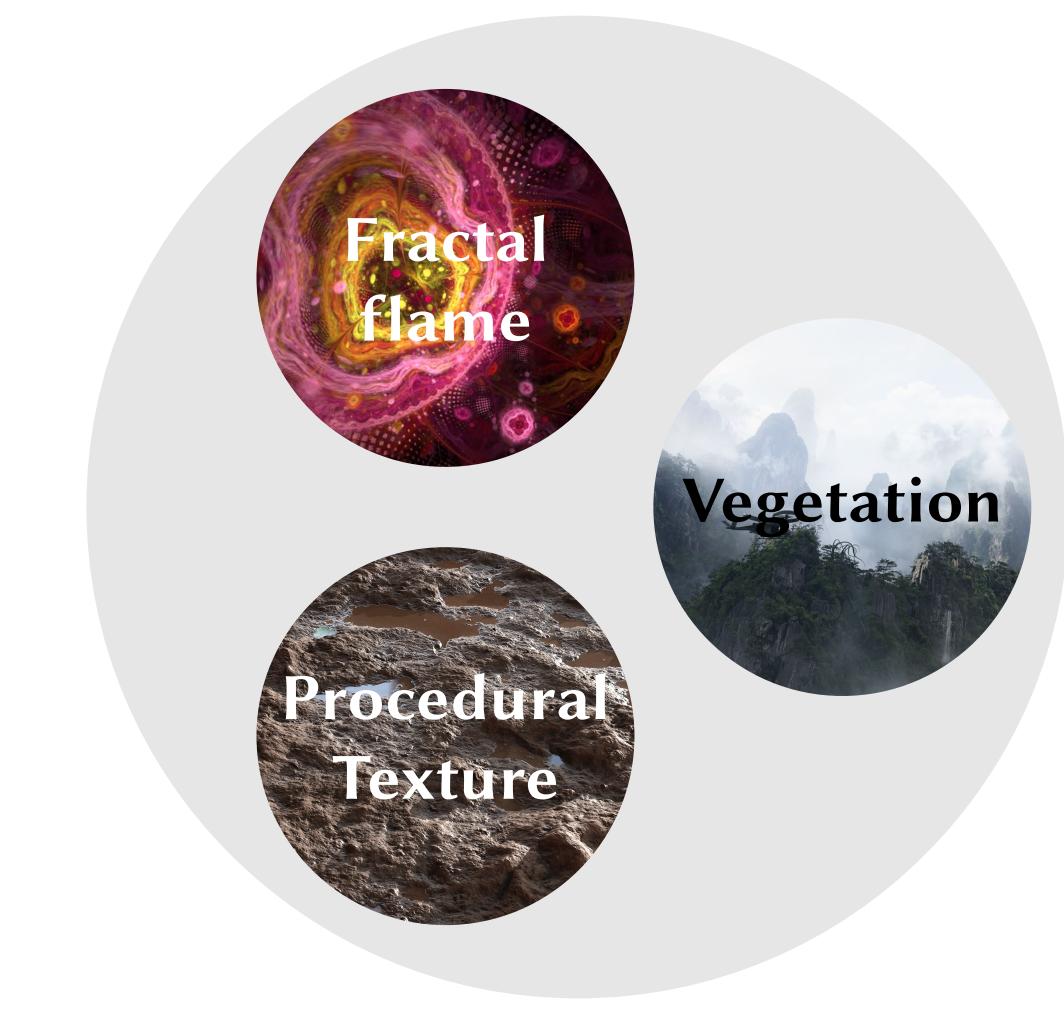
$$S = - \sum_{x \in \mathcal{X}} p(x) \ln p(x) = \langle I \rangle_X$$

$$D_{KL}(p \parallel q) = \sum_{x \in \mathcal{X}} p(x) \ln \frac{p(x)}{q(x)}$$

$$Z = \sum_{x \in \mathcal{X}} e^{-\beta E(x)}$$

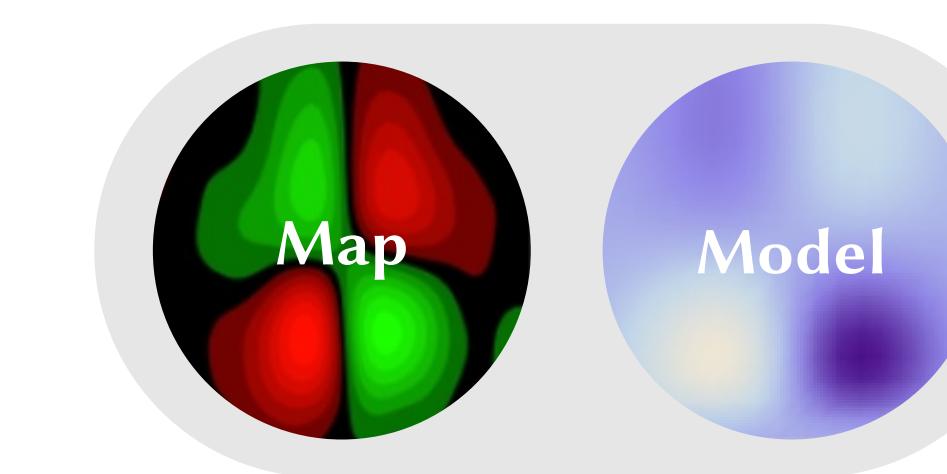


**Visual complexity** arises in many ways. Computational models of these phenomena are the bricks that build up the structure of an image.

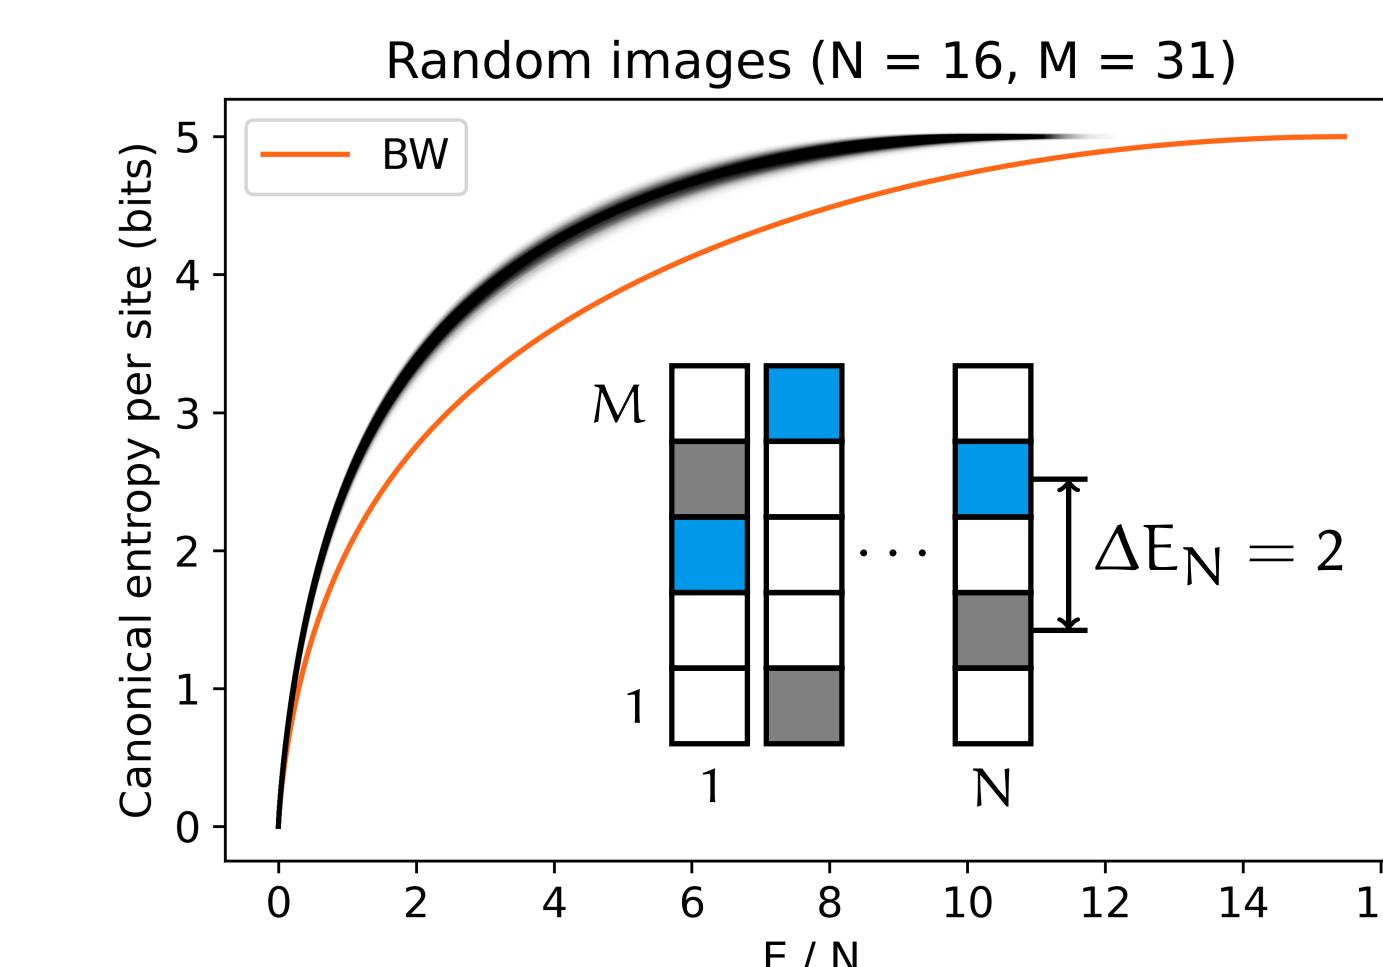
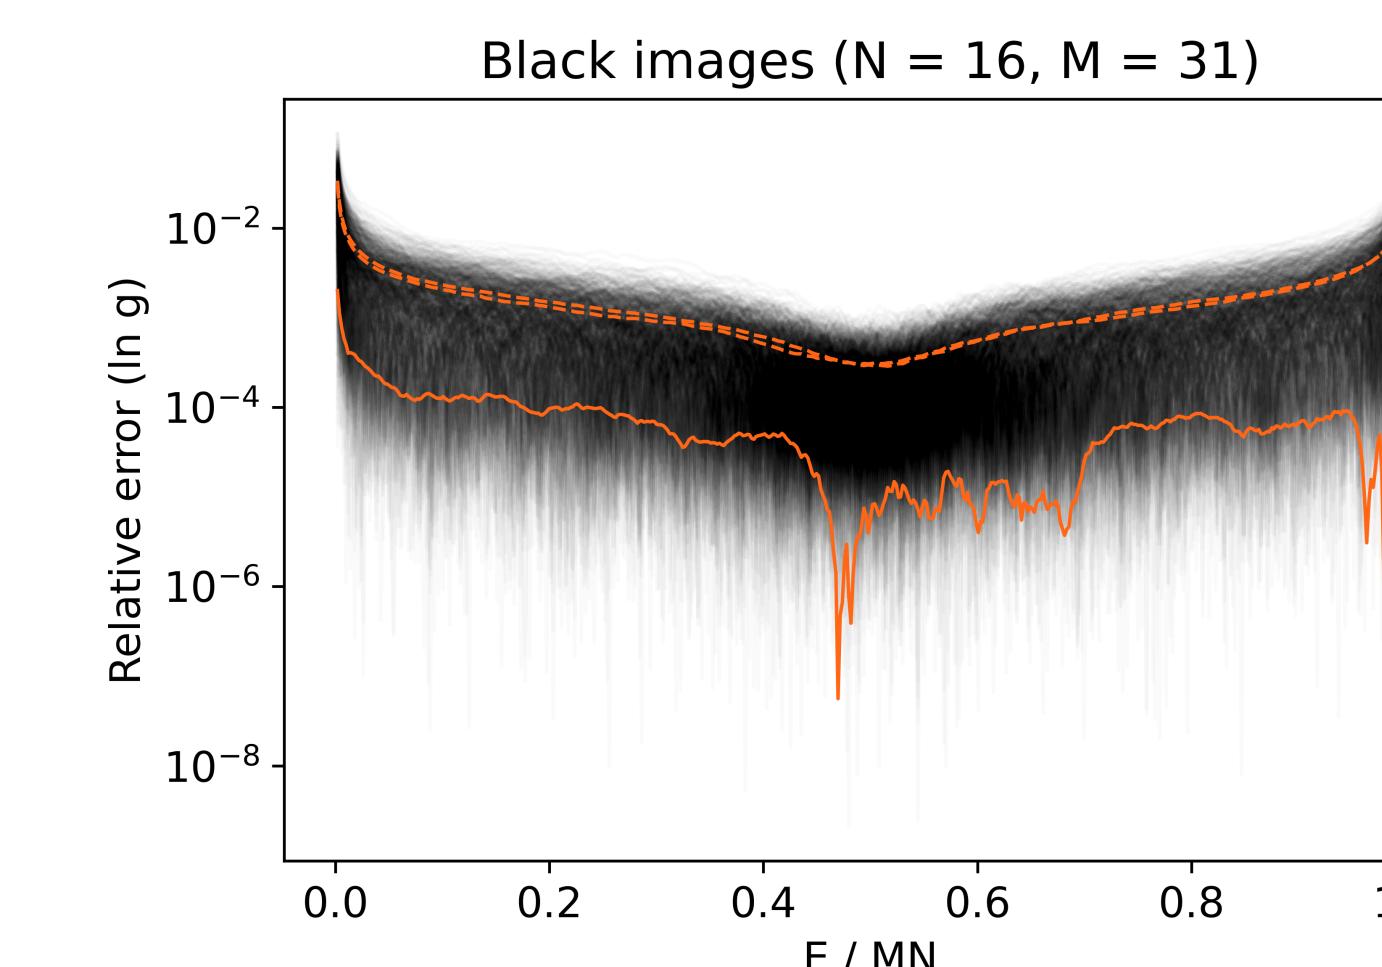


These methods are **widely used** to create graphics.

The electric sheep “fractal flame” screen saver is based on iterated function systems. Noise and turbulence are combined to procedurally generate textures for 3D models. L-systems are a component of *SpeedTree*, which is the software used to create vegetation for movies like *Avatar*.



**Receptive fields** of *simple cells* in area V1 are modeled by space-time derivatives, and are related to smoothed derivatives of the light field entering the eye.



For a notion of lost information in an image, we may perform simulations of **images as discrete-level systems**, where the **energy** is the sum of the differences in pixel values. The *Wang-Landau* algorithm was implemented for images, and confirmed to have the expected convergence properties on **black or white (BW)** images (*left*). Simulations for 1024 random grayscale images (*right*) reveal how entropy scales with how different the pixel values are (average energy).