

1 Natural image statistics

```
1 import numpy as np
2 import numpy.linalg as linalg
3 import matplotlib.pyplot as plt
4 from munkres import Munkres
5 from scipy import stats
6 from PIL import Image, ImageFilter, ImageOps
7 from src.utilities import *
8 from src.intensity_entropy import *
9 from src.kernels import *
10 plt.rcParams['image.cmap'] = 'gray'
11 plt.rcParams['figure.figsize'] = (12.8, 9.6)
```

1.1 The usual histograms

```
1 img = ImageOps.grayscale(Image.open('canyon.jpg'))
2 scale = max(np.shape(img))
3 data = np.array(img)
4 img
```



```
1 def contrast_renormalize(x):
2     mid = np.array(np.shape(x)) // 2
3     std = np.std(x)
4     return (x.item(*mid) - np.mean(x)) / std if std > 0 else 0
5
6 def log_contrast(x):
7     y = np.vectorize(lambda a: np.log(a) if a > 0 else 0)(x)
8     return y - np.mean(y)
9
10 np.array(np.shape(img)) // 5
```

```

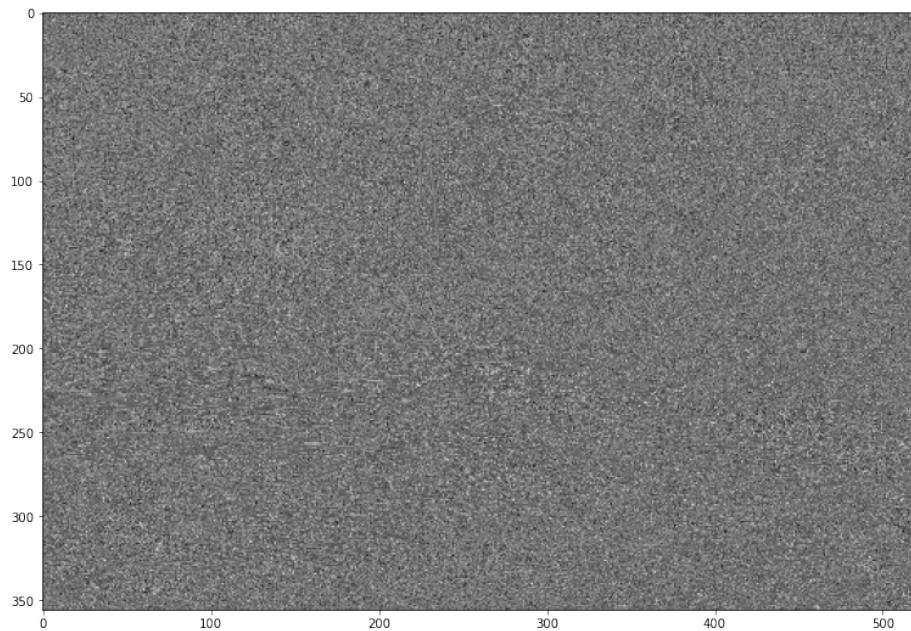
array([356, 518])

1 def renorm_blocks(n, f=contrast_renormalize):
2     return lambda x: mapblocks(*(np.array(np.shape(x)) // n), f, np.array(x))

1 def iterate(f, n):
2     return (lambda x: iterate(f, n-1)(f(x)) if n > 0 else x)

1 rdata = iterate(renorm_blocks(5), 1)(img)
2 plt.imshow(rdata);

```



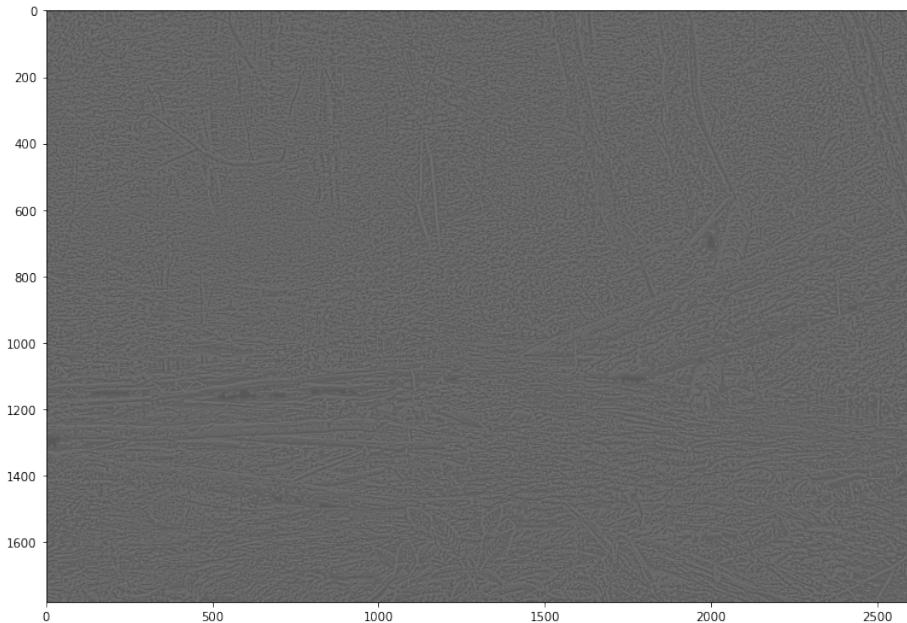
```

1 ldata = log_contrast(1.0*data)
2 plt.imshow(ldata);

```



```
1 l1data = iterate(lambda x: mapbox(5, contrast_renormalize, x), 1)(l1data)
2 l2data = iterate(lambda x: mapbox(5, contrast_renormalize, x), 2)(l1data)
3 plt.imshow(l1data);
```

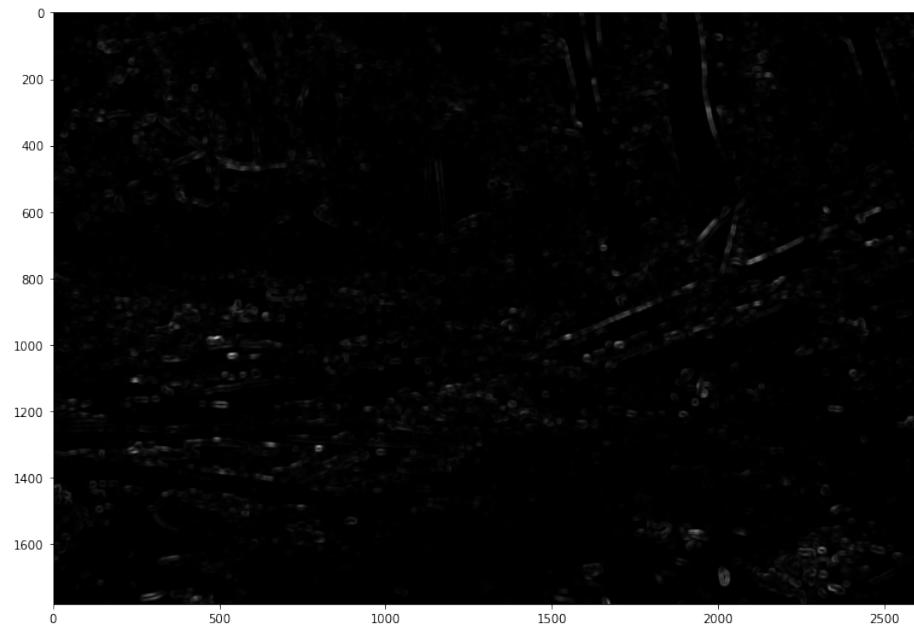


```
1 v1data = iterate(lambda x: mapbox(5, np.var, x), 1)(data)
2 v2data = iterate(lambda x: mapbox(5, np.var, x), 1)(v1data)
```

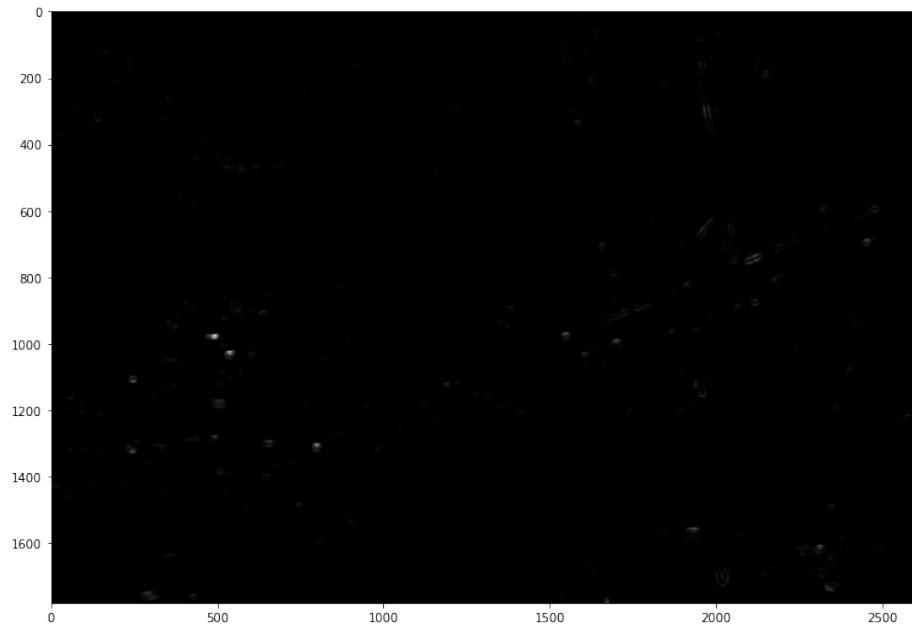
```
3 v3data = iterate(lambda x: mapbox(5, np.var, x), 1)(v2data)
4 plt.imshow(v1data);
```



```
1 plt.imshow(v2data);
```



```
1 plt.imshow(v3data);
```



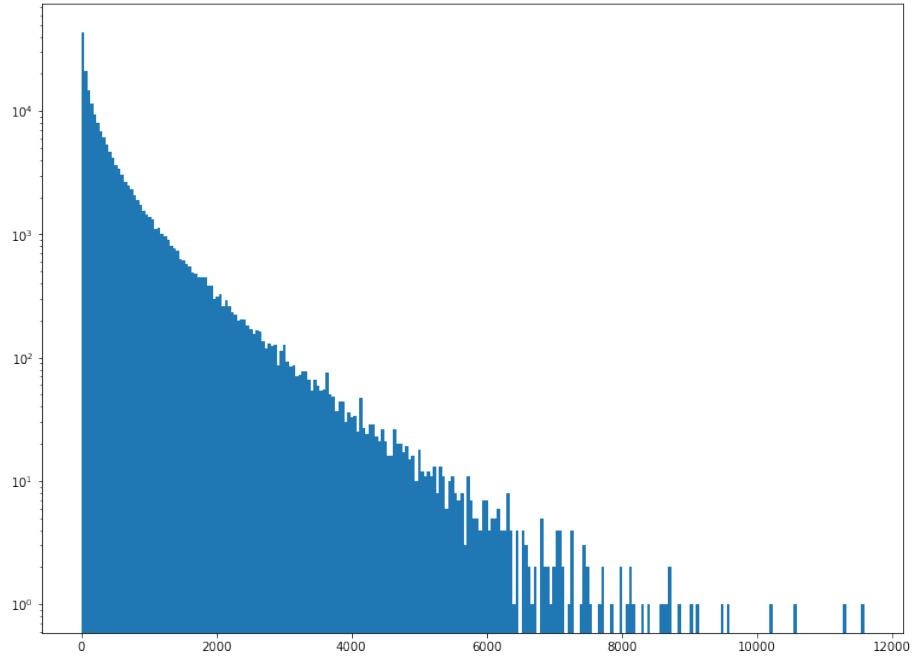
```
1 sep_vblocks = renorm_blocks(5, np.var)(data)
2 plt.imshow(sep_vblocks);
```



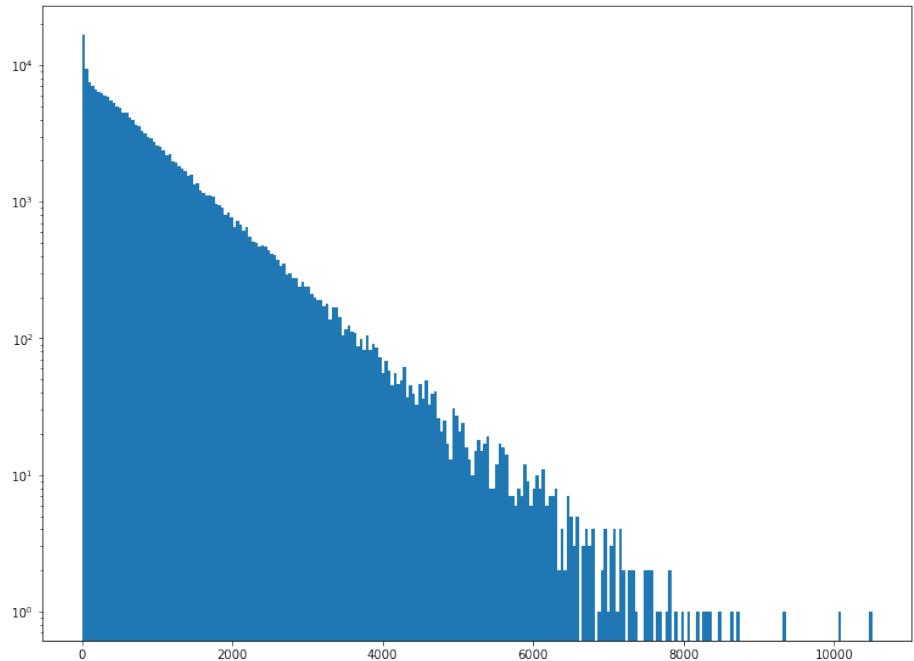
```
1 mean_vblocks = renorm_blocks(5, np.mean)(v1data)
2 plt.imshow(mean_vblocks);
```



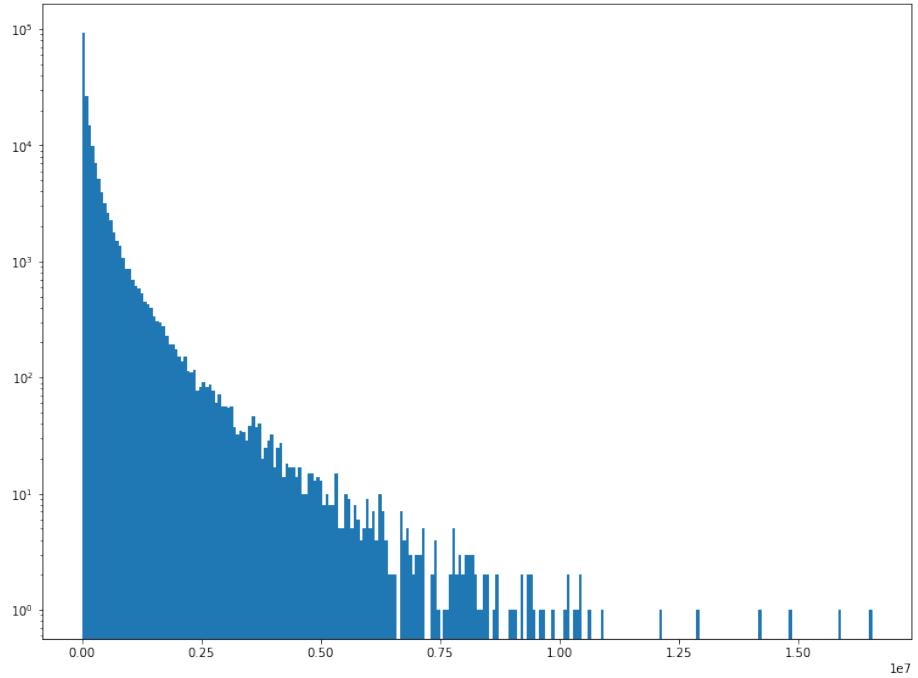
```
1 plt.hist(sep_vblocks.flat, 256)
2 plt.yscale('log');
```



```
1 plt.hist(mean_vblocks.flat, 256)
2 plt.yscale('log');
```



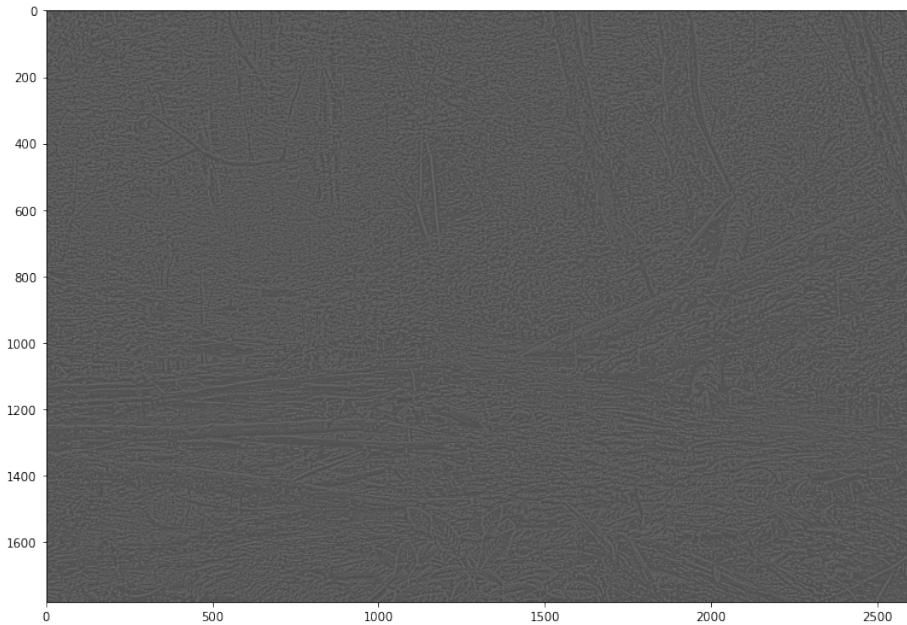
```
1 mean_v2blocks = renorm_blocks(5, np.mean)(v2data)
2 plt.hist(mean_v2blocks.flat, 256)
3 plt.yscale('log');
```



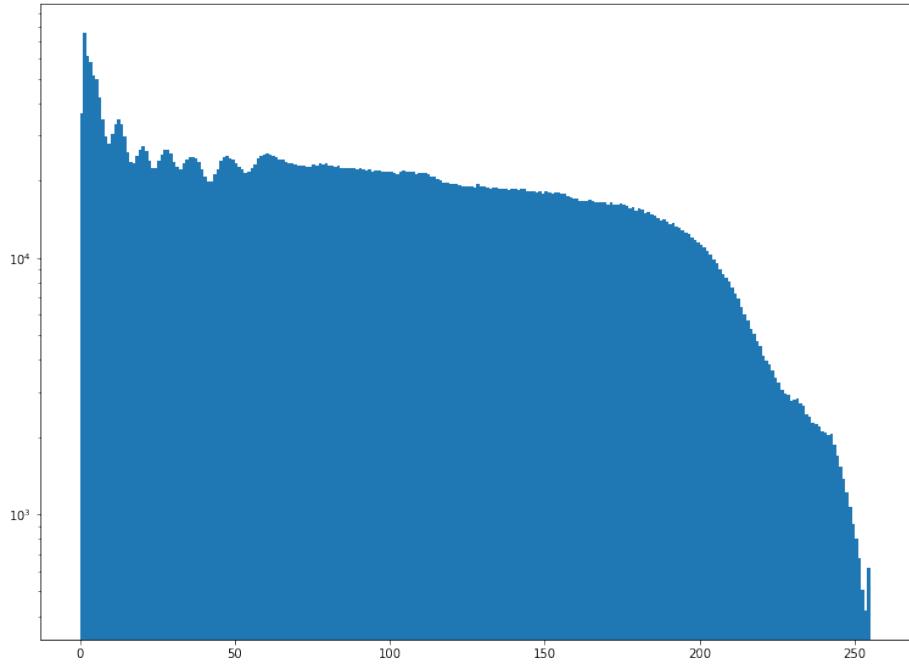
```
1 vdata = iterate(lambda x: mapbox(5, np.var, x), 1)(l1data)
2 plt.imshow(vdata);
```



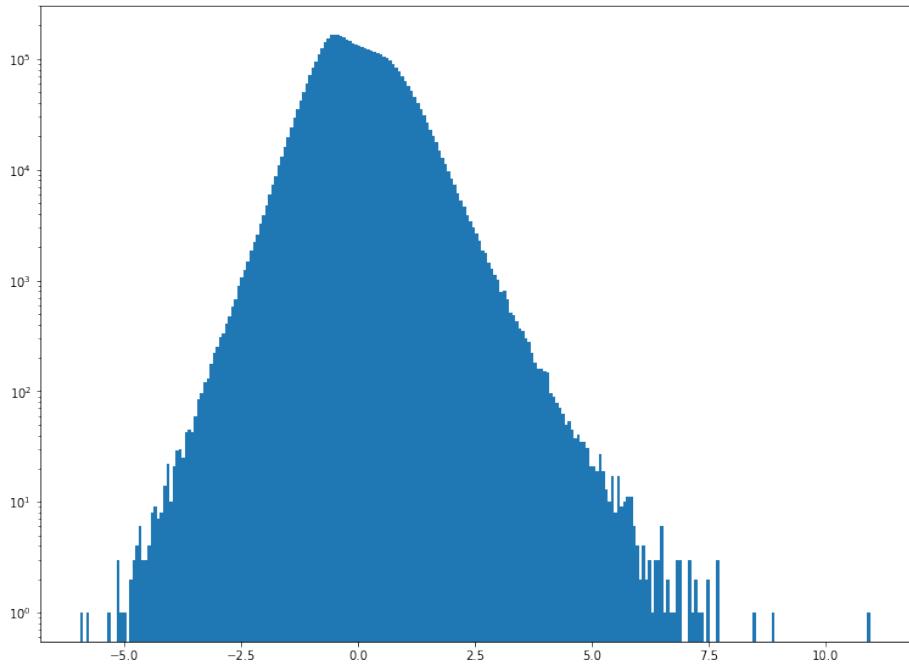
```
1 g1data = iterate(lambda x: mapbox(5, contrast_renormalize, x), 1)(np.array(img))
2 g2data = iterate(lambda x: mapbox(5, contrast_renormalize, x), 2)(np.array(img))
3 plt.imshow(g1data);
```



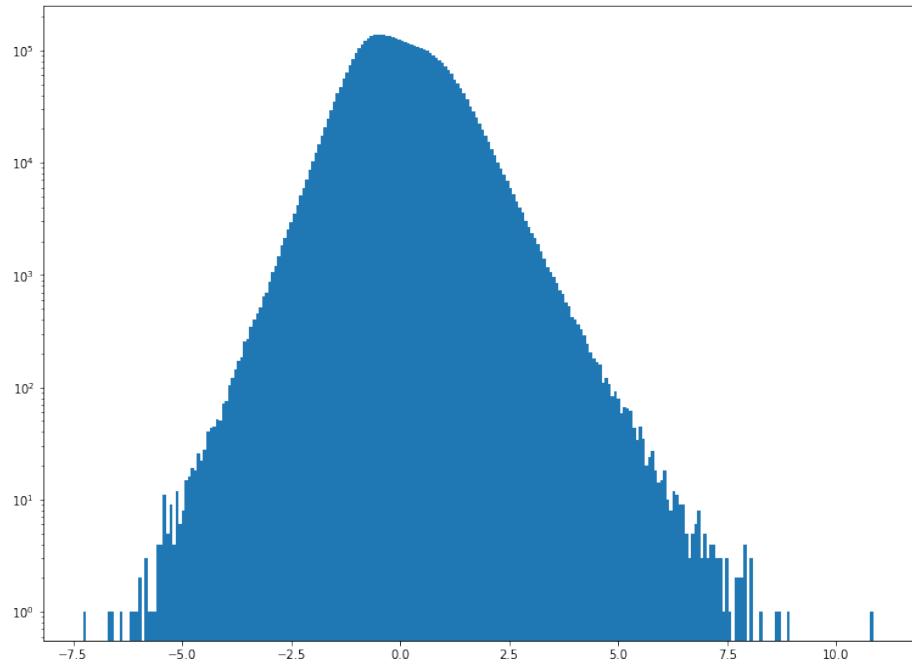
```
1 plt.hist(np.array(img).flat, range(256))
2 plt.yscale('log');
```



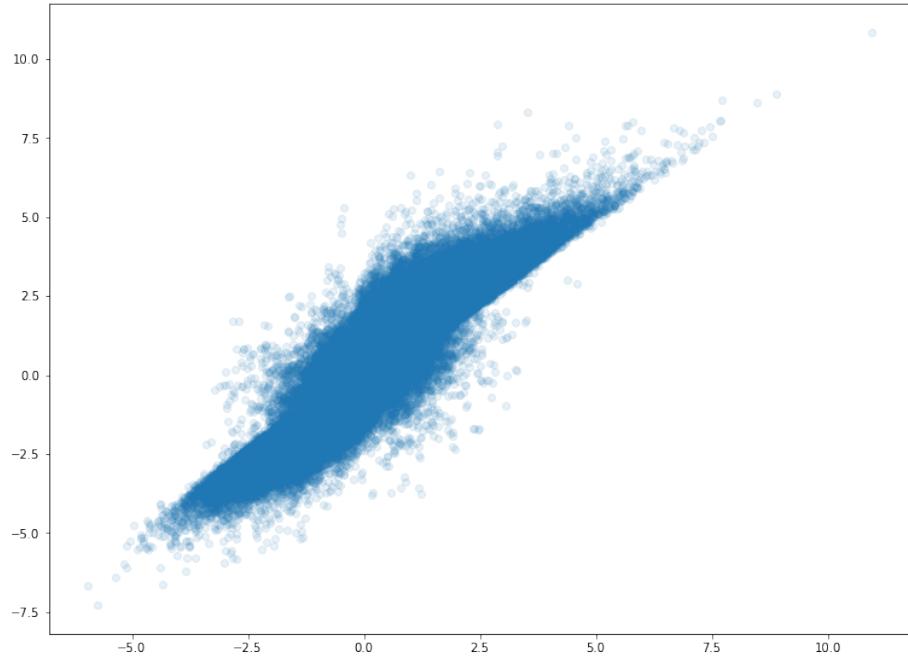
```
1 plt.hist(g1data.flat, 256)
2 plt.yscale('log');
```



```
1 plt.hist(g2data.flat, 256)
2 plt.yscale('log');
```



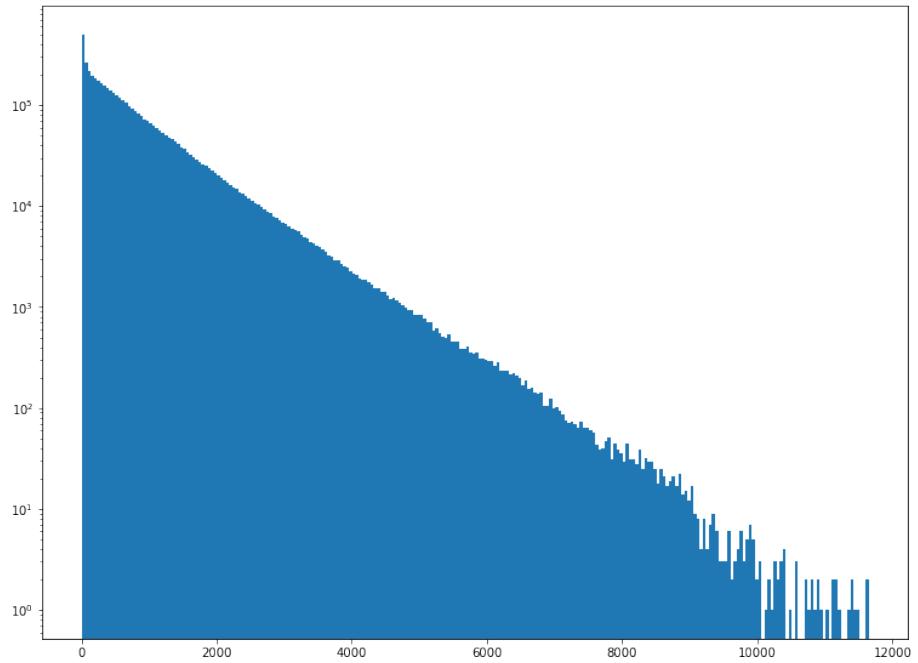
```
1 plt.scatter(g1data.flat, g2data.flat, alpha=0.1);
```



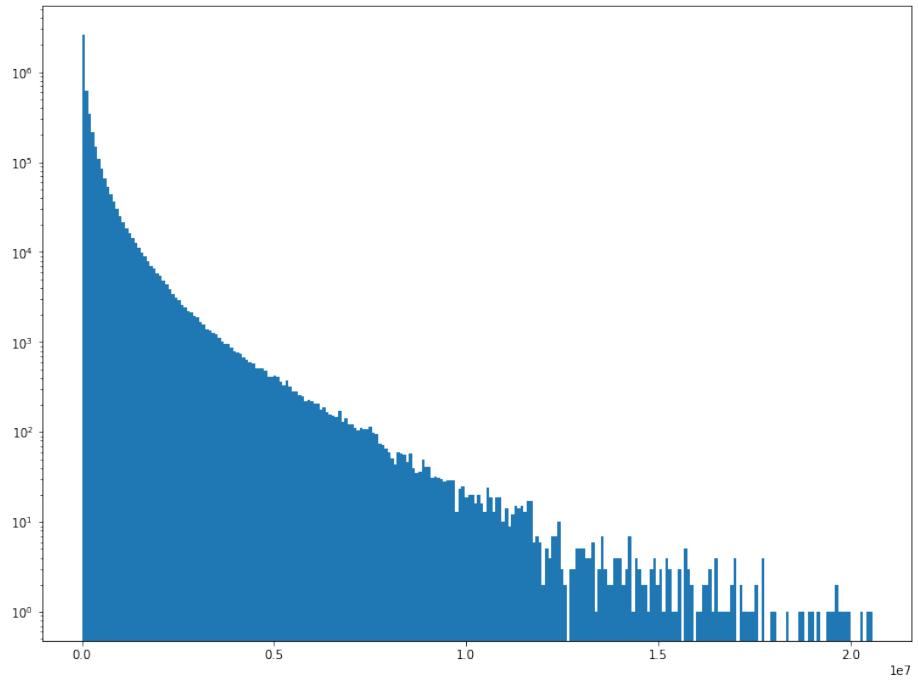
```
1 plt.imshow(vdata);
```



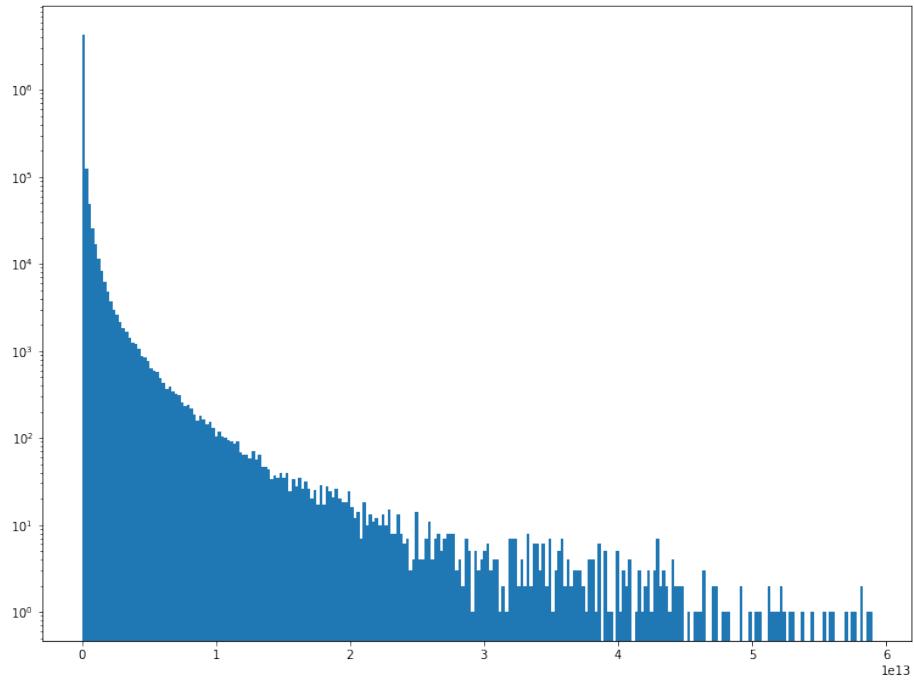
```
1 plt.hist(v1data.flat, 256)
2 plt.yscale('log');
```



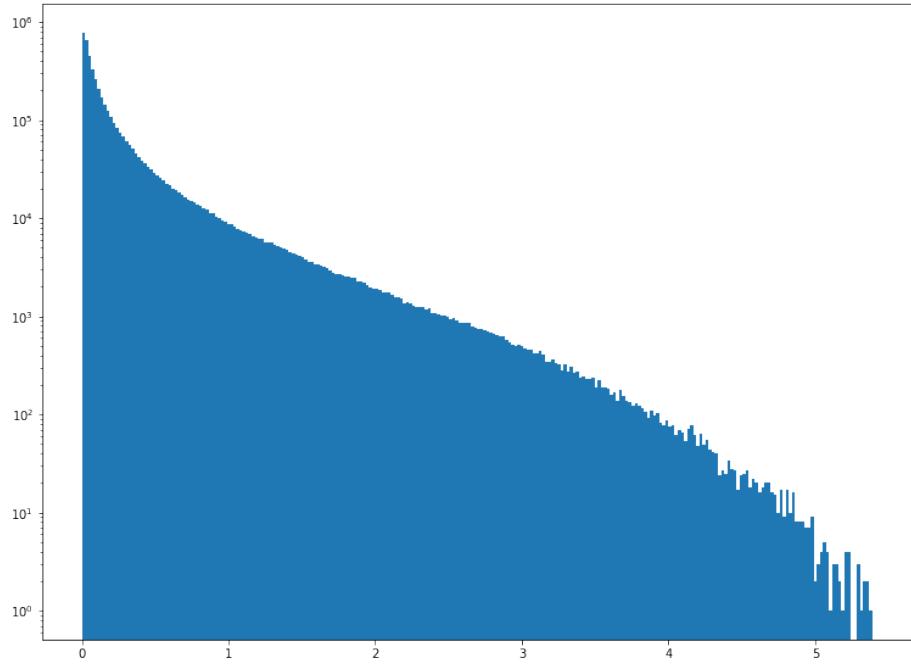
```
1 plt.hist(v2data.flat, 256)
2 plt.yscale('log');
```



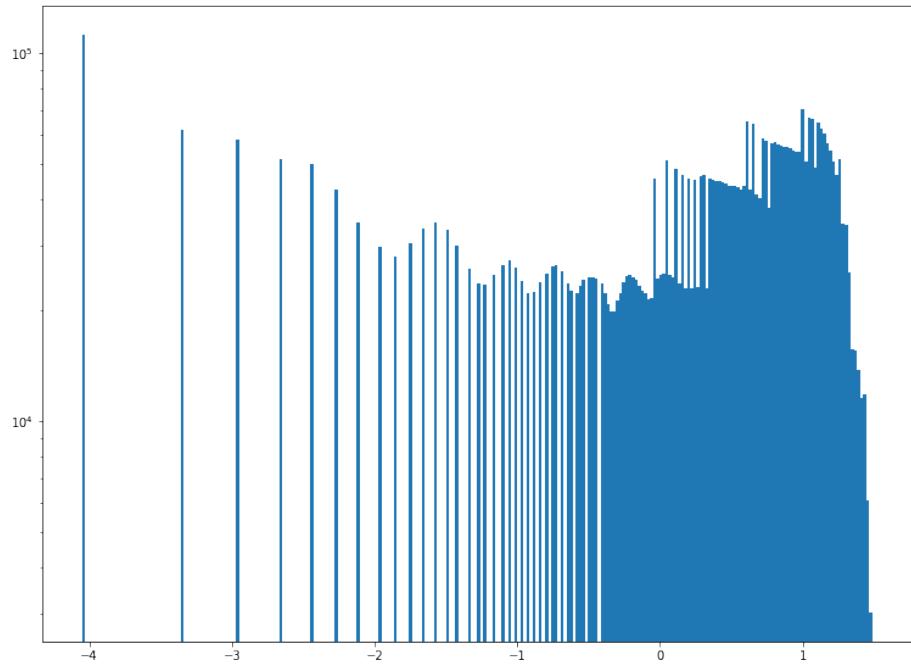
```
1 plt.hist(v3data.flat, 256)
2 plt.yscale('log');
```



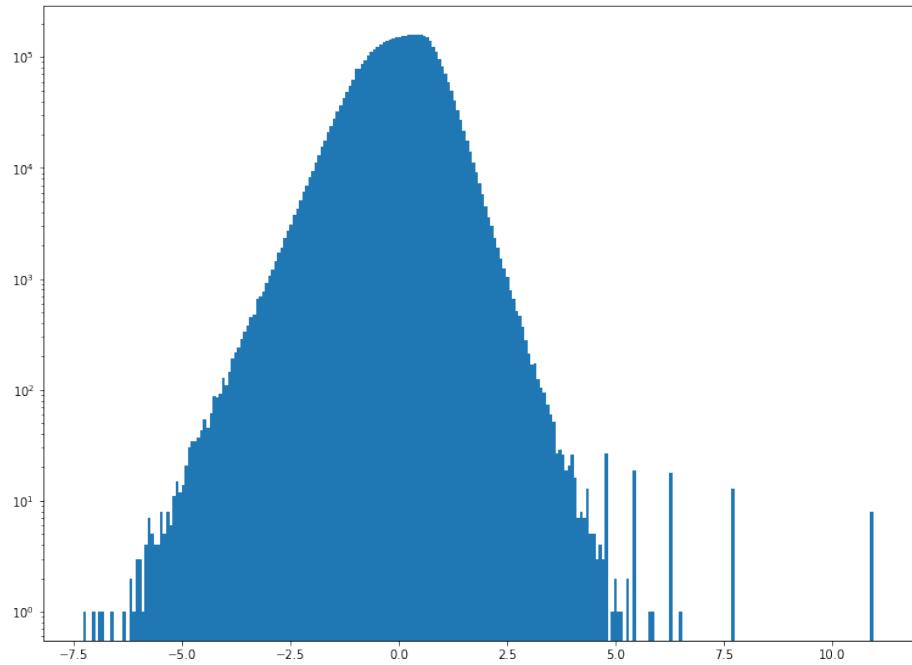
```
1 plt.hist(vdata.flat, 256)
2 plt.yscale('log');
```



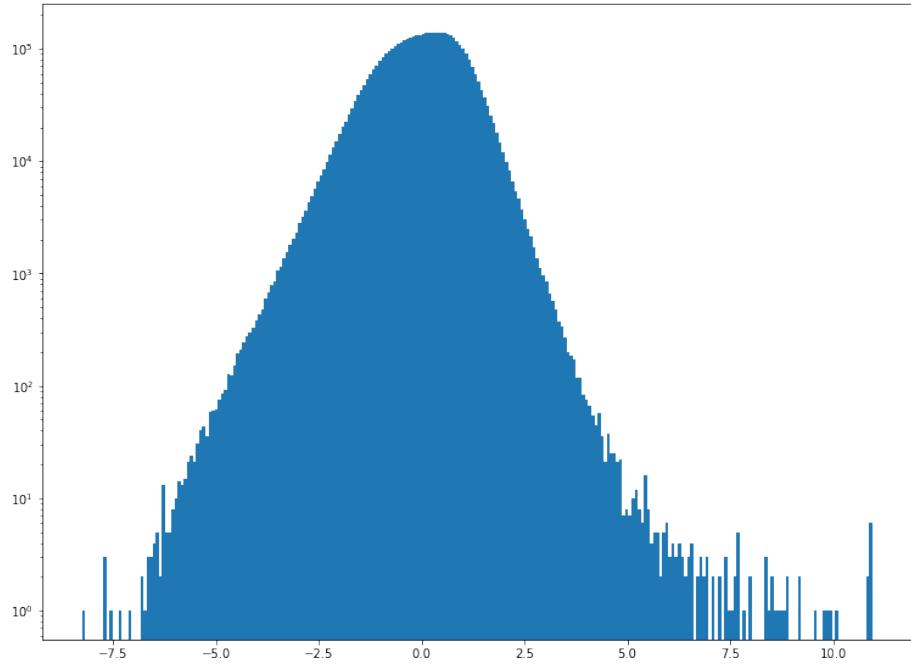
```
1 plt.hist(ldata.flat, 256)
2 plt.yscale('log');
```



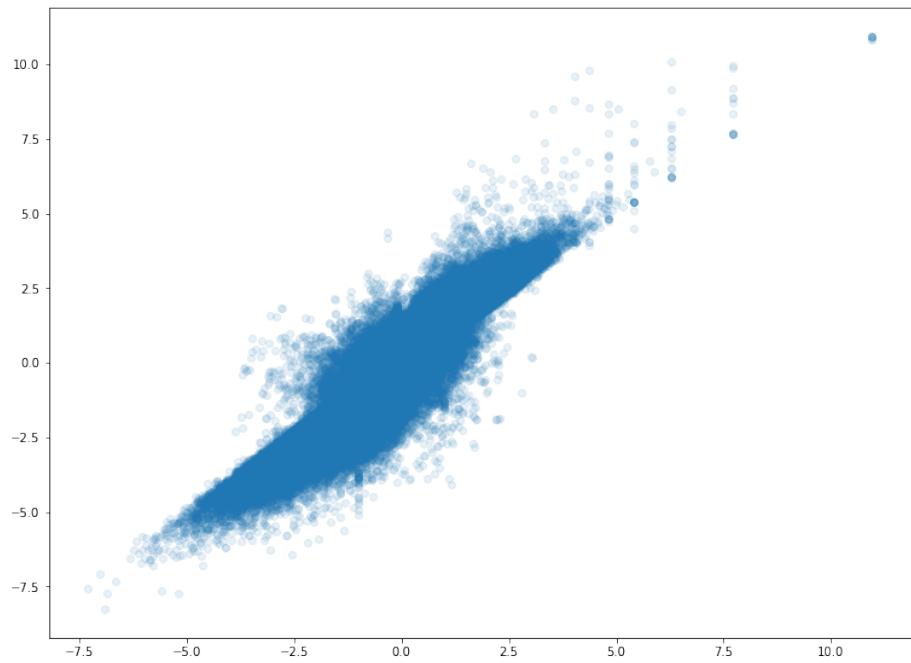
```
1 plt.hist(l1data.flat, 256)
2 plt.yscale('log');
```



```
1 plt.hist(l2data.flat, 256)
2 plt.yscale('log');
```

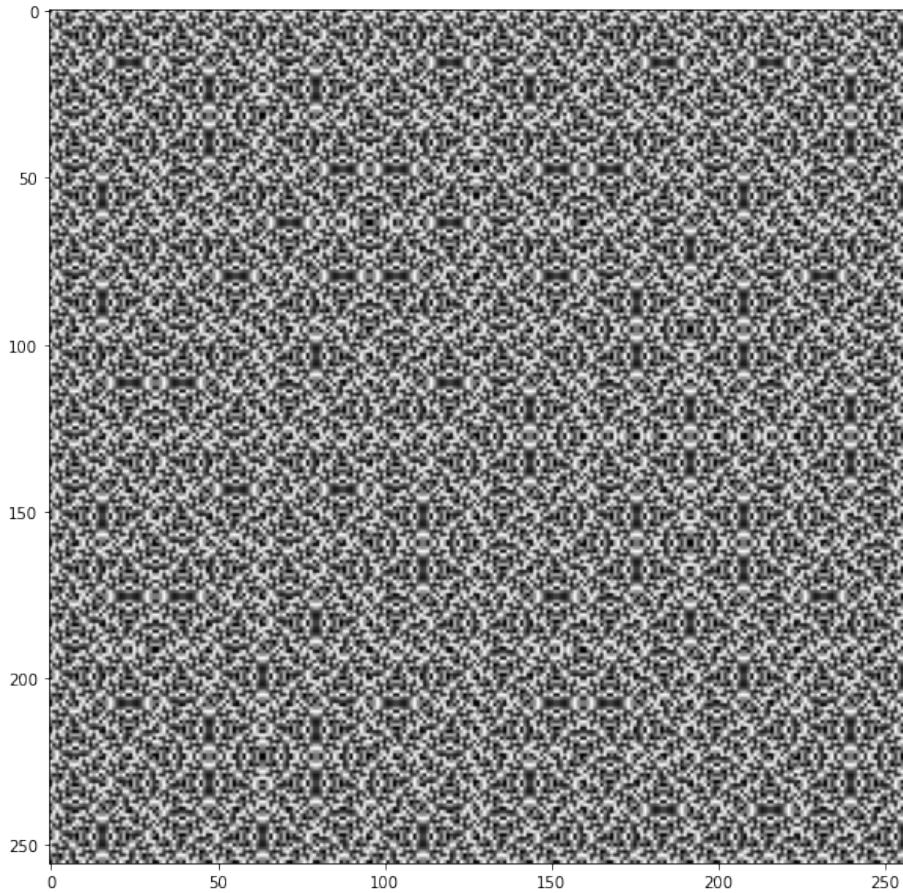


```
plt.scatter(l1data.flat, l2data.flat, alpha=0.1);
```



1.2 Fractal textures

```
1 x = 255 * np.random.rand(2, 2)
2 rep2 = lambda x: np.reshape(np.repeat([x], 4), (2, 2))
3 np.block([[rep2(b) for b in a] for a in x])
array([[ 89.75068064,  89.75068064, 203.5007879 , 203.5007879 ],
       [ 89.75068064,  89.75068064, 203.5007879 , 203.5007879 ],
       [131.58072491, 131.58072491, 76.29484712, 76.29484712],
       [131.58072491, 131.58072491, 76.29484712, 76.29484712]])
1 def grow(f):
2     return lambda x: np.block([[f(b) for b in a] for a in x])
3 def agrow(f):
4     return lambda x: np.block(f(x))
5 def fblock(x):
6     return [[x, np.transpose(x)], [np.transpose(x), np.rot90(x)]]
7 plt.imshow(iterate(agrow(fblock), 5)(np.random.rand(8, 8)));
```



```

1   from PIL import Image
2
3   img = Image.fromarray((255*np.random.rand(2, 2)).astype('uint8'))
4
5   Tinv = np.linalg.inv(np.array([
6       [1, 0, 1],
7       [0, 1, 1],
8       [0, 0, 1]
9   ]))
10  img.transform((2, 2), Image.AFFINE, data=Tinv.flatten()[:6], resample=Image.NEAREST)

```

1.3 Probabalistic inverse neighborhood reductions

Given a list of values y_i for $1 \leq i \leq n$ and a function f on any size list of values, we want to determine a new list X_i of size m for each i so that $f(X_i) = y_i$ and so that all of the generated values $x \in X = \cup_i X_i$ are distributed according to a given $p(x)$. This cannot be done exactly, so we must choose whether to prefer correct y_i values or a correct distribution of X .

1.3.1 Distribution-focused algorithm

It is simplest to sample correctly distributed values and approximate the y_i .

Aside: What value should we add to a set so that it has a specified variance?

```

1  def var_next(var, x, ddof=0):
2      n = len(x) + 1
3      if n <= ddof:
4          return np.nan
5      if n == 1:
6          return 0
7      xmean = np.mean(x)
8      xvarsum = (n - 1) * np.var(x, ddof=0)
9      vardiff = (n-ddof) * var - xvarsum
10     if vardiff > 0:
11         side = -1 if np.random.random_sample() - 1/2 < 0 else 1
12         return xmean + side * np.sqrt(vardiff * n / (n-1))
13     else:
14         return xmean
15
16     x = list(np.random.rand(10))
17     for _ in range(10):
18         x.append(var_next(1, x, ddof=1))
19     np.var(x, ddof=1)

```

1.0

One way to invert a neighborhood reduction function is to sample n values and assign them to the bins in the best way possible. From considering the matrix of all errors in the y -values, we see that this greedy algorithm is a case of the assignment problem, which we may solve with the Kuhn-Munkres algorithm.

```

1 def inverse_nreduce(y, m, f, sampler):
2     n = len(y)
3     xs = -np.ones((n, m))
4     for s in range(m):
5         u = sampler(n)
6         v = [[(y[i] - f(x + [x0]))**2 for x0 in u] for (i, x) in enumerate(xs)]
7         assignments = Munkres().compute(v)
8         for (i, j) in assignments:
9             xs[i, s] = u[j]
10    return xs

```

As an example, we sample from a Laplacian distribution with variance 2, but request sets with variance 1. The result is that the variances of the sets are imperfectly nudged away from 2 and towards 1.

```

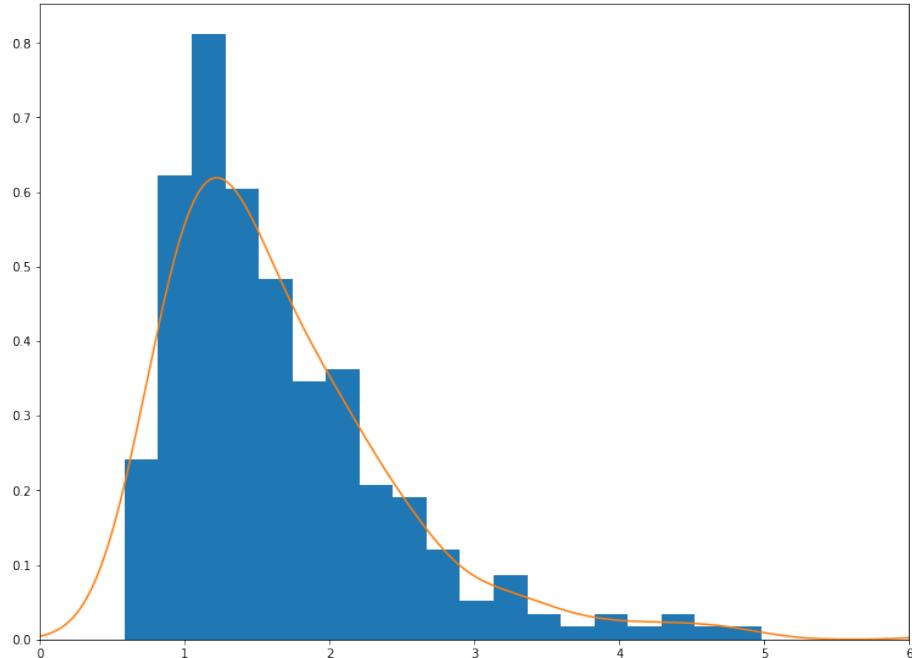
1 inr = inverse_nreduce([1.0 for _ in range(250)], 25, np.var, lambda n: stats.laplace.rvs(size=n))

1 np.var(np.concatenate(inr))

1.8887564926323521

1 nudged_vars = [np.var(r) for r in inr]
2 plt.hist(nudged_vars, 25, density=True)
3 plt.xlim(0, 6)
4 vs = np.linspace(0, 6, 300)
5 nudged_kde = stats.gaussian_kde(nudged_vars)
6 plt.plot(vs, nudged_kde(vs))
7 plt.show()

```

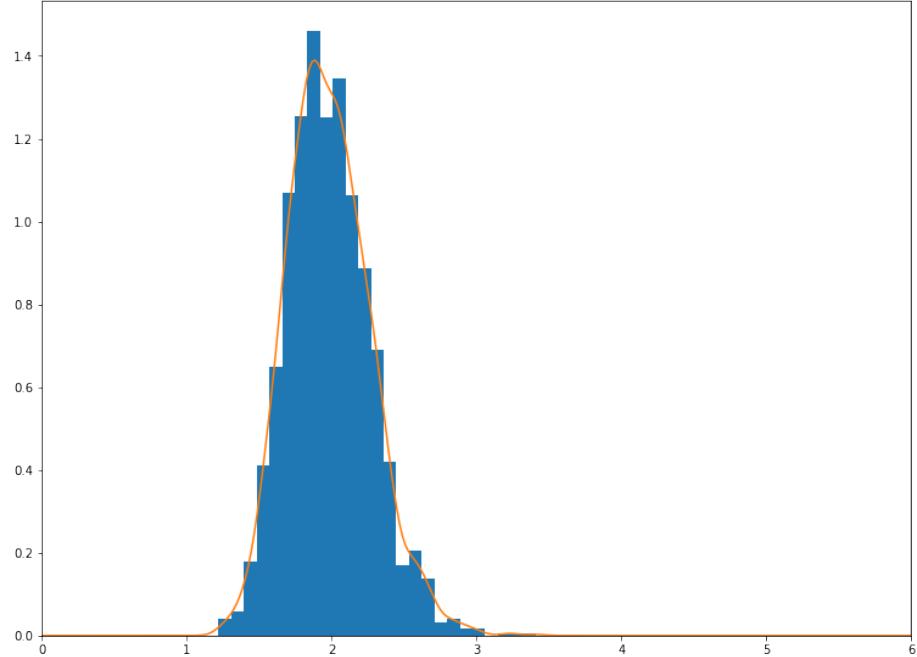


Compare the result of just sampling without redistributing values.

```

1 sample_vars = [np.var(stats.laplace.rvs(size=250)) for _ in range(2500)]
2 plt.hist(sample_vars, 25, density=True)
3 plt.xlim(0, 6)
4 vs = np.linspace(0, 6, 300)
5 sample_kde = stats.gaussian_kde(sample_vars)
6 plt.plot(vs, sample_kde(vs))
7 plt.show()

```



Now what are the distributions of the proposed natural scale-invariant variance images (as in Ruderman's statistics of natural images, doi: 10.1088/0954-898X_5_4_006)? First, we will try an exponential distribution of variances.

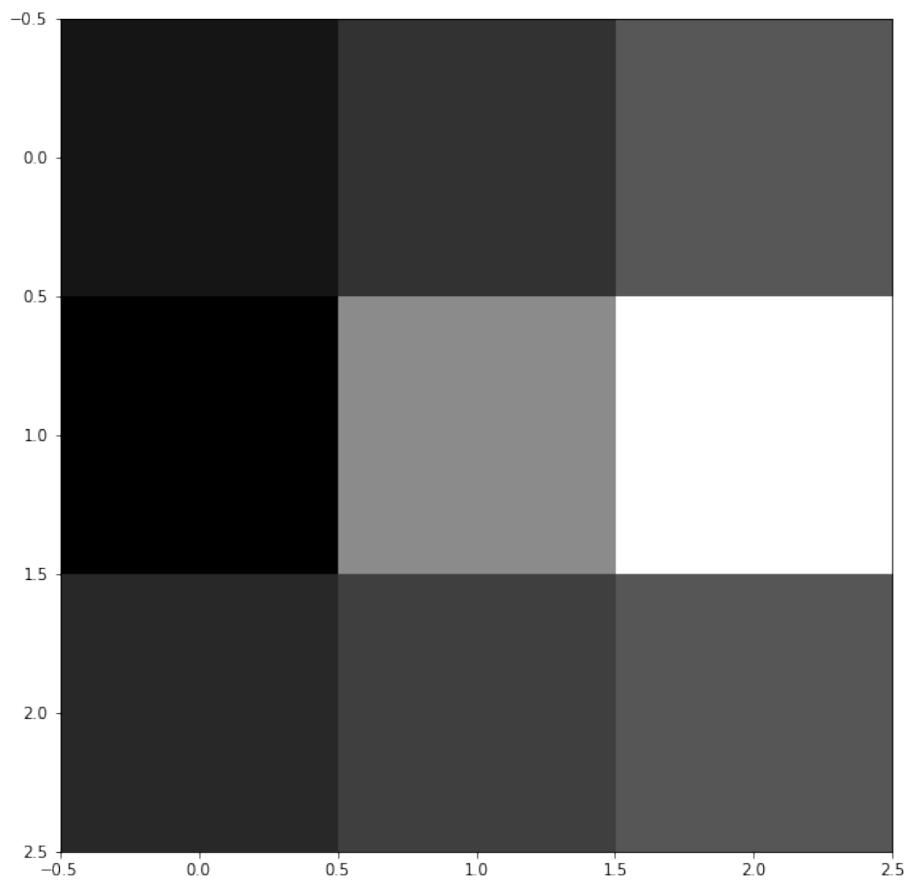
```

1 nθ = 3
2 mθ = 3
3 initial_vars = stats.expon.rvs(size=nθ*mθ)
4 inr_vars = inverse_nreduce(initial_vars, mθ*mθ, np.var, lambda n: (1 / 3) * stats.expon.rvs(size=n))

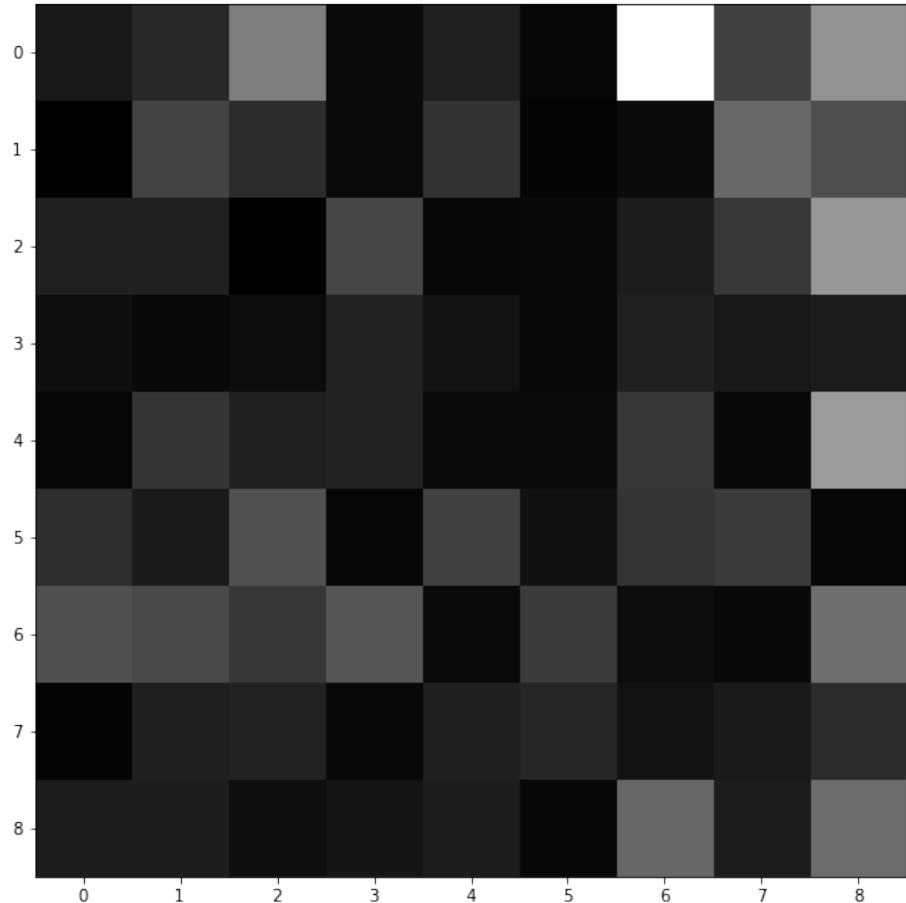
1 inr_mat = np.concatenate([np.concatenate(r, axis=1) for r in np.reshape(inr_vars, (nθ, nθ, mθ, mθ))], axis=0)
2 new_vars = inr_mat.flat

1 plt.imshow(np.reshape(initial_vars, (nθ, nθ)));

```



```
1 plt.imshow(inr_mat);
```

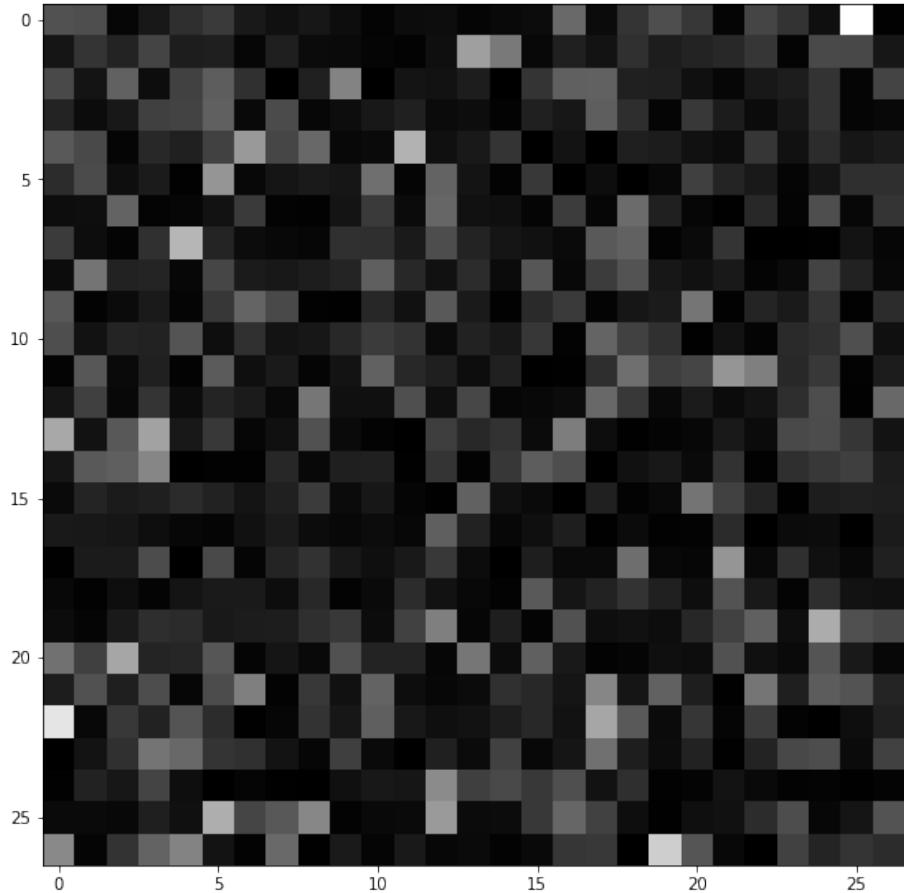


```

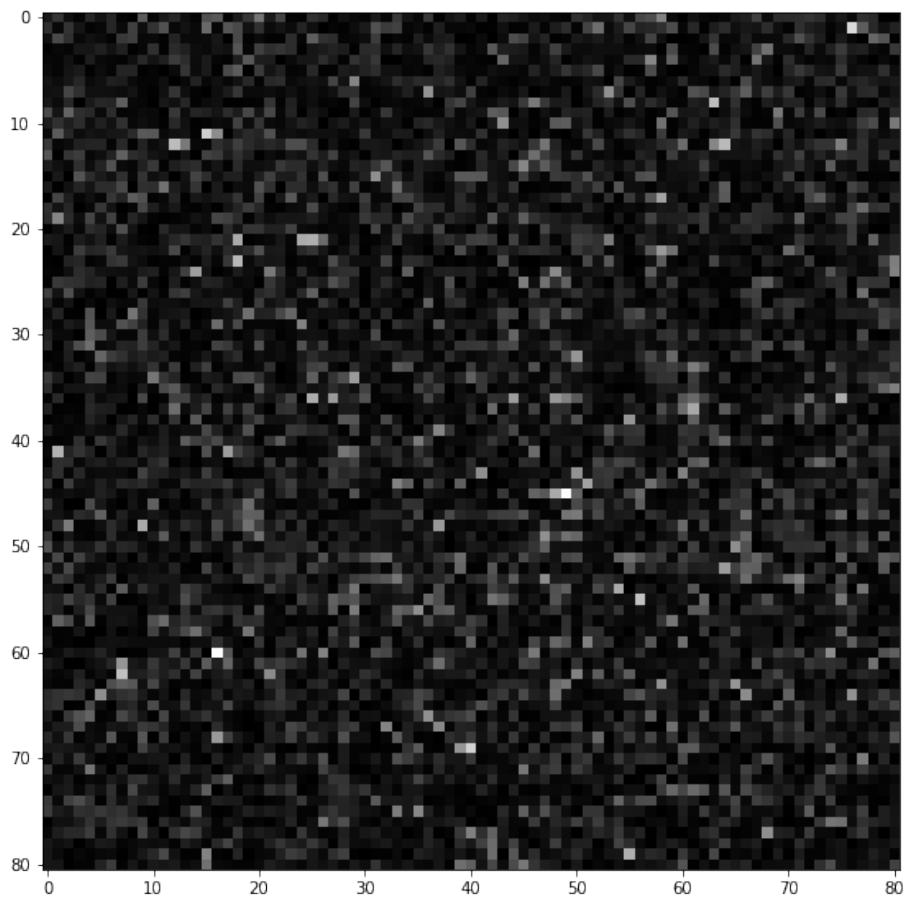
1 inr2_vars = inverse_nreduce(new_vars, m0*m0, np.var, lambda n: (1 / 3) * stats.expon.rvs(size=n))
2 n1 = int(np.sqrt(np.shape(inr2_vars)[0]))
3 inr2_mat = np.concatenate([np.concatenate(r, axis=1) for r in np.reshape(inr2_vars, (n1, n1, m0, m0))]),
4   ↪ axis=0)
5 new2_vars = inr2_mat.flat

6 plt.imshow(inr2_mat);

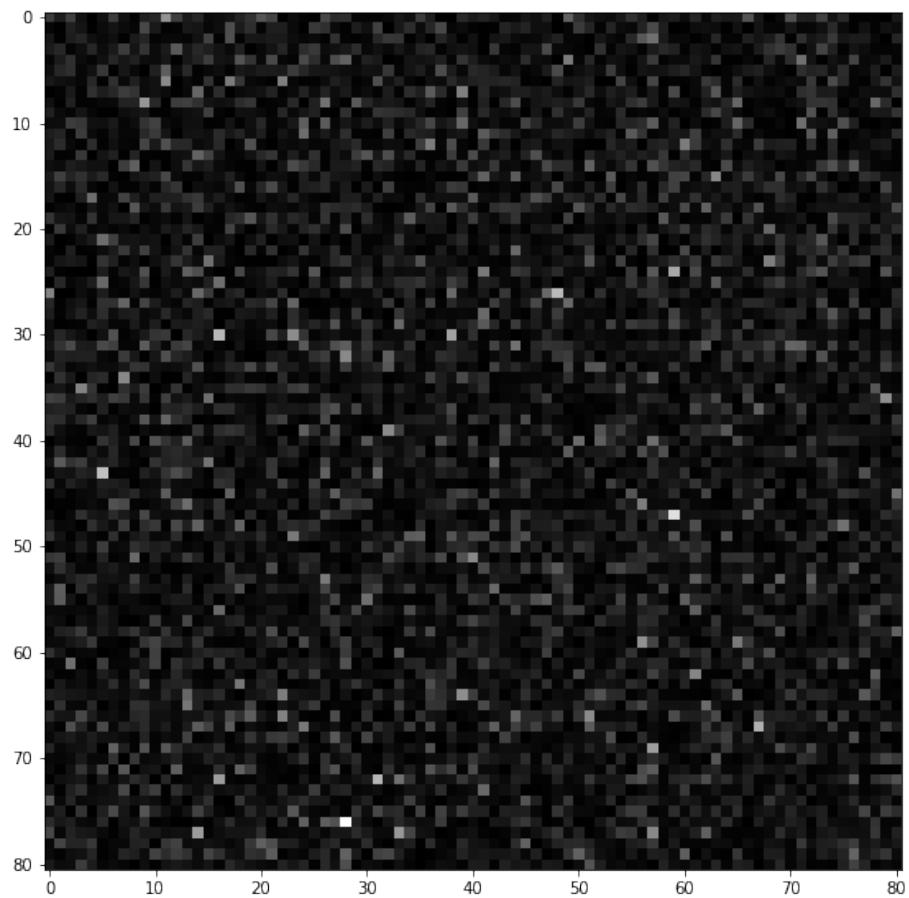
```



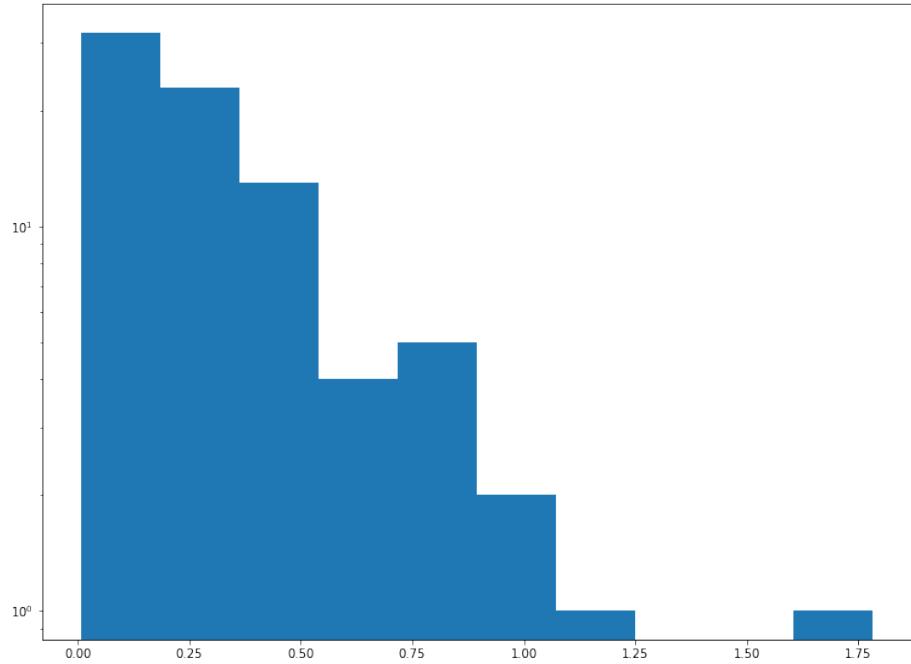
```
1 inr3_vars = inverse_nreduce(new2_vars, m0*m0, np.var, lambda n: (1 / 3) * stats.expon.rvs(size=n))
2 n2 = int(np.sqrt(np.shape(inr3_vars)[0]))
3 inr3_mat = np.concatenate([np.concatenate(r, axis=1) for r in np.reshape(inr3_vars, (n2, n2, m0, m0))],
   ↪ axis=0)
4
5 plt.imshow(inr3_mat);
```



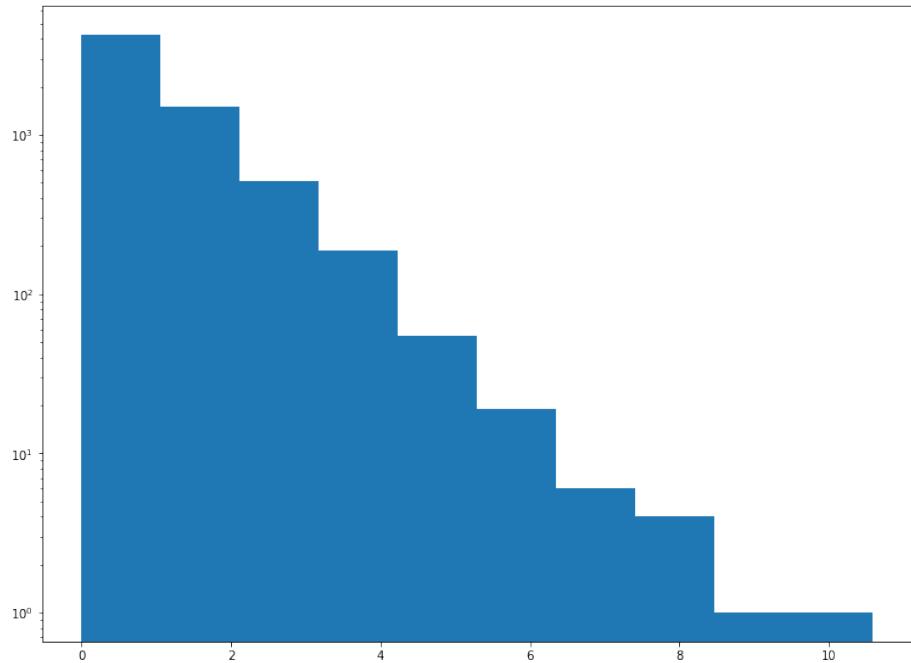
```
1 control_mat = np.reshape(stats.expon.rvs(size=np.size(inr3_mat)), np.shape(inr3_mat))  
1 plt.imshow(control_mat);
```



```
1 plt.hist(inr_mat.flat)
2 plt.yscale('log');
```

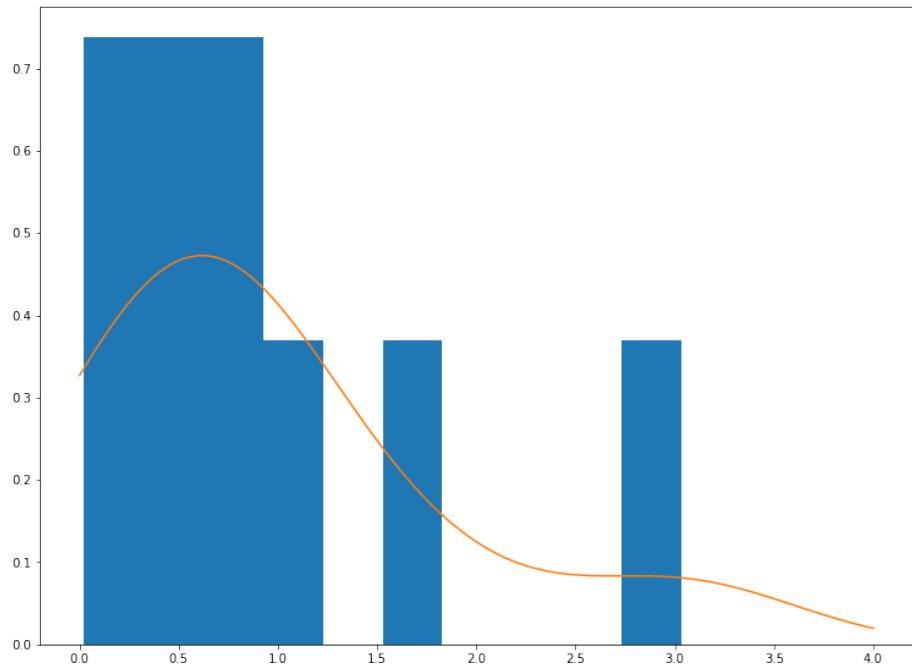


```
1 plt.hist(control_mat.flat)
2 plt.yscale('log');
```

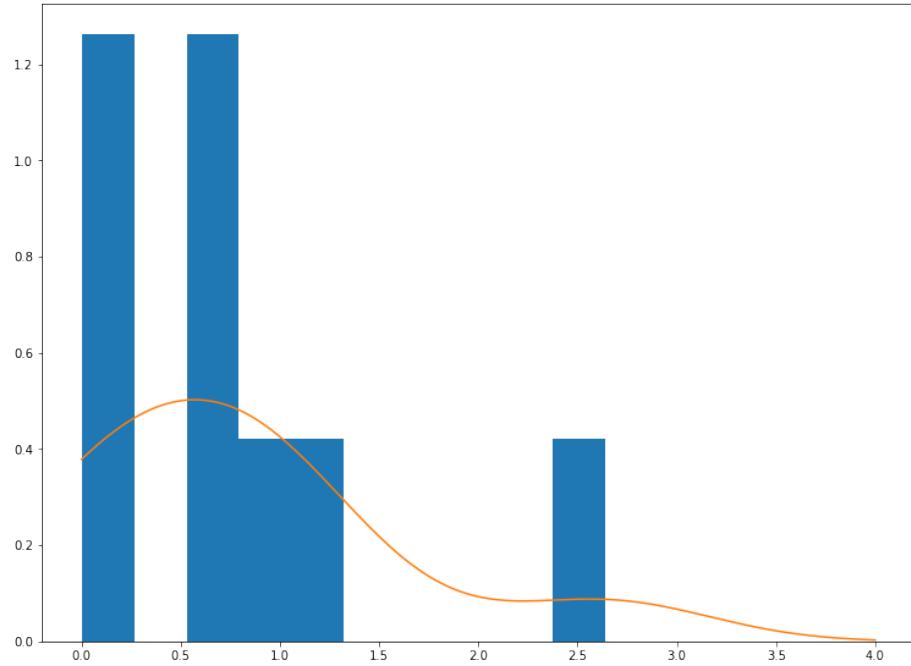


Histograms

```
1 inr_varerrs = [np.abs(initial_vars[i] - np.var(r)) for (i, r) in enumerate(initial_vars)]
2 plt.hist(inr_varerrs, density=True)
3 vs = np.linspace(0, 4, 300)
4 inr_kde = stats.gaussian_kde(inr_varerrs)
5 plt.plot(vs, inr_kde(vs))
6 plt.show()
```



```
1 inr_varerrs = [np.abs(initial_vars[i] - np.var(stats.expon.rvs(size=5*5))) for i in
2     range(len(initial_vars))]
3 plt.hist(inr_varerrs, density=True)
4 vs = np.linspace(0, 4, 300)
5 inr_kde = stats.gaussian_kde(inr_varerrs)
6 plt.plot(vs, inr_kde(vs))
7 plt.show()
```



1.4 Perlin noise

```

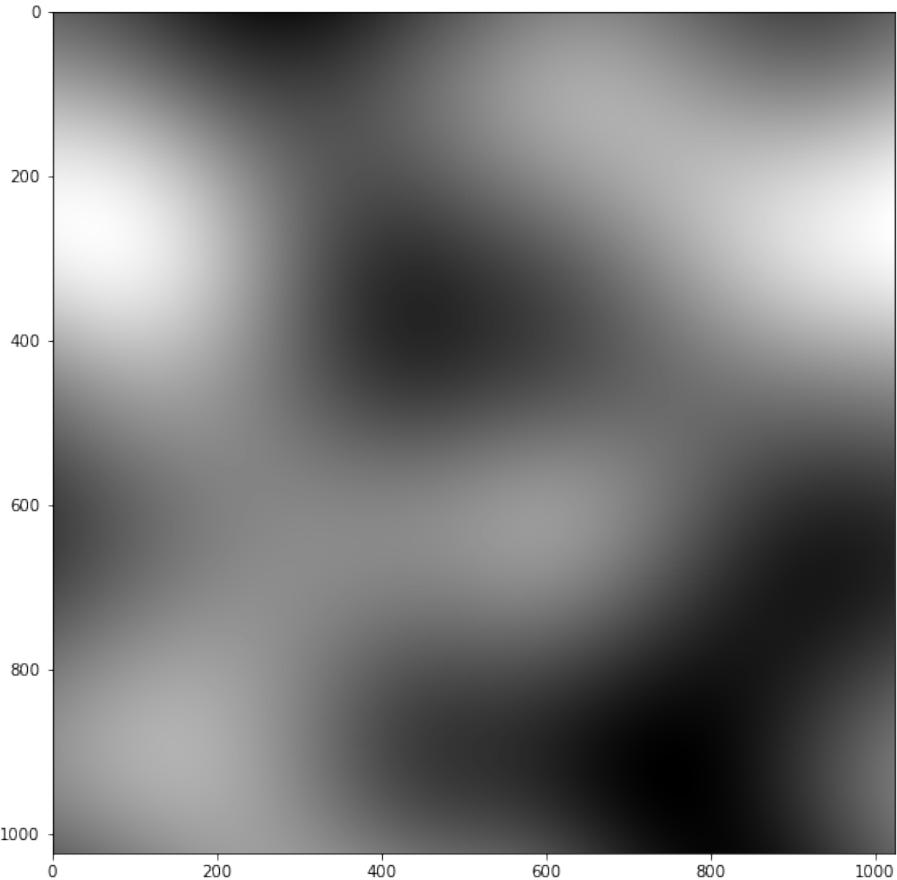
1 # From
2 ↵ https://github.com/pvigier/perlin-numpy/blob/8e3ea24a39e938f631f4101294dcda4ef92bc633/perlin2d.py#L3-L41
3 def generate_perlin_noise_2d(shape, res, tileable=(False, False)):
4     def f(t):
5         return 6*t**5 - 15*t**4 + 10*t**3
6
7     delta = (res[0] / shape[0], res[1] / shape[1])
8     d = (shape[0] // res[0], shape[1] // res[1])
9     grid = np.mgrid[0:res[0]:delta[0],0:res[1]:delta[1]].transpose(1, 2, 0) % 1
10    # Gradients
11    angles = 2*np.pi*np.random.rand(res[0]+1, res[1]+1)
12    gradients = np.dstack((np.cos(angles), np.sin(angles)))
13    if tileable[0]:
14        gradients[-1,:,:] = gradients[0,:,:]
15    if tileable[1]:
16        gradients[:, -1] = gradients[:, 0]
17    gradients = gradients.repeat(d[0], 0).repeat(d[1], 1)
18    g00 = gradients[      :-d[0],      :-d[1]]
19    g10 = gradients[d[0]:,      :-d[1]]
20    g01 = gradients[      :-d[0],d[1]:]
21    g11 = gradients[d[0]:,d[1]:]
22    # Ramps
23    n00 = np.sum(np.dstack((grid[:, :, 0], grid[:, :, 1])) * g00, 2)
24    n10 = np.sum(np.dstack((grid[:, :, 0]-1, grid[:, :, 1])) * g10, 2)
25    n01 = np.sum(np.dstack((grid[:, :, 0], grid[:, :, 1]-1)) * g01, 2)
26    n11 = np.sum(np.dstack((grid[:, :, 0]-1, grid[:, :, 1]-1)) * g11, 2)
27    # Interpolation

```

```

27     t = f(grid)
28     n0 = n00*(1-t[:, :, 0]) + t[:, :, 0]*n10
29     n1 = n01*(1-t[:, :, 0]) + t[:, :, 0]*n11
30     return np.sqrt(2)*((1-t[:, :, 1])*n0 + t[:, :, 1])*n1)
31
32 def generate_fractal_noise_2d(shape, res, octaves=1, persistence=0.5, lacunarity=2, tileable=(False,
33     ↪ False)):
34     noise = np.zeros(shape)
35     frequency = 1
36     amplitude = 1
37     for _ in range(octaves):
38         noise += amplitude * generate_perlin_noise_2d(shape, (frequency*res[0], frequency*res[1]),
39             ↪ tileable)
40         frequency *= lacunarity
41         amplitude *= persistence
42     return noise
43
44 pn = generate_perlin_noise_2d((1024, 1024), (2, 2))
45 plt.imshow(pn, interpolation='lanczos');

```

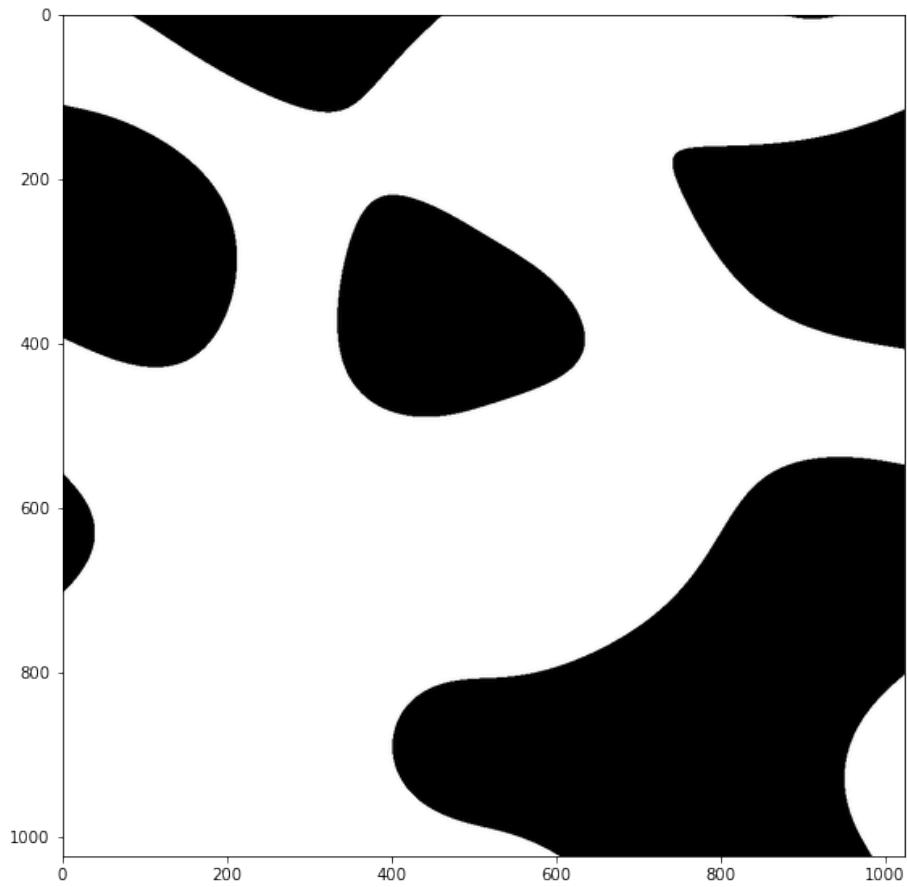


```

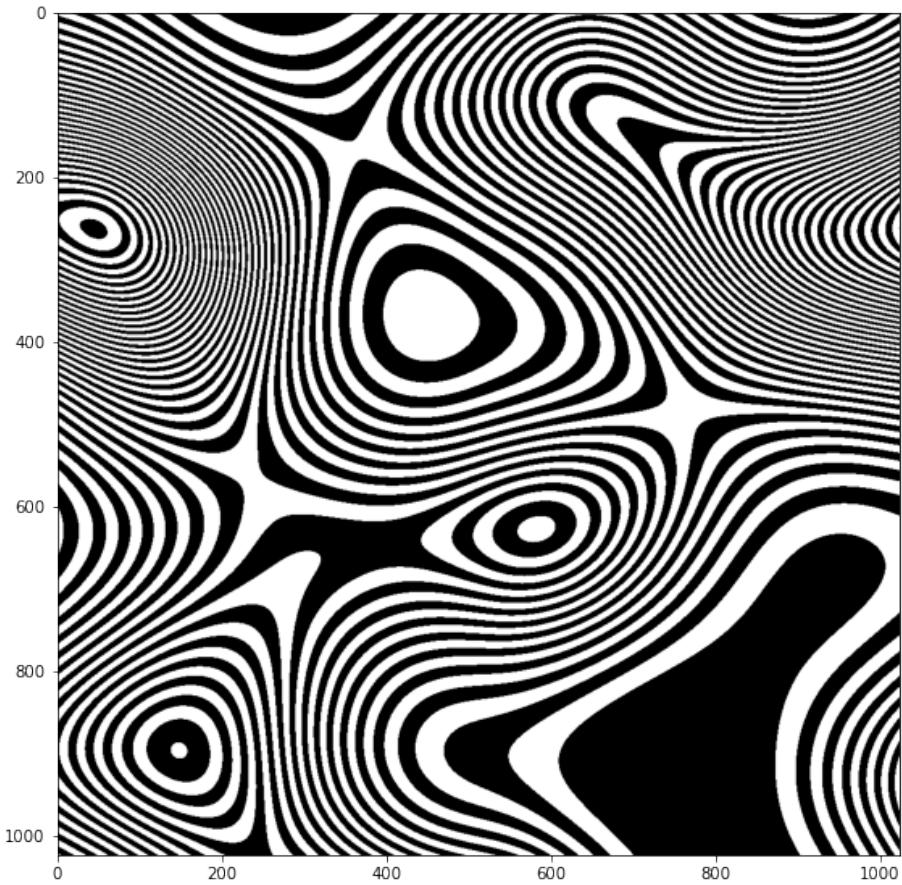
1 pmin, pmax = np.min(pn), np.max(pn)

```

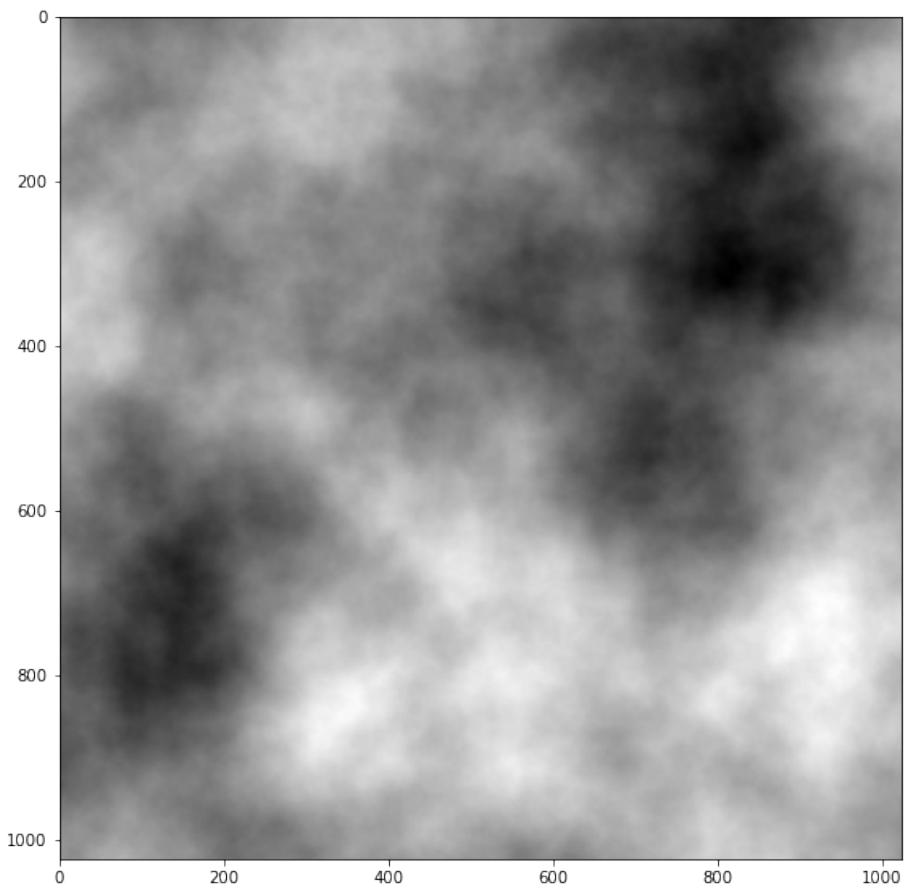
```
2 pn = (pn - pmin) / (pmax - pmin) # Rescale to exactly [0, 1] for thresholding  
3 plt.imshow((0.3 < pn) & (pn < 0.7), interpolation='lanczos');
```



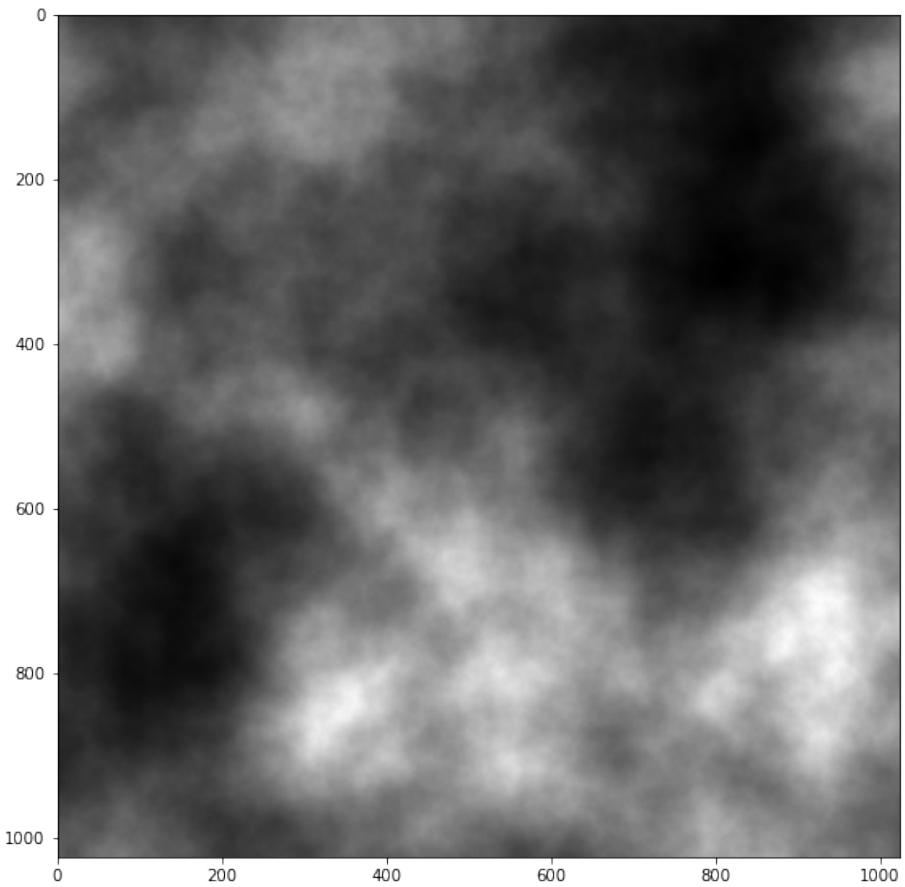
```
1 plt.imshow(np.floor((pn**2)*64) % 2);
```



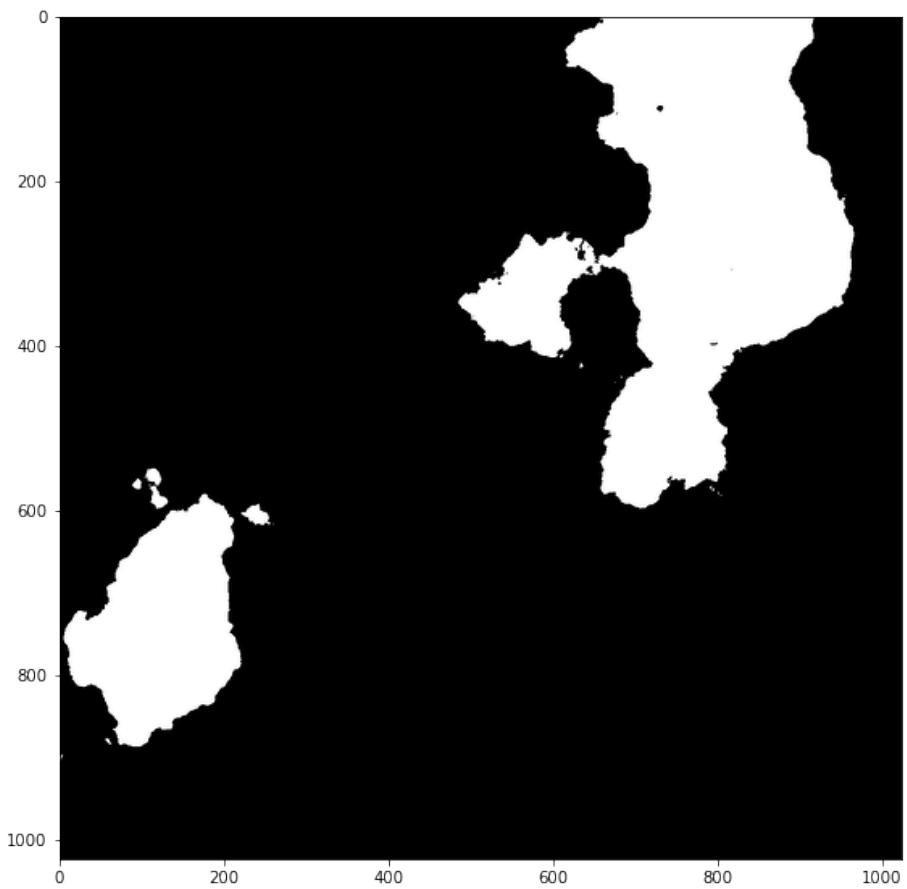
```
1 pfn = generate_fractal_noise_2d((1024, 1024), (2, 2), octaves=10, persistence=0.5)
2 pfn = 0.5 + (pfn / 4) # Rescale to within [0, 1]
3 plt.imshow(pfn, interpolation='lanczos');
```



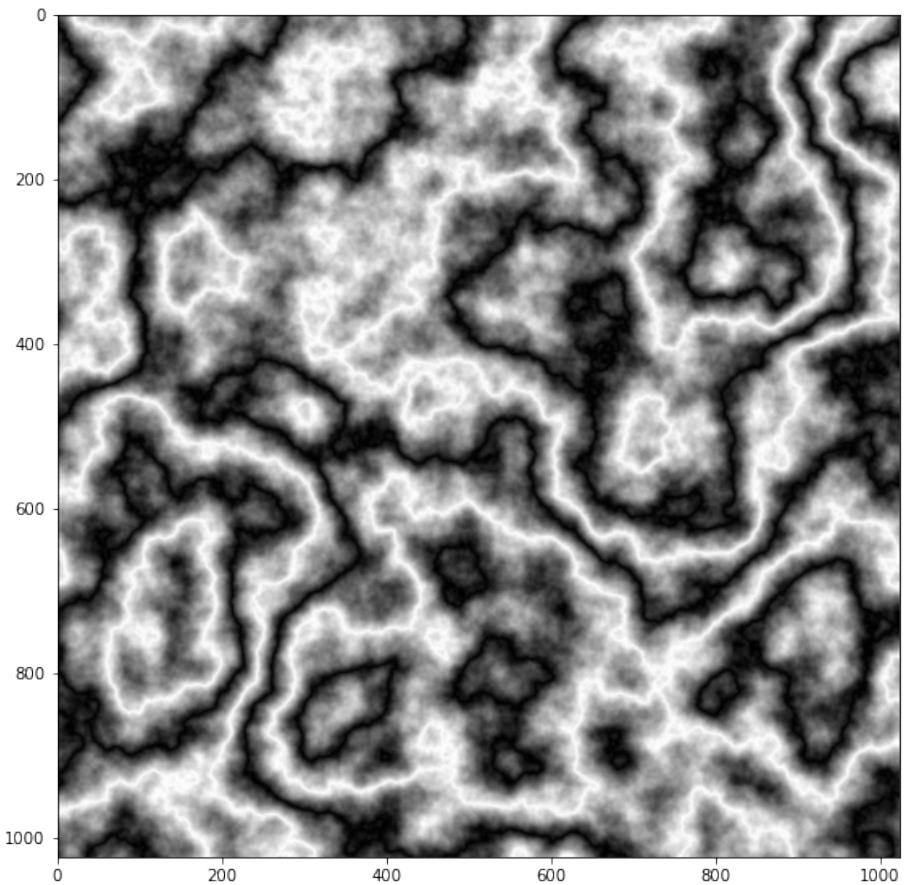
```
    plt.imshow(pfn**3, interpolation='lanczos');
```



```
1 pmin, pmax = np.min(pfn), np.max(pfn)
2 pfn = (pfn - pmin) / (pmax - pmin) # Rescale to exactly [0, 1] for thresholding
3 plt.imshow((0.0 < pfn) & (pfn < 1/3), interpolation='lanczos');
```



```
    plt.imshow(np.abs(((pfn**1)*64) % 16) - 8));
```



```
1 plt.hist(pfn.flat, 200)
2 plt.yscale('log');
```

