

1 Organized parallel simulations

```
1 import numpy as np
2 from multiprocessing import Pool
3 from scipy.signal import windows
4 import sys
5 import time
6 import os, struct # for `urandom`
7 import pprint # for parameters
8 import tempfile
9 import h5py, hickle
```

Since the Numba jitclass objects are not picklable, the relevant parameters to reconstruct a system are passed between processes. We require that all systems be accessible from the systems module.

```
1 if 'src' not in sys.path: sys.path.append('src')
2 import systems

1 def make_params(system):
2     return {system.__class__.__name__:
3             {k: v for k, v in zip(system.state_names(), system.state())}}
4
5 def make_system(system_params, system_prep = lambda x:x):
6     return system_prep([getattr(systems, cls)(**state)
7                         for cls, state in system_params.items()][0])

1 def make_psystems(params, psystem_prep): #:: params → (system → [system]) → [params]
2     log = params.get('log', False)
3     if log:
4         print('Finding parallel bin systems ... ', end='', flush=True)
5     psystems = psystem_prep(make_system(params['system'], **params['parallel']))
6     if log:
7         print('done.')
8     return [(make_params(s), *r) for s, *r in psystems]

1 def urandom_reseed():
2     """Reseeds numpy's RNG from `urandom` and returns the seed."""
3     seed = struct.unpack('I', os.urandom(4))[0]
4     np.random.seed(seed)
5     return seed
6
7 def worker(simulation, psystem, params):
8     log = params.get('log', False)
9     urandom_reseed()
```

```

10     psystem_params, *args = psystem
11     system = make_system(psystem_params)
12     if log:
13         print('(', end='', flush=True)
14         # Individual simulation output is too much when running parallel simulations.
15         params['simulation'].update({'log': False})
16         results = simulation(system, *args, **params['simulation'])
17     if log:
18         print(')', end='', flush=True)
19     return results
20
21 def show_params(params):
22     print('Run parameters')
23     print('-----')
24     pprint.pp(params, sort_dicts=False)
25     print()
26
27 def save_results(results, params, log=False, prefix='simulation-', dir='data'):
28     with tempfile.NamedTemporaryFile( # Note: dir shadows dir()
29         mode='wb', prefix=prefix, suffix='.h5', dir=dir, delete=False) as f:
30         with h5py.File(f, 'w') as hkl:
31             if log:
32                 print('Writing results ... ', end='', flush=True)
33             hickle.dump({
34                 'parameters': params,
35                 'results': results
36             }, hkl)
37             relpath = os.path.relpath(f.name)
38             if log:
39                 print('done: ', relpath)
40         return relpath
41
42 def run(params, simulation, system_prep,
43         psystem_prep = lambda x:x, result_wrapper = lambda x:x, **kwargs):
44     params.update(kwargs)
45     parallel = 'parallel' in params
46     log = params.get('log', False)
47     if log:
48         show_params(params)
49
50     if parallel:
51         psystems = make_psystems(params, psystem_prep)
52     else:
53         psystem = make_system(params['system'], system_prep)
54

```

```

55     if log:
56         if parallel:
57             print('Running || ', end='', flush=True)
58         else:
59             print('Running ...')
60         start_time = time.time()
61
62     if parallel:
63         with Pool() as pool:
64             results = pool.starmap(worker, ((simulation, args, params) for args in psystems))
65             results = [result_wrapper(r) for r in results]
66         else:
67             results = result_wrapper(simulation(*psystem, **params['simulation'], **kwargs))
68
69     if log:
70         seconds = int(time.time() - start_time)
71         if parallel:
72             print(' || done in', seconds, 'seconds.')
73         else:
74             print('... done in', seconds, 'seconds.')
75
76     # Save single-shot results in a singleton list so that we can analyze parallel and
77     # single results the same way.
78     rdict = {'results': results if parallel else [results]}
79     save_params = params.pop('save', False)
80     if save_params:
81         relpath = save_results(results, params, log, **save_params)
82         rdict.update({'file': relpath})
83
84     return rdict

```

We can choose overlapping bins for the parallel processes to negate boundary effects.

```

1  def extend_bin(bins, i, k = 0.05):
2      if len(bins) ≤ 2: # There is only one bin
3          return bins
4      k = max(0, min(1, k))
5      return (bins[i] - (k*(bins[i] - bins[i-1]) if 0 < i else 0),
6              bins[i+1] + (k*(bins[i+2] - bins[i+1]) if i < len(bins) - 2 else 0))

```

Often parallel results are the value of a real function on some grid or list of bins. Given that many of these pieces may overlap, we must combine them back together into a full solution. This requires first transforming the results so that

they are comparable, and then performing the combination. The most common case is repetition of the same real-valued experiment. No transformation is required, and we simply average all the results. Even better, we may assign the values within each piece a varying credence from 0 to 1 and perform weighted sums.

```

1 def join_results(xs, ys, wf = windows.hann):
2     xf = np.array(sorted(set().union(*xs)))
3     xi = [np.intersect1d(xf, x, return_indices=True)[1] for x in xs]
4
5     n, m = len(xf), len(xs)
6     ws = np.zeros((m, n))
7     wc = np.zeros((m, n))
8     for i in range(m):
9         l = len(xs[i])
10        ws[i, xi[i]] = wf(l)
11        wc[i, xi[i]] = np.ones(l)
12    unweighted = np.sum(wc, 0) > 0
13
14    Δys = np.zeros(m)
15    for i in range(m):
16        Σc = Σw = 0
17        for j in range(i):
18            a = Δys[j] * np.ones(n)
19            a[xi[j]] += ys[j]
20            a[xi[i]] -= ys[i]
21            w = ws[i, :] * ws[j, :]
22            Σc += np.dot(a, w)
23            Σw += np.sum(w)
24        Δys[i] = Σc / Σw if i > 0 else 0
25
26    yf = np.zeros(n)
27    for i in range(m):
28        w = ws[i, xi[i]]
29        # The weights are meaningful only as relative weights at overlap points.
30        # We must avoid division by zero at no-overlap points with weight zero.
31        # Note that overlap points with all weights zero will be an issue, as
32        # the weights in that situation are meaningless.
33        w[(w == 0) & unweighted[xi[i]]] = 1
34        yf[xi[i]] += (ys[i] + Δys[i]) * w
35        ws[i, xi[i]] = w
36    Σws = np.sum(ws, 0)
37    yf /= Σws
38
39    return xf, yf, Δys

```

Demonstration of joining overlapping results.

```
1  if __name__ == '__main__':
2      from matplotlib import pyplot as plt
3
4      testn = 50
5      xs = [np.arange(testn), np.arange(testn // 2, testn + testn // 2)]
6      ys = [np.arange(testn) - testn // 2, 4*np.arange(testn // 2, testn + testn // 2) + (testn -
7      ↪ 1)]
8
9      axs, ays, _ = join_results(xs, ys)
10
11     for x, y in zip(xs, ys):
12         plt.plot(x, y, 'black')
13     plt.plot(axs, ays, 'blue');
```

