

1 The Wang-Landau algorithm (density of states)

We determine thermodynamic quantities from the partition function by obtaining the density of states from a simulation.

TODO: * Improve stitching of parallel solutions

```
1 from numba import jit
2 nopython = True

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy import interpolate, special
```

Utility functions.

```
1 from bisect import bisect

1 @jit
2 def binindex(Es, E):
3     return bisect(Es, E, hi=len(Es) - 1) - 1

1 # @jit
2 # def flat(H, tol = 0.2):
3 #     """Determines if an evenly-spaced histogram is approximately flat."""
4 #     Hμ = np.mean(H)
5 #     Hf = np.max(H)
6 #     Hθ = np.min(H)
7 #     return Hf / (1 + tol) < Hμ < Hθ / (1 - tol)
8 @jit
9 def flat(H, tol = 0.2):
10     """Determines if an evenly-spaced histogram is approximately flat."""
11     Hμ = np.mean(H)
12     return not np.any(H < (1 - tol) * Hμ) and np.all(H ≠ 0)
```

1.1 Algorithm

A Wang-Landau algorithm, with quantities as logarithms and with monte-carlo steps proportional to $f^{-1/2}$ (a “Zhou-Bhat schedule”).

We use energy bins encoded by numbers E_i for $i \in [0, N]$, so that there are N bins. The energies E covered by bin i satisfy $E_i \leq E < E_{i+1}$. For the bounded discrete systems that we are considering, we must choose E_N to be an arbitrary number above the maximum energy.

```

1 @jit(forceobj=True)
2 def wanglandau(system,
3     Es,          # The energy bins
4     M = 1_00_000, # Monte carlo step scale
5     ε = 1e-8,     # f tolerance
6     logf0 = 1,    # Initial log f
7     logging = True, # Log progress of f-steps
8     flatness = 0.2 # Desired histogram flatness
9 ):
10     # Initial values
11     E0 = Es[0]
12     Ef = Es[-1]
13     ΔE = Es[1] - E0
14     N = len(Es) - 1
15     logf = logf0
16     logftol = np.log(1 + ε)
17     S = np.zeros(N) # Set all initial g's to 1
18     H = np.zeros(N, dtype=int)
19     i = binindex(Es, system.E)
20
21     if logging:
22         mciters = 0
23         fiter = 0
24         fiters = int(np.ceil(np.log2(logf0) - np.log2(logftol)))
25         print("Wang-Landau START:")
26         print("\t|Es| = {} \n\tM = {} \n\tε = {} \n\tlog f0 = {}".format(len(Es), M, ε, logf0))
27
28     while logftol < logf:
29         H[:] = 0
30         logf /= 2
31         iters = 0
32         niters = int((M + 1) * np.exp(-logf / 2))
33         if logging:
34             fiter += 1
35         while not flat(H, flatness) and iters < niters:
36             system.propose()
37             Ev = system.Ev
38             j = binindex(Es, Ev)
39             # if E0 ≤ Ev ≤ Ef and (
40             if E0 ≤ Ev < Ef and (
41                 S[j] < S[i] or np.random.rand() < np.exp(S[i] - S[j])):
42                 system.accept()
43                 i = j
44             H[i] += 1
45             S[i] += logf

```

```

46         iters += 1
47         if logging:
48             mciters += iters
49             print("f: {} / {} \t({} / {})".format(fiter, fitters, iters, niters))
50
51     if logging:
52         print("Done: {} total MC iterations.".format(mciters))
53     return Es, S, H

```

1.1.1 Parallel construction of the density of states

```

1 from multiprocessing import Pool
2 import copy

```

We can choose overlapping bins for the parallel processes to negate boundary effects.

```

1 def extend_bin(bins, i, k = 0.05):
2     if len(bins) ≤ 2: # There is only one bin
3         return bins
4     k = max(0, min(1, k))
5     return (bins[i] - (k*(bins[i] - bins[i-1]) if 0 < i else 0),
6             bins[i+1] + (k*(bins[i+2] - bins[i+1]) if i < len(bins) - 2 else 0))

```

Try monotonic instead of Wang-Landau steps

```

1 def find_bin_systems(sys, Es, Ebins, N = 1_000_000):
2     """Find systems with energies in the bins given by `Es` by stepping `sys`."""
3     # S = np.zeros(len(Es), dtype=int)
4     systems = [None] * (len(Ebins) - 1)
5     n = 0
6     i = binindex(Es, sys.E)
7     while any(system is None for system in systems) and n < N:
8         for s in range(len(systems)):
9             if systems[s] is None and Ebins[s] ≤ sys.E < Ebins[s + 1]:
10                 systems[s] = copy.deepcopy(sys)
11
12         sys.propose()
13         j = binindex(Es, sys.Ev)
14         if sys.E < sys.Ev:
15             sys.accept()
16         # if S[j] < S[i]:
17         #     i = j
18         # sys.accept()

```

```

19 #         S[i] += 1
20         n += 1
21
22 if N ≤ n:
23     raise ValueError('Could not find bin systems after {} iterations.'.format(N))
24 return systems

```

Now we can construct our parallel systems.

```

1 def parallel_systems(system, Es, n = 8, k = 0.1, N = 1_000_000):
2     Ebins = np.linspace(Es[0], Es[-1], n + 1)
3     systems = find_bin_systems(system, Es, Ebins, N)
4     binEs = [(lambda E0, Ef: Es[(E0 ≤ Es) & (Es ≤ Ef)])(*extend_bin(Ebins, i, k))
5               for i in range(len(Ebins) - 1)]
6     return zip(systems, binEs)

```

We also need a way to reset the random number generator seed in a way that is time-independent and different for each process.

```

1 import os, struct

1 def urandom_reseed():
2     """Reseeds numpy's RNG from `urandom` and returns the seed"""
3     seed = struct.unpack('I', os.urandom(4))[0]
4     np.random.seed(seed)
5     return seed

```

Once we have parallel results, we stitch the pieces of $\ln g(E)$ together.

```

1 def stitch_results(wlresults):
2     E0, S0, _ = wlresults[0]
3     E, S = E0, S0
4     for i in range(1, len(wlresults)):
5         Ev, Sv, _ = wlresults[i]
6         # Assumes overlap is at end regions
7         _, i0s, ivs = np.intersect1d(E0[:-1], Ev[:-1], return_indices=True)
8         # Simplest: join middles of overlap regions
9         l = len(i0s)
10        m = l // 2
11        #         print(l, m, i0s, ivs, i0s[m], S0, Sv)
12        Sv -= Sv[ivs[m]] - S0[i0s[m]]
13        # Simplest: average the overlaps to produce the final value
14        E = np.hstack((E, Ev[1+1:]))
15        S[-1:] = (Sv[ivs] + S0[i0s]) / 2
16        S = np.hstack((S, Sv[1:]))
17        E0, S0 = Ev, Sv
18    return E, S

```

1.2 The 2D Ising model

```
1 class Ising:
2     def __init__(self, n):
3         self.n = n
4         self.spins = np.sign(np.random.rand(n, n) - 0.5)
5         self.E = self.energy()
6         self.Ev = self.E
7     def neighbors(self, i, j):
8         return np.hstack([self.spins[:,j].take([i-1,i+1], mode='wrap'),
9                             self.spins[i,:].take([j-1,j+1], mode='wrap')])
10    def energy(self):
11        return -0.5 * sum(np.sum(s * self.neighbors(i, j))
12                            for (i, j), s in np.ndenumerate(self.spins))
13    def propose(self):
14        i, j = np.random.randint(self.n), np.random.randint(self.n)
15        self.i, self.j = i, j
16        dE = 2 * np.sum(self.spins[i, j] * self.neighbors(i, j))
17        self.dE = dE
18        self.Ev = self.E + dE
19    def accept(self):
20        self.spins[self.i, self.j] *= -1
21        self.E = self.Ev
```

Note that this class-based approach adds some overhead. For speed, instances of Ising should be inlined into the wanglandau.

1.2.1 Simulation

```
1 isingn = 32
2 sys = Ising(isingn)
```

The Ising energies over the full range, with correct end bin. We remove the penultimate energies since $E = 2$ or $E_{\max} - 2$ cannot happen.

```
1 isingE0 = -2 * isingn**2
2 isingEf = 2 * isingn**2
3 isingΔE = 4
4 Es = np.arange(isingE0, isingEf + isingΔE + 1, isingΔE)
5 Es = np.delete(np.delete(Es, -3), 1)

1 psystems = parallel_systems(sys, Es, n = 16, k = 0.5, N = 50_000_000)

1 def parallel_wanglandau(subsystem): # Convenient form for `Pool.map`
2     urandom_reseed()
```

```

3     results = wanglandau(*subsystem, M = 1_000_000, logging=False)
4     print('*', end='', flush=True)
5     return results

1 with Pool() as pool:
2     wlresults = pool.map(parallel_wanglandau, psystems)

1 sEs, sS = stitch_results(wlresults)

1 for Es, S, H in wlresults:
2     plt.plot(Es[:-1], S)

1 plt.plot(sEs[:-1], sS);

1 import os, tempfile, pickle

1 with tempfile.NamedTemporaryFile(mode='wb', prefix='wlresults-ising-', suffix='.pickle',
  ↳ dir='data', delete=False) as f:
2     print(os.path.basename(f.name))
3     pickle.dump(wlresults, f)
4     pickle.dump(sEs, f)
5     pickle.dump(sS, f)

```

1.2.2 Calculating canonical ensemble averages

```

1 gspl = interpolate.splrep(Es, S, s=2*np.sqrt(2))
2 gs = np.exp(interpolate.splev(Es, gspl) - min(S))

1 plt.plot(Es / isingn**2, S)
2 plt.plot(Es / isingn**2, interpolate.splev(Es, gspl))
3 plt.xlabel("E / N")
4 plt.ylabel("log g(E) + C");

```

Translate energies to have minimum zero so that Z is representable.

```

1 nEs = Es - min(Es)

1 Z = lambda beta: np.sum(gs * np.exp(-beta * nEs))

```

Ensemble averages

```

1 beta = [np.exp(k) for k in np.linspace(-3, 1, 200)]
2 Eμ = lambda beta: np.sum(nEs * gs * np.exp(-beta * nEs)) / Z(beta)
3 E2 = lambda beta: np.sum(nEs**2 * gs * np.exp(-beta * nEs)) / Z(beta)
4 CV = lambda beta: (E2(beta) - Eμ(beta)**2) * beta**2
5 F = lambda beta: -np.log(Z(beta)) / beta
6 Sc = lambda beta: beta * Eμ(beta) + np.log(Z(beta))

```

Heat capacity

```
1 plt.plot(np.log(βs), [CV(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("Heat capacity")
4 plt.show()
```

Entropy

```
1 plt.plot(np.log(βs), [Sc(β) for β in βs])
2 plt.xlabel("ln β")
3 plt.ylabel("S(β) + C")
4 plt.show()
```

1.3 Thermal calculations on images


```
1 # @jitclass
2 class StatisticalImage:
3     def __init__(self, I0, M = 2**8 - 1):
4         self.I0 = I0
5         self.I = I0.copy()
6         self.N = len(I0)
7         self.M = M
8         self.E = self.energy()
9         self.Ev = self.E
10    def energy(self):
11        return np.sum(np.abs(self.I - self.I0))
12    def propose(self):
13        i = np.random.randint(self.N)
14        self.i = i
15        x0 = self.I0[i]
16        x = self.I[i]
17        r = np.random.randint(2)
18        if x == 0:
19            dx = r
20        elif x == self.M:
21            dx = -r
22        else:
23            dx = 2*r - 1
24        dE = dx
25        self.dx = dx
26        self.dE = dE
27        self.Ev = self.E + dE
28    def accept(self):
29        self.I[self.i] += self.dx
30        self.E = self.Ev
```

```

1 N = 3
2 M = 3
3 sys = StatisticalImage(np.zeros(N, dtype=int), M)
4 Es = np.arange(0, N*M + 1 + 1)
5 exactS = np.log(exact_bw_gs(N, M)[1])

1 Es, S, H = wanglandau(sys, Es, M = 100_000, ε = 1e-8, flatness = 0.01, logging=False)
2 S -= np.min(S)
3 plt.plot(Es[:-1], S)
4 plt.plot(Es[:-1], exactS);

```



1.3.1 Parallel Simulation

```

1 N = 16
2 M = 2**5 - 1
3 sys = StatisticalImage(np.zeros(N, dtype=int), M) # BW
4 Es = np.arange(N*M + 1 + 1)
5 psystems = parallel_systems(sys, Es, n = 8, k = 0.5, N = 1_000_000)

1 def parallel_wanglandau(subsystem): # Convenient form for `Pool.map`
2     urandom_reseed()
3     results = wanglandau(*subsystem, M = 10_000_000, ε = 1e-10, logging=False)
4     print('*', end='', flush=True)
5     return results

```



```

1 with Pool() as pool:
2     wlresults = pool.map(parallel_wanglandau, psystems)

    *****

1 sEs, sS = stitch_results(wlresults)


1 import os, tempfile, pickle

1 with tempfile.NamedTemporaryFile(mode='wb', prefix='wlresults-image-', suffix='.pickle',
  ↳ dir='data', delete=False) as f:
2     print(os.path.basename(f.name))
3     pickle.dump(N, f)
4     pickle.dump(M, f)
5     pickle.dump(wlresults, f)

    wlresults-image-x7_3j6tp.pickle

1 for Es, S, H in wlresults:
2     plt.plot(Es[:-1], S)

```



```

1 wlEs, S = sEs[:-1], sS

```

Fit a spline to interpolate and optionally clean up noise, giving WL g's up to a normalization constant.

```

1 gspl = interpolate.splrep(wlEs, S, s=0*np.sqrt(2))
2 wlgs = np.exp(interpolate.splev(wlEs, gspl) - min(S))

```

1.3.2 Exact solution

We only compute to halfway since g is symmetric and the other half's large numbers cause numerical instability.

```
1 def reflect(a, center=True):
2     if center:
3         return np.hstack([a[:-1], a[-1], a[-2::-1]])
4     else:
5         return np.hstack([a, a[::-1]])
```

The exact density of states for uniform values. This covers the all gray and all black/white cases. Everything else (normal images) are somewhere between. The gray is a slight approximation: the ground level is not degenerate, but we say it has degeneracy 2 like all the other sites. For the numbers of sites and values we are using, this is insignificant.

```
1 def bw_g(E, N, M, exact=True):
2     return sum((-1)**k * special.comb(N, k, exact=exact) * special.comb(E + N - 1 - k*(M + 1), E
3         ↪ - k*(M + 1), exact=exact)
4         for k in range(int(E / M) + 1))
5 def exact_bw_gs(N, M):
6     Es = np.arange(N*M + 1)
7     gs = np.vectorize(bw_g)(np.arange(1 + N*M // 2), N, M, exact=False)
8     return Es, reflect(gs, len(Es) % 2 == 1)
9
10 def gray_g(E, N, M, exact=True):
11     return 2 * bw_g(E, N, M, exact=exact)
12 def exact_gray_gs(N, M):
13     Es = np.arange(N*M + 1)
14     gs = np.vectorize(gray_g)(np.arange(1 + N*M // 2), N, M, exact=False)
15     return Es, reflect(gs, len(Es) % 2 == 1)
```

Expected results for black/white and gray.

```
1 bw_Es, bw_gs = exact_bw_gs(N=N, M=M)
2 gray_Es, gray_gs = exact_gray_gs(N=N, M=-1 + (M + 1) // 2)
```

Choose what to compare to.


```
1 Es, gs = bw_Es, bw_gs
```

Presumably all of the densities of states for different images fall in the region between the all-gray and all-black/white curves.

```

1 plt.plot(bw_Es / len(bw_Es), np.log(bw_gs), 'black', label='BW')
2 plt.plot(gray_Es / len(gray_Es), np.log(gray_gs), 'gray', label='Gray')
3 plt.plot(wl_Es / len(wl_Es), np.log(wl_gs), label='WL')
4 plt.xlabel('E / MN')
5 plt.ylabel('ln g')
6 plt.title('L = {}'.format(L))
7 plt.legend();

```




```

1 wlgs[[0,-1]]

array([1.          , 1.05022784])

1 plt.plot(wl_Es / len(wl_Es), np.abs(wlgs - bw_gs) / bw_gs)
2 plt.title('Relative error')
3 plt.xlabel('E / MN')
4 plt.ylabel('ε(S)');

```



1.3.3 Calculating canonical ensemble averages

```

1  class CanonicalEnsemble:
2      def __init__(self, Es, gs, name):
3          self.Es = Es
4          self.gs = gs
5          self.name = name
6      def Z(self, β):
7          return np.sum(self.gs * np.exp(-β * self.Es))
8      def average(self, f, β):
9          return np.sum(f(self) * self.gs * np.exp(-β * self.Es)) / self.Z(β)
10     def energy(self, β):
11         return self.average(lambda ens: ens.Es, β)
12     def energy2(self, β):
13         return self.average(lambda ens: ens.Es**2, β)
14     def heat_capacity(self, β):
15         return self.energy2(β) - self.energy(β)**2
16     def free_energy(self, β):
17         return -np.log(self.Z(β)) / β
18     def entropy(self, β):
19         return β * self.energy(β) + np.log(self.Z(β))

```

```

1   $\beta$ s = [np.exp(k) for k in np.linspace(-8, 2, 500)]
2  wlens = CanonicalEnsemble(wlEs, wlgs, 'WL') # Wang-Landau results
3  xens = CanonicalEnsemble(Es, gs, 'Exact') # Exact
4  ensembles = [wlens, xens]


```

Partition function

```

1  for ens in ensembles:
2      plt.plot(-np.log( $\beta$ s), np.log(np.vectorize(ens.Z)( $\beta$ s)), label=ens.name)
3  plt.xlabel("ln kT")
4  plt.ylabel("Log partition function")
5  plt.title('L = {}'.format(L))
6  plt.legend();

```




Helmholtz free energy

```

1  for ens in ensembles:
2      plt.plot(-np.log( $\beta$ s), np.vectorize(ens.free_energy)( $\beta$ s), label=ens.name)
3  plt.xlabel("ln kT")
4  plt.ylabel("Helmholtz free energy")
5  plt.title('L = {}'.format(L))
6  plt.legend();

```




Heat capacity

```

1  for ens in ensembles:
2      plt.plot(-np.log(βs), np.vectorize(ens.heat_capacity)(βs), label=ens.name)
3  plt.xlabel("ln kT")
4  plt.ylabel("Heat capacity")
5  plt.title('L = {}'.format(L))
6  plt.legend();

```



Entropy

```

1  for ens in ensembles:
2      plt.plot(-np.log( $\beta$ s), np.vectorize(ens.entropy)( $\beta$ s), label=ens.name)
3  plt.xlabel("ln kT")
4  plt.ylabel("Canonical entropy")
5  plt.title('L = {}'.format(L))
6  plt.legend();

```

